

Example scripts for User-space, Statically Defined Tracing (USDT)

This directory contains scripts showcasing User-space, Statically Defined Tracing (USDT) support for Bitcoin Core on Linux using. For more information on USDT support in Bitcoin Core see the [USDT documentation](#).

Examples for the two main eBPF front-ends, [bpftrace](#) and [BPF Compiler Collection \(BCC\)](#), with support for USDT, are listed. BCC is used for complex tools and daemons and `bpftrace` is preferred for one-liners and shorter scripts.

To develop and run bpftrace and BCC scripts you need to install the corresponding packages. See [installing bpftrace](#) and [installing BCC](#) for more information. For development there exist a [bpftrace Reference Guide](#), a [BCC Reference Guide](#), and a [bcc Python Developer Tutorial](#).

Examples

The bpftrace examples contain a relative path to the `bitcoind` binary. By default, the scripts should be run from the repository-root and assume a self-compiled `bitcoind` binary. The paths in the examples can be changed, for example, to point to release builds if needed. See the [Bitcoin Core USDT documentation](#) on how to list available tracepoints in your `bitcoind` binary.

WARNING: eBPF programs require root privileges to be loaded into a Linux kernel VM. This means the bpftrace and BCC examples must be executed with root privileges. Make sure to carefully review any scripts that you run with root privileges first!

log_p2p_traffic.bt

A bpftrace script logging information about inbound and outbound P2P network messages. Based on the `net:inbound_message` and `net:outbound_message` tracepoints.

By default, `bpftrace` limits strings to 64 bytes due to the limited stack size in the eBPF VM. For example, Tor v3 addresses exceed the string size limit which results in the port being cut off during logging. The string size limit can be increased with the `BPFTRACE_STRLLEN` environment variable (`BPFTRACE_STRLLEN=70` works fine).

```
$ bpftrace contrib/tracing/log_p2p_traffic.bt
```

Output

```
outbound 'ping' msg to peer 11 (outbound-full-relay,
[2a02:b10c:f747:1:ef:fake:ipv6:addr]:8333) with 8 bytes
inbound 'pong' msg from peer 11 (outbound-full-relay,
[2a02:b10c:f747:1:ef:fake:ipv6:addr]:8333) with 8 bytes
inbound 'inv' msg from peer 16 (outbound-full-relay, XX.XX.XXX.121:8333) with 37 bytes
outbound 'getdata' msg to peer 16 (outbound-full-relay, XX.XX.XXX.121:8333) with 37
bytes
inbound 'tx' msg from peer 16 (outbound-full-relay, XX.XX.XXX.121:8333) with 222 bytes
outbound 'inv' msg to peer 9 (outbound-full-relay,
faketorv3addressa2ufa6odvoi3s77j4uegey0xb10csyfyve2t33curbyd.onion:8333) with 37 bytes
outbound 'inv' msg to peer 7 (outbound-full-relay, XX.XX.XXX.242:8333) with 37 bytes
...
```

p2p_monitor.py

A BCC Python script using cures for an interactive P2P message monitor. Based on the `net:inbound_message` and `net:outbound_message` tracepoints.

Inbound and outbound traffic is listed for each peer together with information about the connection. Peers can be selected individually to view recent P2P messages.

```
$ python3 contrib/tracing/p2p_monitor.py ./src/bitcoind
```

Lists selectable peers and traffic and connection information.

```

P2P Message Monitor
Navigate with UP/DOWN or J/K and select a peer with ENTER or SPACE to see individual
P2P messages

PEER   OUTBOUND                               INBOUND                               TYPE                                ADDR
  0  46                398 byte  61      1407590 byte  block-relay-only
XX.XX.XXX.196:8333
  11 1156            253570 byte  3431    2394924 byte  outbound-full-relay
XXX.X.XX.179:8333
  13 3425            1809620 byte  1236      305458 byte  inbound
XXX.X.X.X:60380
  16 1046            241633 byte  1589    1199220 byte  outbound-full-relay
4faketorv2pbfu7x.onion:8333
  19 577             181679 byte  390      148951 byte  outbound-full-relay
kfake4vctorjv2o2.onion:8333
  20 11              1248 byte  13        1283 byte  block-relay-only
[2600:fake:64d9:b10c:4436:aaaa:fe:bb]:8333
  21 11              1248 byte  13        1299 byte  block-relay-only
XX.XXX.X.155:8333
  22 5               103 byte  1          102 byte  feeler
XX.XX.XXX.173:8333
  23 11              1248 byte  12        1255 byte  block-relay-only
XX.XXX.XXX.220:8333
  24 3               103 byte  1          102 byte  feeler
XXX.XXX.XXX.64:8333
...

```

Showing recent P2P messages between our node and a selected peer.

[illegible]

```

| headers (1 byte) ---> |
|                               <--- feefilter (8 bytes) |
|                               <--- pong (8 bytes) |
|                               <--- headers (82 bytes) |
|                               <--- addr (30003 bytes) |
| inv (1261 byte) ---> |
|                               ... |

```

log_raw_p2p_msgs.py

A BCC Python script showcasing eBPF and USDT limitations when passing data larger than about 32kb. Based on the `net:inbound_message` and `net:outbound_message` tracepoints.

Bitcoin P2P messages can be larger than 32kb (e.g. `tx`, `block`, ...). The eBPF VM's stack is limited to 512 bytes, and we can't allocate more than about 32kb for a P2P message in the eBPF VM. The **message data is cut off** when the message is larger than `MAX_MSG_DATA_LENGTH` (see script). This can be detected in user-space by comparing the data length to the message length variable. The message is cut off when the data length is smaller than the message length. A warning is included with the printed message data.

Data is submitted to user-space (i.e. to this script) via a ring buffer. The throughput of the ring buffer is limited. Each `p2p_message` is about 32kb in size. In- or outbound messages submitted to the ring buffer in rapid succession fill the ring buffer faster than it can be read. Some messages are lost. BCC prints: `Possibly lost 2 samples` on lost messages.

```
$ python3 contrib/tracing/log_raw_p2p_msgs.py ./src/bitcoind
```

```

Logging raw P2P messages.
Messages larger than about 32kb will be cut off!
Some messages might be lost!
  outbound msg 'inv' from peer 4 (outbound-full-relay, XX.XXX.XX.4:8333) with 253
bytes: 0705000000be2245c8f844c9f763748e1a7...
...
Warning: incomplete message (only 32568 out of 53552 bytes)! inbound msg 'tx' from
peer 32 (outbound-full-relay, XX.XXX.XXX.43:8333) with 53552 bytes:
020000000001fd3c01939c85ad6756ed9fc...
...
Possibly lost 2 samples

```

connectblock_benchmark.bt

A `bpftrace` script to benchmark the `ConnectBlock()` function during, for example, a blockchain re-index. Based on the `validation:block_connected` USDT tracepoint.

The script takes three positional arguments. The first two arguments, the start, and end height indicate between which blocks the benchmark should be run. The third acts as a duration threshold in milliseconds. When the `ConnectBlock()` function takes longer than the threshold, information about the block, is printed. For more details, see the header comment in the script.

The following command can be used to benchmark, for example, `ConnectBlock()` between height 20000 and 38000 on SigNet while logging all blocks that take longer than 25ms to connect.

```
$ bpftrace contrib/tracing/connectblock_benchmark.bt 20000 38000 25
```

```
$ ./src/bitcoind -signet -reindex
```

```
Attaching 5 probes...
ConnectBlock Benchmark between height 20000 and 38000 inclusive
Logging blocks taking longer than 25 ms to connect.
Starting Connect Block Benchmark between height 20000 and 38000.
```

| BENCH | 39 blk/s | 59 tx/s | 59 inputs/s | 20 sigops/s | (height 20038) |
|-------------|---|-------------|----------------|----------------|----------------|
| Block 20492 | (000000f555653bb05e2f3c6e79925e01a20dd57033f4dc7c354b46e34735d32b) | | | | 20 |
| tx | 2319 ins | 2318 sigops | took | 38 ms | |
| BENCH | 1840 blk/s | 2117 tx/s | 4478 inputs/s | 2471 sigops/s | (height 21879) |
| BENCH | 1816 blk/s | 4972 tx/s | 4982 inputs/s | 125 sigops/s | (height 23695) |
| BENCH | 2095 blk/s | 2890 tx/s | 2910 inputs/s | 152 sigops/s | (height 25790) |
| BENCH | 1684 blk/s | 3979 tx/s | 4053 inputs/s | 288 sigops/s | (height 27474) |
| BENCH | 1155 blk/s | 3216 tx/s | 3252 inputs/s | 115 sigops/s | (height 28629) |
| BENCH | 1797 blk/s | 2488 tx/s | 2503 inputs/s | 111 sigops/s | (height 30426) |
| BENCH | 1849 blk/s | 6318 tx/s | 6569 inputs/s | 12189 sigops/s | (height 32275) |
| BENCH | 946 blk/s | 20209 tx/s | 20775 inputs/s | 83809 sigops/s | (height 33221) |
| Block 33406 | (0000002adfe4a15cfcd53bd890a89bbae836e5bb7f38bac566f61ad4548c87f6) | | | | 25 |
| tx | 2045 ins | 2090 sigops | took | 29 ms | |
| Block 33687 | (00000073231307a9828e5607ceb8156b402efe56747271a4442e75eb5b77cd36) | | | | 52 |
| tx | 1797 ins | 1826 sigops | took | 26 ms | |
| BENCH | 582 blk/s | 21581 tx/s | 27673 inputs/s | 60345 sigops/s | (height 33803) |
| BENCH | 1035 blk/s | 19735 tx/s | 19776 inputs/s | 51355 sigops/s | (height 34838) |
| Block 35625 | (0000006b00b347390c4768ea9df2655e9ff4b120f29d78594a2a702f8a02c997) | | | | 20 |
| tx | 3374 ins | 3371 sigops | took | 49 ms | |
| BENCH | 887 blk/s | 17857 tx/s | 22191 inputs/s | 24404 sigops/s | (height 35725) |
| Block 35937 | (000000d816d13d6e39b471cd4368db60463a764ba1f29168606b04a22b81ea57) | | | | 75 |
| tx | 3943 ins | 3940 sigops | took | 61 ms | |
| BENCH | 823 blk/s | 16298 tx/s | 21031 inputs/s | 18440 sigops/s | (height 36548) |
| Block 36583 | (000000c3e260556dbf42968aae3f904dba8b8c1fff96a6f6e3aa5365d2e3ad317) | | | | 24 |
| tx | 2198 ins | 2194 sigops | took | 34 ms | |
| Block 36700 | (000000b3b173de9e65a3cfa738d976af6347aaf83fal7ab3f2a4d2ede3ddfacc4) | | | | 73 |
| tx | 1615 ins | 1611 sigops | took | 31 ms | |
| Block 36832 | (0000007859578c02clac37dabd1b9ec19b98f350b56935f5dd3a41e9f79f836e) | | | | 34 |
| tx | 1440 ins | 1436 sigops | took | 26 ms | |
| BENCH | 613 blk/s | 16718 tx/s | 25074 inputs/s | 23022 sigops/s | (height 37161) |
| Block 37870 | (000000f5c1086291ba2d943fb0c3bc82e71c5ee341ee117681d1456fbf6c6c38) | | | | 25 |
| tx | 1517 ins | 1514 sigops | took | 29 ms | |
| BENCH | 811 blk/s | 16031 tx/s | 20921 inputs/s | 18696 sigops/s | (height 37972) |

Took 14055 ms to connect the blocks between height 20000 and 38000.

Histogram of block connection times in milliseconds (ms).

@durations:

| [0] | 16838 | @@@ |
|--------|-------------------|---|
| [1] | 882 <th>@@</th> | @@ |
| [2, 4) | 236 <th> </th> | |
| [4, 8) | 23 <th> </th> | |

| | | | |
|----------|---|--|--|
| [8, 16) | 9 | | |
| [16, 32) | 9 | | |
| [32, 64) | 4 | | |

log_utxocache_flush.py

A BCC Python script to log the UTXO cache flushes. Based on the `utxocache:flush` tracepoint.

```
$ python3 contrib/tracing/log_utxocache_flush.py ./src/bitcoind
```

Logging utxocache flushes. Ctrl-C to end...

| Duration (µs) | Mode | Coins Count | Memory Usage | Prune |
|---------------|-----------|-------------|--------------|-------|
| 730451 | IF_NEEDED | 22990 | 3323.54 kB | True |
| 637657 | ALWAYS | 122320 | 17124.80 kB | False |
| 81349 | ALWAYS | 0 | 1383.49 kB | False |

log_utxos.bt

A `bpftrace` script to log information about the coins that are added, spent, or uncached from the UTXO set.

Based on the `utxocache:add`, `utxocache:spend` and `utxocache:uncache` tracepoints.

```
$ bpftrace contrib/tracing/log_utxos.bt
```

This should produce an output similar to the following. If you see `bpftrace` warnings like `Lost 24 events`, the eBPF perf ring-buffer is filled faster than it is being read. You can increase the ring-buffer size by setting the ENV variable `BPFTRACE_PERF_RB_PAGES` (default 64) at a cost of higher memory usage. See the [bpftrace reference guide](#) for more information.

```
Attaching 4 probes...
OP      Outpoint
Value   Height Coinbase
Added   6ba9ad857e1ef2eb2a2c94f06813c414c7ab273e3d6bd7ad64e000315a887e7c:1
10000  2094512 No
Spent   fa7dc4db56637a151f6649db8f26732956d1c5424c82aae400a83d02b2cc2c87b:0
182264897 2094512 No
Added   eeb2f099b1af6a2a12e6ddd2eeb16fc5968582241d7f08ba202d28b60ac264c7:0
10000  2094512 No
Added   eeb2f099b1af6a2a12e6ddd2eeb16fc5968582241d7f08ba202d28b60ac264c7:1
182254756 2094512 No
Added   a0c7f4ec9cccef2d89672a624a4e6c8237a17572efdd4679eea9e9ee70d2db04:0
10072679 2094513 Yes
Spent   25e0df5cc1aeb1b78e6056bf403e5e8b7e41f138060ca0a50a50134df0549a5e:2
540    2094508 No
Spent   42f383c04e09c26a2378272ec33aa0c1bf4883ca5ab739e8b7e06be5a5787d61:1
3848399 2007724 No
Added   f85e3b4b89270863a389395cc9a4123e417ab19384cef96533c6649abd6b0561:0
3788399 2094513 No
Added   f85e3b4b89270863a389395cc9a4123e417ab19384cef96533c6649abd6b0561:2
540    2094513 No
Spent   a05880b8c77971ed0b9f73062c7c4cdb0ff3856ab14cbf8bc481ed571cd34b83:1
```

5591281046 2094511 No

Added eb689865f7d957938978d6207918748f74e6aa074f47874724327089445b0960:0

5589696005 2094513 No

Added eb689865f7d957938978d6207918748f74e6aa074f47874724327089445b0960:1

1565556 2094513 No