

Intel(R) Dynamic Platform and Thermal Framework Sysfs Interface

Copyright: |copy| 2022 Intel Corporation
Author: Srinivas Pandruvada <srinivas.pandruvada@linux.intel.com>

Introduction

Intel(R) Dynamic Platform and Thermal Framework (DPTF) is a platform level hardware/software solution for power and thermal management.

As a container for multiple power/thermal technologies, DPTF provides a coordinated approach for different policies to effect the hardware state of a system.

Since it is a platform level framework, this has several components. Some parts of the technology is implemented in the firmware and uses ACPI and PCI devices to expose various features for monitoring and control. Linux has a set of kernel drivers exposing hardware interface to user space. This allows user space thermal solutions like "Linux Thermal Daemon" to read platform specific thermal and power tables to deliver adequate performance while keeping the system under thermal limits.

DPTF ACPI Drivers interface

`:file: '/sys/bus/platform/devices/<N>/uuids'`, where <N> =INT3400|INTC1040|INTC1041|INTC10A0

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\linux-master\Documentation\driver-api\thermal\[linux-master] [Documentation] [driver-api] [thermal] intel_dptf.rst, line 33); [backlink](#)

Unknown interpreted text role "file".

`available_uuids (RO)`

A set of UUIDs strings presenting available policies which should be notified to the firmware when the user space can support those policies.

UUID strings:

"42A441D6-AE6A-462b-A84B-4A8CE79027D3": Passive 1
"3A95C389-E4B8-4629-A526-C52C88626BAE": Active
"97C68AE7-15FA-499c-B8C9-5DA81D606E0A": Critical
"63BE270F-1C11-48FD-A6F7-3AF253FF3E2D": Adaptive performance
"5349962F-71E6-431D-9AE8-0A635B710AEE": Emergency call
"9E04115A-AE87-4D1C-9500-0F3E340BFE75": Passive 2
"F5A35014-C209-46A4-993A-EB56DE7530A1": Power Boss
"6ED722A7-9240-48A5-B479-31EEF723D7CF": Virtual Sensor
"16CAF1B7-DD38-40ED-B1C1-1B8A1913D531": Cooling mode
"BE84BABF-C4D4-403D-B495-3128FD44dAC1": HDC

`current_uuid (RW)`

User space can write strings from available UUIDs, one at a time.

`:file: '/sys/bus/platform/devices/<N>/'`, where <N> =INT3400|INTC1040|INTC1041|INTC10A0

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\linux-master\Documentation\driver-api\thermal\[linux-master] [Documentation] [driver-api] [thermal] intel_dptf.rst, line 67); [backlink](#)

Unknown interpreted text role "file".

`imok (WO)`

User space daemon write 1 to respond to firmware event for sending keep alive notification. User space receives THERMAL_EVENT_KEEP_ALIVE kobject uevent notification when firmware calls for user space to respond with imok ACPI method.

odvp* (RO)

Firmware thermal status variable values. Thermal tables calls for different processing based on these variable values.

data_vault (RO)

Binary thermal table. Refer to https://github.com/intel/thermal_daemon for decoding thermal table.

ACPI Thermal Relationship table interface

`:file:`/dev/acpi_thermal_rel``

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\linux-master\Documentation\driver-api\thermal\[linux-master] [Documentation] [driver-api] [thermal]intel_dptf.rst, line 91); [backlink](#)

Unknown interpreted text role "file".

This device provides IOCTL interface to read standard ACPI thermal relationship tables via ACPI methods _TRT and _ART. These IOCTLs are defined in drivers/thermal/intel/int340x_thermal/acpi_thermal_rel.h

IOCTLs:

ACPI_THERMAL_GET_TRT_LEN: Get length of TRT table

ACPI_THERMAL_GET_ART_LEN: Get length of ART table

ACPI_THERMAL_GET_TRT_COUNT: Number of records in TRT table

ACPI_THERMAL_GET_ART_COUNT: Number of records in ART table

ACPI_THERMAL_GET_TRT: Read binary TRT table, length to read is provided via argument to ioctl().

ACPI_THERMAL_GET_ART: Read binary ART table, length to read is provided via argument to ioctl().

DPTF ACPI Sensor drivers

DPTF Sensor drivers are presented as standard thermal sysfs thermal_zone.

DPTF ACPI Cooling drivers

DPTF cooling drivers are presented as standard thermal sysfs cooling_device.

DPTF Processor thermal PCI Driver interface

`:file:`/sys/bus/pci/devices/0000:00:04.0/power_limits``

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\linux-master\Documentation\driver-api\thermal\[linux-master] [Documentation] [driver-api] [thermal]intel_dptf.rst, line 129); [backlink](#)

Unknown interpreted text role "file".

Refer to Documentation/power/powercap/powercap.rst for powercap ABI.

power_limit_0_max_uw (RO)

Maximum powercap sysfs constraint_0 power_limit_uw for Intel RAPL

power_limit_0_step_uw (RO)

Power limit increment/decrements for Intel RAPL constraint 0 power limit

power_limit_0_min_uw (RO)

Minimum powercap sysfs constraint_0 power_limit_uw for Intel RAPL

power_limit_0_tmin_us (RO)

Minimum powercap sysfs constraint_0 time_window_us for Intel RAPL

power_limit_0_tmax_us (RO)

Maximum powercap sysfs constraint_0 time_window_us for Intel RAPL

power_limit_1_max_uw (RO)

Maximum powercap sysfs constraint_1 power_limit_uw for Intel RAPL

power_limit_1_step_uw (RO)

Power limit increment/decrements for Intel RAPL constraint 1 power limit

power_limit_1_min_uw (RO)

Minimum powercap sysfs constraint_1 power_limit_uw for Intel RAPL

power_limit_1_tmin_us (RO)

Minimum powercap sysfs constraint_1 time_window_us for Intel RAPL

power_limit_1_tmax_us (RO)

Maximum powercap sysfs constraint_1_time_window_us for Intel RAPL

`:file:`/sys/bus/pci/devices/0000:00:04.0/``

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\linux-master\Documentation\driver-api\thermal\[linux-master] [Documentation] [driver-api] [thermal]intel_dptf.rst, line 164); [backlink](#)

Unknown interpreted text role "file".

tcc_offset_degree_celsius (RW)

TCC offset from the critical temperature where hardware will throttle CPU.

`:file:`/sys/bus/pci/devices/0000:00:04.0/workload_request``

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\linux-master\Documentation\driver-api\thermal\[linux-master] [Documentation] [driver-api] [thermal]intel_dptf.rst, line 170); [backlink](#)

Unknown interpreted text role "file".

workload_available_types (RO)

Available workload types. User space can specify one of the workload type it is currently executing via workload_type. For example: idle, bursty, sustained etc.

workload_type (RW)

User space can specify any one of the available workload type using this interface.

DPTF Processor thermal RFIM interface

RFIM interface allows adjustment of FIVR (Fully Integrated Voltage Regulator) and DDR (Double Data Rate) frequencies to avoid RF interference with WiFi and 5G.

Switching voltage regulators (VR) generate radiated EMI or RFI at the fundamental frequency and its harmonics. Some harmonics may interfere with very sensitive wireless receivers such as Wi-Fi and cellular that are integrated into host systems like notebook PCs. One of mitigation methods is requesting SOC integrated VR (IVR) switching frequency to a small % and shift away the switching noise harmonic interference from radio channels. OEM or ODMs can use the driver to control SOC IVR operation within the range where it does not impact IVR performance.

DRAM devices of DDR IO interface and their power plane can generate EMI at the data rates. Similar to IVR control mechanism, Intel offers a mechanism by which DDR data rates can be changed if several conditions are met: there is strong RFI interference because of DDR; CPU power management has no other restriction in changing DDR data rates; PC ODMs enable this feature (real time DDR RFI Mitigation referred to as DDR-RFIM) for Wi-Fi from BIOS.

FIVR attributes

`:file:`/sys/bus/pci/devices/0000:00:04.0/fivr/``

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\linux-master\Documentation\driver-api\thermal\[linux-master] [Documentation] [driver-api] [thermal]intel_dptf.rst, line 207); [backlink](#)

Unknown interpreted text role "file".

vco_ref_code_lo (RW)

The VCO reference code is an 11-bit field and controls the FIVR switching frequency. This is the 3-bit LSB field.

vco_ref_code_hi (RW)

The VCO reference code is an 11-bit field and controls the FIVR switching frequency. This is the 8-bit MSB field.

spread_spectrum_pct (RW)

Set the FIVR spread spectrum clocking percentage

spread_spectrum_clk_enable (RW)

Enable/disable of the FIVR spread spectrum clocking feature

rfi_vco_ref_code (RW)

This field is a read only status register which reflects the current FIVR switching frequency

fivr_fffc_rev (RW)

This field indicated the revision of the FIVR HW.

DVFS attributes

`:file:`/sys/bus/pci/devices/0000:00:04.0/dvfs/``

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\linux-master\Documentation\driver-api\thermal\[linux-master] [Documentation] [driver-api] [thermal]intel_dptf.rst, line 233); [backlink](#)

Unknown interpreted text role "file".

rfi_restriction_run_busy (RW)

Request the restriction of specific DDR data rate and set this value 1. Self reset to 0 after operation.

rfi_restriction_err_code (RW)

0 :Request is accepted, 1:Feature disabled, 2: the request restricts more points than it is allowed

rfi_restriction_data_rate_Delta (RW)

Restricted DDR data rate for RFI protection: Lower Limit

rfi_restriction_data_rate_Base (RW)

Restricted DDR data rate for RFI protection: Upper Limit

ddr_data_rate_point_0 (RO)

DDR data rate selection 1st point

ddr_data_rate_point_1 (RO)

DDR data rate selection 2nd point

ddr_data_rate_point_2 (RO)

DDR data rate selection 3rd point

ddr_data_rate_point_3 (RO)

DDR data rate selection 4th point

rfi_disable (RW)

Disable DDR rate change feature

DPTF Power supply and Battery Interface

Refer to Documentation/ABI/testing/sysfs-platform-dptf

DPTF Fan Control

Refer to Documentation/admin-guide/acpi/fan_performance_states.rst

Docutils System Messages

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\linux-master\Documentation\driver-api\thermal\[linux-master] [Documentation] [driver-api] [thermal]intel_dptf.rst, line 7); [backlink](#)

Undefined substitution referenced: "copy".