

Lock Statistics

What

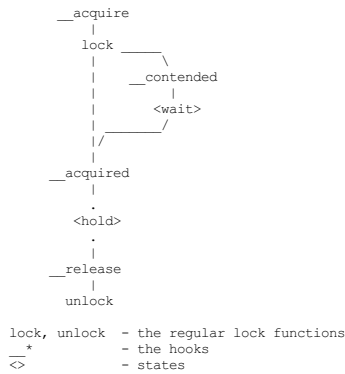
As the name suggests, it provides statistics on locks.

Why

Because things like lock contention can severely impact performance.

How

Lockdep already has hooks in the lock functions and maps lock instances to lock classes. We build on that (see [Documentation/locking/lockdep-design.rst](#)). The graph below shows the relation between the lock functions and the various hooks therein:



With these hooks we provide the following statistics:

- con-bounces
 - number of lock contention that involved x-cpu data
- contentions
 - number of lock acquisitions that had to wait
- wait time
 - min
 - shortest (non-0) time we ever had to wait for a lock
 - max
 - longest time we ever had to wait for a lock
 - total
 - total time we spend waiting on this lock
 - avg
 - average time spent waiting on this lock
- acq-bounces
 - number of lock acquisitions that involved x-cpu data
- acquisitions
 - number of times we took the lock
- hold time
 - min
 - shortest (non-0) time we ever held the lock
 - max
 - longest time we ever held the lock
 - total
 - total time this lock was held
 - avg
 - average time this lock was held

These numbers are gathered per lock class, per read/write state (when applicable).

It also tracks 4 contention points per class. A contention point is a call site that had to wait on lock acquisition.

Configuration

Lock statistics are enabled via CONFIG_LOCK_STAT.

Usage

Enable collection of statistics:

```
# echo 1 >/proc/sys/kernel/lock_stat
```

Disable collection of statistics:

```
# echo 0 >/proc/sys/kernel/lock_stat
```

Look at the current lock statistics:

```
( line numbers not part of actual output, done for clarity in the explanation
  below )

# less /proc/lock_stat

01 lock_stat version 0.4
02-----
03 class name con-bounces contentions waittime-min waittime-max waittime-total waittime-avg
04-----
05
06 6mm->mmap_sem-W: 46 84 0.26 939.10 16371.53 194.90
07 6mm->mmap_sem-R: 37 100 1.31 299502.61 325629.52 3256.30
08-----
09 6mm->mmap_sem 1 [<fffffffff811502a7>] khugepaged_scan_mm_slot+0x57/0x280
10 6mm->mmap_sem 96 [<fffffffff815351c4>] __do_page_fault+0x1d4/0x510
11 6mm->mmap_sem 34 [<fffffffff81113d77>] vm_mmap_pgoff+0x87/0xd0
12 6mm->mmap_sem 17 [<fffffffff81127e71>] vm_munmap+0x41/0x80
13-----
14 6mm->mmap_sem 1 [<fffffffff81046fda>] dup_mmap+0x2a/0x3f0
15 6mm->mmap_sem 60 [<fffffffff81129e29>] Sys_mprotect+0xe9/0x250
16 6mm->mmap_sem 41 [<fffffffff815351c4>] __do_page_fault+0x1d4/0x510
17 6mm->mmap_sem 68 [<fffffffff81113d77>] vm_mmap_pgoff+0x87/0xd0
18-----
19 .....
20
21 unix_table_lock: 110 112 0.21 49.24 163.91 1.46
```

22	-----						
23	unix_table_lock	45	[<fffffffff8150ad8e>]	unix_createl+0x16e/0x1b0			
24	unix_table_lock	47	[<fffffffff8150b111>]	unix_release_sock+0x31/0x250			
25	unix_table_lock	15	[<fffffffff8150ca37>]	unix_find_other+0x117/0x230			
26	unix_table_lock	5	[<fffffffff8150a09f>]	unix_autobind+0x11f/0x1b0			
27	-----						
28	unix_table_lock	39	[<fffffffff8150b111>]	unix_release_sock+0x31/0x250			
29	unix_table_lock	49	[<fffffffff8150ad8e>]	unix_createl+0x16e/0x1b0			
30	unix_table_lock	20	[<fffffffff8150ca37>]	unix_find_other+0x117/0x230			
31	unix_table_lock	4	[<fffffffff8150a09f>]	unix_autobind+0x11f/0x1b0			

This excerpt shows the first two lock class statistics. Line 01 shows the output version - each time the format changes this will be updated. Line 02-04 show the header with column descriptions. Lines 05-18 and 20-31 show the actual statistics. These statistics come in two parts; the actual stats separated by a short separator (line 08, 13) from the contention points.

Lines 09-12 show the first 4 recorded contention points (the code which tries to get the lock) and lines 14-17 show the first 4 recorded contending points (the lock holder). It is possible that the max con-bounces point is missing in the statistics.

The first lock (05-18) is a read/write lock, and shows two lines above the short separator. The contention points don't match the column descriptors, they have two: contentions and [<IP>] symbol. The second set of contention points are the points we're contending with.

The integer part of the time values is in us.

Dealing with nested locks, subclasses may appear:

32						
33							
34	&rq->lock:	13128	13128	0.43	190.53	103881.26	7.91
35	-----						
36	&rq->lock	645	[<fffffffff8103bfc4>]	task_rq_lock+0x43/0x75			
37	&rq->lock	297	[<fffffffff8104ba65>]	try_to_wake_up+0x127/0x25a			
38	&rq->lock	360	[<fffffffff8103c4c5>]	select_task_rq_fair+0x1f0/0x74a			
39	&rq->lock	428	[<fffffffff81045f98>]	scheduler_tick+0x46/0x1fb			
40	-----						
41	&rq->lock	77	[<fffffffff8103bfc4>]	task_rq_lock+0x43/0x75			
42	&rq->lock	174	[<fffffffff8104ba65>]	try_to_wake_up+0x127/0x25a			
43	&rq->lock	4715	[<fffffffff8103ed4b>]	double_rq_lock+0x42/0x54			
44	&rq->lock	893	[<fffffffff81340524>]	schedule+0x157/0x7b8			
45							
46						
47							
48	&rq->lock/1:	1526	11488	0.33	388.73	136294.31	11.86
49	-----						
50	&rq->lock/1	11526	[<fffffffff8103ed58>]	double_rq_lock+0x4f/0x54			
51	-----						
52	&rq->lock/1	5645	[<fffffffff8103ed4b>]	double_rq_lock+0x42/0x54			
53	&rq->lock/1	1224	[<fffffffff81340524>]	schedule+0x157/0x7b8			
54	&rq->lock/1	4336	[<fffffffff8103ed58>]	double_rq_lock+0x4f/0x54			
55	&rq->lock/1	181	[<fffffffff8104ba65>]	try_to_wake_up+0x127/0x25a			

Line 48 shows statistics for the second subclass (/1) of &rq->lock class (subclass starts from 0), since in this case, as line 50 suggests, double_rq_lock actually acquires a nested lock of two spinlocks.

View the top contending locks:

# grep : /proc/lock_stat head							
clockevents_lock:	2926159	2947636	0.15	46882.81	1784540466.34	605.41	
tick_broadcast_lock:	346460	346717	0.18	2257.43	39364622.71	113.54	
&mapping->i_mmap_mutex:	203896	203899	3.36	645530.05	31767507988.39	155800.21	
&rq->lock:	135014	136909	0.18	606.09	842160.68	6.15	
&(&zone->lru_lock)->rlock:	93000	94934	0.16	59.18	188253.78	1.98	
tasklist_lock-W:	40667	41130	0.23	1189.42	428980.51	10.43	
tasklist_lock-R:	21298	21305	0.20	1310.05	215511.12	10.12	
rcu_node_1:	47656	49022	0.16	635.41	193616.41	3.95	
&(&dentry->d_lockref.lock)->rlock:	39791	40179	0.15	1302.08	88851.96	2.21	
rcu_node_0:	29203	30064	0.16	786.55	1555573.00	51.74	

Clear the statistics:

```
# echo 0 > /proc/lock_stat
```