

Universal TUN/TAP device driver

Copyright © 1999-2000 Maxim Krasnyansky <max_mk@yahoo.com>

Linux, Solaris drivers Copyright © 1999-2000 Maxim Krasnyansky <max_mk@yahoo.com>

FreeBSD TAP driver Copyright © 1999-2000 Maksim Yevmenkin <m_evmenkin@yahoo.com>

Revision of this document 2002 by Florian Thiel <florian.thiel@gmx.net>

1. Description

TUN/TAP provides packet reception and transmission for user space programs. It can be seen as a simple Point-to-Point or Ethernet device, which, instead of receiving packets from physical media, receives them from user space program and instead of sending packets via physical media writes them to the user space program.

In order to use the driver a program has to open `/dev/net/tun` and issue a corresponding `ioctl()` to register a network device with the kernel. A network device will appear as `tunXX` or `tapXX`, depending on the options chosen. When the program closes the file descriptor, the network device and all corresponding routes will disappear.

Depending on the type of device chosen the userspace program has to read/write IP packets (with `tun`) or ethernet frames (with `tap`). Which one is being used depends on the flags given with the `ioctl()`.

The package from <http://vtun.sourceforge.net/tun> contains two simple examples for how to use `tun` and `tap` devices. Both programs work like a bridge between two network interfaces. `br_select.c` - bridge based on `select` system call. `br_sigio.c` - bridge based on `async io` and `SIGIO` signal. However, the best example is VTun <http://vtun.sourceforge.net> :))

2. Configuration

Create device node:

```
mkdir /dev/net (if it doesn't exist already)
mknod /dev/net/tun c 10 200
```

Set permissions:

```
e.g. chmod 0666 /dev/net/tun
```

There's no harm in allowing the device to be accessible by non-root users, since `CAP_NET_ADMIN` is required for creating network devices or for connecting to network devices which aren't owned by the user in question. If you want to create persistent devices and give ownership of them to unprivileged users, then you need the `/dev/net/tun` device to be usable by those users.

Driver module autoloading

Make sure that "Kernel module loader" - module auto-loading support is enabled in your kernel. The kernel should load it on first access.

Manual loading

insert the module by hand:

```
modprobe tun
```

If you do it the latter way, you have to load the module every time you need it, if you do it the other way it will be automatically loaded when `/dev/net/tun` is being opened.

3. Program interface

3.1 Network device allocation

`char *dev` should be the name of the device with a format string (e.g. `"tun%d"`), but (as far as I can see) this can be any valid network device name. Note that the character pointer becomes overwritten with the real device name (e.g. `"tun0"`):

```
#include <linux/if.h>
#include <linux/if_tun.h>

int tun_alloc(char *dev)
{
    struct ifreq ifr;
    int fd, err;
```

```

if( (fd = open("/dev/net/tun", O_RDWR)) < 0 )
    return tun_alloc_old(dev);

memset(&ifr, 0, sizeof(ifr));

/* Flags: IFF_TUN   - TUN device (no Ethernet headers)
 *        IFF_TAP   - TAP device
 *
 *        IFF_NO_PI - Do not provide packet information
 */
ifr.ifr_flags = IFF_TUN;
if( *dev )
    strncpy_pad(ifr.ifr_name, dev, IFNAMSIZ);

if( (err = ioctl(fd, TUNSETIFF, (void *) &ifr)) < 0 ){
    close(fd);
    return err;
}
strcpy(dev, ifr.ifr_name);
return fd;
}

```

3.2 Frame format

If flag `IFF_NO_PI` is not set each frame format is:

```

Flags [2 bytes]
Proto [2 bytes]
Raw protocol(IP, IPv6, etc) frame.

```

3.3 Multiqueue tuntap interface

From version 3.8, Linux supports multiqueue tuntap which can use multiple file descriptors (queues) to parallelize packets sending or receiving. The device allocation is the same as before, and if user wants to create multiple queues, `TUNSETIFF` with the same device name must be called many times with `IFF_MULTI_QUEUE` flag.

`char *dev` should be the name of the device, `queues` is the number of queues to be created, `fds` is used to store and return the file descriptors (queues) created to the caller. Each file descriptor were served as the interface of a queue which could be accessed by userspace.

```

#include <linux/if.h>
#include <linux/if_tun.h>

int tun_alloc_mq(char *dev, int queues, int *fds)
{
    struct ifreq ifr;
    int fd, err, i;

    if (!dev)
        return -1;

    memset(&ifr, 0, sizeof(ifr));
    /* Flags: IFF_TUN   - TUN device (no Ethernet headers)
     *        IFF_TAP   - TAP device
     *
     *        IFF_NO_PI - Do not provide packet information
     *        IFF_MULTI_QUEUE - Create a queue of multiqueue device
     */
    ifr.ifr_flags = IFF_TAP | IFF_NO_PI | IFF_MULTI_QUEUE;
    strcpy(ifr.ifr_name, dev);

    for (i = 0; i < queues; i++) {
        if ((fd = open("/dev/net/tun", O_RDWR)) < 0)
            goto err;
        err = ioctl(fd, TUNSETIFF, (void *)&ifr);
        if (err) {
            close(fd);
            goto err;
        }
        fds[i] = fd;
    }

    return 0;
err:
    for (--i; i >= 0; i--)
        close(fds[i]);
    return err;
}

```

A new `ioctl(TUNSETQUEUE)` were introduced to enable or disable a queue. When calling it with `IFF_DETACH_QUEUE` flag, the queue were disabled. And when calling it with `IFF_ATTACH_QUEUE` flag, the queue were enabled. The queue were enabled by default after it was created through `TUNSETIFF`.

`fd` is the file descriptor (queue) that we want to enable or disable, when `enable` is true we enable it, otherwise we disable it:

```
#include <linux/if.h>
#include <linux/if_tun.h>

int tun_set_queue(int fd, int enable)
{
    struct ifreq ifr;

    memset(&ifr, 0, sizeof(ifr));

    if (enable)
        ifr.ifr_flags = IFF_ATTACH_QUEUE;
    else
        ifr.ifr_flags = IFF_DETACH_QUEUE;

    return ioctl(fd, TUNSETQUEUE, (void *)&ifr);
}
```

Universal TUN/TAP device driver Frequently Asked Question

1. What platforms are supported by TUN/TAP driver ?

Currently driver has been written for 3 Unices:

- Linux kernels 2.2.x, 2.4.x
- FreeBSD 3.x, 4.x, 5.x
- Solaris 2.6, 7.0, 8.0

2. What is TUN/TAP driver used for?

As mentioned above, main purpose of TUN/TAP driver is tunneling. It is used by VTun (<http://vtun.sourceforge.net>).

Another interesting application using TUN/TAP is `pipsec` (<http://perso.enst.fr/~beyssac/pipsec/>), a userspace IPSec implementation that can use complete kernel routing (unlike FreeS/WAN).

3. How does Virtual network device actually work ?

Virtual network device can be viewed as a simple Point-to-Point or Ethernet device, which instead of receiving packets from a physical media, receives them from user space program and instead of sending packets via physical media sends them to the user space program.

Let's say that you configured IPv6 on the `tap0`, then whenever the kernel sends an IPv6 packet to `tap0`, it is passed to the application (VTun for example). The application encrypts, compresses and sends it to the other side over TCP or UDP. The application on the other side decompresses and decrypts the data received and writes the packet to the TAP device, the kernel handles the packet like it came from real physical device.

4. What is the difference between TUN driver and TAP driver?

TUN works with IP frames. TAP works with Ethernet frames.

This means that you have to read/write IP packets when you are using tun and ethernet frames when using tap.

5. What is the difference between BPF and TUN/TAP driver?

BPF is an advanced packet filter. It can be attached to existing network interface. It does not provide a virtual network interface. A TUN/TAP driver does provide a virtual network interface and it is possible to attach BPF to this interface.

6. Does TAP driver support kernel Ethernet bridging?

Yes. Linux and FreeBSD drivers support Ethernet bridging.