

Inference and evaluation on the Open Images dataset

This page presents a tutorial for running object detector inference and evaluation measure computations on the [Open Images dataset](#), using tools from the [TensorFlow Object Detection API](#). It shows how to download the images and annotations for the validation and test sets of Open Images; how to package the downloaded data in a format understood by the Object Detection API; where to find a trained object detector model for Open Images; how to run inference; and how to compute evaluation measures on the inferred detections.

Inferred detections will look like the following:

On the validation set of Open Images, this tutorial requires 27GB of free disk space and the inference step takes approximately 9 hours on a single NVIDIA Tesla P100 GPU. On the test set -- 75GB and 27 hours respectively. All other steps require less than two hours in total on both sets.

Installing TensorFlow, the Object Detection API, and Google Cloud SDK

Please run through the [installation instructions](#) to install TensorFlow and all its dependencies. Ensure the Protobuf libraries are compiled and the library directories are added to `PYTHONPATH`. You will also need to `pip install pandas` and `contextlib2`.

Some of the data used in this tutorial lives in Google Cloud buckets. To access it, you will have to [install the Google Cloud SDK](#) on your workstation or laptop.

Preparing the Open Images validation and test sets

In order to run inference and subsequent evaluation measure computations, we require a dataset of images and ground truth boxes, packaged as TFRecords of TFExamples. To create such a dataset for Open Images, you will need to first download ground truth boxes from the [Open Images website](#):

```
# From tensorflow/models/research
mkdir oid
cd oid
wget
https://storage.googleapis.com/openimages/2017_07/annotations_human_bbox_2017_07.tar.gz

tar -xvf annotations_human_bbox_2017_07.tar.gz
```

Next, download the images. In this tutorial, we will use lower resolution images provided by [CVDF](#). Please follow the instructions on [CVDF's Open Images repository page](#) in order to gain access to the cloud bucket with the images. Then run:

```
# From tensorflow/models/research/oid
SPLIT=validation # Set SPLIT to "test" to download the images in the test set
mkdir raw_images_${SPLIT}
gsutil -m rsync -r gs://open-images-dataset/${SPLIT} raw_images_${SPLIT}
```

Another option for downloading the images is to follow the URLs contained in the [image URLs and metadata CSV files](#) on the Open Images website.

At this point, your `tensorflow/models/research/oid` directory should appear as follows:

```
|-- 2017_07
|   |-- test
|   |   |-- annotations-human-bbox.csv
|   |-- train
|   |   |-- annotations-human-bbox.csv
|   |-- validation
|       |-- annotations-human-bbox.csv
|-- raw_images_validation (if you downloaded the validation split)
|   |-- ... (41,620 files matching regex "[0-9a-f]{16}.jpg")
|-- raw_images_test (if you downloaded the test split)
|   |-- ... (125,436 files matching regex "[0-9a-f]{16}.jpg")
`-- annotations_human_bbox_2017_07.tar.gz
```

Next, package the data into TFRecords of TFExamples by running:

```
# From tensorflow/models/research/oid
SPLIT=validation # Set SPLIT to "test" to create TFRecords for the test split
mkdir ${SPLIT}_tfrecords

PYTHONPATH=$PYTHONPATH:$(readlink -f ..) \
python -m object_detection/dataset_tools/create_oid_tf_record \
  --input_box_annotations_csv 2017_07/${SPLIT}/annotations-human-bbox.csv \
  --input_images_directory raw_images_${SPLIT} \
  --input_label_map ../object_detection/data/oid_bbox_trainable_label_map.pbtxt \
  --output_tf_record_path_prefix ${SPLIT}_tfrecords/${SPLIT}.tfrecord \
  --num_shards=100
```

To add image-level labels, use the `--input_image_label_annotations_csv` flag.

This results in 100 TFRecord files (shards), written to `oid/${SPLIT}_tfrecords`, with filenames matching `${SPLIT}.tfrecord-000[0-9][0-9]-of-00100`. Each shard contains approximately the same number of images and is defacto a representative random sample of the input data. [This enables](#) a straightforward work division scheme for distributing inference and also approximate measure computations on subsets of the validation and test sets.

Inferring detections

Inference requires a trained object detection model. In this tutorial we will use a model from the [detections model zoo](#), which can be downloaded and unpacked by running the commands below. More information about the model, such as its architecture and how it was trained, is available in the [model zoo page](#).

```
# From tensorflow/models/research/oid
wget
http://download.tensorflow.org/models/object_detection/faster_rcnn_inception_resnet_v2_

tar -zxvf faster_rcnn_inception_resnet_v2_atrous_oid_14_10_2017.tar.gz
```

At this point, data is packed into TFRecords and we have an object detector model. We can run inference using:

```
# From tensorflow/models/research/oid
SPLIT=validation # or test
TF_RECORD_FILES=$(ls -l ${SPLIT}_tfrecords/* | tr '\n' ',')

PYTHONPATH=$PYTHONPATH:$(readlink -f ..) \
python -m object_detection/inference/infer_detections \
  --input_tfrecord_paths=$TF_RECORD_FILES \
  --output_tfrecord_path=${SPLIT}_detections.tfrecord-00000-of-00001 \
  --
inference_graph=faster_rcnn_inception_resnet_v2_atrous_oid/frozen_inference_graph.pb \
  --discard_image_pixels
```

Inference preserves all fields of the input TFExamples, and adds new fields to store the inferred detections. This allows [computing evaluation measures](#) on the output TFRecord alone, as groundtruth boxes are preserved as well. Since measure computations don't require access to the images, `infer_detections` can optionally discard them with the `--discard_image_pixels` flag. Discarding the images drastically reduces the size of the output TFRecord.

Accelerating inference

Running inference on the whole validation or test set can take a long time to complete due to the large number of images present in these sets (41,620 and 125,436 respectively). For quick but approximate evaluation, inference and the subsequent measure computations can be run on a small number of shards. To run for example on 2% of all the data, it is enough to set `TF_RECORD_FILES` as shown below before running `infer_detections`:

```
TF_RECORD_FILES=$(ls ${SPLIT}_tfrecords/${SPLIT}.tfrecord-0000[0-1]-of-00100 | tr
'\n' ',')
```

Please note that computing evaluation measures on a small subset of the data introduces variance and bias, since some classes of objects won't be seen during evaluation. In the example above, this leads to 13.2% higher mAP on the first two shards of the validation set compared to the mAP for the full set ([see mAP results](#)).

Another way to accelerate inference is to run it in parallel on multiple TensorFlow devices on possibly multiple machines. The script below uses [tmux](#) to run a separate `infer_detections` process for each GPU on different partition of the input data.

```
# From tensorflow/models/research/oid
SPLIT=validation # or test
NUM_GPUS=4
NUM_SHARDS=100

tmux new-session -d -s "inference"
function tmux_start { tmux new-window -d -n "inference:GPU$1" "${*:2}; exec bash"; }
for gpu_index in $(seq 0 $((NUM_GPUS-1))); do
  start_shard=$(( $gpu_index * NUM_SHARDS / NUM_GPUS ))
  end_shard=$(( ($gpu_index + 1) * NUM_SHARDS / NUM_GPUS - 1))
  TF_RECORD_FILES=$(seq -s, -f "${SPLIT}_tfrecords/${SPLIT}.tfrecord-%05.0f-of-${printf '%05d' NUM_SHARDS}" $start_shard $end_shard)
  tmux_start $gpu_index \
```

```

PYTHONPATH=$PYTHONPATH:$(readlink -f ..) CUDA_VISIBLE_DEVICES=$gpu_index \
python -m object_detection/inference/infer_detections \
    --input_tfrecord_paths=$TF_RECORD_FILES \
    --output_tfrecord_path=${SPLIT}_detections.tfrecord-$(printf "%05d" $gpu_index)-
of-$(printf "%05d" $NUM_GPUS) \
    --
inference_graph=faster_rcnn_inception_resnet_v2_atrous_oid/frozen_inference_graph.pb
\
    --discard_image_pixels
done

```

After all `infer_detections` processes finish, `tensorflow/models/research/oid` will contain one output TFRecord from each process, with name matching `validation_detections.tfrecord-0000[0-3]-of-00004`.

Computing evaluation measures

To compute evaluation measures on the inferred detections you first need to create the appropriate configuration files:

```

# From tensorflow/models/research/oid
SPLIT=validation # or test
NUM_SHARDS=1 # Set to NUM_GPUS if using the parallel evaluation script above

mkdir -p ${SPLIT}_eval_metrics

echo "
label_map_path: '../object_detection/data/oid_bbox_trainable_label_map.pbtxt'
tf_record_input_reader: { input_path: '${SPLIT}_detections.tfrecord@${NUM_SHARDS}' }
" > ${SPLIT}_eval_metrics/${SPLIT}_input_config.pbtxt

echo "
metrics_set: 'oid_V2_detection_metrics'
" > ${SPLIT}_eval_metrics/${SPLIT}_eval_config.pbtxt

```

And then run:

```

# From tensorflow/models/research/oid
SPLIT=validation # or test

PYTHONPATH=$PYTHONPATH:$(readlink -f ..) \
python -m object_detection/metrics/offline_eval_map_corloc \
    --eval_dir=${SPLIT}_eval_metrics \
    --eval_config_path=${SPLIT}_eval_metrics/${SPLIT}_eval_config.pbtxt \
    --input_config_path=${SPLIT}_eval_metrics/${SPLIT}_input_config.pbtxt

```

The first configuration file contains an `object_detection.protos.InputReader` message that describes the location of the necessary input files. The second file contains an `object_detection.protos.EvalConfig` message that describes the evaluation metric. For more information about these protos see the corresponding source files.

Expected mAPs

The result of running `offline_eval_map_corloc` is a CSV file located at `${SPLIT}_eval_metrics/metrics.csv`. With the above configuration, the file will contain average precision at $\text{IoU} \geq 0.5$ for each of the classes present in the dataset. It will also contain the $\text{mAP@IoU} \geq 0.5$. Both the per-class average precisions and the mAP are computed according to the [Open Images evaluation protocol](#). The expected mAPs for the validation and test sets of Open Images in this case are:

Set	Fraction of data	Images	mAP@IoU ≥ 0.5
validation	everything	41,620	39.2%
validation	first 2 shards	884	52.4%
test	everything	125,436	37.7%
test	first 2 shards	2,476	50.8%