

# Quick Start: Distributed Training on the Oxford-IIIT Pets Dataset on Google Cloud



This page is a walkthrough for training an object detector using the TensorFlow Object Detection API. In this tutorial, we'll be training on the Oxford-IIIT Pets dataset to build a system to detect various breeds of cats and dogs. The output of the detector will look like the following:

## Setting up a Project on Google Cloud

To accelerate the process, we'll run training and evaluation on [Google Cloud ML Engine](#) to leverage multiple GPUs. To begin, you will have to set up Google Cloud via the following steps (if you have already done this, feel free to skip to the next section):

1. [Create a GCP project](#).
2. [Install the Google Cloud SDK](#) on your workstation or laptop. This will provide the tools you need to upload files to Google Cloud Storage and start ML training jobs.
3. [Enable the ML Engine APIs](#). By default, a new GCP project does not enable APIs to start ML Engine training jobs. Use the above link to explicitly enable them.
4. [Set up a Google Cloud Storage \(GCS\) bucket](#). ML Engine training jobs can only access files on a Google Cloud Storage bucket. In this tutorial, we'll be required to upload our dataset and configuration to GCS.

Please remember the name of your GCS bucket, as we will reference it multiple times in this document. Substitute `${YOUR_GCS_BUCKET}` with the name of your bucket in this document. For your convenience, you should define the environment variable below:

```
export YOUR_GCS_BUCKET=${YOUR_GCS_BUCKET}
```

It is also possible to run locally by following [the running locally instructions](#).

## Installing TensorFlow and the TensorFlow Object Detection API

Please run through the [installation instructions](#) to install TensorFlow and all its dependencies. Ensure the Protobuf libraries are compiled and the library directories are added to `PYTHONPATH`.

## Getting the Oxford-IIIT Pets Dataset and Uploading it to Google Cloud Storage

In order to train a detector, we require a dataset of images, bounding boxes and classifications. For this demo, we'll use the Oxford-IIIT Pets dataset. The raw dataset for Oxford-IIIT Pets lives [here](#). You will need to download both the image dataset [images.tar.gz](#) and the groundtruth data [annotations.tar.gz](#) to the `tensorflow/models/research/` directory and unzip them. This may take some time.

```
# From tensorflow/models/research/
wget http://www.robots.ox.ac.uk/~vgg/data/pets/data/images.tar.gz
wget http://www.robots.ox.ac.uk/~vgg/data/pets/data/annotations.tar.gz
tar -xvf images.tar.gz
tar -xvf annotations.tar.gz
```

After downloading the tarballs, your `tensorflow/models/research/` directory should appear as follows:

```
- images.tar.gz
- annotations.tar.gz
+ images/
+ annotations/
+ object_detection/
... other files and directories
```

The TensorFlow Object Detection API expects data to be in the TFRecord format, so we'll now run the `create_pet_tf_record` script to convert from the raw Oxford-IIIT Pet dataset into TFRecords. Run the following commands from the `tensorflow/models/research/` directory:

```
# From tensorflow/models/research/
python object_detection/dataset_tools/create_pet_tf_record.py \
  --label_map_path=object_detection/data/pet_label_map.pbtxt \
  --data_dir=`pwd` \
  --output_dir=`pwd`
```

Note: It is normal to see some warnings when running this script. You may ignore them.

Two 10-sharded TFRecord files named `pet_faces_train.record-*` and `pet_faces_val.record-*` should be generated in the `tensorflow/models/research/` directory.

Now that the data has been generated, we'll need to upload it to Google Cloud Storage so the data can be accessed by ML Engine. Run the following command to copy the files into your GCS bucket (substituting `${YOUR_GCS_BUCKET}`):

```
# From tensorflow/models/research/
gsutil cp pet_faces_train.record-* gs://${YOUR_GCS_BUCKET}/data/
gsutil cp pet_faces_val.record-* gs://${YOUR_GCS_BUCKET}/data/
gsutil cp object_detection/data/pet_label_map.pbtxt
gs://${YOUR_GCS_BUCKET}/data/pet_label_map.pbtxt
```

Please remember the path where you upload the data to, as we will need this information when configuring the pipeline in a following step.

## Downloading a COCO-pretrained Model for Transfer Learning

Training a state of the art object detector from scratch can take days, even when using multiple GPUs! In order to speed up training, we'll take an object detector trained on a different dataset (COCO), and reuse some of it's parameters to initialize our new model.

Download our [COCO-pretrained Faster R-CNN with Resnet-101 model](#). Unzip the contents of the folder and copy the `model.ckpt*` files into your GCS Bucket.

```
wget
http://storage.googleapis.com/download.tensorflow.org/models/object_detection/faster_rcnn_resnet101_coco_11_06_2017.tar.gz

tar -xvf faster_rcnn_resnet101_coco_11_06_2017.tar.gz
gsutil cp faster_rcnn_resnet101_coco_11_06_2017/model.ckpt.*
gs://${YOUR_GCS_BUCKET}/data/
```

Remember the path where you uploaded the model checkpoint to, as we will need it in the following step.

## Configuring the Object Detection Pipeline

In the TensorFlow Object Detection API, the model parameters, training parameters and eval parameters are all defined by a config file. More details can be found [here](#). For this tutorial, we will use some predefined templates provided with the source code. In the `object_detection/samples/configs` folder, there are skeleton object\_detection configuration files. We will use `faster_rcnn_resnet101_pets.config` as a starting point for configuring the pipeline. Open the file with your favourite text editor.

We'll need to configure some paths in order for the template to work. Search the file for instances of `PATH_TO_BE_CONFIGURED` and replace them with the appropriate value (typically `gs://${YOUR_GCS_BUCKET}/data/`). Afterwards upload your edited file onto GCS, making note of the path it was uploaded to (we'll need it when starting the training/eval jobs).

```
# From tensorflow/models/research/

# Edit the faster_rcnn_resnet101_pets.config template. Please note that there
# are multiple places where PATH_TO_BE_CONFIGURED needs to be set.
sed -i "s|PATH_TO_BE_CONFIGURED|gs://${YOUR_GCS_BUCKET}/data|g" \
    object_detection/samples/configs/faster_rcnn_resnet101_pets.config

# Copy edited template to cloud.
gsutil cp object_detection/samples/configs/faster_rcnn_resnet101_pets.config \
    gs://${YOUR_GCS_BUCKET}/data/faster_rcnn_resnet101_pets.config
```

## Checking Your Google Cloud Storage Bucket

At this point in the tutorial, you should have uploaded the training/validation datasets (including label map), our COCO trained FasterRCNN finetune checkpoint and your job configuration to your Google Cloud Storage Bucket. Your bucket should look like the following:

```
+ ${YOUR_GCS_BUCKET}/
+ data/
+   - faster_rcnn_resnet101_pets.config
+   - model.ckpt.index
+   - model.ckpt.meta
+   - model.ckpt.data-00000-of-00001
+   - pet_label_map.pbtxt
```

```
- pet_faces_train.record-*  
- pet_faces_val.record-*
```

You can inspect your bucket using the [Google Cloud Storage browser](#).

## Starting Training and Evaluation Jobs on Google Cloud ML Engine

Before we can start a job on Google Cloud ML Engine, we must:

1. Package the TensorFlow Object Detection code.
2. Write a cluster configuration for our Google Cloud ML job.

To package the TensorFlow Object Detection code, run the following commands from the

`tensorflow/models/research/` directory:

```
# From tensorflow/models/research/  
bash object_detection/dataset_tools/create_pycocotools_package.sh /tmp/pycocotools  
python setup.py sdist  
(cd slim && python setup.py sdist)
```

This will create python packages `dist/object_detection-0.1.tar.gz`, `slim/dist/slim-0.1.tar.gz`, and `/tmp/pycocotools/pycocotools-2.0.tar.gz`.

For running the training Cloud ML job, we'll configure the cluster to use 5 training jobs and three parameters servers.

The configuration file can be found at `object_detection/samples/cloud/cloud.yml`.

Note: The code sample below is supported for use with 1.12 runtime version.

To start training and evaluation, execute the following command from the `tensorflow/models/research/` directory:

```
# From tensorflow/models/research/  
gcloud ml-engine jobs submit training `whoami`_object_detection_pets_`date`  
+%m_%d_%Y_%H_%M_%S` \  
  --runtime-version 1.12 \  
  --job-dir=gs://${YOUR_GCS_BUCKET}/model_dir \  
  --packages dist/object_detection-0.1.tar.gz,slim/dist/slim-  
0.1.tar.gz,/tmp/pycocotools/pycocotools-2.0.tar.gz \  
  --module-name object_detection.model_main \  
  --region us-central1 \  
  --config object_detection/samples/cloud/cloud.yml \  
  -- \  
  --model_dir=gs://${YOUR_GCS_BUCKET}/model_dir \  
  --  
pipeline_config_path=gs://${YOUR_GCS_BUCKET}/data/faster_rcnn_resnet101_pets.config
```

Users can monitor and stop training and evaluation jobs on the [ML Engine Dashboard](#).

## Monitoring Progress with Tensorboard

You can monitor progress of the training and eval jobs by running Tensorboard on your local machine:

```
# This command needs to be run once to allow your local machine to access your
# GCS bucket.
gcloud auth application-default login

tensorboard --logdir=gs://{YOUR_GCS_BUCKET}/model_dir
```

Once Tensorboard is running, navigate to `localhost:6006` from your favourite web browser. You should see something similar to the following:

Make sure your Tensorboard version is the same minor version as your TensorFlow (1.x)

You will also want to click on the images tab to see example detections made by the model while it trains. After about an hour and a half of training, you can expect to see something like this:

Note: It takes roughly 10 minutes for a job to get started on ML Engine, and roughly an hour for the system to evaluate the validation dataset. It may take some time to populate the dashboards. If you do not see any entries after half an hour, check the logs from the [ML Engine Dashboard](#). Note that by default the training jobs are configured to go for much longer than is necessary for convergence. To save money, we recommend killing your jobs once you've seen that they've converged.

## Exporting the TensorFlow Graph

After your model has been trained, you should export it to a TensorFlow graph proto. First, you need to identify a candidate checkpoint to export. You can search your bucket using the [Google Cloud Storage Browser](#). The file should be stored under `{YOUR_GCS_BUCKET}/model_dir`. The checkpoint will typically consist of three files:

- `model.ckpt- $\{CHECKPOINT\_NUMBER\}$ .data-00000-of-00001`
- `model.ckpt- $\{CHECKPOINT\_NUMBER\}$ .index`
- `model.ckpt- $\{CHECKPOINT\_NUMBER\}$ .meta`

After you've identified a candidate checkpoint to export, run the following command from

```
tensorflow/models/research/ :
```

```
# From tensorflow/models/research/
gsutil cp gs://{YOUR_GCS_BUCKET}/model_dir/model.ckpt- $\{CHECKPOINT\_NUMBER\}$ .* .
python object_detection/export_inference_graph.py \
    --input_type image_tensor \
    --pipeline_config_path
object_detection/samples/configs/faster_rcnn_resnet101_pets.config \
    --trained_checkpoint_prefix model.ckpt- $\{CHECKPOINT\_NUMBER\}$  \
    --output_directory exported_graphs
```

Afterwards, you should see a directory named `exported_graphs` containing the SavedModel and frozen graph.

## Configuring the Instance Segmentation Pipeline

Mask prediction can be turned on for an object detection config by adding `predict_instance_masks: true` within the `MaskRCNNBoxPredictor`. Other parameters such as mask size, number of convolutions in the mask layer, and the convolution hyper parameters can be defined. We will use `mask_rcnn_resnet101_pets.config` as a starting point for configuring the instance segmentation pipeline. Everything above that was mentioned about object detection holds true for instance segmentation. Instance segmentation consists of an object detection model with an additional head that predicts the object mask inside each predicted box once we remove the training and other details. Please refer to the section on [Running an Instance Segmentation Model](#) for instructions on how to configure a model that predicts masks in addition to object bounding boxes.

## What's Next

Congratulations, you have now trained an object detector for various cats and dogs! There different things you can do now:

1. [Test your exported model using the provided Jupyter notebook.](#)
2. [Experiment with different model configurations.](#)
3. Train an object detector using your own data.