

# Quantization

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\pytorch-master\docs\source\quantization.rst, line 6)**

Unknown directive type "automodule".

```
.. automodule:: torch.quantization
```

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\pytorch-master\docs\source\quantization.rst, line 7)**

Unknown directive type "automodule".

```
.. automodule:: torch.quantization.fx
```

## Warning

Quantization is in beta and subject to change.

## Introduction to Quantization

Quantization refers to techniques for performing computations and storing tensors at lower bitwidths than floating point precision. A quantized model executes some or all of the operations on tensors with integers rather than floating point values. This allows for a more compact model representation and the use of high performance vectorized operations on many hardware platforms. PyTorch supports INT8 quantization compared to typical FP32 models allowing for a 4x reduction in the model size and a 4x reduction in memory bandwidth requirements. Hardware support for INT8 computations is typically 2 to 4 times faster compared to FP32 compute. Quantization is primarily a technique to speed up inference and only the forward pass is supported for quantized operators.

PyTorch supports multiple approaches to quantizing a deep learning model. In most cases the model is trained in FP32 and then the model is converted to INT8. In addition, PyTorch also supports quantization aware training, which models quantization errors in both the forward and backward passes using fake-quantization modules. Note that the entire computation is carried out in floating point. At the end of quantization aware training, PyTorch provides conversion functions to convert the trained model into lower precision.

At lower level, PyTorch provides a way to represent quantized tensors and perform operations with them. They can be used to directly construct models that perform all or part of the computation in lower precision. Higher-level APIs are provided that incorporate typical workflows of converting FP32 model to lower precision with minimal accuracy loss.

Quantization requires users to be aware of three concepts:

1. Quantization Config (Qconfig): Specifies how weights and activations are to be quantized. Qconfig is needed to create a quantized model.
2. Backend: Refers to kernels that support quantization, usually with different numerics.
3. Quantization engine (torch.backends.quantization.engine): When a quantized model is executed, the engine specifies which backend is to be used for execution. It is important to ensure that the engine is consistent with the Qconfig.

## Quantization API Summary

PyTorch provides two different modes of quantization: Eager Mode Quantization and FX Graph Mode Quantization.

Eager Mode Quantization is a beta feature. User needs to do fusion and specify where quantization and dequantization happens manually, also it only supports modules and not functionals.

FX Graph Mode Quantization is a new automated quantization framework in PyTorch, and currently it's a prototype feature. It improves upon Eager Mode Quantization by adding support for functionals and automating the quantization process, although people might need to refactor the model to make the model compatible with FX Graph Mode Quantization (symbolically traceable with `torch.fx`). Note that FX Graph Mode Quantization is not expected to work on arbitrary models since the model might not be symbolically traceable, we will integrate it into domain libraries like torchvision and users will be able to quantize models similar to the ones in supported domain libraries with FX Graph Mode Quantization. For arbitrary models we'll provide general guidelines, but to actually make it work, users might need to be familiar with `torch.fx`, especially on how to make a model symbolically traceable.

New users of quantization are encouraged to try out FX Graph Mode Quantization first, if it does not work, user may try to follow the guideline of [using FX Graph Mode Quantization](#) or fall back to eager mode quantization.

The following table compares the differences between Eager Mode Quantization and FX Graph Mode Quantization:

	Eager Mode Quantization	FX Graph Mode Quantization
Release Status	beta	prototype
Operator Fusion	Manual	Automatic
Quant/DeQuant Placement	Manual	Automatic
Quantizing Modules	Supported	Supported
Quantizing Functionals/Torch Ops	Manual	Automatic
Support for Customization	Limited Support	Fully Supported
Quantization Mode Support	Post Training Quantization: Static, Dynamic, Weight Only	Post Training Quantization: Static, Dynamic, Weight Only
	Quantization Aware Training: Static	Quantization Aware Training: Static
Input/Output Model Type	<code>torch.nn.Module</code>	<code>torch.nn.Module</code> (May need some refactors to make the model compatible with FX Graph Mode Quantization)

There are three types of quantization supported:

1. dynamic quantization (weights quantized with activations read/stored in floating point and quantized for compute)
2. static quantization (weights quantized, activations quantized, calibration required post training)

- static quantization aware training (weights quantized, activations quantized, quantization numerics modeled during training)

Please see our [Introduction to Quantization on Pytorch](#) blog post for a more comprehensive overview of the tradeoffs between these quantization types.

Operator coverage varies between dynamic and static quantization and is captured in the table below. Note that for FX quantization, the corresponding functionals are also supported.

	Static Quantization	Dynamic Quantization
nn.Linear	Y	Y
nn.Conv1d/2d/3d	Y	N
nn.LSTM	N	Y
nn.GRU	N	Y
nn.RNNCell	N	Y
nn.GRUCell	N	Y
nn.LSTMCell	N	Y
nn.EmbeddingBag	Y (activations are in fp32)	Y
nn.Embedding	Y	N
nn.MultiheadAttention	Not Supported	Not supported
Activations	Broadly supported	Un-changed, computations stay in fp32

## Eager Mode Quantization

### Dynamic Quantization

This is the simplest to apply form of quantization where the weights are quantized ahead of time but the activations are dynamically quantized during inference. This is used for situations where the model execution time is dominated by loading weights from memory rather than computing the matrix multiplications. This is true for LSTM and Transformer type models with small batch size.

Diagram:

```
# original model
# all tensors and computations are in floating point
previous_layer_fp32 -- linear_fp32 -- activation_fp32 -- next_layer_fp32
/
linear_weight_fp32

# dynamically quantized model
# linear and LSTM weights are in int8
previous_layer_fp32 -- linear_int8_w_fp32_inp -- activation_fp32 -- next_layer_fp32
/
linear_weight_int8
```

API example:

```
import torch

# define a floating point model
class M(torch.nn.Module):
    def __init__(self):
        super(M, self).__init__()
        self.fc = torch.nn.Linear(4, 4)

    def forward(self, x):
        x = self.fc(x)
        return x

# create a model instance
model_fp32 = M()
# create a quantized model instance
model_int8 = torch.quantization.quantize_dynamic(
    model_fp32, # the original model
    {torch.nn.Linear}, # a set of layers to dynamically quantize
    dtype=torch.qint8) # the target dtype for quantized weights

# run the model
input_fp32 = torch.randn(4, 4, 4, 4)
res = model_int8(input_fp32)
```

To learn more about dynamic quantization please see our [dynamic quantization tutorial](#).

### Static Quantization

Static quantization quantizes the weights and activations of the model. It fuses activations into preceding layers where possible. It requires calibration with a representative dataset to determine optimal quantization parameters for activations. Post Training Quantization is typically used when both memory bandwidth and compute savings are important with CNNs being a typical use case. Static quantization is also known as Post Training Quantization or PTQ.

Diagram:

```
# original model
# all tensors and computations are in floating point
previous_layer_fp32 -- linear_fp32 -- activation_fp32 -- next_layer_fp32
/
linear_weight_fp32

# statically quantized model
# weights and activations are in int8
previous_layer_int8 -- linear_with_activation_int8 -- next_layer_int8
/
linear_weight_int8
```

API Example:

```

import torch

# define a floating point model where some layers could be statically quantized
class M(torch.nn.Module):
    def __init__(self):
        super(M, self).__init__()
        # QuantStub converts tensors from floating point to quantized
        self.quant = torch.quantization.QuantStub()
        self.conv = torch.nn.Conv2d(1, 1, 1)
        self.relu = torch.nn.ReLU()
        # DeQuantStub converts tensors from quantized to floating point
        self.dequant = torch.quantization.DeQuantStub()

    def forward(self, x):
        # manually specify where tensors will be converted from floating
        # point to quantized in the quantized model
        x = self.quant(x)
        x = self.conv(x)
        x = self.relu(x)
        # manually specify where tensors will be converted from quantized
        # to floating point in the quantized model
        x = self.dequant(x)
        return x

# create a model instance
model_fp32 = M()

# model must be set to eval mode for static quantization logic to work
model_fp32.eval()

# attach a global qconfig, which contains information about what kind
# of observers to attach. Use 'fbgemm' for server inference and
# 'qnnpack' for mobile inference. Other quantization configurations such
# as selecting symmetric or asymmetric quantization and MinMax or L2Norm
# calibration techniques can be specified here.
model_fp32.qconfig = torch.quantization.get_default_qconfig('fbgemm')

# Fuse the activations to preceding layers, where applicable.
# This needs to be done manually depending on the model architecture.
# Common fusions include 'conv + relu' and 'conv + batchnorm + relu'
model_fp32_fused = torch.quantization.fuse_modules(model_fp32, [['conv', 'relu']])

# Prepare the model for static quantization. This inserts observers in
# the model that will observe activation tensors during calibration.
model_fp32_prepared = torch.quantization.prepare(model_fp32_fused)

# calibrate the prepared model to determine quantization parameters for activations
# in a real world setting, the calibration would be done with a representative dataset
input_fp32 = torch.randn(4, 1, 4, 4)
model_fp32_prepared(input_fp32)

# Convert the observed model to a quantized model. This does several things:
# quantizes the weights, computes and stores the scale and bias value to be
# used with each activation tensor, and replaces key operators with quantized
# implementations.
model_int8 = torch.quantization.convert(model_fp32_prepared)

# run the model, relevant calculations will happen in int8
res = model_int8(input_fp32)

```

To learn more about static quantization, please see the [static quantization tutorial](#).

## Quantization Aware Training

Quantization Aware Training models the effects of quantization during training allowing for higher accuracy compared to other quantization methods. During training, all calculations are done in floating point, with `fake_quant` modules modeling the effects of quantization by clamping and rounding to simulate the effects of INT8. After model conversion, weights and activations are quantized, and activations are fused into the preceding layer where possible. It is commonly used with CNNs and yields a higher accuracy compared to static quantization. Quantization Aware Training is also known as QAT.

Diagram:

```

# original model
# all tensors and computations are in floating point
previous_layer_fp32 -- linear_fp32 -- activation_fp32 -- next_layer_fp32
                        /
                    linear_weight_fp32

# model with fake_quant for modeling quantization numerics during training
previous_layer_fp32 -- fq -- linear_fp32 -- activation_fp32 -- fq -- next_layer_fp32
                        /
                    linear_weight_fp32 -- fq

# quantized model
# weights and activations are in int8
previous_layer_int8 -- linear_with_activation_int8 -- next_layer_int8
                        /
                    linear_weight_int8

```

API Example:

```

import torch

# define a floating point model where some layers could benefit from QAT
class M(torch.nn.Module):
    def __init__(self):
        super(M, self).__init__()
        # QuantStub converts tensors from floating point to quantized
        self.quant = torch.quantization.QuantStub()
        self.conv = torch.nn.Conv2d(1, 1, 1)
        self.bn = torch.nn.BatchNorm2d(1)

```

```

self.relu = torch.nn.ReLU()
# DeQuantStub converts tensors from quantized to floating point
self.dequant = torch.quantization.DeQuantStub()

def forward(self, x):
    x = self.quant(x)
    x = self.conv(x)
    x = self.bn(x)
    x = self.relu(x)
    x = self.dequant(x)
    return x

# create a model instance
model_fp32 = M()

# model must be set to train mode for QAT logic to work
model_fp32.train()

# attach a global qconfig, which contains information about what kind
# of observers to attach. Use 'fbgemm' for server inference and
# 'qnnpack' for mobile inference. Other quantization configurations such
# as selecting symmetric or asymmetric quantization and MinMax or L2Norm
# calibration techniques can be specified here.
model_fp32.qconfig = torch.quantization.get_default_qat_qconfig('fbgemm')

# fuse the activations to preceding layers, where applicable
# this needs to be done manually depending on the model architecture
model_fp32_fused = torch.quantization.fuse_modules(model_fp32,
    [['conv', 'bn', 'relu']])

# Prepare the model for QAT. This inserts observers and fake_quants in
# the model that will observe weight and activation tensors during calibration.
model_fp32_prepared = torch.quantization.prepare_qat(model_fp32_fused)

# run the training loop (not shown)
training_loop(model_fp32_prepared)

# Convert the observed model to a quantized model. This does several things:
# quantizes the weights, computes and stores the scale and bias value to be
# used with each activation tensor, fuses modules where appropriate,
# and replaces key operators with quantized implementations.
model_fp32_prepared.eval()
model_int8 = torch.quantization.convert(model_fp32_prepared)

# run the model, relevant calculations will happen in int8
res = model_int8(input_fp32)

```

To learn more about quantization aware training, please see the [QAT tutorial](#).

## (Prototype) FX Graph Mode Quantization

Quantization types supported by FX Graph Mode can be classified in two ways:

1. Post Training Quantization (apply quantization after training, quantization parameters are calculated based on sample calibration data)
2. Quantization Aware Training (simulate quantization during training so that the quantization parameters can be learned together with the model using training data)

And then each of these two may include any or all of the following types:

- Weight Only Quantization (only weight is statically quantized)
- Dynamic Quantization (weight is statically quantized, activation is dynamically quantized)
- Static Quantization (both weight and activations are statically quantized)

These two ways of classification are independent, so theoretically we can have 6 different types of quantization.

The supported quantization types in FX Graph Mode Quantization are:

- Post Training Quantization
  - Weight Only Quantization
  - Dynamic Quantization
  - Static Quantization
- Quantization Aware Training
  - Static Quantization

There are multiple quantization types in post training quantization (weight only, dynamic and static) and the configuration is done through *qconfig\_dict* (an argument of the *prepare\_fx* function).

API Example:

```

import torch.quantization.quantize_fx as quantize_fx
import copy

model_fp = UserModel(...)

#
# post training dynamic/weight_only quantization
#

# we need to deepcopy if we still want to keep model_fp unchanged after quantization since quantization apis change t
model_to_quantize = copy.deepcopy(model_fp)
model_to_quantize.eval()
qconfig_dict = {"": torch.quantization.default_dynamic_qconfig}
# prepare
model_prepared = quantize_fx.prepare_fx(model_to_quantize, qconfig_dict)
# no calibration needed when we only have dynamic/weight_only quantization
# quantize
model_quantized = quantize_fx.convert_fx(model_prepared)

#
# post training static quantization

```

```
#

model_to_quantize = copy.deepcopy(model_fp)
qconfig_dict = {"": torch.quantization.get_default_qconfig('qnnpack')}
model_to_quantize.eval()
# prepare
model_prepared = quantize_fx.prepare_fx(model_to_quantize, qconfig_dict)
# calibrate (not shown)
# quantize
model_quantized = quantize_fx.convert_fx(model_prepared)

#
# quantization aware training for static quantization
#

model_to_quantize = copy.deepcopy(model_fp)
qconfig_dict = {"": torch.quantization.get_default_qat_qconfig('qnnpack')}
model_to_quantize.train()
# prepare
model_prepared = quantize_fx.prepare_qat_fx(model_to_quantize, qconfig_dict)
# training loop (not shown)
# quantize
model_quantized = quantize_fx.convert_fx(model_prepared)

#
# fusion
#
model_to_quantize = copy.deepcopy(model_fp)
model_fused = quantize_fx.fuse_fx(model_to_quantize)
```

Please see the following tutorials for more information about FX Graph Mode Quantization:

- [User Guide on Using FX Graph Mode Quantization](#)
- [FX Graph Mode Post Training Static Quantization](#)
- [FX Graph Mode Post Training Dynamic Quantization](#)

## Quantization API Reference

The `:doc:Quantization API Reference <quantization-support>` contains documentation of quantization APIs, such as quantization passes, quantized tensor operations, and supported quantized modules and functions.

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\pytorch-master [docs] [source]quantization.rst, line 481); [backlink](#)**

Unknown interpreted text role "doc".

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\pytorch-master [docs] [source]quantization.rst, line 485)**

Unknown directive type "toctree".

```
.. toctree::
    :hidden:

    quantization-support
    torch.ao.ns._numeric_suite
    torch.ao.ns._numeric_suite_fx
```

## Quantization Backend Configuration

The `:doc:Quantization Backend Configuration <quantization-backend-configuration>` contains documentation on how to configure the quantization workflows for various backends.

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\pytorch-master [docs] [source]quantization.rst, line 495); [backlink](#)**

Unknown interpreted text role "doc".

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\pytorch-master [docs] [source]quantization.rst, line 498)**

Unknown directive type "toctree".

```
.. toctree::
    :hidden:

    quantization-backend-configuration
```

## Quantized Tensors

PyTorch supports both per tensor and per channel asymmetric linear quantization. Per tensor means that all the values within the tensor are scaled the same way. Per channel means that for each dimension, typically the channel dimension of a tensor, the values in the tensor are scaled and offset by a different value (effectively the scale and offset become vectors). This allows for lesser error in converting tensors to quantized values.

The mapping is performed by converting the floating point tensors using



Note that, we ensure that zero in floating point is represented with no error after quantization, thereby ensuring that operations like padding do not cause additional quantization error.

In order to do quantization in PyTorch, we need to be able to represent quantized data in Tensors. A Quantized Tensor allows for storing quantized data (represented as int8/uint8/int32) along with quantization parameters like scale and zero\_point. Quantized Tensors allow for many useful operations making quantized arithmetic easy, in addition to allowing for serialization of data in a quantized format.

## Natively supported backends

Today, PyTorch supports the following backends for running quantized operators efficiently:

- x86 CPUs with AVX2 support or higher (without AVX2 some operations have inefficient implementations), via *fbgemm* (<https://github.com/pytorch/FBGEMM>).
- ARM CPUs (typically found in mobile/embedded devices), via *qnnpack* (<https://github.com/pytorch/pytorch/tree/master/aten/src/ATen/native/quantized/cpu/qnnpack>).

The corresponding implementation is chosen automatically based on the PyTorch build mode, though users have the option to override this by setting `torch.backends.quantization.engine` to *fbgemm* or *qnnpack*.

### Note

At the moment PyTorch doesn't provide quantized operator implementations on CUDA - this is the direction for future work. Move the model to CPU in order to test the quantized functionality.

Quantization-aware training (through `class:~torch.quantization.FakeQuantize`, which emulates quantized numerics in fp32) supports both CPU and CUDA.

**System Message: ERROR/3 (D: \onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\ [pytorch-master] [docs] [source] quantization.rst, line 549); backlink**

Unknown interpreted text role "class".

When preparing a quantized model, it is necessary to ensure that qconfig and the engine used for quantized computations match the backend on which the model will be executed. The qconfig controls the type of observers used during the quantization passes. The qengine controls whether *fbgemm* or *qnnpack* specific packing function is used when packing weights for linear and convolution functions and modules. For example:

Default settings for fbgemm:

```
# set the qconfig for PTQ
qconfig = torch.quantization.get_default_qconfig('fbgemm')
# or, set the qconfig for QAT
qconfig = torch.quantization.get_default_qat_qconfig('fbgemm')
# set the qengine to control weight packing
torch.backends.quantized.engine = 'fbgemm'
```

Default settings for qnnpack:

```
# set the qconfig for PTQ
qconfig = torch.quantization.get_default_qconfig('qnnpack')
# or, set the qconfig for QAT
qconfig = torch.quantization.get_default_qat_qconfig('qnnpack')
# set the qengine to control weight packing
torch.backends.quantized.engine = 'qnnpack'
```

## Quantization Customizations

While default implementations of observers to select the scale factor and bias based on observed tensor data are provided, developers can provide their own quantization functions. Quantization can be applied selectively to different parts of the model or configured differently for different parts of the model.

We also provide support for per channel quantization for `conv2d()`, `conv3d()` and `linear()`

Quantization workflows work by adding (e.g. adding observers as `.observer` submodule) or replacing (e.g. converting `nn.Conv2d` to `nn.quantized.Conv2d`) submodules in the model's module hierarchy. It means that the model stays a regular `nn.Module`-based instance throughout the process and thus can work with the rest of PyTorch APIs.

### Quantization Custom Module API

Both Eager mode and FX graph mode quantization APIs provide a hook for the user to specify module quantized in a custom way, with user defined logic for observation and quantization. The user needs to specify:

1. The Python type of the source fp32 module (existing in the model)
2. The Python type of the observed module (provided by user). This module needs to define a `from_float` function which defines how the observed module is created from the original fp32 module.
3. The Python type of the quantized module (provided by user). This module needs to define a `from_observed` function which defines how the quantized module is created from the observed module.
4. A configuration describing (1), (2), (3) above, passed to the quantization APIs.

The framework will then do the following:

1. during the *prepare* module swaps, it will convert every module of type specified in (1) to the type specified in (2), using the `from_float` function of the class in (2).
2. during the *convert* module swaps, it will convert every module of type specified in (2) to the type specified in (3), using the `from_observed` function of the class in (3).

Currently, there is a requirement that *ObservedCustomModule* will have a single Tensor output, and an observer will be added by the framework (not by the user) on that output. The observer will be stored under the `activation_post_process` key as an attribute of the custom module instance. Relaxing these restrictions may be done at a future time.

Example:

```

import torch
import torch.nn.quantized as nnq
import torch.quantization.quantize_fx

# original fp32 module to replace
class CustomModule(torch.nn.Module):
    def __init__(self):
        super().__init__()
        self.linear = torch.nn.Linear(3, 3)

    def forward(self, x):
        return self.linear(x)

# custom observed module, provided by user
class ObservedCustomModule(torch.nn.Module):
    def __init__(self, linear):
        super().__init__()
        self.linear = linear

    def forward(self, x):
        return self.linear(x)

    @classmethod
    def from_float(cls, float_module):
        assert hasattr(float_module, 'qconfig')
        observed = cls(float_module.linear)
        observed.qconfig = float_module.qconfig
        return observed

# custom quantized module, provided by user
class StaticQuantCustomModule(torch.nn.Module):
    def __init__(self, linear):
        super().__init__()
        self.linear = linear

    def forward(self, x):
        return self.linear(x)

    @classmethod
    def from_observed(cls, observed_module):
        assert hasattr(observed_module, 'qconfig')
        assert hasattr(observed_module, 'activation_post_process')
        observed_module.linear.activation_post_process = \
            observed_module.activation_post_process
        quantized = cls(nnq.Linear.from_float(observed_module.linear))
        return quantized

#
# example API call (Eager mode quantization)
#

m = torch.nn.Sequential(CustomModule()).eval()

prepare_custom_config_dict = {
    "float_to_observed_custom_module_class": {
        CustomModule: ObservedCustomModule
    }
}
convert_custom_config_dict = {
    "observed_to_quantized_custom_module_class": {
        ObservedCustomModule: StaticQuantCustomModule
    }
}

m.qconfig = torch.quantization.default_qconfig
mp = torch.quantization.prepare(
    m, prepare_custom_config_dict=prepare_custom_config_dict)
# calibration (not shown)
mq = torch.quantization.convert(
    mp, convert_custom_config_dict=convert_custom_config_dict)

#
# example API call (FX graph mode quantization)
#

m = torch.nn.Sequential(CustomModule()).eval()

qconfig_dict = {'': torch.quantization.default_qconfig}
prepare_custom_config_dict = {
    "float_to_observed_custom_module_class": {
        "static": {
            CustomModule: ObservedCustomModule,
        }
    }
}
convert_custom_config_dict = {
    "observed_to_quantized_custom_module_class": {
        "static": {
            ObservedCustomModule: StaticQuantCustomModule,
        }
    }
}

mp = torch.quantization.quantize_fx.prepare_fx(
    m, qconfig_dict, prepare_custom_config_dict=prepare_custom_config_dict)
# calibration (not shown)
mq = torch.quantization.quantize_fx.convert_fx(
    mp, convert_custom_config_dict=convert_custom_config_dict)

```

## Model Preparation for Quantization (Eager Mode)

It is necessary to currently make some modifications to the model definition prior to Eager mode quantization. This is because

currently quantization works on a module by module basis. Specifically, for all quantization techniques, the user needs to:

1. Convert any operations that require output requantization (and thus have additional parameters) from functionals to module form (for example, using `torch.nn.ReLU` instead of `torch.nn.functional.relu`).
2. Specify which parts of the model need to be quantized either by assigning `.qconfig` attributes on submodules or by specifying `qconfig_dict`. For example, setting `model.conv1.qconfig = None` means that the `model.conv` layer will not be quantized, and setting `model.linear1.qconfig = custom_qconfig` means that the quantization settings for `model.linear1` will be using `custom_qconfig` instead of the global `qconfig`.

For static quantization techniques which quantize activations, the user needs to do the following in addition:

1. Specify where activations are quantized and de-quantized. This is done using `:class:`~torch.quantization.QuantStub`` and `:class:`~torch.quantization.DeQuantStub`` modules.

```
System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\[pytorch-master] [docs]
[source]quantization.rst, line 748); backlink
Unknown interpreted text role "class".
```

```
System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\[pytorch-master] [docs]
[source]quantization.rst, line 748); backlink
Unknown interpreted text role "class".
```

2. Use `:class:`~torch.nn.quantized.FloatFunctional`` to wrap tensor operations that require special handling for quantization into modules. Examples are operations like `add` and `cat` which require special handling to determine output quantization parameters.

```
System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\[pytorch-master] [docs]
[source]quantization.rst, line 751); backlink
Unknown interpreted text role "class".
```

3. Fuse modules: combine operations/modules into a single module to obtain higher accuracy and performance. This is done using the `:func:`~torch.quantization.fuse_modules`` API, which takes in lists of modules to be fused. We currently support the following fusions: `[Conv, ReLU]`, `[Conv, BatchNorm]`, `[Conv, BatchNorm, ReLU]`, `[Linear, ReLU]`

```
System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\[pytorch-master] [docs]
[source]quantization.rst, line 755); backlink
Unknown interpreted text role "func".
```

## Best Practices

1. Set the `reduce_range` argument on observers to `True` if you are using the `fbgemm` backend. This argument prevents overflow on some `int8` instructions by reducing the range of quantized data type by 1 bit.

## Common Errors

### Passing a non-quantized Tensor into a quantized kernel

If you see an error similar to:

```
RuntimeError: Could not run 'quantized::some_operator' with arguments from the 'CPU' backend...
```

This means that you are trying to pass a non-quantized Tensor to a quantized kernel. A common workaround is to use `torch.quantization.QuantStub` to quantize the tensor. This needs to be done manually in Eager mode quantization. An e2e example:

```
class M(torch.nn.Module):
    def __init__(self):
        super().__init__()
        self.quant = torch.quantization.QuantStub()
        self.conv = torch.nn.Conv2d(1, 1, 1)

    def forward(self, x):
        # during the convert step, this will be replaced with a
        # `quantize_per_tensor` call
        x = self.quant(x)
        x = self.conv(x)
        return x
```

### Passing a quantized Tensor into a non-quantized kernel

If you see an error similar to:

```
RuntimeError: Could not run 'aten::thnn_conv2d_forward' with arguments from the 'QuantizedCPU' backend.
```

This means that you are trying to pass a quantized Tensor to a non-quantized kernel. A common workaround is to use `torch.quantization.DeQuantStub` to dequantize the tensor. This needs to be done manually in Eager mode quantization. An e2e example:

```
class M(torch.nn.Module):
    def __init__(self):
        super().__init__()
```



```

self.quant = torch.quantization.QuantStub()
self.conv1 = torch.nn.Conv2d(1, 1, 1)
# this module will not be quantized (see `qconfig = None` logic below)
self.conv2 = torch.nn.Conv2d(1, 1, 1)
self.dequant = torch.quantization.DeQuantStub()

def forward(self, x):
    # during the convert step, this will be replaced with a
    # `quantize_per_tensor` call
    x = self.quant(x)
    x = self.conv1(x)
    # during the convert step, this will be replaced with a
    # `dequantize` call
    x = self.dequant(x)
    x = self.conv2(x)
    return x

m = M()
m.qconfig = some_qconfig
# turn off quantization for conv2
m.conv2.qconfig = None

```

## Saving and Loading Quantized models

When calling `torch.load` on a quantized model, if you see an error like:

```
AttributeError: 'LinearPackedParams' object has no attribute '_modules'
```

This is because directly saving and loading a quantized model using `torch.save` and `torch.load` is not supported. To save/load quantized models, the following ways can be used:

### 1. Saving/Loading the quantized model state\_dict

An example:

```

class M(torch.nn.Module):
    def __init__(self):
        super().__init__()
        self.linear = nn.Linear(5, 5)
        self.relu = nn.ReLU()

    def forward(self, x):
        x = self.linear(x)
        x = self.relu(x)
        return x

m = M().eval()
prepare_orig = prepare_fx(m, {'': default_qconfig})
prepare_orig(torch.rand(5, 5))
quantized_orig = convert_fx(prepare_orig)

# Save/load using state_dict
b = io.BytesIO()
torch.save(quantized_orig.state_dict(), b)

m2 = M().eval()
prepared = prepare_fx(m2, {'': default_qconfig})
quantized = convert_fx(prepared)
b.seek(0)
quantized.load_state_dict(torch.load(b))

```

### 2. Saving/Loading scripted quantized models using `torch.jit.save` and `torch.jit.load`

An example:

```

# Note: using the same model M from previous example
m = M().eval()
prepare_orig = prepare_fx(m, {'': default_qconfig})
prepare_orig(torch.rand(5, 5))
quantized_orig = convert_fx(prepare_orig)

# save/load using scripted model
scripted = torch.jit.script(quantized_orig)
b = io.BytesIO()
torch.jit.save(scripted, b)
b.seek(0)
scripted_quantized = torch.jit.load(b)

```

## Numerical Debugging (prototype)

### Warning

Numerical debugging tooling is early prototype and subject to change.

- `ref:torch_ao_ns_numeric_suite` Eager mode numeric suite

**System Message: ERROR/3** (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\pytorch-master) [docs]  
`[source] quantization.rst, line 896); backlink`

Unknown interpreted text role "ref".

- `ref:torch_ao_ns_numeric_suite_fx` FX numeric suite

**System Message: ERROR/3** (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\pytorch-master) [docs]  
`[source] quantization.rst, line 898); backlink`

Unknown interpreted text role "ref".

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\pytorch-master [docs] [source] quantization.rst, line 904)**

Unknown directive type "py:module".

```
.. py:module:: torch.ao
```

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\pytorch-master [docs] [source] quantization.rst, line 905)**

Unknown directive type "py:module".

```
.. py:module:: torch.ao.nn
```

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\pytorch-master [docs] [source] quantization.rst, line 906)**

Unknown directive type "py:module".

```
.. py:module:: torch.ao.nn.sparse
```

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\pytorch-master [docs] [source] quantization.rst, line 907)**

Unknown directive type "py:module".

```
.. py:module:: torch.ao.nn.sparse.quantized
```

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\pytorch-master [docs] [source] quantization.rst, line 908)**

Unknown directive type "py:module".

```
.. py:module:: torch.ao.nn.sparse.quantized.dynamic
```

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\pytorch-master [docs] [source] quantization.rst, line 909)**

Unknown directive type "py:module".

```
.. py:module:: torch.ao.ns
```

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\pytorch-master [docs] [source] quantization.rst, line 910)**

Unknown directive type "py:module".

```
.. py:module:: torch.ao.ns.fx
```

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\pytorch-master [docs] [source] quantization.rst, line 911)**

Unknown directive type "py:module".

```
.. py:module:: torch.ao.quantization
```

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\pytorch-master [docs] [source] quantization.rst, line 912)**

Unknown directive type "py:module".

```
.. py:module:: torch.ao.quantization.fx
```

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\pytorch-master [docs] [source] quantization.rst, line 913)**

Unknown directive type "py:module".

```
.. py:module:: torch.ao.quantization.fx.backend_config
```

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\pytorch-master [docs] [source] quantization.rst, line 914)**

Unknown directive type "py:module".

```
.. py:module:: torch.ao.sparsity
```

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\[pytorch-master] [docs] [source]quantization.rst, line 915)**

Unknown directive type "py:module".

```
.. py:module:: torch.ao.sparsity.experimental
```

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\[pytorch-master] [docs] [source]quantization.rst, line 916)**

Unknown directive type "py:module".

```
.. py:module:: torch.ao.sparsity.experimental.pruner
```

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\[pytorch-master] [docs] [source]quantization.rst, line 917)**

Unknown directive type "py:module".

```
.. py:module:: torch.ao.sparsity.scheduler
```

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\[pytorch-master] [docs] [source]quantization.rst, line 918)**

Unknown directive type "py:module".

```
.. py:module:: torch.ao.sparsity.sparsifier
```