

Swift Compiler Performance

This document is a guide to understanding, diagnosing and reporting compilation-performance problems in the swift compiler. That is: the speed at which the compiler compiles code, not the speed at which that code runs.

While this guide is lengthy, it should all be relatively straightforward. Performance analysis is largely a matter of patience, thoroughness and perseverance, measuring carefully and consistently, and gradually eliminating noise and focusing on a signal.

Table of Contents

- [Swift Compiler Performance](#)
 - [Table of Contents](#)
 - [Outline of processes and factors affecting compilation performance](#)
 - [Compilation modes](#)
 - [Primary-file vs. WMO](#)
 - [Amount of optimization](#)
 - [Workload variability, approximation and laziness](#)
 - [Incremental compilation](#)
 - [Lazy resolution](#)
 - [Summing up: high level picture of compilation performance](#)
 - [Known problem areas](#)
 - [How to diagnose compilation performance problems](#)
 - [Tools and options](#)
 - [Profilers](#)
 - [Instruments.app](#)
 - [Perf](#)
 - [Diagnostic options](#)
 - [Post-processing tools for diagnostics](#)
 - [Artifact-analysis tools](#)
 - [Minimizers](#)
 - [Git bisect](#)
 - [Creduce](#)
 - [General bisection](#)
 - [Isolating a regression](#)
 - [Driver diagnosis](#)
 - [Finding areas in need of general improvement](#)
 - [Compiler counters](#)
 - [Scale-test](#)
 - [How to report bugs most usefully](#)
 - [How else to help out if you want to improve compile times](#)

Outline of processes and factors affecting compilation performance

This section is intended to provide a high-level orientation around what the compiler is doing when it's run -- beyond the obvious "compiling" -- and what major factors influence how much time it spends.

When you compile or run a Swift program, either with Xcode or on the command line, you typically invoke `swift` or `swiftc` (the latter is a symbolic link to the former), which is a program that can behave in very different ways depending on its arguments.

It may compile or execute code directly, but it will usually instead turn around and run one or more copies of `swift` or `swiftc` as subprocesses. In typical batch compilation, the first copy of `swiftc` runs as a so-called **driver** process, and it then executes a number of so-called **frontend** subprocesses, in a process tree. It's essential, when interpreting Swift compilation, to have a clear picture of which processes are run and what they're doing:

- **Driver:** the top-level `swiftc` process in a tree of subprocesses. Responsible for deciding which files need compiling or recompiling and running child processes — so-called **jobs** — to perform compilation and linking steps. For most of its execution, it is idle, waiting for subprocesses to complete.
- **Frontend Jobs:** subprocesses launched by the driver, running `swift -frontend ...` and performing compilation, generating PCH files, merging modules, etc. These are the jobs that incur the bulk of the costs of compiling.
- **Other Jobs:** subprocesses launched by the driver, running `ld`, `swift -modulewrap`, `swift-autolink-extract`, `dsymutil`, `dwarfdump` and similar tools involved in finishing off a batch of work done by the frontend jobs. Some of these will be the `swift` program too, but they're not "doing frontend jobs" and so will have completely different profiles.

The set of jobs that are run, and the way they spend their time, is itself highly dependent on **compilation modes**. Information concerning those modes that's relevant to compilation performance is recounted in the following section; for more details on the driver, see [the driver docs](#), as well as docs on [driver internals](#) and [driver parseable output](#).

After discussing compilation modes in the following section, we'll also touch on large-scale variation in workload that can occur *without* obvious hotspots, in terms of laziness strategies and approximations.

Compilation modes

There are many different options for controlling the driver and frontend jobs, but the two dimensions that cause the most significant variation in behaviour are often referred to as *modes*. These modes make the biggest difference, and it's important when looking at compilation to be clear on which mode `swiftc` is running in, and often to perform separate analysis for each mode. The significant modes are:

- **Primary-file vs. whole-module:** this varies depending on whether the driver is run with the flag `-wmo` (a.k.a. `-whole-module-optimization`).
 - **Batch vs. single-file** primary-file mode. This distinction refines the behaviour of primary-file mode, with the new batch mode added in the Swift 4.2 release cycle. Batching eliminates much of the overhead of primary-file mode, and will eventually become the default way of running primary-file mode, but until that time it is explicitly enabled by passing the `-enable-batch-mode` flag.
- **Optimizing vs. non-optimizing:** this varies depending on whether the driver (and thus each frontend) is run with the flags `-O`, `-Osize`, or `-Ounchecked` (each of which turn on one or more sets of optimizations), or the default (no-optimization) which is synonymous with `-Onone` or `-Oplayground`.

When you build a program in Xcode or using `xcodebuild`, often there is a *configuration* parameter that will switch both of these modes simultaneously. That is, typical code has two configurations:

- **Debug** which combines primary-file mode with `-Onone`
- **Release** which combines WMO mode with `-O`

But these parameters can be varied independently and the compiler will spend its time very differently depending on their settings, so it's worth understanding both dimensions in a bit more detail.

Primary-file (with and without batching) vs. WMO

This is the most significant variable in how the compiler behaves, so it's worth getting perfectly clear:

- In **primary-file mode**, the driver divides the work it has to do between multiple frontend processes, emitting partial results and merging those results when all the frontends finish. Each frontend job itself reads *all* the files in the module, and focuses on one or more *primary* file(s) among the set it read, which it compiles, lazily analyzing other referenced definitions from the module as needed. This mode has two sub-modes:
 - In the **single-file** sub-mode, it runs *one frontend job per file*, with each job having a single primary.
 - In the **batch** sub-mode, it runs *one frontend job per CPU*, identifying an equal-sized "batch" of the module's files as primaries.
- In **whole-module optimization (WMO) mode**, the driver runs one frontend job for the entire module, no matter what. That frontend reads all the files in the module *once* and compiles them all at once.

For example: if your module has 100 files in it:

- Running `swiftc *.swift` will compile in **single-file mode**, and will thus run 100 frontend subprocesses, each of which will parse all 100 inputs (for a total of 10,000 parses), and then each subprocess will (in parallel) compile the definitions in its single primary file.
- Running `swiftc -enable-batch-mode *.swift` will compile in **batch** mode, and on a system with 4 CPUs will run 4 frontend subprocesses, each of which will parse all 100 inputs (for a total of 400 parses), and then each subprocess will (in parallel) compile the definitions of 25 primary files (one quarter of the module in each process).
- Running `swiftc -wmo *.swift` will compile in **whole-module** mode, and will thus run *one* frontend subprocess, which then reads all 100 files *once* (for a total of 100 parses) and compiles the definitions in all of them, in order (serially).

Why do multiple modes exist? Because they have different strengths and weaknesses; neither is perfect:

- Primary-file mode's advantages are that the driver can do **incremental compilation** by only running frontends for files that it thinks are out of date, as well as running multiple frontend jobs **in parallel**, making use of multiple cores. Its disadvantage is that each frontend job has to read *all the source files* in the module before focusing on its primary-files of interest, which means that a *portion* of the frontend job's work is being done *quadratically* in the number of jobs. Usually this portion is relatively small and fast, but because it's quadratic, it can easily go wrong. The addition of **batch mode** was specifically to eliminate this quadratic increase in early work.
- WMO mode's advantages are that it can do certain optimizations that only work when they are sure they're looking at the entire module, and it avoids the quadratic work in the early phases of primary-file mode. Its disadvantages are that it always rebuilds everything, and that it exploits parallelism worse (at least before LLVM IR code-generation, which is always multithreaded).

Whole-module mode does enable a set of optimizations that are not possible when compiling in primary-file mode. In particular, in modules with a lot of private dead code, whole-module mode can eliminate the dead code earlier and avoid needless work compiling it, making for both smaller output and faster compilation.

It is therefore possible that, in certain cases (such as with limited available parallelism / many modules built in parallel), building in whole-module mode with optimization disabled can complete in less time than batched primary-file mode. This scenario depends on many factors seldom gives a significant advantage, and since using it trades-away support for incremental compilation entirely, it is not a recommended configuration.

Amount of optimization

This document isn't the right place to give a detailed overview of the compiler architecture, but it's important to keep in mind that the compiler deals with Swift code in memory in 3 major representations, and can therefore be conceptually divided into 3 major stages, the latter 2 of which behave differently depending on optimization mode:

- **ASTs** (Abstract Syntax Trees): this is the representation (defined in the `lib/AST` directory) closest to what's in a source file, produced from Swift source code, Swift modules and Clang modules (in `lib/Parse`, `lib/Serialization` and `lib/ClangImporter` respectively) and interpreted by resolution, typechecking and high-level semantics functions (in `lib/Sema`) early-on in compilation.
- **SIL** (Swift Intermediate Language): this is a form that's private to the Swift compiler, lower-level and more-explicit than the AST representation, but still higher-level and more Swift-specific than a machine-oriented representation like LLVM. It's defined in `lib/SIL`, produced by code in `lib/SILGen` and *optionally* optimized by code in `lib/SILOptimizer`.
- **LLVM IR** (Low Level Virtual Machine Intermediate Representation): this is a form that's an abstract representation of the machine language being compiled for; it doesn't contain any Swift-specific knowledge, rather it's a form the Swift compiler *generates from SIL* (in `lib/IRGen`) and then hands off as input to the [LLVM backend](#), a library upon which the Swift compiler depends. LLVM has its own *optional optimizations* that apply to LLVM IR before it's lowered to machine code.

When running the Swift compiler in optimizing mode, many SIL and LLVM optimizations are turned on, making those phases of compilation (in each frontend job) take significantly more time and memory. When running in non-optimizing mode, SIL and LLVM IR are still produced and consumed along the way, but only as part of lowering, with comparatively few "simple" optimizations applied.

Additionally, the IRGen and LLVM phases can operate (and usually are operated) in parallel, using multiple threads in each frontend job, as controlled by the `-num-threads` flag. This option only applies to the latter phases, however: the AST and SIL-related phases never run multithreaded.

The amount of work done to the AST representation (in particular: importing, resolving and typechecking ASTs) does not vary between different optimization modes. However, it does vary *significantly* between different projects and among seemingly-minor changes to code, depending on the amount of laziness the frontend is able to exploit.

Workload variability, approximation and laziness

While some causes of slow compilation have definite hotspots (which we will get to shortly), one final thing to keep in mind when doing performance analysis is that the compiler tries to be *lazy* in a variety of ways, and that laziness does not always work: it is driven by certain approximations and assumptions that often err on the side of doing more work than strictly necessary.

The outcome of a failure in laziness is not usually a visible hotspot in a profile: rather, it's the appearance of doing "too much work altogether" across a generally-flat profile. Two areas in particular where this occurs — and where there are significant, ongoing improvements to be made — are in incremental compilation and lazy resolution.

Incremental compilation

As mentioned in the section on primary-file mode, the driver has an *incremental mode* that can be used to attempt to avoid running frontend jobs entirely. When successful, this is the most effective form of time-saving possible: nothing is faster than a process that doesn't even run.

Unfortunately judgements about when a file "needs recompiling" are themselves driven by an auxiliary data structure that summarizes the dependencies between files, and this data structure is necessarily a conservative approximation. The approximation is weaker than it should be, and as a result the driver often runs more frontend jobs than it should.

Lazy resolution

Swift source files contain names that refer to definitions outside the enclosing file, and frequently outside of the enclosing module. These "external" definitions are resolved lazily from two very different locations (both called "modules"):

- C/ObjC modules, provided by the Clang importer
- Serialized Swift modules

Despite their differences, both kinds of modules support laziness in the Swift compiler in one crucial way: they are both kinds of *indexed* binary file formats that permit loading *single definitions* out of by name, without having to load the entire contents of the module.

When the Swift compiler manages to be lazy and limit the number of definitions it tries to load from modules, it can be very fast; the file formats support very cheap access. But often the logic in the Swift compiler is unnecessarily conservative about exploiting this potential laziness, and so it loads more definitions than it should.

Summing up: high level picture of compilation performance

Swift compilation performance varies *significantly* by at least the following parameters:

- WMO vs. primary-file (non-WMO) mode, including batching thereof
- Optimizing vs. non-optimizing mode
- Quantity of incremental work avoided (if in non-WMO)
- Quantity of external definitions lazily loaded

When approaching Swift compilation performance, it's important to be aware of these parameters and keep them in mind, as they tend to frame the problem you're analyzing: changing one (or any of the factors influencing them, in a project) will likely completely change the resulting profile.

Known problem areas

These are areas where we know the compiler has room for improvement, performance-wise, where it's worth searching for existing bugs on the topic, finding an existing team member who knows the area, and trying to relate the problem you're seeing to some of the existing strategies and plans for improvement:

- Incremental mode is over-approximate, runs too many subprocesses.
- Too many referenced (non-primary-file) definitions are type-checked beyond the point they need to be, during the quadratic phase.
- Expression type inference solves constraints inefficiently, and can sometimes behave super-linearly or even exponentially.
- Periodically the analysis phase of a SIL optimization fails to cache overlapping subproblems, causing a super-linear slowdown.
- Some SIL-to-IR lowerings (eg. large value types) can generate too much LLVM IR, increasing the time spent in LLVM.

(Subsystem experts: please add further areas of concern here.)

How to diagnose compilation performance problems

Compiler performance analysis breaks down into two broad categories of work, depending on what you're trying to do:

- Isolating a regression
- Finding areas that need general improvement

In all cases, it's important to be familiar with several tools and compiler options we have at our disposal. If you know about all these tools, you can skip the following section.

Tools and options

You'll use several tools along the way. These come in 5 main categories:

- Profilers
- Diagnostic options built-in to the compiler (timers, counters)
- Post-processing tools to further analyze diagnostic output
- Tools to generally analyze the output artifacts of the compiler
- Tools to minimize the regression range or testcases

Profilers

The basic tool of performance analysis is a profiler, and you *will need* to learn to use at least one profiler for the purposes of this work. The main two profilers we use are `Instruments.app` on macOS, and `perf(1)` on Linux. Both are freely available and extremely powerful; this document will barely scratch the surface of what they can do.

Instruments.app

Instruments is a tool on macOS that ships as part of Xcode. It contains graphical and batch interfaces to a very wide variety of profiling services; see [here](#) for more documentation.

The main way we will use `Instruments.app` is in "Counter" mode, to record and analyze a single run of swiftc. We will also use it in simple push-button interactive mode, as a normal application. While it's possible to run Instruments in batch mode on the command-line, the batch interface is less reliable than running it as an interactive application, and frequently causes lockups or fails to collect data.

Before starting, you should also be sure you are going to profile a version of Swift *without* DWARF debuginfo; while in theory debuginfo will give a higher-resolution, more-detailed profile, in practice Instruments will often stall out and become unresponsive trying to process the additional detail.

Similarly, be sure that as many applications as possible (especially those with debuginfo themselves!) are closed, so that Instruments has little additional material to symbolicate as possible. It collects a *whole system profile* at very high resolution, so you want to make its life easy by profiling on a quiet machine doing little beyond the task you're interested in.

Once you're ready, follow these steps:

- Open `Xcode.app`
- Click `Xcode => Open Developer Tool => Instruments` (Once it's open, you might want to pin `Instruments.app` to the dock for ease of access)
- Select the `Counters` profiling template
- Open a terminal and get prepared to run your test-case
- Switch back to `Instruments.app`
- Press the red `record` button in the top-left of the instruments panel

- Quickly switch to your terminal, run the test-case you wish to profile, and as soon as it's finished switch back to `Instruments.app` and press the stop button.

That's it! You should have a profile gathered.

Ideally you want to get to a situation that looks like this:



Instruments Profile with terminal

In the main panel you can see a time-sorted set of process and call-frame samples, which you can filter to show only swift processes by typing `swift` in the `Input Filter` box at the bottom of the window. Each line in the main panel can be expanded by clicking the triangle at its left, showing the callees as indented sub-frames.

If you hover over the line corresponding to a specific `swift` process, you'll see a small arrow enclosed in a grey circle to the right of the line. Click on it and instruments will shift focus of the main panel to just that process' subtree (and recalculate time-percentages accordingly). Once you're focused on a specific `swift` process, you can begin looking at its individual stack-frame profile.

In the panel to the right of the main panel, you can see the heaviest stack trace within the currently-selected line of the main panel. If you click on one of the frames in that stack, the main panel will automatically expand every level between the current frame and the frame you clicked on. For example, clicking 11 frames down the hottest stack, on the frame called `swift::ModuleFile::getModule`, will expand the main panel to show something like this:



Instruments Profile with terminal

Click around a profile by expanding and contracting nodes in the stack tree, and you'll pretty quickly get a feeling for where the program is spending its time. Each line in the main display shows both the cumulative sample count and running time of its subtree (including all of its children), as well as its own frame-specific `Self` time.

In the example above, it's pretty clear that the compiler is spending 66% of its time in `Sema`, and the heaviest stack inside there is the time spent deserializing external definitions (which matches a known problem area, mentioned earlier).

If you want to keep notes on what you're seeing while exploring a profile, you can expand and collapse frames until you see a meaningful pattern, then select the displayed set of stack frames and copy them as text (using `⌘-C` as usual) and paste it into a text file; whitespace indentation will be inserted in the copied text, to keep the stack structure readable.

If you have two profiles and want to compare them, Instruments does have a mode for direct diffing between profiles, but it doesn't work when the profiles are gathered from different binaries, so for purposes of comparing different swift compilers, you'll typically have to do manual comparison of the profiles.

Perf

Perf is a Linux profiler that runs on the command line. In many Linux distributions it's included in a package called `linux-tools` that needs to be separately installed. It's small, fast, robust, flexible, and can be easily scripted; the main disadvantages are that it lacks any sort of GUI and only runs on Linux, so you can't use it to diagnose problems in builds that need macOS or iOS frameworks or run under `xcodebuild`.

Perf is documented [on the kernel wiki](#) as well as on [Brendan Gregg's website](#).

Using `perf` requires access to hardware performance counters, so you cannot use it in most virtual machines (unless they virtualize access to performance counters). Further, you will need root access to give yourself permission to use the profiling interface of the kernel.

The simplest use of `perf` just involves running your command under `perf stat`. This gives high level performance counters including an instructions-executed count, which is a comparatively-stable approximation of total execution cost, and is often enough to pick out a regression when bisecting (see below):

```
$ perf stat swiftc t.swift

Performance counter stats for 'swiftc t.swift':

      2140.543052      task-clock (msec)      #    0.966 CPUs utilized
           17          context-switches      #    0.008 K/sec
           6          cpu-migrations        #    0.003 K/sec
      52,084          page-faults            #    0.024 M/sec
 5,373,530,212        cycles                 #    2.510 GHz
 9,709,304,679        instructions          #    1.81  insn per cycle
1,812,011,233         branches              #   846.519 M/sec
   22,026,587         branch-misses        #    1.22% of all branches

2.216754787 seconds time elapsed
```


The fact that `perf` gives relatively stable and precise cost measurements means that it can be made into a useful subroutine when doing other performance-analysis tasks, such as bisecting (see section on `git bisect`) or reducing (see section on `creduce`). A shell function like the following is very useful:

```
count_instructions() {
    perf stat -x , --log-fd 3 \
        -e instructions -r 10 "$@" \
        3>&1 2>/dev/null 1>&2 | cut -d , -f 1
}
```

To gather a full profile with `perf` -- when not just using it as a batch counter -- use the `perf record` and `perf report` commands; depending on configuration you might need to play with the `--call-graph` and `-e` parameters to get a clear picture:

```
$ perf record -e cycles -c 10000 --call-graph=lbr swiftc t.swift
[ perf record: Woken up 5 times to write data ]
[ perf record: Captured and wrote 1.676 MB perf.data (9731 samples) ]
```

Once recorded, data will be kept in a file called `perf.data`, which is the default file acted-upon by `perf report`. Running it should give you something like the following textual user interface, which operates similarly to `Instruments.app`, only using cursor keys:

 `perf report`

Diagnostic options

The Swift compiler has a variety of built-in diagnostic options. Some are interpreted by the driver, others are interpreted by the frontend jobs that the driver runs: these have to be passed on the driver command-line with `-Xfrontend` to get passed through to the frontends. In a multi-frontend, primary-file-mode compilation, any such `-Xfrontend` option will be passed to *all* frontend job, which means the diagnostic output from *all* frontend jobs will be sent to standard output in sequence. This makes diagnosing problems with these options somewhat challenging; they work better if you can reduce the problem to a single frontend process before using these options.

Further, some of the diagnostic options depend on instrumentation that's only present in assert compilers; if you pass them to release compilers, you will get limited or no output. It's often useful to have both release and assert compilers on hand while you're working.

- `-driver-time-compilation` : interpreted by the driver; emits a high-level timing of the frontend and other jobs that the driver executes. It can guide early investigation to see which file in a primary-file-mode compilation is taking the majority of time, or is taking more or less time than when comparing compilations. Its output looks like this:

```
=====
                        Driver Compilation Time
=====
Total Execution Time: 0.0001 seconds (0.0490 wall clock)

---User Time---  --System Time--  --User+System--  ---Wall Time---  ---
Name ---
  0.0000 ( 82.0%)   0.0001 ( 59.5%)   0.0001 ( 69.0%)   0.0284 ( 58.0%)
{compile: t-177627.o <= t.swift}
  0.0000 ( 18.0%)   0.0000 ( 40.5%)   0.0000 ( 31.0%)   0.0206 ( 42.0%) {link:
t <= t-177627.o}
  0.0001 (100.0%)   0.0001 (100.0%)   0.0001 (100.0%)   0.0490 (100.0%) Total
```

- `-Xfrontend -debug-time-function-bodies` : asks each frontend to print out the time spent typechecking *every function* in the program, sorted by time taken. The output is therefore voluminous, but can help when reducing a testcase to the "one bad function" that causes it. The output looks like this:

```
9.16ms test.swift:15:6 func find<R>(_ range: R, value: R.Element) -> R where R
: IteratorProtocol, R.Element : Eq
0.28ms test.swift:27:6 func findIf<R>(_ range: R, predicate: (R.Element) ->
Bool) -> R where R : IteratorProtocol
2.81ms test.swift:40:6 func count<R>(_ range: R, value: R.Element) -> Int where
R : IteratorProtocol, R.Element : Eq
0.64ms test.swift:51:6 func countIf<R>(_ range: R, predicate: (R.Element) ->
Bool) -> Int where R : IteratorProtocol
...
```

- `-Xfrontend -debug-time-expression-type-checking` : similar to `-debug-time-function-bodies` , but prints a separate timer for *every expression* in the program, much more detail than just the functions. The output looks like this:

```
0.20ms test.swift:17:16
1.82ms test.swift:18:12
6.35ms test.swift:19:8
0.11ms test.swift:22:5
0.02ms test.swift:24:10
0.02ms test.swift:30:16
...
```

- `-Xfrontend -print-stats` : activates statistic counters inside each frontend, printing them out when the frontend exits. By default, most statistics are enabled only in assert builds, so in a release build this option will do nothing. In an assert build, its output will look like this:

```

=====
... Statistics Collected ...
=====

 4 BasicCalleeAnalysis          - computeMethodCallees
 6 Clang module importer         - # of imported clang entities
11 Constraint solver largest system - # of connected components split
176 Constraint solver largest system - # of disjunction terms explored
 9 Constraint solver largest system - # of disjunctions explored
59 Constraint solver largest system - # of constraints simplified
 0 Constraint solver largest system - # of simplification iterations
232 Constraint solver largest system - # of solution states explored
42 Constraint solver largest system - # of type variable bindings attempted
38 Constraint solver largest system - # of type variables bound
79 Constraint solver largest system - # of constraints not simplified
94 Constraint solver largest system - # of the largest solution attempt
 6 Constraint solver overall      - Number of solutions discarded
361 Constraint solver overall      - # of solution attempts
130 Constraint solver overall      - # of connected components split
1898 Constraint solver overall     - # of disjunction terms explored
119 Constraint solver overall      - # of disjunctions explored
844 Constraint solver overall      - # of constraints simplified
...

```

- `-Xfrontend -print-clang-stats` : prints counters associated with the clang AST reader, which is operated as a subsystem of the swift compiler when importing definitions from C/ObjC. Its output is added to the end of whatever output comes from `-print-stats` , and looks like this:

```

*** AST File Statistics:
1/194 source location entries read (0.515464%)
5/182 types read (2.747253%)
7/318 declarations read (2.201258%)
6/251 identifiers read (2.390438%)
0/22 macros read (0.000000%)
0/70 statements read (0.000000%)
0/22 macros read (0.000000%)
1/22 lexical declcontexts read (4.545455%)
0/16 visible declcontexts read (0.000000%)
6 / 6 identifier table lookups succeeded (100.000000%)
...

```

- `-Xfrontend -print-stats -Xfrontend -print-inst-counts` : an extended form of `-print-stats` that activates a separate statistic counter for every kind of SIL instruction generated during compilation. Its output looks like this:

```

...
163 sil-instcount          - Number of AllocStackInst
 90 sil-instcount          - Number of ApplyInst
 92 sil-instcount          - Number of BeginAccessInst
212 sil-instcount          - Number of BranchInst
 80 sil-instcount          - Number of BuiltinInst

```

37 sil-instcount	- Number of CondBranchInst
6 sil-instcount	- Number of CondFailInst
136 sil-instcount	- Number of CopyAddrInst
177 sil-instcount	- Number of DeallocStackInst
21 sil-instcount	- Number of DebugValueAddrInst
6 sil-instcount	- Number of DebugValueInst
128 sil-instcount	- Number of DestroyAddrInst
...	

Unified stats reporter

In an attempt to unify collection and reporting of the various statistic-gathering options, recent versions of the compiler support a partly redundant command `-stats-output-dir <directory>` that writes *all* driver and primary frontend counters and timers (though not per-function timers) to JSON files in `<directory>`.

This option also provides *some* high-level counters that are "always available" regardless of whether you're using an assert or release build, though assert builds still get *more* counters (all of those available through `-print-stats`). If you are using a new-enough compiler, `-stats-output-dir` often simplifies analysis, since its output is machine-readable and aggregates all the jobs in a multi-job compilation, and there's a post-processing script `utils/process-stats-dir.py` to work with these files in aggregate.

For example, to compile a file with the unified stats reporter enabled, first make a directory in which to output the stats, then compile with the `-stats-output-dir` flag:

```
$ mkdir /tmp/stats
$ swiftc -c test.swift -stats-output-dir /tmp/stats
$ ls /tmp/stats
stats-1518219149045080-swift-frontend-test-test.swift-x86_64_apple_macosx10.13-o-Onone-531621672.json
$ cat /tmp/stats/*.json
{
  "AST.NumSourceBuffers": 1,
  "AST.NumSourceLines": 1,
  "AST.NumSourceLinesPerSecond": 3,
  "AST.NumLinkLibraries": 0,
  "AST.NumLoadedModules": 4,
  "AST.NumTotalClangImportedEntities": 0,
  ...
  "time.swift.Parsing.wall": 5.038023e-03,
  "time.swift.Parsing.user": 7.200000e-05,
  "time.swift.Parsing.sys": 4.794000e-03,
  "time.swift-frontend.test-test.swift-x86_64_apple_macosx10.13-o-Onone.wall":
3.239949e-01,
  "time.swift-frontend.test-test.swift-x86_64_apple_macosx10.13-o-Onone.user":
2.152100e-02,
  "time.swift-frontend.test-test.swift-x86_64_apple_macosx10.13-o-Onone.sys":
2.897520e-01
}
```

TRACING STATS EVENTS

Furthermore, recent versions `-stats-output-dir` have a secondary, experimental (and much more voluminous mode) called `-trace-stats-events`, that writes *trace files* in CSV to the stats output directory. These trace files

show -- in quite verbose detail, declaration and expression at a time -- the costs incurred by various phases of the compiler, both in terms of absolute time and in terms of any changers to statistics being tracked by the unified stats reporter.

For example, to compile a small file with `-trace-stats-events`, pass it as an extra argument to a compilation already using `-stats-output-dir`:

```
$ mkdir /tmp/stats
$ swiftc -c test.swift -stats-output-dir /tmp/stats -trace-stats-events
$ ls /tmp/stats
stats-1518219460129565-swift-frontend-test-test.swift-x86_64_apple_macosx10.13-o-Onone-1576107381.json
trace-1518219460129597-swift-frontend-test-test.swift-x86_64_apple_macosx10.13-o-Onone-1471252712.csv
$ head /tmp/stats/trace-1518219460129597-swift-frontend-test-test.swift-x86_64_apple_macosx10.13-o-Onone-1471252712.csv
Time,Live,IsEntry,EventName,CounterName,CounterDelta,CounterValue,EntityName,EntityRange

40032,0,"entry","typecheck-decl","Sema.NumDeclsDeserialized",91,91,"foo","
[test.swift:1:1 - line:1:32]"
40032,0,"entry","typecheck-decl","Sema.NumLazyGenericEnvironments",40,40,"foo","
[test.swift:1:1 - line:1:32]"
40032,0,"entry","typecheck-decl","Sema.NumLazyIterableDeclContexts",40,40,"foo","
[test.swift:1:1 - line:1:32]"
40032,0,"entry","typecheck-decl","Sema.NumTypesDeserialized",106,106,"foo","
[test.swift:1:1 - line:1:32]"
40032,0,"entry","typecheck-decl","Sema.NumUnloadedLazyIterableDeclContexts",40,40,"foo","[test.swift:1:1 -
line:1:32]"
40135,0,"entry","typecheck-decl","Sema.InterfaceTypeRequest",1,1,"","[test.swift:1:13
- line:1:29]"
...
```

The data volume in these trace files can be quite overwhelming, and the contents a little hard to read without formatting; for extraction and analysis it can be helpful to load them into a separate tool such as an [SQLite database](#) or a command line CSV processor such as [xsv](#).

```
$ cat /tmp/stats/trace-1518219460129597-swift-frontend-test-test.swift-x86_64_apple_macosx10.13-o-Onone-1471252712.csv \
| xsv search --select CounterName DeclsDeserialized \
| xsv sort --reverse --numeric --select CounterDelta \
| xsv table
Time   Live  IsEntry  EventName          CounterName          CounterDelta
CounterValue  EntityName  EntityRange
43279  0     entry   emit-SIL           Sema.NumDeclsDeserialized  360          517
_      [test.swift:1:17 - line:1:17]
40032  0     entry   typecheck-decl     Sema.NumDeclsDeserialized  91           91
foo    [test.swift:1:1 - line:1:32]
41324  735   exit    typecheck-decl     Sema.NumDeclsDeserialized  40           156
[test.swift:1:13 - line:1:29]
40432  0     entry   typecheck-decl     Sema.NumDeclsDeserialized  25           116
_      [test.swift:1:17 - line:1:17]
```

```

43712 206 exit emit-SIL Sema.NumDeclsDeserialized 18 535
_ [test.swift:1:17 - line:1:17]
41448 97 exit typecheck-fn Sema.NumDeclsDeserialized 1 157
_ [test.swift:1:17 - line:1:17]

```

Post-processing tools for diagnostics

If you dump diagnostic output using `-stats-output-dir <dir>`, the resulting files in `<dir>` will be simple JSON files that can be processed with any JSON-reading program or library, such as `jq`. Alternatively, a bulk-analysis script also exists in `utils/process-stats-dir.py`, which permits a variety of aggregation and analysis tasks.

Here is an example of how to use `-stats-output-dir` together with `utils/process-stats-dir.py` to analyze the difference in compilation performance between two compilers, say `${OLD}/swiftc` and

```
${NEW}/swiftc :
```

```

$ mkdir stats-old stats-new
$ ${OLD}/swiftc -stats-output-dir stats-old test.swift
$ ${NEW}/swiftc -stats-output-dir stats-new test.swift
$ utils/process-stats-dir.py --compare-stats-dirs stats-old stats-new
old      new      delta_pct      name
1402939 1430732 1.98      AST.NumASTBytesAllocated
7        0        -100.0     AST.NumUsedConformances
232      231      -0.43     Constraint solver largest system.LargestNumStatesExplored
42       41      -2.38     Constraint solver largest
system.LargestNumTypeVariableBindings
38       37      -2.63     Constraint solver largest system.LargestNumTypeVariablesBound
79       78      -1.27     Constraint solver largest
system.LargestNumUnsimplifiedConstraints
2593     2592     -0.04     Constraint solver overall.OverallNumStatesExplored
589      588     -0.17     Constraint solver overall.OverallNumTypeVariableBindings
482      481     -0.21     Constraint solver overall.OverallNumTypeVariablesBound
972      971     -0.1      Constraint solver overall.OverallNumUnsimplifiedConstraints
2593     2592     -0.04     ConstraintSystem.incrementScopeCounter
2948     2947     -0.03     Generic signature builder.NumArchetypeAnchorCacheHits
14767    15924    7.84     Generic signature builder.NumConformanceConstraints
9701     10858    11.93    Generic signature builder.NumConformances
5013     4241     -15.4     Generic signature builder.NumPotentialArchetypes
5776     4084     -29.29    Generic signature builder.NumSameTypeConstraints
...

```

When comparing two stats directories that contain the combined results of multiple projects, it can be helpful to select a single project with `--select-module` and/or group counters by module with `--group-by-module`.

Artifact-analysis tools

Many performance issues manifest in the object files or module files produced by the Swift compiler -- say, by generating too much code -- so it can sometimes be helpful to look at the files the compiler outputs directly. The following tools are helpful in such cases:

- `llvm-objdump`, `llvm-otool` and `llvm-size`, `llvm-nm` (which are LLVM-project implementations of the `objdump`, `otool`, `size` and `nm` tools) permit analysis of object files: their sizes, their headers, the set of symbols within them, and even their complete disassembled contents.

- `c++filt` and `swift-demangle` are commands that read from stdin and write to stdout, transforming the text they read by *demangling names* in C++ and Swift, respectively. If you ever seen long, ugly symbol names in diagnostic output from a tool reading a binary artifact, it may read much better after being piped through one or another of these tools.
- `llvm-bcanalyzer` can print (in rough form) the contents of LLVM bitcode streams, such as Swift module files and the PCH/PCM files clang stores its serialized ASTs in. The latter requires combining `llvm-objdump` and `llvm-bcanalyzer` in the following fashion: `llvm-objdump --raw-clang-ast file.pcm | llvm-bcanalyzer -dump`
- `llvm-dwarfdump` and `llvm-dis` can be used to print textual representations of DWARF debug information and LLVM bitcode, respectively. These are usually a bit lower-level than necessary when doing performance analysis, but can be helpful in certain cases.
- `utils/cmpcodesize/cmpcodesize.py` provides a detailed, organized set of size comparisons between the artifacts in a pair of object files emitted by the Swift compiler.

Minimizers

Git bisect

The `git` command has a sub-command called `bisect` that helps with -- and can even completely automate -- the process of reducing a regression range from a large period (weeks to months, hundreds of revisions) down to a single revision that was the cause of a problem. As the name implies, bisect performs bisection -- repeatedly dividing a range in half, then in quarters, and so forth -- so usually within a matter of a dozen steps it can isolate a problem. It is documented in full [here](#) as well as in the `git-bisect` man page.

The `bisect` command can be run in manual mode or automatic mode. In manual mode, you follow these steps:

1. Start with a command like `git bisect start <bad> <good>` that sets up the initial regression range to between the revisions `<bad>` and `<good>` (replace those with revisions you know to be bad and good).
2. Git will repeatedly check out the revision in the middle of the current regression range, asking you to evaluate it.
3. Build the revision it checked out, evaluate it using whatever method you like, and then either run `git bisect good` or `git bisect bad` (literally those words: `bad` or `good`). If the revision it checked out can't be evaluated, run `git bisect skip` and it will shift to an adjacent revision, adjusting its search strategy accordingly.
4. When it's reduced to a single culprit, git will print out the "first bad revision" log entry. At this point it's worth manually re-confirming that the identified revision is indeed bad, and the revision immediately-previous to it is indeed good. Periodically some confounding issue will confuse the bisection search, and you'll be directed to the wrong candidate.

In automatic mode, one simply runs the command `git bisect run test.sh` on some shell script you've written `test.sh`, and `git-bisect` performs the entire bisection search automatically. The shell script needs to exit with a return value that tells `git bisect` whether the revision it's run on is good (`exit 0`), bad (`exit 1`), or should be skipped (`exit 125`).

How the test-script determines the presence of a regression is up to you: a typical approach is to measure against a baseline (eg. using `utils/process-stats-dir.py --compare-to-csv-baseline`, if your regression range covers compilers that all support `-stats-output-dir`). Alternatively, just measure raw time or instruction counts.

An example script that uses the `perf`-based `count_instructions` shell function (see the section on `perf`) to judge whether a revision contains a bug looks something like this:

```
#!/bin/sh
THRESHOLD=500000000
CURR=$(git describe)
utils/update-checkout --scheme main --reset-to-remote --clone --clean
git checkout ${CURR}
utils/update-checkout --scheme main --match-timestamp
git checkout ${CURR}
if utils/build-script -r
then
    V=$(count_instructions ../build/Ninja-ReleaseAssert/swift-linux-x86_64/bin/swiftpc
-c test.swift)
    if [ ${V} -gt ${THRESHOLD} ]
    then
        # Bad
        exit 1
    else
        # Good
        exit 0
    fi
else
    # Skip
    exit 125
fi
```

Note that in the example, the `utils/update-checkout` script is called twice, once to reset the adjacent repositories to their head state, and once with the `--match-timestamp` argument to match the adjacent repositories to the latest point in their history before the timestamp of the primary Swift repository being bisected. This mechanism is necessary if the regression range includes incompatible changes to `clang`, `llvm` or similar adjacent repositories.

Creduce

The `creduce` tool takes an input program and a script, and repeatedly cuts pieces out of the input program and re-runs the script on the program to see if the script still considers the residual input "interesting". It is documented [here](#) and is available through Homebrew on macOS, or in most Linux package ecosystems.

You can use `creduce` to automate the otherwise-tedious task of taking a failing or slow input and reducing it to "just the bad part" that triggers the problem you're after.

For performance analysis, using `creduce` requires that the script can reliably tell when it's "still seeing a regression" and when it's not. This means having a reliable timer or cost-measurement tool; the simplest and most reliable we know of is `perf` so we'll assume you're using it here, via the `count_instructions` shell function described in the section on `perf`; but other measurement tools also work, for example using `utils/process-stats-dir.py` to test changes in performance counters. To use `creduce` you need to write a script something like the following:

```
#!/bin/sh
INPUT=test.swift
OLD=${HOME}/old-toolchain/usr/bin/swift
NEW=${HOME}/new-toolchain/usr/bin/swift
```

```

THRESHOLD=50000000
VOLD=$(count_instructions ${OLD} -frontend -c ${INPUT})
VNEW=$(count_instructions ${NEW} -frontend -c ${INPUT})
VDIFF=$(( ${VNEW} - ${VOLD} ))
if [ ${VDIFF} -gt ${THRESHOLD} ]
then
    # Interesting
    exit 0
else
    # Not interesting
    exit 1
fi

```

Note that, as with `git-bisect`, any measurement tool will work in place of `count_instructions`; if you are looking at a regression-range in which all the compilers have `-stats-output-dir` support, for example, you may well prefer to use `utils/process-stats-dir.py --compare-to-csv-baseline`, for example.

General bisection

When all else fails, coding up a manual bisection is often possible given a numbered set of testcases. The LLVM project ships with a very generic helper script for this, `llvm/util/bisect`, that takes a numeric range and a general subprocess and bisects the range until it finds the place the process changes from success to failure.

Isolating a regression

Follow these steps if you've observed (or think you're seeing) the compiler getting slower between versions:

1. Make sure the before-and-after compilers you're comparing are as close to identical as possible.
 - Ensure both compilers are built the same way: same configuration options, same host toolchains, same optimizations enabled.
 - Ensure both compilers are *release* compilers *without assertions* (Note: nightly snapshots from swift.org have assertions turned on.) You may also want to build (or download) assertion-enabled compilers for finer-grained counter analysis (see below) but keep in mind that they run strictly slower and do significantly different work than release (non-assert) compilers, so are not representative of what users will be using.
2. Measure the high-level timing of the compilation, using `time(1)` or `utils/rusage.py`. Run the compilation a few times to be sure you're not seeing noise.
3. Determine the **compilation mode** in which the regression is occurring. Check to see if changing the compilation mode keeps the regression happening.
 - If the regression occurs regardless of optimization setting, you have reason to believe the cause is in `Sema` or a non-optimizing part of `SILGen`, `IRGen` or `LLVM`. Alternatively if it only occurs in optimizing builds, you have reason to believe the cause is in `SILOptimizer`, `IRGen` or `LLVM`.
 - If the regression occurs regardless of compilation mode, you have reason to believe the cause is localized to frontend jobs. Alternatively if it occurs only in primary-file mode, and goes away in whole-module mode, you have reason to believe the driver may also be implicated.
4. Try to isolate your analysis to a specific frontend process in the compilation.

- If the regression only occurs in whole-module build then you're already dealing with a single process (the sole frontend job).
 - If the regression *also* (or *only*) happens in primary-file mode, you can compile with `-driver-time-compilation` to see which if any frontend job is slower.
 - If *all* frontend jobs are slower, pick one at random to focus on for purposes of analysis.
 - If *no* frontend jobs are slower but the overall compilation seems to be, or if a different number of frontend jobs are running depending on which compiler is run, then the problem is likely in the driver. Skip to the section on driver diagnosis.
5. Assuming you're looking at a frontend process: extract the command-line for the single process (of the form `swift -frontend ...`) by running the build in verbose mode, and put the command-line in a shell script so you can re-run it on its own, without the interference of the driver or other processes. Make a copy of the script that runs the old compiler and a different copy that runs the new compiler (or make the compiler version a parameter of the script). Reconfirm that *just those two isolated frontend processes* still show the regression you're interested in isolating.
 6. Check the value of performance counters between the two compilers via the unified stats reporter (`-stats-output-dir`).
 7. Run both frontend processes under a profiler and compare the profiles in detail. At this point there ought to be *some* sign of a difference, either in counters or profile; if everything looks identical, you either have a deeper mystery than this document will cover, or you've lost the signal by accidentally perturbing the environment / the input / the processes under study. If the problem is *blindingly obvious* at this point, stop and fix the bug, otherwise proceed to narrowing the problem.
 8. Reduce the testcase. That is, figure out the smallest input file that causes the regression you're observing. If the problem you're looking at occurs in the frontend, you might be able to do this manually by running the input with `-Xfrontend -debug-time-function-bodies` and deleting all but the most expensive function, or reduce even further from a function to a single expression via `-Xfrontend -debug-time-expression-type-checking` ; but keep in mind these options only track the time spent *typechecking* a given function or expression; they do not help in reduction of testcases for problems that occur in other subsystems of the compiler. For general reduction, the normal approach is "bisection", also called "delta-debugging": repeatedly delete half the remaining contents of the file, and see if the regression remains. If so, repeat the process in the remaining half; if not, restore the half you deleted and switch your attention to it. This process -- along with several other reduction heuristics that are sensitive to the structure of the code -- can also be automated with the tool `creduce` . See the section on `creduce` for details.
 9. Bisect the regression range. That is, figure out the smallest range of changes to the compiler (typically a single revision in the git history) that caused the regression. If you have more network bandwidth than compute power available, you might want to begin this part by downloading snapshots of the compiler from swift.org. While only a handful of recent snapshots are linked on the swift.org webpage, all historical snapshots remain available to download by substituting the appropriate timestamp into the snapshot URL. For example, the main-branch, macOS snapshot from June 9 2017 is available at <https://swift.org/builds/development/xcode/swift-DEVELOPMENT-SNAPSHOT-2017-06-09-a/swift-DEVELOPMENT-SNAPSHOT-2017-06-09-a-osx.pkg>, and the July 10 2017, swift-4.0-branch Linux snapshot is at <https://swift.org/builds/swift-4.0-branch/ubuntu1604/swift-4.0-DEVELOPMENT-SNAPSHOT-2017-07-10-a/swift-4.0-DEVELOPMENT-SNAPSHOT-2017-07-10-a-ubuntu16.04.tar.gz>. While such snapshots have asserts enabled -- so they do not entirely match the performance characteristics of release compilers -- it is often the case that a regression in a release compiler will still show up in an assert compiler, and downloading snapshots to narrow a regression range can often be much faster than building multiple compilers. Once

you've narrowed a regression range to within a few days (or however far you can get with snapshots alone), you will likely also need to switch to bisection using `git bisect`, which can semi-automate or totally-automate the remaining search, depending on how much shell scripting you want to do and how precisely you're able to measure the difference. See the section on `git-bisect` for details.

10. File or fix the bug! You have everything necessary at this point, and if you can't see how to fix the bug yourself, you should have enough information to hand it off to the person whose change caused the regression; they can usually make short work of fixing it (or at worst, consider whether to just revert the culprit revision). You should include in the bug report:

- A reduced single-file testcase to feed to a single frontend job
- A bisected culprit revision
- Evidence (in the form of a profile and/or change in counters) of a definite performance problem caused by that revision, on that testcase.

Driver diagnosis

One special area where things rarely go wrong, but can be quite serious when they do, is in the driver process orchestrating the frontend jobs. A sure sign that something is amiss in the driver is when running the same compilation scenario, from the same incremental state, but on different compilers, produces a different set of frontend jobs (as seen by `-driver-time-compilation`).

To diagnose a driver problem, several additional diagnostic flags exist: `-driver-print-actions`, `-driver-print-bindings`, `-driver-print-jobs`, `-driver-show-incremental`, and `-driver-show-job-lifecycle`. By carefully comparing these, one can sometimes determine the difference in dependency-analysis and job-execution logic, between one compilation and another.

It is usually also helpful to look at the `.swiftdeps` files generated by the driver. These files contain the driver's summary-view of the dependencies between entities defined and referenced in each source file; it is from these files that the driver decides when a file "needs" to be rebuilt because it depends on another file that needs to be rebuilt, and so on transitively. The file format is [documented here](#).

Finding areas in need of general improvement

If you're not hunting a particular regression, but just trying to improve some problem area in the compiler in general, you usually want to focus on *sorting* costs rather than *isolating*. That is: you'll want to use the tools discussed above (profilers, timers, counters, etc.) to pick off the *largest* cost-centers in a problem area, and work on those first.

Sorting only tells you where a cost center is, though, not whether that cost center is *reasonable*. That is, it's often helpful to try to differentiate cost-centers that "should" be expensive (because they're compiling a lot of code) from cost-centers that "shouldn't" be expensive (but are, because they have a bug or design flaw). One tool that's useful for differentiating these cases is the `utils/scale-test` script.

Scale-test runs on counters, so it's worth taking a short digression into the set of counters that exist in the Swift compiler and how they fit together.

Compiler counters

The Swift compiler has two separate (though related) subsystems for counting the work it does.

1. The LLVM `STATISTIC()` system included by `#include "llvm/ADT/Statistic.h"` and lightly wrapped by `#include "swift/Basic/Statistic.h"`. This system consists of macros and helper structures for atomic, static counters that self-register in a global list. This subsystem is shared with Clang and LLVM in general, and so decisions about whether to enable or disable it (or indeed, conditionally compile it out) are typically shared across all three projects. In practice, most of these definitions are

compiled-out of a non-assert build, because the level of atomic counting and memory-fencing is considered inappropriate and potentially too expensive to count inner loops in production builds. When present, these counters are reported by `-Xfrontend -print-stats`

2. The Swift-specific `UnifiedStatsReporter` system also included by `#include "swift/Basic/Statistic.h"`. This (newer) system consists of a Swift-specific struct full of counters passed around between subsystems of interest. These counters are *always compiled-in* to a Swift build, regardless of build setting. As such, should have *negligible cost* when not counting/reporting: as much as possible, access is arranged to either involve a non-atomic operation (in an inner loop) or a single high-level check before a batch of measurements (outside a loop). These counters are reported by `-stats-output-dir <dir>`

The `UnifiedStatsReporter` system has `Unified` in its name partly because it *subsumes* the other statistic and timer reporting systems in LLVM: it merges its counters with any LLVM `STATISTIC()` counters that existed in the current build, as well as any timers in the compiler, when reporting. Thus whenever possible, you should rely on the output from this subsystem by passing `-stats-output-dir <dir>` and parsing the resulting JSON files; it will always be as good as the output from the `STATISTIC()` counters, if they're present.

Scale-test

This script works with parametrized templates of code (in the same `.gyb` format as other generated code in the compiler), repeatedly instantiating the template with larger and larger parameter values, and measuring the relationship between the input parameter value and the outputs of the compiler counters.

The `utils/scale-test` script works with the same counters inside the compiler as the `-Xfrontend -print-stats` option mentioned above, so it should be used with an assert compiler. It works in concert with the `utils/gyb` boilerplate code-generation tool. The process is straightforward:

1. Write a small `.gyb` file that contains a code pattern that can be varied by varying a numeric input parameter called `N`. For example, the following file (taken from the testsuite) varies the number of nested dictionaries:

```
private let _: [Any?] = [[
  %for i in range(0, N):
    "A": [
      "B" : "C",
    %end
  %for i in range(0, N):
    ]
  %end
]]
```

2. Run the file under the `utils/scale-test` script. You will at least need to pass a `--swiftc-binary` and `.gyb` template filename; by default, it will test on values of `N` ranging from 10 to 100 in steps of 10, and fit scaling curves to *all* counters that it measures, printing any that scale worse than $O(n^{1.2})$. For example, the following will give an initial survey of the script above:

```
$ utils/scale-test --swiftc-binary=../../usr/bin/swiftc test.swift.gyb
O(n^0.0) : BasicCalleeAnalysis.computeMethodCallees
O(n^0.0) : Clang module importer.NumTotalImportedEntities
O(n^0.0) : Constraint solver largest system.LargestNumDisjunctionTerms
```

```
O(n^0.0) : Constraint solver largest system.LargestNumDisjunctions
...
O(n^1.0) : regalloc.NumCopies
O(n^1.0) : isel.NumFastIselSuccessIndependent
O(n^1.0) : regalloc.NumLoads
O(n^1.0) : Constraint solver overall.NumSolutionAttempts
O(n^1.9) : Constraint solver overall.TotalNumTypeVariables
O(n^2.0) : Constraint solver overall.OverallNumComponentsSplit
O(n^2.0) : Constraint solver largest system.LargestNumComponentsSplit
...
```

3. Focus in on a single problematic counter if you see any, and make a regression scale-test -- a file with a `// RUN: %scale-test line --` that runs in as little time as possible while still capturing the scaling pattern. Adjust the scaling ranges with `--begin`, `--end`, `--step` and the counters selected with `--select`. It can be helpful to pass `--values` to see the values being extracted. Scale-test uses a numerical nonlinear optimizer that needs at least 4 or 5 data points to fit a curve properly, and will generally do better with 6 or 7 points. By default, `utils/scale-test` will exit 1 (in error) if it fits a scaling curve worse than `O(n^1.2)`, but this threshold can be adjusted up or down to fit the scenario using `--polynomial-threshold` or `--exponential-threshold`.
4. Add the test to the regression suite. If you're testing a case that you have a fix for, commit the case along with the fix as a positive case, that passes (i.e. shows the bug is fixed). If you're testing a case you do *not* yet have a fix for, add a case with `// RUN: %scale-test --invert-result` and leave it as essentially an XFAIL for others to look at: a test that shows how the current code *scales badly*.
5. If you did *not* find a problematic counter, it's still important to write and commit a testcase, since it will be able to catch *future* regressions if the compiler ever starts scaling poorly in the area you've tested. In general, the more scale tests, the more likely performance regressions will be prevented.
6. If you do not see a counter covering a function in the compiler you're interested in the scaling behaviour of, by all means add a new one! Statistics are easy to add to a file, it takes only a few simple steps:
 - Add a define like `#define DEBUG_TYPE "subsystem-name"` to the file
 - Add an include like `#include <swift/Basic/Statistic.h>` to the file
 - Add `SWIFT_FUNC_STAT;` to the first line of a function to measure
 - Optionally, add separate `STATISTIC(SomeStat, "description");` definitions and manually increment `SomeStat++;` where you like; `SWIFT_FUNC_STAT;` is just a short-form of declaring and incrementing a local `STATISTIC()` named by the function.

How to report bugs most usefully

The three components that make a performance bug easiest to act on are:

1. A reduced testcase that reliably reproduces the problem
2. A revision range where it started going wrong
3. Measurement of how it's going wrong, numerically

If you can't get to #2 or #3, at least try for #1. The smaller the testcase the better: a single file or even a single function is ideal, if you can run a reducer. Much of the work a compiler engineer has to do when diagnosing a bug is just reproduction and reduction, which does not take much knowledge of the internals of the compiler, just time and patience.

How else to help out if you want to improve compile times

- File bugs with reduced testcases and culprit revisions, if you see a regression.
- If you can understand such a bug well enough to fix it, open a pull request with the fix! Straightforward fixes to performance regressions are likely to be merged straight away.
- Add `STATISTIC()` or `SWIFT_FUNC_STAT`-type counters to the compiler, as described in the scale-tests section. Alternatively, if you want a counter that will be "always available" in production builds (and potentially tracked by Apple's performance-tracking CI system), add a counter to `UnifiedStatsReporter`.
- Add scale-tests, to eliminate the possibility of the compiler performing quadratic-or-worse work where it's expected to be linear, or linear-or-worse work where it's expected to be constant.
- Add Open Source projects to the [source-compatibility testsuite](#). Apple's internal CI infrastructure is now tracking selected non-assert-build `UnifiedStatsReporter` counters on those projects, and the team is far more likely to catch a regression if it's shown by a project in the testsuite.
- If you're comfortable making changes to the compiler itself, and don't have a specific testcase you're concerned with, consider working on some of the known problem areas listed above. Get in touch with the people regularly working on those areas and see if there's some set of tasks you can look at. Improvements add up!