

Nodes are created by calling the [createNode](#) action. Nodes can be any object.

A node is stored in Redux under the `nodes` namespace, whose state is a map of the node ID to the actual node object.

## Sourcing Nodes

The creation of nodes occurs primarily in the [sourceNodes](#) bootstrap phase. Nodes created during this phase are top level nodes. I.e, they have no parent. This is represented by source plugins setting the node's `parent` field to `null`. Nodes created via transform plugins (who implement [onCreateNode](#)) will have source nodes as their parents, or other transformed nodes. For a rough overview of what happens when source nodes run, see the [traceId illustration](#).

## Parent/Child/Refs

There are a few different scenarios for creating parent/child relationships.

### Node relationship storage model

All nodes in Gatsby are stored in a flat structure in the Redux `nodes` namespace. A node's `children` field is an array of node IDs, whose nodes are also at the top level of the Redux namespace. Here's an example of the `nodes` namespace.

```
{
  `id1`: { type: `File`, children: [`id2`, `id3`], ...other_fields },
  `id2`: { type: `markdownRemark`, ...other_fields },
  `id3`: { type: `postsJson`, ...other_fields }
}
```

An important note here is that we do not store a distinct collection of each type of child. Rather we store a single collection that they're all packed into. This has some implications on [child field inference](#) in the Schema Generation phase.

### Explicitly recording a parent/child relationship

This occurs when a transformer plugin implements [onCreateNode](#) in order to create some child of the originally created node. In this case, the transformer plugin will call [createParentChildLink](#), with the original node, and the newly created node. All this does is push the child's node ID onto the parent's `children` collection and resave the parent to Redux.

This does **not** automatically create a `parent` field on the child node. If a plugin author wishes to allow child nodes to navigate to their parents in GraphQL queries, they must explicitly set `childNode.parent: 'parent.id'` when creating the child node.

### Foreign Key reference ( `__NODE` )

We've established that child nodes are stored at the top level in Redux, and are referenced via ids in their parent's `children` collection. The same mechanism drives foreign key relationships. Foreign key fields have a `__NODE` suffix on the field name. At query time, Gatsby will take the field's value as an ID, and search Redux for a matching node. This is explained in more detail in [schema gqlTypes](#).

## Plain objects at creation time

Let's say you create the following node by passing it to `createNode`

```
{
  foo: 'bar',
  baz: {
    car: 10
  }
}
```

The value for `baz` is itself an object. That value's parent is the top level object. In this case, Gatsby saves the top level node as is to Redux. It doesn't attempt to extract `baz` into its own node. It does however track the subobject's root NodeID using [Node Tracking](#).

During schema compilation, Gatsby will infer the sub object's type while [creating the gqlType](#).

## Fresh/stale nodes

Every time a build is re-run, there is a chance that a node that exists in the Redux store no longer exists in the original data source. E.g. a file might be deleted from disk between runs. We need a way to indicate that fact to Gatsby.

To track this, there is a Redux `nodesTouched` namespace that tracks whether a particular node ID has been touched. This occurs whenever a node is created (handled by [CREATE\\_NODE](#)), or an explicit call to [touchNode](#).

When a `source-nodes` plugin runs again, it generally recreates nodes (which automatically touches them too). But in some cases, such as [transformer-screenshot](#), a node might not change, but we still want to keep it around for the build. In these cases, we must explicitly call `touchNode`.

Any nodes that aren't touched by the end of the `source-nodes` phase, are deleted. This is performed via a diff between the `nodesTouched` and `nodes` Redux namespaces, in [source-nodes.ts](#)

## Changing a node's fields

From a site developer's point of view, nodes are immutable. In the sense that if you change a node object, those changes will not be seen by other parts of Gatsby. To make a change to a node, it must be persisted to Redux via an action.

So, how do you add a field to an existing node? E.g. perhaps in `onCreateNode`, you want to add a transformer specific field? You can call [createNodeField](#) and this will add your field to the node's `node.fields` object and then persists it to Redux. This can then be referenced by other parts of your plugin at later stages of the build.

## Node Tracking

When a node is created, `createNode` will track all its fields against its nodeid. See [Node Tracking Docs](#) for more.