

BPF Type Format (BTF)

1. Introduction

BTF (BPF Type Format) is the metadata format which encodes the debug info related to BPF program/map. The name BTF was used initially to describe data types. The BTF was later extended to include function info for defined subroutines, and line info for source/line information.

The debug info is used for map pretty print, function signature, etc. The function signature enables better bpf program/function kernel symbol. The line info helps generate source annotated translated byte code, jited code and verifier log.

The BTF specification contains two parts,

- BTF kernel API
- BTF ELF file format

The kernel API is the contract between user space and kernel. The kernel verifies the BTF info before using it. The ELF file format is a user space contract between ELF file and libbpf loader.

The type and string sections are part of the BTF kernel API, describing the debug info (mostly types related) referenced by the bpf program. These two sections are discussed in details in [ref: BTF_Type_String](#).

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\linux-master\Documentation\bpf\[linux-master] [Documentation] [bpf]btf.rst, line 26); [backlink](#)

Unknown interpreted text role "ref".

2. BTF Type and String Encoding

The file `include/uapi/linux/btf.h` provides high-level definition of how types/strings are encoded.

The beginning of data blob must be:

```
struct btf_header {
    __u16    magic;
    __u8     version;
    __u8     flags;
    __u32    hdr_len;

    /* All offsets are in bytes relative to the end of this header */
    __u32    type_off;      /* offset of type section      */
    __u32    type_len;      /* length of type section      */
    __u32    str_off;       /* offset of string section     */
    __u32    str_len;       /* length of string section     */
};
```

The magic is `0xeB9F`, which has different encoding for big and little endian systems, and can be used to test whether BTF is generated for big- or little-endian target. The `btf_header` is designed to be extensible with `hdr_len` equal to `sizeof(struct btf_header)` when a data blob is generated.

2.1 String Encoding

The first string in the string section must be a null string. The rest of string table is a concatenation of other null-terminated strings.

2.2 Type Encoding

The type id 0 is reserved for `void` type. The type section is parsed sequentially and type id is assigned to each recognized type starting from id 1. Currently, the following types are supported:

```
#define BTF_KIND_INT          1      /* Integer          */
#define BTF_KIND_PTR          2      /* Pointer          */
#define BTF_KIND_ARRAY        3      /* Array            */
#define BTF_KIND_STRUCT        4      /* Struct           */
#define BTF_KIND_UNION         5      /* Union            */
#define BTF_KIND_ENUM          6      /* Enumeration      */
#define BTF_KIND_FWD           7      /* Forward          */
#define BTF_KIND_TYPEDEF       8      /* Typedef          */
#define BTF_KIND_VOLATILE      9      /* Volatile         */
#define BTF_KIND_CONST         10     /* Const            */
#define BTF_KIND_RESTRICT      11     /* Restrict         */
#define BTF_KIND_FUNC          12     /* Function         */
#define BTF_KIND_FUNC_PROTO    13     /* Function Proto   */
#define BTF_KIND_VAR           14     /* Variable         */
#define BTF_KIND_DATASEC       15     /* Section          */
```

```
#define BTF_KIND_FLOAT      16      /* Floating point      */
#define BTF_KIND_DECL_TAG  17      /* Decl Tag           */
#define BTF_KIND_TYPE_TAG  18      /* Type Tag           */
```

Note that the type section encodes debug info, not just pure types. `BTF_KIND_FUNC` is not a type, and it represents a defined subprogram.

Each type contains the following common data:

```
struct btf_type {
    u32 name_off;
    /* "info" bits arrangement
     * bits 0-15: vlen (e.g. # of struct's members)
     * bits 16-23: unused
     * bits 24-28: kind (e.g. int, ptr, array...etc)
     * bits 29-30: unused
     * bit 31: kind_flag, currently used by
     *         struct, union and fwd
     */
    u32 info;
    /* "size" is used by INT, ENUM, STRUCT and UNION.
     * "size" tells the size of the type it is describing.
     *
     * "type" is used by PTR, TYPEDEF, VOLATILE, CONST, RESTRICT,
     * FUNC, FUNC_PROTO, DECL_TAG and TYPE_TAG.
     * "type" is a type_id referring to another type.
     */
    union {
        u32 size;
        u32 type;
    };
};
```

For certain kinds, the common data are followed by kind-specific data. The `name_off` in `struct btf_type` specifies the offset in the string table. The following sections detail encoding of each kind.

2.2.1 BTF_KIND_INT

`struct btf_type` encoding requirement:

- `name_off`: any valid offset
- `info.kind_flag`: 0
- `info.kind`: `BTF_KIND_INT`
- `info.vlen`: 0
- `size`: the size of the int type in bytes.

`btf_type` is followed by a `u32` with the following bits arrangement:

```
#define BTF_INT_ENCODING(VAL) ((VAL) & 0xf0000000) >> 24
#define BTF_INT_OFFSET(VAL) ((VAL) & 0x00ff0000) >> 16
#define BTF_INT_BITS(VAL) ((VAL) & 0x000000ff)
```

The `BTF_INT_ENCODING` has the following attributes:

```
#define BTF_INT_SIGNED (1 << 0)
#define BTF_INT_CHAR (1 << 1)
#define BTF_INT_BOOL (1 << 2)
```

The `BTF_INT_ENCODING()` provides extra information: signedness, char, or bool, for the int type. The char and bool encoding are mostly useful for pretty print. At most one encoding can be specified for the int type.

The `BTF_INT_BITS()` specifies the number of actual bits held by this int type. For example, a 4-bit bitfield encodes `BTF_INT_BITS()` equals to 4. The `btf_type.size * 8` must be equal to or greater than `BTF_INT_BITS()` for the type. The maximum value of `BTF_INT_BITS()` is 128.

The `BTF_INT_OFFSET()` specifies the starting bit offset to calculate values for this int. For example, a bitfield struct member has:

- btf member bit offset 100 from the start of the structure,
- btf member pointing to an int type,
- the int type has `BTF_INT_OFFSET() = 2` and `BTF_INT_BITS() = 4`

Then in the struct memory layout, this member will occupy 4 bits starting from bits $100 + 2 = 102$.

Alternatively, the bitfield struct member can be the following to access the same bits as the above:

- btf member bit offset 102,
- btf member pointing to an int type,
- the int type has `BTF_INT_OFFSET() = 0` and `BTF_INT_BITS() = 4`

The original intention of `BTF_INT_OFFSET()` is to provide flexibility of bitfield encoding. Currently, both llvm and pahole generate `BTF_INT_OFFSET() = 0` for all int types.

2.2.2 BTF_KIND_PTR

struct `btf_type` encoding requirement:

- `name_off`: 0
- `info.kind_flag`: 0
- `info.kind`: `BTF_KIND_PTR`
- `info.vlen`: 0
- `type`: the pointee type of the pointer

No additional type data follow `btf_type`.

2.2.3 BTF_KIND_ARRAY

struct `btf_type` encoding requirement:

- `name_off`: 0
- `info.kind_flag`: 0
- `info.kind`: `BTF_KIND_ARRAY`
- `info.vlen`: 0
- `size/type`: 0, not used

`btf_type` is followed by one struct `btf_array`:

```
struct btf_array {
    __u32    type;
    __u32    index_type;
    __u32    nelems;
};
```

The struct `btf_array` encoding:

- `type`: the element type
- `index_type`: the index type
- `nelems`: the number of elements for this array (0 is also allowed).

The `index_type` can be any regular int type (u8, u16, u32, u64, unsigned `__int128`). The original design of including `index_type` follows DWARF, which has an `index_type` for its array type. Currently in BTF, beyond type verification, the `index_type` is not used.

The struct `btf_array` allows chaining through element type to represent multidimensional arrays. For example, for `int a[5][6]`, the following type information illustrates the chaining:

- [1]: int
- [2]: array, `btf_array.type = [1], btf_array.nelems = 6`
- [3]: array, `btf_array.type = [2], btf_array.nelems = 5`

Currently, both pahole and llvm collapse multidimensional array into one-dimensional array, e.g., for `a[5][6]`, the `btf_array.nelems` is equal to 30. This is because the original use case is map pretty print where the whole array is dumped out so one-dimensional array is enough. As more BTF usage is explored, pahole and llvm can be changed to generate proper chained representation for multidimensional arrays.

2.2.4 BTF_KIND_STRUCT

2.2.5 BTF_KIND_UNION

struct `btf_type` encoding requirement:

- `name_off`: 0 or offset to a valid C identifier
- `info.kind_flag`: 0 or 1
- `info.kind`: `BTF_KIND_STRUCT` or `BTF_KIND_UNION`
- `info.vlen`: the number of struct/union members
- `info.size`: the size of the struct/union in bytes

`btf_type` is followed by `info.vlen` number of struct `btf_member`:

```
struct btf_member {
    __u32    name_off;
    __u32    type;
    __u32    offset;
};
```

struct `btf_member` encoding:

- `name_off`: offset to a valid C identifier

- type: the member type
- offset: <see below>

If the type info `kind_flag` is not set, the offset contains only bit offset of the member. Note that the base type of the bitfield can only be int or enum type. If the bitfield size is 32, the base type can be either int or enum type. If the bitfield size is not 32, the base type must be int, and int type `BTF_INT_BITS()` encodes the bitfield size.

If the `kind_flag` is set, the `btf_member.offset` contains both member bitfield size and bit offset. The bitfield size and bit offset are calculated as below.:

```
#define BTF_MEMBER_BITFIELD_SIZE(val)    ((val) >> 24)
#define BTF_MEMBER_BIT_OFFSET(val)      ((val) & 0xffffffff)
```

In this case, if the base type is an int type, it must be a regular int type:

- `BTF_INT_OFFSET()` must be 0.
- `BTF_INT_BITS()` must be equal to $\{1, 2, 4, 8, 16\} * 8$.

The following kernel patch introduced `kind_flag` and explained why both modes exist:

<https://github.com/torvalds/linux/commit/9d5f9f701b1891466fb3dbb1806ad97716f95cc3#diff-fa650a64fdd3968396883d2fe8215ff3>

2.2.6 BTF_KIND_ENUM

struct `btf_type` encoding requirement:

- name_off: 0 or offset to a valid C identifier
- info.kind_flag: 0
- info.kind: `BTF_KIND_ENUM`
- info.vlen: number of enum values
- size: 4

`btf_type` is followed by `info.vlen` number of struct `btf_enum`.:

```
struct btf_enum {
    __u32    name_off;
    __s32    val;
};
```

The `btf_enum` encoding:

- name_off: offset to a valid C identifier
- val: any value

2.2.7 BTF_KIND_FWD

struct `btf_type` encoding requirement:

- name_off: offset to a valid C identifier
- info.kind_flag: 0 for struct, 1 for union
- info.kind: `BTF_KIND_FWD`
- info.vlen: 0
- type: 0

No additional type data follow `btf_type`.

2.2.8 BTF_KIND_TYPEDEF

struct `btf_type` encoding requirement:

- name_off: offset to a valid C identifier
- info.kind_flag: 0
- info.kind: `BTF_KIND_TYPEDEF`
- info.vlen: 0
- type: the type which can be referred by name at `name_off`

No additional type data follow `btf_type`.

2.2.9 BTF_KIND_VOLATILE

struct `btf_type` encoding requirement:

- name_off: 0
- info.kind_flag: 0
- info.kind: `BTF_KIND_VOLATILE`
- info.vlen: 0
- type: the type with volatile qualifier

No additional type data follow `btf_type`.

2.2.10 BTF_KIND_CONST

struct `btf_type` encoding requirement:

- `name_off`: 0
- `info.kind_flag`: 0
- `info.kind`: `BTF_KIND_CONST`
- `info.vlen`: 0
- `type`: the type with `const` qualifier

No additional type data follow `btf_type`.

2.2.11 BTF_KIND_RESTRICT

struct `btf_type` encoding requirement:

- `name_off`: 0
- `info.kind_flag`: 0
- `info.kind`: `BTF_KIND_RESTRICT`
- `info.vlen`: 0
- `type`: the type with `restrict` qualifier

No additional type data follow `btf_type`.

2.2.12 BTF_KIND_FUNC

struct `btf_type` encoding requirement:

- `name_off`: offset to a valid C identifier
- `info.kind_flag`: 0
- `info.kind`: `BTF_KIND_FUNC`
- `info.vlen`: 0
- `type`: a `BTF_KIND_FUNC_PROTO` type

No additional type data follow `btf_type`.

A `BTF_KIND_FUNC` defines not a type, but a subprogram (function) whose signature is defined by `type`. The subprogram is thus an instance of that type. The `BTF_KIND_FUNC` may in turn be referenced by a `func_info` in the [ref`BTF_Ext_Section`](#) (ELF) or in the arguments to [ref`BPF_Prog_Load`](#) (ABI).

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\linux-master\Documentation\bpf\[linux-master] [Documentation] [bpf]btf.rst, line 372); [backlink](#)

Unknown interpreted text role "ref".

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\linux-master\Documentation\bpf\[linux-master] [Documentation] [bpf]btf.rst, line 372); [backlink](#)

Unknown interpreted text role "ref".

2.2.13 BTF_KIND_FUNC_PROTO

struct `btf_type` encoding requirement:

- `name_off`: 0
- `info.kind_flag`: 0
- `info.kind`: `BTF_KIND_FUNC_PROTO`
- `info.vlen`: # of parameters
- `type`: the return type

`btf_type` is followed by `info.vlen` number of struct `btf_param`:

```
struct btf_param {
    __u32    name_off;
    __u32    type;
};
```

If a `BTF_KIND_FUNC_PROTO` type is referred by a `BTF_KIND_FUNC` type, then `btf_param.name_off` must point to a valid C identifier except for the possible last argument representing the variable argument. The `btf_param.type` refers to parameter type.

If the function has variable arguments, the last parameter is encoded with `name_off` = 0 and `type` = 0.

2.2.14 BTF_KIND_VAR

struct btf_type encoding requirement:

- name_off: offset to a valid C identifier
- info.kind_flag: 0
- info.kind: BTF_KIND_VAR
- info.vlen: 0
- type: the type of the variable

btf_type is followed by a single struct btf_variable with the following data:

```
struct btf_var {
    __u32 linkage;
};
```

struct btf_var encoding:

- linkage: currently only static variable 0, or globally allocated variable in ELF sections 1

Not all type of global variables are supported by LLVM at this point. The following is currently available:

- static variables with or without section attributes
- global variables with section attributes

The latter is for future extraction of map key/value type id's from a map definition.

2.2.15 BTF_KIND_DATASEC

struct btf_type encoding requirement:

- name_off: offset to a valid name associated with a variable or one of .data/.bss/.rodata
- info.kind_flag: 0
- info.kind: BTF_KIND_DATASEC
- info.vlen: # of variables
- size: total section size in bytes (0 at compilation time, patched to actual size by BPF loaders such as libbpf)

btf_type is followed by info.vlen number of struct btf_var_secinfo.:

```
struct btf_var_secinfo {
    __u32 type;
    __u32 offset;
    __u32 size;
};
```

struct btf_var_secinfo encoding:

- type: the type of the BTF_KIND_VAR variable
- offset: the in-section offset of the variable
- size: the size of the variable in bytes

2.2.16 BTF_KIND_FLOAT

struct btf_type encoding requirement:

- name_off: any valid offset
- info.kind_flag: 0
- info.kind: BTF_KIND_FLOAT
- info.vlen: 0
- size: the size of the float type in bytes: 2, 4, 8, 12 or 16.

No additional type data follow btf_type.

2.2.17 BTF_KIND_DECL_TAG

struct btf_type encoding requirement:

- name_off: offset to a non-empty string
- info.kind_flag: 0
- info.kind: BTF_KIND_DECL_TAG
- info.vlen: 0
- type: struct, union, func, var or typedef

btf_type is followed by struct btf_decl_tag.:

```
struct btf_decl_tag {
    __u32 component_idx;
```

```
};
```

The `name_off` encodes `btf_decl_tag` attribute string. The type should be `struct`, `union`, `func`, `var` or `typedef`. For `var` or `typedef` type, `btf_decl_tag.component_idx` must be `-1`. For the other three types, if the `btf_decl_tag` attribute is applied to the `struct`, `union` or `func` itself, `btf_decl_tag.component_idx` must be `-1`. Otherwise, the attribute is applied to a `struct/union` member or a `func` argument, and `btf_decl_tag.component_idx` should be a valid index (starting from 0) pointing to a member or an argument.

2.2.17 BTF_KIND_TYPE_TAG

struct `btf_type` encoding requirement:

- `name_off`: offset to a non-empty string
- `info.kind_flag`: 0
- `info.kind`: `BTF_KIND_TYPE_TAG`
- `info.vlen`: 0
- `type`: the type with `btf_type_tag` attribute

Currently, `BTF_KIND_TYPE_TAG` is only emitted for pointer types. It has the following btf type chain:

```
ptr -> [type_tag]*
    -> [const | volatile | restrict | typedef]*
    -> base_type
```

Basically, a pointer type points to zero or more `type_tag`, then zero or more `const/volatile/restrict/typedef` and finally the base type. The base type is one of `int`, `ptr`, `array`, `struct`, `union`, `enum`, `func_proto` and `float` types.

3. BTF Kernel API

The following `bpf syscall` command involves BTF:

- `BPF_BTF_LOAD`: load a blob of BTF data into kernel
- `BPF_MAP_CREATE`: map creation with btf key and value type info.
- `BPF_PROG_LOAD`: prog load with btf function and line info.
- `BPF_BTF_GET_FD_BY_ID`: get a btf fd
- `BPF_OBJ_GET_INFO_BY_FD`: btf, func_info, line_info and other btf related info are returned.

The workflow typically looks like:

```
Application:
  BPF_BTF_LOAD
    |
    v
  BPF_MAP_CREATE and BPF_PROG_LOAD
    |
    v
  .....

Introspection tool:
  .....
  BPF_{PROG,MAP}_GET_NEXT_ID (get prog/map id's)
    |
    v
  BPF_{PROG,MAP}_GET_FD_BY_ID (get a prog/map fd)
    |
    v
  BPF_OBJ_GET_INFO_BY_FD (get bpf_prog_info/bpf_map_info with btf_id)
    |
    v
  BPF_BTF_GET_FD_BY_ID (get btf_fd)
    |
    v
  BPF_OBJ_GET_INFO_BY_FD (get btf)
    |
    v
  pretty print types, dump func signatures and line info, etc.
```

3.1 BPF_BTF_LOAD

Load a blob of BTF data into kernel. A blob of data, described in `ref`BTF_Type_String``, can be directly loaded into the kernel. A `btf_fd` is returned to a userspace.

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\linux-master\Documentation\bpf\[linux-master] [Documentation] [bpf]btf.rst, line 565); [backlink](#)

Unknown interpreted text role "ref".

3.2 BPF_MAP_CREATE

A map can be created with `btf_fd` and specified key/value type id.:

```
__u32    btf_fd;           /* fd pointing to a BTF type data */
__u32    btf_key_type_id;  /* BTF type_id of the key */
__u32    btf_value_type_id; /* BTF type_id of the value */
```

In `libbpf`, the map can be defined with extra annotation like below:

```
struct {
    __uint(type, BPF_MAP_TYPE_ARRAY);
    __type(key, int);
    __type(value, struct ipv_counts);
    __uint(max_entries, 4);
} btf_map SEC(".maps");
```

During ELF parsing, `libbpf` is able to extract key/value type_id's and assign them to `BPF_MAP_CREATE` attributes automatically.

3.3 BPF_PROG_LOAD

During `prog_load`, `func_info` and `line_info` can be passed to kernel with proper values for the following attributes:

```
__u32    insn_cnt;
__aligned_u64 insns;
.....
__u32     prog_btf_fd;    /* fd pointing to BTF type data */
__u32     func_info_rec_size; /* userspace bpf_func_info size */
__aligned_u64 func_info; /* func info */
__u32     func_info_cnt; /* number of bpf_func_info records */
__u32     line_info_rec_size; /* userspace bpf_line_info size */
__aligned_u64 line_info; /* line info */
__u32     line_info_cnt; /* number of bpf_line_info records */
```

The `func_info` and `line_info` are an array of below, respectively.:

```
struct bpf_func_info {
    __u32 insn_off; /* [0, insn_cnt - 1] */
    __u32 type_id; /* pointing to a BTF_KIND_FUNC type */
};
struct bpf_line_info {
    __u32 insn_off; /* [0, insn_cnt - 1] */
    __u32 file_name_off; /* offset to string table for the filename */
    __u32 line_off; /* offset to string table for the source line */
    __u32 line_col; /* line number and column number */
};
```

`func_info_rec_size` is the size of each `func_info` record, and `line_info_rec_size` is the size of each `line_info` record. Passing the record size to kernel make it possible to extend the record itself in the future.

Below are requirements for `func_info`:

- `func_info[0].insn_off` must be 0.
- the `func_info` `insn_off` is in strictly increasing order and matches bpf func boundaries.

Below are requirements for `line_info`:

- the first insn in each func must have a `line_info` record pointing to it.
- the `line_info` `insn_off` is in strictly increasing order.

For `line_info`, the line number and column number are defined as below:

```
#define BPF_LINE_INFO_LINE_NUM(line_col) ((line_col) >> 10)
#define BPF_LINE_INFO_LINE_COL(line_col) ((line_col) & 0x3ff)
```

3.4 BPF_{PROG,MAP}_GET_NEXT_ID

In kernel, every loaded program, map or btf has a unique id. The id won't change during the lifetime of a program, map, or btf.

The `bpf` syscall command `BPF_{PROG,MAP}_GET_NEXT_ID` returns all id's, one for each command, to user space, for `bpf` program or maps, respectively, so an inspection tool can inspect all programs and maps.

3.5 BPF_{PROG,MAP}_GET_FD_BY_ID

An introspection tool cannot use id to get details about program or maps. A file descriptor needs to be obtained first for reference-counting purpose.

3.6 BPF_OBJ_GET_INFO_BY_FD

Once a program/map fd is acquired, an introspection tool can get the detailed information from kernel about this fd, some of which are BTF-related. For example, `bpf_map_info` returns `btf_id` and key/value type ids. `bpf_prog_info` returns `btf_id`, `func_info`,

and line info for translated bpf byte codes, and jited_line_info.

3.7 BPF_BTFF_GET_FD_BY_ID

With `btf_id` obtained in `bpf_map_info` and `bpf_prog_info`, bpf syscall command `BPF_BTFF_GET_FD_BY_ID` can retrieve a btf fd. Then, with command `BPF_OBJ_GET_INFO_BY_FD`, the btf blob, originally loaded into the kernel with `BPF_BTFF_LOAD`, can be retrieved.

With the btf blob, `bpf_map_info`, and `bpf_prog_info`, an introspection tool has full btf knowledge and is able to pretty print map key/values, dump func signatures and line info, along with byte/jit codes.

4. ELF File Format Interface

4.1 .BTF section

The .BTF section contains type and string data. The format of this section is same as the one describe in [ref`BTF_Type_String`](#).

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\linux-master\Documentation\bpf\[linux-master] [Documentation] [bpf]btf.rst, line 686); [backlink](#)

Unknown interpreted text role "ref".

4.2 .BTF.ext section

The .BTF.ext section encodes `func_info` and `line_info` which needs loader manipulation before loading into the kernel.

The specification for .BTF.ext section is defined at `tools/lib/bpf/btf.h` and `tools/lib/bpf/btf.c`.

The current header of .BTF.ext section:

```
struct btf_ext_header {
    __u16    magic;
    __u8     version;
    __u8     flags;
    __u32    hdr_len;

    /* All offsets are in bytes relative to the end of this header */
    __u32    func_info_off;
    __u32    func_info_len;
    __u32    line_info_off;
    __u32    line_info_len;
};
```

It is very similar to .BTF section. Instead of type/string section, it contains `func_info` and `line_info` section. See [ref`BPF_Prog_Load`](#) for details about `func_info` and `line_info` record format.

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\linux-master\Documentation\bpf\[linux-master] [Documentation] [bpf]btf.rst, line 715); [backlink](#)

Unknown interpreted text role "ref".

The `func_info` is organized as below.:

```
func_info_rec_size
btf_ext_info_sec for section #1 /* func_info for section #1 */
btf_ext_info_sec for section #2 /* func_info for section #2 */
...
```

`func_info_rec_size` specifies the size of `bpf_func_info` structure when .BTF.ext is generated. `btf_ext_info_sec`, defined below, is a collection of `func_info` for each specific ELF section.:

```
struct btf_ext_info_sec {
    __u32    sec_name_off; /* offset to section name */
    __u32    num_info;
    /* Followed by num_info * record_size number of bytes */
    __u8     data[0];
};
```

Here, `num_info` must be greater than 0.

The `line_info` is organized as below.:

```
line_info_rec_size
btf_ext_info_sec for section #1 /* line_info for section #1 */
btf_ext_info_sec for section #2 /* line_info for section #2 */
...
```

`line_info_rec_size` specifies the size of `bpf_line_info` structure when `.BTF.ext` is generated.

The interpretation of `bpf_func_info->insn_off` and `bpf_line_info->insn_off` is different between kernel API and ELF API. For kernel API, the `insn_off` is the instruction offset in the unit of `struct bpf_insn`. For ELF API, the `insn_off` is the byte offset from the beginning of section (`btf_ext_info_sec->sec_name_off`).

4.2 .BTF_ids section

The `.BTF_ids` section encodes BTF ID values that are used within the kernel.

This section is created during the kernel compilation with the help of macros defined in `include/linux/btf_ids.h` header file. Kernel code can use them to create lists and sets (sorted lists) of BTF ID values.

The `BTF_ID_LIST` and `BTF_ID` macros define unsorted list of BTF ID values, with following syntax:

```
BTF_ID_LIST(list)
BTF_ID(type1, name1)
BTF_ID(type2, name2)
```

resulting in following layout in `.BTF_ids` section:

```
__BTF_ID__type1__name1__1:
.zero 4
__BTF_ID__type2__name2__2:
.zero 4
```

The `u32 list[];` variable is defined to access the list.

The `BTF_ID_UNUSED` macro defines 4 zero bytes. It's used when we want to define unused entry in `BTF_ID_LIST`, like:

```
BTF_ID_LIST(bpf_skb_output_btf_ids)
BTF_ID(struct, sk_buff)
BTF_ID_UNUSED
BTF_ID(struct, task_struct)
```

The `BTF_SET_START/END` macros pair defines sorted list of BTF ID values and their count, with following syntax:

```
BTF_SET_START(set)
BTF_ID(type1, name1)
BTF_ID(type2, name2)
BTF_SET_END(set)
```

resulting in following layout in `.BTF_ids` section:

```
__BTF_ID__set__set:
.zero 4
__BTF_ID__type1__name1__3:
.zero 4
__BTF_ID__type2__name2__4:
.zero 4
```

The `struct btf_id_set set;` variable is defined to access the list.

The `typeX` name can be one of following:

```
struct, union, typedef, func
```

and is used as a filter when resolving the BTF ID value.

All the BTF ID lists and sets are compiled in the `.BTF_ids` section and resolved during the linking phase of kernel build by `resolve_btfids` tool.

5. Using BTF

5.1 bpftool map pretty print

With BTF, the map key/value can be printed based on fields rather than simply raw bytes. This is especially valuable for large structure or if your data structure has bitfields. For example, for the following map,:

```
enum A { A1, A2, A3, A4, A5 };
typedef enum A __A;
struct tmp_t {
    char a1:4;
    int a2:4;
    int :4;
    __u32 a3:4;
    int b;
    __A b1:4;
    enum A b2:4;
};
```

```

struct {
    __uint(type, BPF_MAP_TYPE_ARRAY);
    __type(key, int);
    __type(value, struct tmp_t);
    __uint(max_entries, 1);
} tmpmap SEC(".maps");

```

bpftool is able to pretty print like below:

```

[ {
  "key": 0,
  "value": {
    "a1": 0x2,
    "a2": 0x4,
    "a3": 0x6,
    "b": 7,
    "b1": 0x8,
    "b2": 0xa
  }
}
]

```

5.2 bpftool prog dump

The following is an example showing how `func_info` and `line_info` can help prog dump with better kernel symbol names, function prototypes and line information.:

```

$ bpftool prog dump jited pinned /sys/fs/bpf/test_btf_haskv
[...]
int test_long_fname_2(struct dummy_tracepoint_args * arg):
bpf_prog_44a040bf25481309_test_long_fname_2:
; static int test_long_fname_2(struct dummy_tracepoint_args *arg)
0:  push    %rbp
1:  mov     %rsp,%rbp
4:  sub     $0x30,%rsp
b:  sub     $0x28,%rbp
f:  mov     %rbx,0x0(%rbp)
13: mov     %r13,0x8(%rbp)
17: mov     %r14,0x10(%rbp)
1b: mov     %r15,0x18(%rbp)
1f: xor     %eax,%eax
21: mov     %rax,0x20(%rbp)
25: xor     %esi,%esi
; int key = 0;
27: mov     %esi,-0x4(%rbp)
; if (!arg->sock)
2a: mov     0x8(%rdi),%rdi
; if (!arg->sock)
2e: cmp     $0x0,%rdi
32: je      0x0000000000000070
34: mov     %rbp,%rsi
; counts = bpf_map_lookup_elem(&btf_map, &key);
[...]

```

5.3 Verifier Log

The following is an example of how `line_info` can help debugging verification failure.:

```

/* The code at tools/testing/selftests/bpf/test_xdp_noinline.c
 * is modified as below.
 */
data = (void *) (long)xdp->data;
data_end = (void *) (long)xdp->data_end;
/*
if (data + 4 > data_end)
    return XDP_DROP;
*/
*(u32 *)data = dst->dst;

$ bpftool prog load ./test_xdp_noinline.o /sys/fs/bpf/test_xdp_noinline type xdp
; data = (void *) (long)xdp->data;
224: (79) r2 = *(u64 *) (r10 -112)
225: (61) r2 = *(u32 *) (r2 +0)
; *(u32 *)data = dst->dst;
226: (63) *(u32 *) (r2 +0) = r1
invalid access to packet, off=0 size=4, R2(id=0,off=0,r=0)
R2 offset is outside of the packet

```

6. BTF Generation

You need latest pahole

<https://git.kernel.org/pub/scm/devel/pahole/pahole.git/>

or llvm (8.0 or later). The pahole acts as a dwarf2btf converter. It doesn't support .BTF.ext and btfBTF_KIND_FUNC type yet. For example,:

```
-bash-4.4$ cat t.c
struct t {
    int a:2;
    int b:3;
    int c:2;
} g;
-bash-4.4$ gcc -c -O2 -g t.c
-bash-4.4$ pahole -JV t.o
File t.o:
[1] STRUCT t kind_flag=1 size=4 vlen=3
    a type_id=2 bitfield_size=2 bits_offset=0
    b type_id=2 bitfield_size=3 bits_offset=2
    c type_id=2 bitfield_size=2 bits_offset=5
[2] INT int size=4 bit_offset=0 nr_bits=32 encoding=SIGNED
```

The llvm is able to generate .BTF and .BTF.ext directly with -g for bpf target only. The assembly code (-S) is able to show the BTF encoding in assembly format.:

```
-bash-4.4$ cat t2.c
typedef int __int32;
struct t2 {
    int a2;
    int (*f2)(char q1, __int32 q2, ...);
    int (*f3)();
} g2;
int main() { return 0; }
int test() { return 0; }
-bash-4.4$ clang -c -g -O2 -target bpf t2.c
-bash-4.4$ readelf -S t2.o
.....
[ 8] .BTF                                PROGBITS                0000000000000000 00000247
                                0000000000000016e 0000000000000000 0 0 1
[ 9] .BTF.ext                          PROGBITS                0000000000000000 000003b5
                                0000000000000060 0000000000000000 0 0 1
[10] .rel.BTF.ext                      REL                     0000000000000000 000007e0
                                0000000000000040 0000000000000010 16 9 8
.....
-bash-4.4$ clang -S -g -O2 -target bpf t2.c
-bash-4.4$ cat t2.s
.....
        .section          .BTF,"",@progbits
        .short            60319
        .byte             1
        .byte             0
        .long             24
        .long             0
        .long             220
        .long             220
        .long             122
        .long             0
        .long             218103808
        .long             2
        .long             83
        .long             16777216
        .long             4
        .long             16777248
        .byte             0
        .ascii            ".text"
        .byte             0
        .ascii            "/home/yhs/tmp-pahole/t2.c"
        .byte             0
        .ascii            "int main() { return 0; }"
        .byte             0
        .ascii            "int test() { return 0; }"
        .byte             0
        .ascii            "int"
        .section          .BTF.ext,"",@progbits
        .short            60319
        .byte             1
        .byte             0
        .long             24
```

```
.long 0
.long 28
.long 28
.long 44
.long 8          # FuncInfo
.long 1          # FuncInfo section string offset=1
.long 2
.long .Lfunc_begin0
.long 3
.long .Lfunc_begin1
.long 5
.long 16         # LineInfo
.long 1          # LineInfo section string offset=1
.long 2
.long .Ltmp0
.long 7
.long 33
.long 7182       # Line 7 Col 14
.long .Ltmp3
.long 7
.long 58
.long 8206       # Line 8 Col 14
```

7. Testing

Kernel bpf selftest *test_btf.c* provides extensive set of BTF-related tests.