

x86 Feature Flags

Introduction

On x86, flags appearing in `/proc/cpuinfo` have an `X86_FEATURE` definition in `arch/x86/include/asm/cpufeatures.h`. If the kernel cares about a feature or KVM want to expose the feature to a KVM guest, it can and should have an `X86_FEATURE_*` defined. These flags represent hardware features as well as software features.

If users want to know if a feature is available on a given system, they try to find the flag in `/proc/cpuinfo`. If a given flag is present, it means that the kernel supports it and is currently making it available. If such flag represents a hardware feature, it also means that the hardware supports it.

If the expected flag does not appear in `/proc/cpuinfo`, things are murkier. Users need to find out the reason why the flag is missing and find the way how to enable it, which is not always easy. There are several factors that can explain missing flags: the expected feature failed to enable, the feature is missing in hardware, platform firmware did not enable it, the feature is disabled at build or run time, an old kernel is in use, or the kernel does not support the feature and thus has not enabled it. In general, `/proc/cpuinfo` shows features which the kernel supports. For a full list of CPUID flags which the CPU supports, use `tools/arch/x86/kcpuid`.

How are feature flags created?

a: Feature flags can be derived from the contents of CPUID leaves.

These feature definitions are organized mirroring the layout of CPUID leaves and grouped in words with offsets as mapped in `enum cpuid_leafs` in `cpufeatures.h` (see `arch/x86/include/asm/cpufeatures.h` for details). If a feature is defined with a `X86_FEATURE_<name>` definition in `cpufeatures.h`, and if it is detected at run time, the flags will be displayed accordingly in `/proc/cpuinfo`. For example, the flag "avx2" comes from `X86_FEATURE_AVX2` in `cpufeatures.h`.

b: Flags can be from scattered CPUID-based features.

Hardware features enumerated in sparsely populated CPUID leaves get software-defined values. Still, CPUID needs to be queried to determine if a given feature is present. This is done in `init_scattered_cpuid_features()`. For instance, `X86_FEATURE_CQM_LLC` is defined as `11*32 + 0` and its presence is checked at runtime in the respective CPUID leaf [`EAX=f`, `ECX=0`] bit `EDX[1]`.

The intent of scattering CPUID leaves is to not bloat `struct cpuinfo_x86.x86_capability[]` unnecessarily. For instance, the CPUID leaf [`EAX=7`, `ECX=0`] has 30 features and is dense, but the CPUID leaf [`EAX=7`, `EAX=1`] has only one feature and would waste 31 bits of space in the `x86_capability[]` array. Since there is a `struct cpuinfo_x86` for each possible CPU, the wasted memory is not trivial.

c: Flags can be created synthetically under certain conditions for hardware features.

Examples of conditions include whether certain features are present in `MSR_IA32_CORE_CAPS` or specific CPU models are identified. If the needed conditions are met, the features are enabled by the `set_cpu_cap` or `setup_force_cpu_cap` macros. For example, if bit 5 is set in `MSR_IA32_CORE_CAPS`, the feature `X86_FEATURE_SPLIT_LOCK_DETECT` will be enabled and "split_lock_detect" will be displayed. The flag "ring3mwait" will be displayed only when running on `INTEL_FAM6_XEON_PHI_KNL` processors.

d: Flags can represent purely software features.

These flags do not represent hardware features. Instead, they represent a software feature implemented in the kernel. For example, Kernel Page Table Isolation is purely software feature and its feature flag `X86_FEATURE_PTI` is also defined in `cpufeatures.h`.

Naming of Flags

The script `arch/x86/kernel/cpu/mkcapflags.sh` processes the `#define X86_FEATURE_<name>` from `cpufeatures.h` and generates the `x86_cap/bug_flags[]` arrays in `kernel/cpu/capflags.c`. The names in the resulting `x86_cap/bug_flags[]` are used to populate `/proc/cpuinfo`. The naming of flags in the `x86_cap/bug_flags[]` are as follows:

a: The name of the flag is from the string in `X86_FEATURE_<name>` by default.

By default, the flag `<name>` in `/proc/cpuinfo` is extracted from the respective `X86_FEATURE_<name>` in `cpufeatures.h`. For example, the flag "avx2" is from `X86_FEATURE_AVX2`.

b: The naming can be overridden.

If the comment on the line for the `#define X86_FEATURE_*` starts with a double-quote character (`"`), the string inside the double-quote characters will be the name of the flags. For example, the flag "sse4_1" comes from the comment "sse4_1" following the `X86_FEATURE_XMM4_1` definition.

There are situations in which overriding the displayed name of the flag is needed. For instance, `/proc/cpuinfo` is a userspace interface and must remain constant. If, for some reason, the naming of `X86_FEATURE_<name>` changes, one shall override the new naming with the name already used in `/proc/cpuinfo`.

c: The naming override can be "", which means it will not appear in `/proc/cpuinfo`.

The feature shall be omitted from `/proc/cpuinfo` if it does not make sense for the feature to be exposed to userspace. For example, `X86_FEATURE_ALWAYS` is defined in `cpufeatures.h` but that flag is an internal kernel feature used in the alternative runtime patching functionality. So, its name is overridden with "". Its flag will not appear in `/proc/cpuinfo`.

Flags are missing when one or more of these happen

a: The hardware does not enumerate support for it.

For example, when a new kernel is running on old hardware or the feature is not enabled by boot firmware. Even if the hardware is new, there might be a problem enabling the feature at run time, the flag will not be displayed.

b: The kernel does not know about the flag.

For example, when an old kernel is running on new hardware.

c: The kernel disabled support for it at compile-time.

For example, if 5-level-paging is not enabled when building (i.e., `CONFIG_X86_5LEVEL` is not selected) the flag "la57" will not show up [1]. Even though the feature will still be detected via `CPUID`, the kernel disables it by clearing via `setup_clear_cpu_cap(X86_FEATURE_LA57)`.

d: The feature is disabled at boot-time.

A feature can be disabled either using a command-line parameter or because it failed to be enabled. The command-line parameter `clearcpuid=` can be used to disable features using the feature number as defined in `/arch/x86/include/asm/cpufeatures.h`. For instance, User Mode Instruction Protection can be disabled using `clearcpuid=514`. The number 514 is calculated from `#define X86_FEATURE_UMIP (16*32 + 2)`.

In addition, there exists a variety of custom command-line parameters that disable specific features. The list of parameters includes, but is not limited to, `nofsbase`, `nosmap`, and `nosmep`. 5-level paging can also be disabled using "no5lvl". `SMAP` and `SMEP` are disabled with the aforementioned parameters, respectively.

e: The feature was known to be non-functional.

The feature was known to be non-functional because a dependency was missing at runtime. For example, `AVX` flags will not show up if `XSAVE` feature is disabled since they depend on `XSAVE` feature. Another example would be broken CPUs and them missing microcode patches. Due to that, the kernel decides not to enable a feature.

[1] 5-level paging uses linear address of 57 bits.