

# MobileNet

This folder contains building code for [MobileNetV2](#) and [MobileNetV3](#) networks. The architectural definition for each model is located in [mobilenet\\_v2.py](#) and [mobilenet\\_v3.py](#) respectively.


For MobilenetV1 please refer to [this page](#)

We have also introduced a family of MobileNets customized for the Edge TPU accelerator found in [Google Pixel4](#) devices. The architectural definition for MobileNetEdgeTPU is located in [mobilenet\\_v3.py](#).

## Performance

### Mobilenet V3 latency

This is the timing of [MobileNetV2](#) vs [MobileNetV3](#) using TF-Lite on the large core of Pixel 1 phone.

 Mobilenet V2 and V3 Latency for Pixel 1.png

### MACs

MACs, also sometimes known as MADDs - the number of multiply-accumulates needed to compute an inference on a single image is a common metric to measure the efficiency of the model. Full size Mobilenet V3 on image size 224 uses ~215 Million MADDs (MMadds) while achieving accuracy 75.1%, while Mobilenet V2 uses ~300MMadds and achieving accuracy 72%. By comparison ResNet-50 uses approximately 3500 MMAdds while achieving 76% accuracy.

Below is the graph comparing Mobilenets and a few selected networks. The size of each blob represents the number of parameters. Note for [ShuffleNet](#) there are no published size numbers. We estimate it to be comparable to MobileNetV2 numbers.

 madds\_top1\_accuracy

### Mobilenet EdgeTPU latency

The figure below shows the Pixel 4 Edge TPU latency of int8-quantized Mobilenet EdgeTPU compared with MobilenetV2 and the minimalistic variants of MobilenetV3 (see below).

 Mobilenet Edge TPU latency for Pixel 4 Edge TPU.png

## Pretrained models

### Mobilenet V3 Imagenet Checkpoints

All mobilenet V3 checkpoints were trained with image resolution 224x224. All phone latencies are in milliseconds, measured on large core. In addition to large and small models this page also contains so-called minimalistic models, these models have the same per-layer dimensions characteristic as MobilenetV3 however, they don't utilize any of the advanced blocks (squeeze-and-excite units, hard-swish, and 5x5 convolutions). While these models are less efficient on CPU, we find that they are much more performant on GPU/DSP.

Imagenet Checkpoint	MACs (M)	Params (M)	Top1	Pixel 1	Pixel 2	Pixel 3

<a href="#">Large dm=1 (float)</a>	217	5.4	75.2	51.2	61	44
<a href="#">Large dm=1 (8-bit)</a>	217	5.4	73.9	44	42.5	32
<a href="#">Large dm=0.75 (float)</a>	155	4.0	73.3	39.8	48	34
<a href="#">Small dm=1 (float)</a>	66	2.9	67.5	15.8	19.4	14.4
<a href="#">Small dm=1 (8-bit)</a>	66	2.9	64.9	15.5	15	10.7
<a href="#">Small dm=0.75 (float)</a>	44	2.4	65.4	12.8	15.9	11.6

#### Minimalistic checkpoints:

Imagenet Checkpoint	MACs (M)	Params (M)	Top1	Pixel 1	Pixel 2	Pixel 3
<a href="#">Large minimalistic (float)</a>	209	3.9	72.3	44.1	51	35
<a href="#">Large minimalistic (8-bit)</a>	209	3.9	71.3	37	35	27
<a href="#">Small minimalistic (float)</a>	65	2.0	61.9	12.2	15.1	11

#### Edge TPU checkpoints:

Imagenet Checkpoint	MACs (M)	Params (M)	Top1	Pixel 4 Edge TPU	Pixel 4 CPU
<a href="#">MobilenetEdgeTPU dm=0.75 (8-bit)</a>	624	2.9	73.5	3.1	13.8
<a href="#">MobilenetEdgeTPU dm=1 (8-bit)</a>	990	4.0	75.6	3.6	20.6

Note: 8-bit quantized versions of the MobilenetEdgeTPU models were obtained using Tensorflow Lite's [post training quantization](#) tool.

#### Mobilenet V2 Imagenet Checkpoints

Classification Checkpoint	Quantized	MACs (M)	Parameters (M)	Top 1 Accuracy	Top 5 Accuracy	Mobile CPU (ms) Pixel 1
<a href="#">float v2 1.4 224</a>	<a href="#">uint8</a>	582	6.06	75.0	92.5	138.0
<a href="#">float v2 1.3 224</a>	<a href="#">uint8</a>	509	5.34	74.4	92.1	123.0
<a href="#">float v2 1.0 224</a>	<a href="#">uint8</a>	300	3.47	71.8	91.0	73.8
<a href="#">float v2 1.0 192</a>	<a href="#">uint8</a>	221	3.47	70.7	90.1	55.1
<a href="#">float v2 1.0 160</a>	<a href="#">uint8</a>	154	3.47	68.8	89.0	40.2
<a href="#">float v2 1.0 128</a>	<a href="#">uint8</a>	99	3.47	65.3	86.9	27.6
<a href="#">float v2 1.0 96</a>	<a href="#">uint8</a>	56	3.47	60.3	83.2	17.6
<a href="#">float v2 0.75 224</a>	<a href="#">uint8</a>	209	2.61	69.8	89.6	55.8
<a href="#">float v2 0.75 192</a>	<a href="#">uint8</a>	153	2.61	68.7	88.9	41.6

<a href="#">float v2 0.75 160</a>	<a href="#">uint8</a>	107	2.61	66.4	87.3	30.4
<a href="#">float v2 0.75 128</a>	<a href="#">uint8</a>	69	2.61	63.2	85.3	21.9
<a href="#">float v2 0.75 96</a>	<a href="#">uint8</a>	39	2.61	58.8	81.6	14.2
<a href="#">float v2 0.5 224</a>	<a href="#">uint8</a>	97	1.95	65.4	86.4	28.7
<a href="#">float v2 0.5 192</a>	<a href="#">uint8</a>	71	1.95	63.9	85.4	21.1
<a href="#">float v2 0.5 160</a>	<a href="#">uint8</a>	50	1.95	61.0	83.2	14.9
<a href="#">float v2 0.5 128</a>	<a href="#">uint8</a>	32	1.95	57.7	80.8	9.9
<a href="#">float v2 0.5 96</a>	<a href="#">uint8</a>	18	1.95	51.2	75.8	6.4
<a href="#">float v2 0.35 224</a>	<a href="#">uint8</a>	59	1.66	60.3	82.9	19.7
<a href="#">float v2 0.35 192</a>	<a href="#">uint8</a>	43	1.66	58.2	81.2	14.6
<a href="#">float v2 0.35 160</a>	<a href="#">uint8</a>	30	1.66	55.7	79.1	10.5
<a href="#">float v2 0.35 128</a>	<a href="#">uint8</a>	20	1.66	50.8	75.0	6.9
<a href="#">float v2 0.35 96</a>	<a href="#">uint8</a>	11	1.66	45.5	70.4	4.5

## Training

### V3

The following configuration, achieves 74.6% using 8 GPU setup and 75.2% using 2x2 TPU setup.

Final Top 1 Accuracy	74.6	
learning_rate	0.16	Total learning rate. (Per clone learning rate is 0.02)
rmsprop_momentum	0.9	
rmsprop_decay	0.9	
rmsprop_epsilon	0.002	
learning_rate_decay_factor	0.99	
optimizer	RMSProp	
warmup_epochs	5	Slim uses per clone epoch, so the the flag value is 0.6
num_epochs_per_decay	3	Slim uses per clone epoch, so the flag value is 0.375
batch_size (per chip)	192	
moving_average_decay	0.9999	
weight_decay	1e-5	
init_stddev	0.008	
dropout_keep_prob	0.8	

bn_moving_average_decay	0.997	
bn_epsilon	0.001	
label_smoothing	0.1	

## V2

The numbers above can be reproduced using slim's [train\\_image\\_classifier](#) . Below is the set of parameters that achieves 72.0% for full size MobileNetV2, after about 700K when trained on 8 GPU. If trained on a single GPU the full convergence is after 5.5M steps. Also note that learning rate and num\_epochs\_per\_decay both need to be adjusted depending on how many GPUs are being used due to slim's internal averaging.

```
--model_name="mobilenet_v2"
--learning_rate=0.045 * NUM_GPUS #slim internally averages clones so we compensate
--preprocessing_name="inception_v2"
--label_smoothing=0.1
--moving_average_decay=0.9999
--batch_size= 96
--num_clones = NUM_GPUS # you can use any number here between 1 and 8 depending on
your hardware setup.
--learning_rate_decay_factor=0.98
--num_epochs_per_decay = 2.5 / NUM_GPUS # train_image_classifier does per clone
epochs
```

## Example

See this [ipython notebook](#) or open and run the network directly in [Colaboratory](#).