

[Intel's Deep Learning Inference Engine](#) (DL IE) is a part of [Intel® OpenVINO™ toolkit](#). You can use it as a computational backend for OpenCV deep learning module.

To use OpenCV with Inference Engine, choose one of the options:

- [Intel® OpenVINO™ toolkit](#) - includes ready to use build of OpenCV
- [OpenCV+OpenVINO Windows package \(community version\)](#).
- [Build from source](#)
 - [Linux](#)
 - [Microsoft Windows](#)
 - [Raspbian](#)

Intel® OpenVINO™ toolkit

- Download and install [Intel® OpenVINO™ toolkit](#).

Important note: if you want to transfer the installed Inference Engine binaries to another machine w/o running OpenVINO installer there, you need [the redistributable files of Intel C++ compiler](#) (use the latest update, 64-bit version), otherwise the Inference Engine or some of its essential plugins will refuse to load and run, which may result in an app crash.

- To perform deep neural networks inference on ARM CPUs, build ARM CPU Plugin from [openvino contrib](#). ARM CPU Plugin isn't distributed in the package. This way requires building OpenVINO and OpenCV from source. ARM CPU Plugin build [instructions](#).
 - OpenVINO release tag [2021.3](#)
 - OpenVINO contrib release tag [2021.3](#)

OpenCV+OpenVINO Windows package (community version)

This community package uses open source version of Inference Engine from [Deep Learning Deployment Toolkit repository](#) (distributed under Apache 2 license).

Hardware requirements:

- CPU with support of AVX2 instruction set
- *[optional]* Intel® Integrated Graphics
- *[optional]* Intel® Neural Compute Stick 2

Software requirements:

- Windows* 10, 64-bit
- *[recommended]* Microsoft* Visual Studio* 2019
- or MSVS* 2019 redistributables: [vc_redist.x64.exe](#)
- Python 3.x (64-bit) to use OpenCV Python bindings
- CMake 3.5.1+ (MSVS 2015+) or CMake 3.14+ (MSVS 2019+)

OpenCV 4.5.5

OpenVINO version	GitHub releases	SourceForge
2021.4.2 (release binaries)	.7z (164.5 Mb) / .zip (253.4 Mb)	.7z (164.5 Mb) / .zip (253.4 Mb)
2021.4.2 (debug binaries)	.7z (196.8 Mb)	.7z (196.8 Mb)

DLDLT components included in this release:

- IE MKLDNN plugin (CPU)
- IE cIDNN plugin (GPU-OpenCL)
- IE Myriad plugin (VPU)
- IE Hetero plugin
- Intel® Threading Building Blocks (TBB)
- nGraph

Both package versions include .PDB files for better debugging experience (generated by MSVS 2019). Package with debug binaries is intended for development of C++ applications only (missing Python support, etc).

► Archive

First steps

Steps below require package with release binaries. Steps for debug binaries might differ.

1. Extract package contents

- Unpack .7z archive using 7-Zip: <https://www.7-zip.org/download.html>
- .zip archive can be extracted using Windows 10 builtin tools (check content menu in explorer)

2. Run OpenCV Python samples (can be launched from explorer):

```
<package_root>\src\opencv\samples\python\_run_winpack_demo.cmd
```

Note: OpenCV DNN Python samples are located in `samples\dnn` directory

3. Open console terminal and configure OpenCV environment:

```
CALL <package_root>\build\setup_vars_opencv4.cmd  
opencv_version
```

For better experience consider running MSVS instance from the terminal with configured environment (type `start MyProject.sln`)

4. Run opencv_version:

```
<package_root>\build\bin\opencv_version_win32.exe
```

or execute in the terminal:

```
<package_root>\build\setup_vars_opencv4.cmd opencv_version -v
```

5. Run Python interpreter (python.exe should be in `PATH`):

```
<package_root>\build\setup_vars_opencv4.cmd python  
> import cv2 as cv  
> print(cv.getBuildInformation())
```

Note: Python support is missing in the package with "Debug" binaries. Use package with "Release" binaries instead.

6. Build and run C++ sample:

- open `src\opencv\samples` in explorer (unmaximize window)
- open `src\opencv\samples\cpp` in explorer (unmaximize window)
- drag&drop `drawing.cpp` file from `cpp` directory onto `_winpack_build_sample.cmd` script
- or execute in the terminal:

```
<package_root>\src\opencv\samples\_winpack_build_sample.cmd
<package_root>\src\opencv\samples\cpp\drawing.cpp
```

Build OpenCV from source

Build OpenCV with pre-built Inference Engine binaries from OpenVINO toolkit.

Linux

Setup environment variables to detect Inference Engine:

```
source /opt/intel/opencvino/bin/setupvars.sh
export ngraph_DIR=/opt/intel/opencvino/deployment_tools/ngraph/cmake/
```

Build OpenCV with extra flags:

```
cmake \
  -DWITH_INF_ENGINE=ON \
  -DENABLE_CXX11=ON \
  ...
```

Microsoft Windows

Setup environment variables to detect Inference Engine:

```
"C:\Program Files (x86)\IntelSWTools\opencvino\bin\setupvars.bat"
```

Build OpenCV with extra flags:

```
cmake ^
  -DWITH_INF_ENGINE=ON ^
  -DENABLE_CXX11=ON ^
  ...
```

Raspbian Buster

Use Docker to cross-compile OpenCV for Raspberry Pi. Check that `uname -m` detects `armv7l` CPU architecture (starts from Raspberry Pi 2 model B).

1. Create a folder named `debian_armhf` with a file `Dockerfile` with the following content:

```
FROM debian:buster
```

USER root

```
RUN dpkg --add-architecture armhf && \
apt-get update && \
apt-get install -y --no-install-recommends \
crossbuild-essential-armhf \
cmake \
pkg-config \
wget \
xz-utils \
libgtk2.0-dev:armhf \
libpython-dev:armhf \
libpython3-dev:armhf \
python-numpy \
python3-numpy \
libgstreamer1.0-dev:armhf \
libgstreamer-plugins-base1.0-dev:armhf

# Install Inference Engine
RUN wget --no-check-certificate
https://download.01.org/opencv/2019/openvinotoolkit/R2/l_openvino_toolkit_runtime_rasp
&& \
tar -xf l_openvino_toolkit_runtime_raspbian_p_2019.2.242.tgz
```

2. Build a Docker image

```
docker image build -t debian_armhf debian_armhf
```

3. Run Docker container mounting source code folder from host.

```
docker run -it -v /absolute/path/to/opencv:/opencv debian_armhf /bin/bash
```

4. Build

```
cd opencv && mkdir opencv_build && mkdir opencv_install && cd opencv_build
cmake -DCMAKE_BUILD_TYPE=Release \
-DCMAKE_INSTALL_PREFIX="../../opencv_install" \
-DOPENCV_CONFIG_INSTALL_PATH="cmake" \
-DCMAKE_TOOLCHAIN_FILE="../../platforms/linux/arm-gnueabi.toolchain.cmake" \
-DWITH_IPP=OFF \
-DBUILD_TESTS=OFF \
-DBUILD_PERF_TESTS=OFF \
-DOPENCV_ENABLE_PKG_CONFIG=ON \
-DPKG_CONFIG_EXECUTABLE="/usr/bin/arm-linux-gnueabi-pkg-config" \
-DPYTHON2_INCLUDE_PATH="/usr/include/python2.7" \
-DPYTHON2_NUMPY_INCLUDE_DIRS="/usr/local/lib/python2.7/dist-
packages/numpy/core/include" \
-DPYTHON3_INCLUDE_PATH="/usr/include/python3.7" \
-DPYTHON3_NUMPY_INCLUDE_DIRS="/usr/local/lib/python3.7/dist-
packages/numpy/core/include" \
```

```

-DPYTHON3_CVPY_SUFFIX=".cpython-37m-arm-linux-gnueabi.hf.so" \
-DENABLE_NEON=ON \
-DCPU_BASELINE="NEON" \
-DWITH_INF_ENGINE=ON \
-
DINF_ENGINE_LIB_DIRS="/l_openvino_toolkit_runtime_raspbian_p_2019.2.242/inference_engine
\
-
DINF_ENGINE_INCLUDE_DIRS="/l_openvino_toolkit_runtime_raspbian_p_2019.2.242/inference_e
\
-DCMAKE_FIND_ROOT_PATH="/l_openvino_toolkit_runtime_raspbian_p_2019.2.242" \
-DENABLE_CXX11=ON ..
make -j4 && make install

```

5. Copy `opencv_install` to the board. Follow

http://docs.opencv.org/latest/docs_install_guides_installing_opencv_raspbian.html to install OpenVINO distribution for Raspberry Pi. Then type the following commands to specify new location of OpenCV:

```

export PYTHONPATH=/path/to/opencv_install/lib/python2.7/dist-packages/:$PYTHONPATH
export PYTHONPATH=/path/to/opencv_install/lib/python3.7/dist-packages/:$PYTHONPATH
export LD_LIBRARY_PATH=/path/to/opencv_install/lib/:$LD_LIBRARY_PATH

```

Usage

- Enable Intel's Inference Engine backend right after `cv::dnn::readNet` invocation:

```
net.setPreferableBackend(DNN_BACKEND_INFERENCE_ENGINE);
```

NOTE: Starts from OpenCV 3.4.2 (OpenVINO 2018.R2) this backend is used by default if OpenCV is built with the Inference Engine support. To switch to origin implementation, use `DNN_BACKEND_OPENCV`. Also, the Inference engine backend is the only available option when the loaded model is represented in OpenVINO™ Model Optimizer format (`.bin` and `.xml`).

- Then, optionally you can also set the device to use for the inference (by default it will use CPU):

```

net.setPreferableTarget(DNN_TARGET_OPENCL);
// the possible options are
// DNN_TARGET_CPU,
// DNN_TARGET_OPENCL,
// DNN_TARGET_OPENCL_FP16
// (fall back to OPENCL if the hardware does not support FP16),
// DNN_TARGET_MYRIAD

```

- You may also import [pre-trained models](#) from [Open Model Zoo](#) passing paths to `.bin` and `.xml` files to `cv::dnn::readNet` function.

- Other names and brands may be claimed as the property of others.