

Distributed Data Parallel Benchmark

This tool is used to measure distributed training iteration time. This is helpful for evaluating the performance impact of code changes to `torch.nn.parallel.DistributedDataParallel`, `torch.distributed`, or anything in between.

It optionally produces a JSON file with all measurements, allowing for an easy A/B comparison of code, configuration, or environment. This comparison can be produced by `diff.py`.

Requirements

This benchmark depends on PyTorch and torchvision.

How to run

Run as many copies of this script as you have model replicas.

If you launch a single task per machine with multiple GPUs, consider using [torch.distributed.launch](#) to spawn multiple processes per machine.

Example output (only on rank 0):

```
-----
PyTorch distributed benchmark suite
-----

* PyTorch version: 1.4.0a0+05140f0
* CUDA version: 10.0
* Distributed backend: nccl

--- nvidia-smi topo -m ---

      GPU0   GPU1   GPU2   GPU3   GPU4   GPU5   GPU6   GPU7   mlx5_2   mlx5_0
mlx5_3  mlx5_1  CPU Affinity
GPU0    X      NV1    NV1    NV2    NV2    SYS    SYS    SYS    SYS    PIX
SYS     PHB    0-19,40-59
GPU1    NV1    X      NV2    NV1    SYS    NV2    SYS    SYS    SYS    PIX
SYS     PHB    0-19,40-59
GPU2    NV1    NV2    X      NV2    SYS    SYS    NV1    SYS    SYS    PHB
SYS     PIX    0-19,40-59
GPU3    NV2    NV1    NV2    X      SYS    SYS    SYS    NV1    SYS    PHB
SYS     PIX    0-19,40-59
GPU4    NV2    SYS    SYS    SYS    X      NV1    NV1    NV2    PIX    SYS
PHB     SYS    0-19,40-59
GPU5    SYS    NV2    SYS    SYS    NV1    X      NV2    NV1    PIX    SYS
PHB     SYS    0-19,40-59
GPU6    SYS    SYS    NV1    SYS    NV1    NV2    X      NV2    PHB    SYS
PIX     SYS    0-19,40-59
GPU7    SYS    SYS    SYS    NV1    NV2    NV1    NV2    X      PHB    SYS
PIX     SYS    0-19,40-59
mlx5_2  SYS    SYS    SYS    SYS    PIX    PIX    PHB    PHB    X      SYS
PHB     SYS
```

mlx5_0	PIX	PIX	PHB	PHB	SYS	SYS	SYS	SYS	SYS	X
SYS	PHB									
mlx5_3	SYS	SYS	SYS	SYS	PHB	PHB	PIX	PIX	PHB	SYS
X	SYS									
mlx5_1	PHB	PHB	PIX	PIX	SYS	SYS	SYS	SYS	SYS	PHB
SYS	X									

Legend:

X = Self
 SYS = Connection traversing PCIe as well as the SMP interconnect between NUMA nodes (e.g., QPI/UPI)
 NODE = Connection traversing PCIe as well as the interconnect between PCIe Host Bridges within a NUMA node
 PHB = Connection traversing PCIe as well as a PCIe Host Bridge (typically the CPU)
 PXB = Connection traversing multiple PCIe switches (without traversing the PCIe Host Bridge)
 PIX = Connection traversing a single PCIe switch
 NV# = Connection traversing a bonded set of # NVLinks

Benchmark: resnet50 with batch size 32

ex/sec	sec/iter	ex/sec	sec/iter	ex/sec	sec/iter	ex/sec	sec/iter
1 GPUs -- no ddp:		p50: 0.097s	329/s	p75: 0.097s	329/s	p90: 0.097s	
329/s p95: 0.097s		329/s					
1 GPUs -- 1M/1G:		p50: 0.100s	319/s	p75: 0.100s	318/s	p90: 0.100s	
318/s p95: 0.100s		318/s					
2 GPUs -- 1M/2G:		p50: 0.103s	310/s	p75: 0.103s	310/s	p90: 0.103s	
310/s p95: 0.103s		309/s					
4 GPUs -- 1M/4G:		p50: 0.103s	310/s	p75: 0.103s	310/s	p90: 0.103s	
310/s p95: 0.103s		310/s					
8 GPUs -- 1M/8G:		p50: 0.104s	307/s	p75: 0.104s	307/s	p90: 0.104s	
306/s p95: 0.104s		306/s					
16 GPUs -- 2M/8G:		p50: 0.104s	306/s	p75: 0.104s	306/s	p90: 0.104s	
306/s p95: 0.104s		306/s					

Benchmark: resnet101 with batch size 32

ex/sec	sec/iter	ex/sec	sec/iter	ex/sec	sec/iter	ex/sec	sec/iter
1 GPUs -- no ddp:		p50: 0.162s	197/s	p75: 0.162s	197/s	p90: 0.162s	
197/s p95: 0.162s		197/s					
1 GPUs -- 1M/1G:		p50: 0.171s	187/s	p75: 0.171s	186/s	p90: 0.171s	
186/s p95: 0.172s		185/s					
2 GPUs -- 1M/2G:		p50: 0.176s	182/s	p75: 0.176s	181/s	p90: 0.176s	
181/s p95: 0.176s		181/s					
4 GPUs -- 1M/4G:		p50: 0.176s	182/s	p75: 0.176s	181/s	p90: 0.176s	
181/s p95: 0.176s		181/s					

```

 8 GPUs --    1M/8G: p50: 0.179s    179/s p75: 0.179s    178/s p90: 0.180s
178/s p95: 0.180s    177/s
 16 GPUs --    2M/8G: p50: 0.179s    178/s p75: 0.180s    177/s p90: 0.183s
174/s p95: 0.188s    170/s

```

Benchmark: resnext50_32x4d with batch size 32

ex/sec	sec/iter	ex/sec	sec/iter	ex/sec	sec/iter
1 GPUs -- no ddp:	p50: 0.145s	220/s	p75: 0.145s	220/s	p90: 0.145s
220/s p95: 0.145s	220/s				
1 GPUs -- 1M/1G:	p50: 0.147s	217/s	p75: 0.147s	217/s	p90: 0.148s
216/s p95: 0.148s	216/s				
2 GPUs -- 1M/2G:	p50: 0.153s	209/s	p75: 0.153s	209/s	p90: 0.153s
209/s p95: 0.153s	209/s				
4 GPUs -- 1M/4G:	p50: 0.153s	208/s	p75: 0.153s	208/s	p90: 0.154s
208/s p95: 0.154s	208/s				
8 GPUs -- 1M/8G:	p50: 0.157s	204/s	p75: 0.157s	204/s	p90: 0.157s
203/s p95: 0.157s	203/s				
16 GPUs -- 2M/8G:	p50: 0.157s	203/s	p75: 0.157s	203/s	p90: 0.158s
203/s p95: 0.158s	202/s				

Benchmark: resnext101_32x8d with batch size 32

ex/sec	sec/iter	ex/sec	sec/iter	ex/sec	sec/iter
1 GPUs -- no ddp:	p50: 0.415s	77/s	p75: 0.415s	77/s	p90: 0.416s
76/s p95: 0.417s	76/s				
1 GPUs -- 1M/1G:	p50: 0.425s	75/s	p75: 0.426s	75/s	p90: 0.426s
75/s p95: 0.426s	75/s				
2 GPUs -- 1M/2G:	p50: 0.438s	73/s	p75: 0.439s	72/s	p90: 0.439s
72/s p95: 0.439s	72/s				
4 GPUs -- 1M/4G:	p50: 0.439s	72/s	p75: 0.439s	72/s	p90: 0.440s
72/s p95: 0.440s	72/s				
8 GPUs -- 1M/8G:	p50: 0.447s	71/s	p75: 0.447s	71/s	p90: 0.448s
71/s p95: 0.448s	71/s				
16 GPUs -- 2M/8G:	p50: 0.450s	71/s	p75: 0.451s	70/s	p90: 0.451s
70/s p95: 0.451s	70/s				

How to diff

Run the benchmark with the `--json PATH_TO_REPORT_FILE` argument to produce the JSON file that the diff script can consume.

Then, run the diff script as follows:

```

$ python3 diff.py PATH_TO_BASELINE_FILE PATH_TO_TEST_FILE

```

	baseline		test
bucket_size:	25	vs	1
cuda_version:	10.0	vs	10.0
distributed_backend:	nccl	vs	nccl

pytorch_version: 1.4.0a0+05140f0 vs 1.4.0a0+05140f0

Benchmark: resnet50 with batch size 32

		sec/iter	ex/sec	diff		sec/iter	ex/sec	diff
1 GPUs:	p75:	0.101s	317/s	-0.3%	p95:	0.101s	317/s	-0.4%
2 GPUs:	p75:	0.104s	306/s	-1.0%	p95:	0.104s	306/s	-1.0%
4 GPUs:	p75:	0.105s	305/s	-1.6%	p95:	0.105s	304/s	-1.8%
8 GPUs:	p75:	0.107s	299/s	-2.6%	p95:	0.107s	298/s	-2.7%
16 GPUs:	p75:	0.108s	294/s	-3.8%	p95:	0.122s	262/s	-16.4%

Benchmark: resnet101 with batch size 32

		sec/iter	ex/sec	diff		sec/iter	ex/sec	diff
1 GPUs:	p75:	0.172s	185/s	-1.2%	p95:	0.172s	185/s	-1.3%
2 GPUs:	p75:	0.179s	178/s	-2.1%	p95:	0.179s	178/s	-2.0%
4 GPUs:	p75:	0.180s	177/s	-2.6%	p95:	0.180s	177/s	-2.6%
8 GPUs:	p75:	0.184s	173/s	-3.5%	p95:	0.184s	173/s	-3.5%
16 GPUs:	p75:	0.187s	170/s	-0.1%	p95:	0.204s	157/s	-7.9%

Benchmark: resnext50_32x4d with batch size 32

		sec/iter	ex/sec	diff		sec/iter	ex/sec	diff
1 GPUs:	p75:	0.149s	214/s	-1.0%	p95:	0.149s	214/s	-0.9%
2 GPUs:	p75:	0.156s	205/s	-1.5%	p95:	0.156s	205/s	-1.6%
4 GPUs:	p75:	0.156s	204/s	-1.6%	p95:	0.157s	204/s	-1.8%
8 GPUs:	p75:	0.159s	200/s	-1.5%	p95:	0.159s	200/s	-1.5%
16 GPUs:	p75:	0.161s	198/s	-1.9%	p95:	0.162s	197/s	-2.3%

Benchmark: resnext101_32x8d with batch size 32

		sec/iter	ex/sec	diff		sec/iter	ex/sec	diff
1 GPUs:	p75:	0.427s	74/s	-0.8%	p95:	0.428s	74/s	-0.7%
2 GPUs:	p75:	0.444s	72/s	-1.3%	p95:	0.445s	71/s	-0.7%
4 GPUs:	p75:	0.444s	72/s	-1.1%	p95:	0.445s	71/s	-0.8%
8 GPUs:	p75:	0.452s	70/s	-1.3%	p95:	0.452s	70/s	-1.3%
16 GPUs:	p75:	0.455s	70/s	-0.7%	p95:	0.456s	70/s	-0.6%

This compares throughput between `bucket_cap_mb=25` (the default) and `bucket_cap_mb=1` on 8 DGX machines with V100 GPUs. It confirms that even for a relatively small model on machines with a very fast interconnect (4x 100Gb InfiniBand per machine), it still pays off to batch allreduce calls.