

orphan:

The Swift Array Design

Author: Dave Abrahams

Date: 2014-04-10

Goals

1. Performance equivalent to C arrays for subscript get/set of non-class element types is the most important performance goal.
2. It should be possible to receive an `NSArray` from Cocoa, represent it as an `Array<AnyObject>`, and pass it right back to Cocoa as an `NSArray` in $O(1)$ and with no memory allocations.
3. Arrays should be usable as stacks, so we want amortized $O(1)$ append and $O(1)$ popBack. Together with goal #1, this implies a `std::vector`-like layout, with a reserved tail memory capacity that can exceed the number of actual stored elements.

To achieve goals 1 and 2 together, we use static knowledge of the element type: when it is statically known that the element type is not a class, code and checks accounting for the possibility of wrapping an `NSArray` are eliminated. An `Array` of Swift value types always uses the most efficient possible representation, identical to that of `ContiguousArray`.

Components

Swift provides three generic array types, all of which have amortized $O(1)$ growth. In this document, statements about `ArrayType` apply to all three of the components.

- `ContiguousArray<Element>` is the fastest and simplest of the three--use this when you need "C array" performance. The elements of a `ContiguousArray` are always stored contiguously in memory.



- `Array<Element>` is like `ContiguousArray<Element>`, but optimized for efficient conversions from Cocoa and back--when `Element` can be a class type, `Array<Element>` can be backed by the (potentially non-contiguous) storage of an arbitrary `NSArray` rather than by a Swift `ContiguousArray`. `Array<Element>` also supports up- and downcasts between arrays of related class types. When `Element` is known to be a non-class type, the performance of `Array<Element>` is identical to that of `ContiguousArray<Element>`.



- `ArraySlice<Element>` is a subrange of some `Array<Element>` or `ContiguousArray<Element>`; it's the result of using slice notation, e.g. `a[7...21]` on any Swift array `a`. A slice always has contiguous storage and "C array" performance. Slicing an `ArrayType` is $O(1)$ unless the source is an `Array<Element>` backed by an `NSArray` that doesn't supply contiguous storage.

`ArraySlice` is recommended for transient computations but not for long-term storage. Since it references a sub-range of some shared backing buffer, a `ArraySlice` may artificially prolong the lifetime of elements outside the `ArraySlice` itself.



Mutation Semantics

The `ArrayTypes` have full value semantics via copy-on-write (COW):

```
var a = [1, 2, 3]
let b = a
a[1] = 42
print(b[1]) // prints "2"
```

Bridging Rules and Terminology for all Types

- Every class type or `@objc` existential (such as `AnyObject`) is **bridged** to Objective-C and **bridged back** to Swift via the identity transformation, i.e. it is **bridged verbatim**.
- A type `T` that is not **bridged verbatim** can conform to `BridgedToObjectiveC`, which specifies its conversions to and from ObjectiveC:

```
protocol _BridgedToObjectiveC {
    typealias _ObjectiveCType: AnyObject
    func _bridgeToObjectiveC() -> _ObjectiveCType
    class func _forceBridgeFromObjectiveC(_: _ObjectiveCType) -> Self
}
```

Note

Classes and `@objc` existentials shall not conform to `_BridgedToObjectiveC`, a restriction that's not currently enforceable at compile-time.

- Some generic types (*ArrayType*<T> in particular) bridge to Objective-C only if their element types bridge. These types conform to `_ConditionallyBridgedToObjectiveC`:

```
protocol _ConditionallyBridgedToObjectiveC : _BridgedToObjectiveC {
    class func _isBridgedToObjectiveC() -> Bool
    class func _conditionallyBridgeFromObjectiveC(_: _ObjectiveCType) -> Self?
}
```

Bridging from, or *bridging back* to, a type `T` conforming to `_ConditionallyBridgedToObjectiveC` when `T._isBridgedToObjectiveC()` is false is a user programming error that may be diagnosed at runtime. `_conditionallyBridgeFromObjectiveC` can be used to attempt to bridge back, and return `nil` if the entire object cannot be bridged.

Implementation Note

There are various ways to move this detection to compile-time

- For a type `T` that is not **bridged verbatim**,
 - if `T` conforms to `BridgedToObjectiveC` and either
 - `T` does not conform to `_ConditionallyBridgedToObjectiveC`
 - or, `T._isBridgedToObjectiveC()`
 then a value `x` of type `T` is **bridged** as `T._ObjectiveCType` via `x._bridgeToObjectiveC()`, and an object `y` of `T._ObjectiveCType` is **bridged back** to `T` via `T._forceBridgeFromObjectiveC(y)`
 - Otherwise, `T` **does not bridge** to Objective-C

Array Type Conversions

From here on, this document deals only with `Array` itself, and not `Slice` or `ContiguousArray`, which support a subset of `Array`'s conversions. Future revisions will add descriptions of `Slice` and `ContiguousArray` conversions.

Kinds of Conversions

In these definitions, `Base` is `AnyObject` or a trivial subtype thereof, `Derived` is a trivial subtype of `Base`, and `X` conforms to `_BridgedToObjectiveC`:

- **Trivial bridging** implicitly converts `[Base]` to `NSArray` in $O(1)$. This is simply a matter of returning the `Array`'s internal buffer, which is-a `NSArray`.
- **Trivial bridging back** implicitly converts `NSArray` to `[AnyObject]` in $O(1)$ plus the cost of calling `copy()` on the `NSArray`.
[1]
- **Implicit conversions** between `Array` types
 - **Implicit upcasting** implicitly converts `[Derived]` to `[Base]` in $O(1)$.
 - **Implicit bridging back** implicitly converts `[X]` to `[X._ObjectiveCType]` in $O(N)$.

Note

Either type of implicit conversion may be combined with **trivial bridging** in an implicit conversion to `NSArray`.

- **Checked conversions** convert `[T]` to `[U]?` in $O(N)$ via `a as [U]`.
 - **Checked downcasting** converts `[Base]` to `[Derived]?`.
 - **Checked bridging back** converts `[T]` to `[X]?` where `X._ObjectiveCType` is `T` or a trivial subtype thereof.
- **Forced conversions** convert `[AnyObject]` or `NSArray` to `[T]` implicitly, in bridging thanks between Swift and Objective-C.

For example, when a user writes a Swift method taking `[NSView]`, it is exposed to Objective-C as a method taking `NSArray`, which is force-converted to `[NSView]` when called from Objective-C.

- **Forced downcasting** converts `[AnyObject]` to `[Derived]` in $O(1)$
- **Forced bridging back** converts `[AnyObject]` to `[X]` in $O(N)$.

A forced conversion where any element fails to convert is considered a user programming error that may trap. In the case of forced downcasts, the trap may be **deferred** to the point where an offending element is accessed.

Note

Both checked and forced downcasts may be combined with **trivial bridging back** in conversions from `NSArray`.

Maintaining Type-Safety

Both upcasts and forced downcasts raise type-safety issues.

Upcasts

TODO: this section is outdated.

When up-casting an `[Derived]` to `[Base]`, a buffer of `Derived` object can simply be `unsafeBitCast`'ed to a buffer of elements of type `Base`--as long as the resulting buffer is never mutated. For example, we cannot allow a `Base` element to be inserted in the buffer, because the buffer's destructor will destroy the elements with the (incorrect) static presumption that they have `Derived` type.

Furthermore, we can't (logically) copy the buffer just prior to mutation, since the `[Base]` may be copied prior to mutation, and our shared subscript assignment semantics imply that all copies must observe its subscript assignments.

Therefore, converting `[T]` to `[U]` is akin to resizing: the new `Array` becomes logically independent. To avoid an immediate $O(N)$ conversion cost, and preserve shared subscript assignment semantics, we use a layer of indirection in the data structure. Further, when `T` is a subclass of `U`, the intermediate object is marked to prevent in-place mutation of the buffer; it will be copied upon its first mutation:



Deferred Checking for Forced Downcasts

In forced downcasts, if any element fails to have dynamic type `Derived`, it is considered a programming error that may cause a trap. Sometimes we can do this check in $O(1)$ because the source holds a known buffer type. Rather than incur $O(N)$ checking for the other cases, the new intermediate object is marked for deferred checking, and all element accesses through that object are dynamically typechecked, with a trap upon failure (except in `-Ounchecked` builds).

When the resulting array is later up-cast (other than to a type that can be validated in $O(1)$ by checking the type of the underlying buffer), the result is also marked for deferred checking.

[1] This `copy()` may amount to a retain if the `NSArray` is already known to be immutable. We could eventually optimize out the copy if we can detect that the `NSArray` is uniquely referenced. Our current unique-reference detection applies only to Swift objects, though.