

```
+++ title = "Provisioning" description = "" keywords = ["grafana", "provisioning"] aliases =  
["/docs/grafana/latest/installation/provisioning"] weight = 800 +++
```

## Provisioning Grafana

In previous versions of Grafana, you could only use the API for provisioning data sources and dashboards. But that required the service to be running before you started creating dashboards and you also needed to set up credentials for the HTTP API. In v5.0 we decided to improve this experience by adding a new active provisioning system that uses config files. This will make GitOps more natural as data sources and dashboards can be defined via files that can be version controlled. We hope to extend this system to later add support for users, orgs and alerts as well.

### Config File

Check out the [\[configuration\]\({{< relref "configuration.md" >}}\)](#) page for more information on what you can configure in `grafana.ini`

#### Config File Locations

- Default configuration from `$WORKING_DIR/conf/defaults.ini`
- Custom configuration from `$WORKING_DIR/conf/custom.ini`
- The custom configuration file path can be overridden using the `--config` parameter

**Note:** If you have installed Grafana using the `deb` or `rpm` packages, then your configuration file is located at `/etc/grafana/grafana.ini`. This path is specified in the Grafana `init.d` script using `--config` file parameter.

#### Using Environment Variables

It is possible to use environment variable interpolation in all 3 provisioning configuration types. Allowed syntax is either `$ENV_VAR_NAME` or `${ENV_VAR_NAME}` and can be used only for values not for keys or bigger parts of the configurations. It is not available in the dashboard's definition files just the dashboard provisioning configuration. Example:

```
datasources:  
- name: Graphite  
  url: http://localhost:$PORT  
  user: $USER  
  secureJsonData:  
    password: $PASSWORD
```

If you have a literal `$` in your value and want to avoid interpolation, `$$` can be used.

---

## Configuration Management Tools

Currently we do not provide any scripts/manifests for configuring Grafana. Rather than spending time learning and creating scripts/manifests for each tool, we think our time is better spent making Grafana easier to provision. Therefore, we heavily rely on the expertise of the community.

Tool	Project

Puppet	<a href="https://forge.puppet.com/puppet/grafana">https://forge.puppet.com/puppet/grafana</a>
Ansible	<a href="https://github.com/cloudalchemy/ansible-grafana">https://github.com/cloudalchemy/ansible-grafana</a>
Chef	<a href="https://github.com/JonathanTron/chef-grafana">https://github.com/JonathanTron/chef-grafana</a>
Saltstack	<a href="https://github.com/salt-formulas/salt-formula-grafana">https://github.com/salt-formulas/salt-formula-grafana</a>
Jsonnet	<a href="https://github.com/grafana/grafonnet-lib/">https://github.com/grafana/grafonnet-lib/</a>

## Data sources

*This feature is available from v5.0*

It's possible to manage data sources in Grafana by adding one or more YAML config files in the `provisioning/datasources` directory. Each config file can contain a list of `datasources` that will get added or updated during start up. If the data source already exists, then Grafana updates it to match the configuration file. The config file can also contain a list of data sources that should be deleted. That list is called `deleteDatasources`. Grafana will delete data sources listed in `deleteDatasources` before inserting/updating those in the `datasource` list.

## Running Multiple Grafana Instances

If you are running multiple instances of Grafana you might run into problems if they have different versions of the `datasource.yaml` configuration file. The best way to solve this problem is to add a version number to each `datasource` in the configuration and increase it when you update the config. Grafana will only update `datasources` with the same or lower version number than specified in the config. That way, old configs cannot overwrite newer configs if they restart at the same time.

## Example data source Config File

```
# config file version
apiVersion: 1

# list of datasources that should be deleted from the database
deleteDatasources:
  - name: Graphite
    orgId: 1

# list of datasources to insert/update depending
# what's available in the database
datasources:
  # <string, required> name of the datasource. Required
  - name: Graphite
    # <string, required> datasource type. Required
    type: graphite
    # <string, required> access mode. proxy or direct (Server or Browser in the UI).
    # Required
    access: proxy
    # <int> org id. will default to orgId 1 if not specified
    orgId: 1
    # <string> custom UID which can be used to reference this datasource in other
```

```

parts of the configuration, if not specified will be generated automatically
uid: my_unique_uid
# <string> url
url: http://localhost:8080
# <string> Deprecated, use secureJsonData.password
password:
# <string> database user, if used
user:
# <string> database name, if used
database:
# <bool> enable/disable basic auth
basicAuth:
# <string> basic auth username
basicAuthUser:
# <string> Deprecated, use secureJsonData.basicAuthPassword
basicAuthPassword:
# <bool> enable/disable with credentials headers
withCredentials:
# <bool> mark as default datasource. Max one per org
isDefault:
# <map> fields that will be converted to json and stored in jsonData
jsonData:
  graphiteVersion: '1.1'
  tlsAuth: true
  tlsAuthWithCACert: true
# <string> json object of data that will be encrypted.
secureJsonData:
  tlsCACert: '...'
  tlsClientCert: '...'
  tlsClientKey: '...'
# <string> database password, if used
password:
# <string> basic auth password
basicAuthPassword:
version: 1
# <bool> allow users to edit datasources from the UI.
editable: false

```

### Custom Settings per Datasource

Please refer to each datasource documentation for specific provisioning examples.

Datasource	Misc
Elasticsearch	Elasticsearch uses the <code>database</code> property to configure the index for a datasource

### JSON Data

Since not all datasources have the same configuration settings we only have the most common ones as fields. The rest should be stored as a json blob in the `jsonData` field. Here are the most common settings that the core datasources use.

**Note:** Datasources tagged with HTTP\* below denotes any data source which communicates using the HTTP protocol, e.g. all core data source plugins except MySQL, PostgreSQL and MSSQL.

Name	Type	Datasource	Description
tlsAuth	boolean	HTTP*, MySQL	Enable TLS authentication using client cert configured in secure json data
tlsAuthWithCACert	boolean	HTTP*, MySQL, PostgreSQL	Enable TLS authentication using CA cert
tlsSkipVerify	boolean	HTTP*, MySQL, PostgreSQL, MSSQL	Controls whether a client verifies the server's certificate chain and host name.
serverName	string	HTTP*, MSSQL	Optional. Controls the server name used for certificate common name/subject alternative name verification. Defaults to using the data source URL.
timeout	string	HTTP*	Request timeout in seconds. Overrides dataproxy.timeout option
graphiteVersion	string	Graphite	Graphite version
timeInterval	string	Prometheus, Elasticsearch, InfluxDB, MySQL, PostgreSQL and MSSQL	Lowest interval/step value that should be used for this data source.
httpMode	string	Influxdb	HTTP Method. 'GET', 'POST', defaults to GET
maxSeries	number	Influxdb	Max number of series/tables that Grafana processes
httpMethod	string	Prometheus	HTTP Method. 'GET', 'POST', defaults to POST
customQueryParameters	string	Prometheus	Query parameters to add, as a URL-encoded string.
manageAlerts	boolean	Prometheus and Loki	Manage alerts via Alerting UI
esVersion	string	Elasticsearch	Elasticsearch version (E.g. 7.0.0, 7.6.1)
timeField	string	Elasticsearch	Which field that should be used as timestamp
interval	string	Elasticsearch	Index date time format. nil(No Pattern), 'Hourly', 'Daily', 'Weekly', 'Monthly' or 'Yearly'
logMessageField	string	Elasticsearch	Which field should be used as the log message

logLevelField	string	Elasticsearch	Which field should be used to indicate the priority of the log message
maxConcurrentShardRequests	number	Elasticsearch	Maximum number of concurrent shard requests that each sub-search request executes per node. Defaults to 5 if esVersion is greater than or equals 7.0.0. When the esVersion is less than 7.0.0 and greater than or equals 5.6.0, then the default value is 256. Option is ignored when esVersion is less than 5.6.0.
sigV4Auth	boolean	Elasticsearch and Prometheus	Enable usage of SigV4
sigV4AuthType	string	Elasticsearch and Prometheus	SigV4 auth provider. default/credentials/keys
sigV4ExternalId	string	Elasticsearch and Prometheus	Optional SigV4 External ID
sigV4AssumeRoleArn	string	Elasticsearch and Prometheus	Optional SigV4 ARN role to assume
sigV4Region	string	Elasticsearch and Prometheus	SigV4 AWS region
sigV4Profile	string	Elasticsearch and Prometheus	Optional SigV4 credentials profile
authType	string	Cloudwatch	Auth provider. default/credentials/keys
externalId	string	Cloudwatch	Optional External ID
assumeRoleArn	string	Cloudwatch	Optional ARN role to assume
defaultRegion	string	Cloudwatch	Optional default AWS region
customMetricsNamespaces	string	Cloudwatch	Namespaces of Custom Metrics
profile	string	Cloudwatch	Optional credentials profile
tsdbVersion	string	OpenTSDB	Version
tsdbResolution	string	OpenTSDB	Resolution
sslmode	string	PostgreSQL	SSLmode. 'disable', 'require', 'verify-ca' or 'verify-full'
tlsConfigurationMethod	string	PostgreSQL	SSL Certificate configuration, either by 'file-path' or 'file-content'

sslRootCertFile	string	PostgreSQL, MSSQL	SSL server root certificate file, must be readable by the Grafana user
sslCertFile	string	PostgreSQL	SSL client certificate file, must be readable by the Grafana user
sslKeyFile	string	PostgreSQL	SSL client key file, must be readable by <i>only</i> the Grafana user
encrypt	string	MSSQL	Connection SSL encryption handling. 'disable', 'false' or 'true'
postgresVersion	number	PostgreSQL	Postgres version as a number (903/904/905/906/1000) meaning v9.3, v9.4, ..., v10
timescaledb	boolean	PostgreSQL	Enable usage of TimescaleDB extension
maxOpenConns	number	MySQL, PostgreSQL and MSSQL	Maximum number of open connections to the database (Grafana v5.4+)
maxIdleConns	number	MySQL, PostgreSQL and MSSQL	Maximum number of connections in the idle connection pool (Grafana v5.4+)
connMaxLifetime	number	MySQL, PostgreSQL and MSSQL	Maximum amount of time in seconds a connection may be reused (Grafana v5.4+)
keepCookies	array	HTTP*	Cookies that needs to be passed along while communicating with datasources

## Secure Json Data

```
{"authType": "keys", "defaultRegion": "us-west-2", "timeField": "@timestamp"}
```

Secure json data is a map of settings that will be encrypted with [secret key]({{< relref "configuration.md#secret-key" >}}) from the Grafana config. The purpose of this is only to hide content from the users of the application. This should be used for storing TLS Cert and password that Grafana will append to the request on the server side. All of these settings are optional.

**Note:** Datasources tagged with HTTP\* below denotes any data source which communicates using the HTTP protocol, e.g. all core data source plugins except MySQL, PostgreSQL and MSSQL.

Name	Type	Datasource	Description
tlsCACert	string	HTTP*, MySQL, PostgreSQL	CA cert for out going requests
tlsClientCert	string	HTTP*, MySQL, PostgreSQL	TLS Client cert for outgoing requests
tlsClientKey	string	HTTP*, MySQL, PostgreSQL	TLS Client key for outgoing requests
password	string	HTTP*, MySQL, PostgreSQL, MSSQL	password

basicAuthPassword	string	HTTP*	password for basic authentication
accessKey	string	Cloudwatch	Access key for connecting to Cloudwatch
secretKey	string	Cloudwatch	Secret key for connecting to Cloudwatch
sigV4AccessKey	string	Elasticsearch and Prometheus	SigV4 access key. Required when using keys auth provider
sigV4SecretKey	string	Elasticsearch and Prometheus	SigV4 secret key. Required when using keys auth provider

### Custom HTTP headers for datasources

Data sources managed by Grafanas provisioning can be configured to add HTTP headers to all requests going to that datasource. The header name is configured in the `jsonData` field and the header value should be configured in `secureJsonData` .

```
apiVersion: 1

datasources:
- name: Graphite
  jsonData:
    httpHeaderName1: 'HeaderName'
    httpHeaderName2: 'Authorization'
  secureJsonData:
    httpHeaderValue1: 'HeaderValue'
    httpHeaderValue2: 'Bearer XXXXXXXXXX'
```

## Plugins

*This feature is available from v7.1*

You can manage plugins in Grafana by adding one or more YAML config files in the [ `provisioning/plugins` ] (`{{< relref "configuration.md#provisioning" >}}`) directory. Each config file can contain a list of `apps` that will be updated during start up. Grafana updates each app to match the configuration file.

### Example plugin configuration file

```
apiVersion: 1

apps:
# <string> the type of app, plugin identifier. Required
- type: raintank-worldping-app
# <int> Org ID. Default to 1, unless org_name is specified
  org_id: 1
# <string> Org name. Overrides org_id unless org_id not specified
  org_name: Main Org.
# <bool> disable the app. Default to false.
  disabled: false
# <map> fields that will be converted to json and stored in jsonData. Custom per
```

```

app.
  jsonData:
    # key/value pairs of string to object
    key: value
  # <map> fields that will be converted to json, encrypted and stored in
  secureJsonData. Custom per app.
  secureJsonData:
    # key/value pairs of string to string
    key: value

```

## Dashboards

You can manage dashboards in Grafana by adding one or more YAML config files in the `[ provisioning/dashboards ]` directory. Each config file can contain a list of dashboards providers that load dashboards into Grafana from the local filesystem.

The dashboard provider config file looks somewhat like this:

```

apiVersion: 1

providers:
  # <string> an unique provider name. Required
  - name: 'a unique provider name'
    # <int> Org id. Default to 1
    orgId: 1
    # <string> name of the dashboard folder.
    folder: ''
    # <string> folder UID. will be automatically generated if not specified
    folderUid: ''
    # <string> provider type. Default to 'file'
    type: file
    # <bool> disable dashboard deletion
    disableDeletion: false
    # <int> how often Grafana will scan for changed dashboards
    updateIntervalSeconds: 10
    # <bool> allow updating provisioned dashboards from the UI
    allowUiUpdates: false
    options:
      # <string, required> path to dashboard files on disk. Required when using the
      'file' type
      path: /var/lib/grafana/dashboards
      # <bool> use folder names from filesystem to create folders in Grafana
      foldersFromFilesStructure: true

```

When Grafana starts, it will update/insert all dashboards available in the configured path. Then later on poll that path every **updateIntervalSeconds** and look for updated json files and update/insert those into the database.

**Note:** Dashboards are provisioned to the General folder if the `folder` option is missing or empty.

### Making changes to a provisioned dashboard



It's possible to make changes to a provisioned dashboard in the Grafana UI. However, it is not possible to automatically save the changes back to the provisioning source. If `allowUiUpdates` is set to `true` and you make changes to a provisioned dashboard, you can `Save` the dashboard then changes will be persisted to the Grafana database.

**Note:** If a provisioned dashboard is saved from the UI and then later updated from the source, the dashboard stored in the database will always be overwritten. The `version` property in the JSON file will not affect this, even if it is lower than the existing dashboard.

If a provisioned dashboard is saved from the UI and the source is removed, the dashboard stored in the database will be deleted unless the configuration option `disableDeletion` is set to `true`.

If `allowUiUpdates` is configured to `false`, you are not able to make changes to a provisioned dashboard. When you click `Save`, Grafana brings up a *Cannot save provisioned dashboard* dialog. The screenshot below illustrates this behavior.

Grafana offers options to export the JSON definition of a dashboard. Either `Copy JSON to Clipboard` or `Save JSON to file` can help you synchronize your dashboard changes back to the provisioning source.

Note: The JSON definition in the input field when using `Copy JSON to Clipboard` or `Save JSON to file` will have the `id` field automatically removed to aid the provisioning workflow.

{{< figure src="/static/img/docs/v51/provisioning\_cannot\_save\_dashboard.png" max-width="500px" class="docs-image--no-shadow" >}}

## Reusable Dashboard URLs

If the dashboard in the JSON file contains an [UID]({{< relref "../dashboards/json-model.md" >}}), Grafana forces insert/update on that UID. This allows you to migrate dashboards between Grafana instances and provisioning Grafana from configuration without breaking the URLs given because the new dashboard URL uses the UID as identifier. When Grafana starts, it updates/inserts all dashboards available in the configured folders. If you modify the file, then the dashboard is also updated. By default, Grafana deletes dashboards in the database if the file is removed. You can disable this behavior using the `disableDeletion` setting.

**Note:** Provisioning allows you to overwrite existing dashboards which leads to problems if you re-use settings that are supposed to be unique. Be careful not to re-use the same `title` multiple times within a folder or `uid` within the same installation as this will cause weird behaviors.

## Provision folders structure from filesystem to Grafana

If you already store your dashboards using folders in a git repo or on a filesystem, and also you want to have the same folder names in the Grafana menu, you can use `foldersFromFilesStructure` option.

For example, to replicate these dashboards structure from the filesystem to Grafana,

```
/etc/dashboards
├── /server
│   ├── /common_dashboard.json
│   └── /network_dashboard.json
└── /application
    ├── /requests_dashboard.json
    └── /resources_dashboard.json
```

you need to specify just this short provision configuration file.

```

apiVersion: 1

providers:
- name: dashboards
  type: file
  updateIntervalSeconds: 30
  options:
    path: /etc/dashboards
    foldersFromFilesStructure: true

```

`server` and `application` will become new folders in Grafana menu.

**Note:** `folder` and `folderUid` options should be empty or missing to make `foldersFromFilesStructure` work.

**Note:** To provision dashboards to the General folder, store them in the root of your `path`.

## Alert Notification Channels

Alert Notification Channels can be provisioned by adding one or more YAML config files in the `provisioning/notifiers` directory.

Each config file can contain the following top-level fields:

- `notifiers`, a list of alert notifications that will be added or updated during start up. If the notification channel already exists, Grafana will update it to match the configuration file.
- `delete_notifiers`, a list of alert notifications to be deleted before inserting/updating those in the `notifiers` list.

Provisioning looks up alert notifications by uid, and will update any existing notification with the provided uid.

By default, exporting a dashboard as JSON will use a sequential identifier to refer to alert notifications. The field `uid` can be optionally specified to specify a string identifier for the alert name.

```

{
  ...
  "alert": {
    ...,
    "conditions": [...],
    "frequency": "24h",
    "noDataState": "ok",
    "notifications": [
      {"uid": "notifier1"},
      {"uid": "notifier2"},
    ]
  }
  ...
}

```

### Example Alert Notification Channels Config File

```

notifiers:
- name: notification-channel-1
  type: slack
  uid: notifier1
  # either
  org_id: 2
  # or
  org_name: Main Org.
  is_default: true
  send_reminder: true
  frequency: 1h
  disable_resolve_message: false
  # See `Supported Settings` section for settings supported for each
  # alert notification type.
  settings:
    recipient: 'XXX'
    uploadImage: true
    token: 'xoxb' # legacy setting since Grafana v7.2 (stored non-encrypted)
    url: https://slack.com # legacy setting since Grafana v7.2 (stored non-
encrypted)
    # Secure settings that will be encrypted in the database (supported since
Grafana v7.2). See `Supported Settings` section for secure settings supported for
each notifier.
    secure_settings:
      token: 'xoxb'
      url: https://slack.com

delete_notifiers:
- name: notification-channel-1
  uid: notifier1
  # either
  org_id: 2
  # or
  org_name: Main Org.
- name: notification-channel-2
  # default org_id: 1

```

## Supported Settings

The following sections detail the supported settings and secure settings for each alert notification type. Secure settings are stored encrypted in the database and you add them to `secure_settings` in the YAML file instead of `settings`.

**Note:** Secure settings is supported since Grafana v7.2.

### Alert notification pushover

Name	Secure setting
apiToken	yes
userKey	yes

device	
priority	
okPriority	
retry	
expire	
sound	
okSound	

#### Alert notification discord

Name	Secure setting
url	yes
avatar_url	
content	
use_discord_username	

#### Alert notification slack

Name	Secure setting
url	yes
recipient	
username	
icon_emoji	
icon_url	
uploadImage	
mentionUsers	
mentionGroups	
mentionChannel	
token	yes

#### Alert notification victorops

Name
url
autoResolve

**Alert notification** kafka

Name
kafkaRestProxy
kafkaTopic

**Alert notification** LINE

Name	Secure setting
token	yes

**Alert notification** pagerduty

Name	Secure setting
integrationKey	yes
autoResolve	

**Alert notification** sensu

Name	Secure setting
url	
source	
handler	
username	
password	yes

**Alert notification** sensugo

Name	Secure setting
url	
apikey	yes
entity	
check	
handler	
namespace	

**Alert notification** prometheus-alertmanager

Name	Secure setting
url	

basicAuthUser	
basicAuthPassword	yes

#### Alert notification `teams`

Name
url

#### Alert notification `dingding`

Name
url

#### Alert notification `email`

Name
singleEmail
addresses

#### Alert notification `hipchat`

Name
url
apikey
roomid

#### Alert notification `opsgenie`

Name	Secure setting
apiKey	yes
apiUrl	
autoClose	
overridePriority	
sendTagsAs	

#### Alert notification `telegram`

Name	Secure setting
bottoken	yes
chatid	
uploadImage	

#### Alert notification `threema`

Name	Secure setting
gateway_id	
recipient_id	
api_secret	yes

#### Alert notification `webhook`

Name	Secure setting
url	
httpMethod	
username	
password	yes

#### Alert notification `googlechat`

Name
url

## Grafana Enterprise

Grafana Enterprise supports provisioning for the following resources:

- [Access Control Provisioning]({{< relref "../enterprise/access-control/provisioning.md" >}})