

Kernel Support for miscellaneous Binary Formats (binfmt_misc)

This Kernel feature allows you to invoke almost (for restrictions see below) every program by simply typing its name in the shell. This includes for example compiled Java(TM), Python or Emacs programs.

To achieve this you must tell `binfmt_misc` which interpreter has to be invoked with which binary. `Binfmt_misc` recognises the binary-type by matching some bytes at the beginning of the file with a magic byte sequence (masking out specified bits) you have supplied. `Binfmt_misc` can also recognise a filename extension aka `.com` or `.exe`.

First you must mount `binfmt_misc`:

```
mount binfmt_misc -t binfmt_misc /proc/sys/fs/binfmt_misc
```

To actually register a new binary type, you have to set up a string looking like

```
:name:type:offset:magic:mask:interpreter:flags (where you can choose the : upon your needs) and echo it to /proc/sys/fs/binfmt_misc/register.
```

Here is what the fields mean:

- name**
is an identifier string. A new `/proc` file will be created with this name below `/proc/sys/fs/binfmt_misc`; cannot contain slashes `/` for obvious reasons.
- type**
is the type of recognition. Give `M` for magic and `E` for extension.
- offset**
is the offset of the magic/mask in the file, counted in bytes. This defaults to 0 if you omit it (i.e. you write `:name:type::magic...`). Ignored when using filename extension matching.
- magic**
is the byte sequence `binfmt_misc` is matching for. The magic string may contain hex-encoded characters like `\x0a` or `\xA4`. Note that you must escape any NUL bytes; parsing halts at the first one. In a shell environment you might have to write `\\x0a` to prevent the shell from eating your `\`. If you chose filename extension matching, this is the extension to be recognised (without the `.`, the `\x0a` specials are not allowed). Extension matching is case sensitive, and slashes `/` are not allowed!
- mask**
is an (optional, defaults to all 0xff) mask. You can mask out some bits from matching by supplying a string like magic and as long as magic. The mask is anded with the byte sequence of the file. Note that you must escape any NUL bytes; parsing halts at the first one. Ignored when using filename extension matching.
- interpreter**
is the program that should be invoked with the binary as first argument (specify the full path)
- flags**
is an optional field that controls several aspects of the invocation of the interpreter. It is a string of capital letters, each controls a certain aspect. The following flags are supported:

P - preserve-argv[0]

Legacy behavior of `binfmt_misc` is to overwrite the original `argv[0]` with the full path to the binary. When this flag is included, `binfmt_misc` will add an argument to the argument vector for this purpose, thus preserving the original `argv[0]`. e.g. If your `interp` is set to `/bin/foo` and you run `blah` (which is in `/usr/local/bin`), then the kernel will execute `/bin/foo` with `argv[]` set to `["/bin/foo", "/usr/local/bin/blah", "blah"]`. The `interp` has to be aware of this so it can execute `/usr/local/bin/blah` with `argv[]` set to `["blah"]`.

O - open-binary

Legacy behavior of `binfmt_misc` is to pass the full path of the binary to the interpreter as an argument. When this flag is included, `binfmt_misc` will open the file for reading and pass its descriptor as an argument, instead of the full path, thus allowing the interpreter to execute non-readable binaries. This feature should be used with care - the interpreter has to be trusted not to emit the contents of the non-readable binary.

C - credentials

Currently, the behavior of `binfmt_misc` is to calculate the credentials and security token of the new process according to the interpreter. When this flag is included, these attributes are calculated according to the binary. It also implies the `O` flag. This feature should be used with care as the interpreter will run with root permissions when a setuid binary owned by root is run with `binfmt_misc`.

F - fix binary

The usual behaviour of `binfmt_misc` is to spawn the binary lazily when the `misc` format file is invoked. However, this doesn't work very well in the face of mount namespaces and `changelroots`, so the `F` mode opens the binary as soon as the emulation is installed and uses the opened image to spawn the emulator, meaning it is always available once installed, regardless of how the environment changes.

There are some restrictions:

- the whole register string may not exceed 1920 characters
- the magic must reside in the first 128 bytes of the file, i.e. `offset+size(magic)` has to be less than 128
- the interpreter string may not exceed 127 characters

To use `binfmt_misc` you have to mount it first. You can mount it with `mount -t binfmt_misc none`

`/proc/sys/fs/binfmt_misc` command, or you can add a line `none /proc/sys/fs/binfmt_misc binfmt_misc defaults 0 0` to your `/etc/fstab` so it auto mounts on boot.

You may want to add the binary formats in one of your `/etc/rc` scripts during boot-up. Read the manual of your init program to figure out how to do this right.

Think about the order of adding entries! Later added entries are matched first!

A few examples (assumed you are in `/proc/sys/fs/binfmt_misc`):

- enable support for `em86` (like `binfmt_em86`, for Alpha AXP only):

```
echo ':i386:M::\x7fELF\x01\x00\x00\x00\x00\x00\x00\x00\x00\x00\x02\x00\x03:\xff\xff\xff\xff\xff\xff\xe0\xe0\xff\xff\xff\xff' > register
echo ':i486:M::\x7fELF\x01\x00\x00\x00\x00\x00\x00\x00\x00\x00\x02\x00\x06:\xff\xff\xff\xff\xff\xff\xe0\xe0\xff\xff\xff\xff' > register
```

- enable support for packed DOS applications (pre-configured `dosexu` hdmages):

```
echo ':DEXE:M::\x0eDEX::/usr/bin/dosexec:' > register
```

- enable support for Windows executables using wine:

```
echo ':DOSWin:M:MZ::/usr/local/bin/wine:' > register
```

For java support see [Documentation/admin-guide/java.rst](#)

You can enable/disable `binfmt_misc` or one binary type by echoing 0 (to disable) or 1 (to enable) to

`/proc/sys/fs/binfmt_misc/status` or `/proc/.../the_name`. Catting the file tells you the current status of `binfmt_misc/the_entry`.

You can remove one entry or all entries by echoing -1 to `/proc/.../the_name` or `/proc/sys/fs/binfmt_misc/status`.

Hints

If you want to pass special arguments to your interpreter, you can write a wrapper script for it. See `doc:Documentation/admin-guide/java.rst` `<./java>` for an example.

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\linux-master\Documentation\admin-guide\linux-master [Documentation] [admin-guide]binfmt-misc.rst, line 142); [backlink](#)
Unknown interpreted text role "doc".

Your interpreter should NOT look in the PATH for the filename; the kernel passes it the full filename (or the file descriptor) to use. Using `$PATH` can cause unexpected behaviour and can be a security hazard.

Richard GÄnther <rguenth@tat.physik.uni-tuebingen.de>