

Process Number Controller

Abstract

The process number controller is used to allow a cgroup hierarchy to stop any new tasks from being fork()'d or clone()'d after a certain limit is reached.

Since it is trivial to hit the task limit without hitting any kmemcg limits in place, PIDs are a fundamental resource. As such, PID exhaustion must be preventable in the scope of a cgroup hierarchy by allowing resource limiting of the number of tasks in a cgroup.

Usage

In order to use the *pids* controller, set the maximum number of tasks in *pids.max* (this is not available in the root cgroup for obvious reasons). The number of processes currently in the cgroup is given by *pids.current*.

Organisational operations are not blocked by cgroup policies, so it is possible to have *pids.current* > *pids.max*. This can be done by either setting the limit to be smaller than *pids.current*, or attaching enough processes to the cgroup such that *pids.current* > *pids.max*. However, it is not possible to violate a cgroup policy through *fork()* or *clone()*. *fork()* and *clone()* will return -EAGAIN if the creation of a new process would cause a cgroup policy to be violated.

To set a cgroup to have no limit, set *pids.max* to "max". This is the default for all new cgroups (N.B. that PID limits are hierarchical, so the most stringent limit in the hierarchy is followed).

pids.current tracks all child cgroup hierarchies, so *parent/pids.current* is a superset of *parent/child/pids.current*.

The *pids.events* file contains event counters:

- *max*: Number of times fork failed because limit was hit.

Example

First, we mount the *pids* controller:

```
# mkdir -p /sys/fs/cgroup/pids
# mount -t cgroup -o pids none /sys/fs/cgroup/pids
```

Then we create a hierarchy, set limits and attach processes to it:

```
# mkdir -p /sys/fs/cgroup/pids/parent/child
# echo 2 > /sys/fs/cgroup/pids/parent/pids.max
# echo $$ > /sys/fs/cgroup/pids/parent/cgroup.procs
# cat /sys/fs/cgroup/pids/parent/pids.current
2
#
```

It should be noted that attempts to overcome the set limit (2 in this case) will fail:


```
# cat /sys/fs/cgroup/pids/parent/pids.current
2
# ( /bin/echo "Here's some processes for you." | cat )
sh: fork: Resource temporary unavailable
#
```

Even if we migrate to a child cgroup (which doesn't have a set limit), we will not be able to overcome the most stringent limit in the hierarchy (in this case, parent's):

```
# echo $$ > /sys/fs/cgroup/pids/parent/child/cgroup.procs
# cat /sys/fs/cgroup/pids/parent/pids.current
2
# cat /sys/fs/cgroup/pids/parent/child/pids.current
2
# cat /sys/fs/cgroup/pids/parent/child/pids.max
max
# ( /bin/echo "Here's some processes for you." | cat )
sh: fork: Resource temporary unavailable
#
```

We can set a limit that is smaller than *pids.current*, which will stop any new processes from being forked at all (note that the shell itself counts towards *pids.current*):

```
# echo 1 > /sys/fs/cgroup/pids/parent/pids.max
# /bin/echo "We can't even spawn a single process now."
sh: fork: Resource temporary unavailable
# echo 0 > /sys/fs/cgroup/pids/parent/pids.max
```



```
# /bin/echo "We can't even spawn a single process now."  
sh: fork: Resource temporary unavailable  
#
```

