

BPF drgn tools

drgn scripts is a convenient and easy to use mechanism to retrieve arbitrary kernel data structures. drgn is not relying on kernel UAPI to read the data. Instead it's reading directly from `/proc/kcore` or `vmcore` and pretty prints the data based on DWARF debug information from `vnlinux`.

This document describes BPF related drgn tools.

See [drgn/tools](#) for all tools available at the moment and [drgn/doc](#) for more details on drgn itself.

bpf_inspect.py

Description

[bpf_inspect.py](#) is a tool intended to inspect BPF programs and maps. It can iterate over all programs and maps in the system and print basic information about these objects, including id, type and name.

The main use-case [bpf_inspect.py](#) covers is to show BPF programs of types `BPF_PROG_TYPE_EXT` and `BPF_PROG_TYPE_TRACING` attached to other BPF programs via `freplace/fentry/fexit` mechanisms, since there is no user-space API to get this information.

Getting started

List BPF programs (full names are obtained from BTf):

```
% sudo bpf_inspect.py prog
 27: BPF_PROG_TYPE_TRACEPOINT      tracepoint_tcp_tcp_send_reset
4632: BPF_PROG_TYPE_CGROUP_SOCK_ADDR tw_ip_t_bind
49464: BPF_PROG_TYPE_RAW_TRACEPOINT raw_tracepoint_sched_process_exit
```

List BPF maps:

```
% sudo bpf_inspect.py map
2577: BPF_MAP_TYPE_HASH          tw_ip_t_vips
4050: BPF_MAP_TYPE_STACK_TRACE   stack_traces
4069: BPF_MAP_TYPE_PERCPU_ARRAY   ned_dctcp_cntr
```

Find BPF programs attached to BPF program `test_pkt_access`:

```
% sudo bpf_inspect.py p | grep test_pkt_access
650: BPF_PROG_TYPE_SCHED_CLS      test_pkt_access
654: BPF_PROG_TYPE_TRACING         test_main
655: BPF_PROG_TYPE_TRACING         test_subprog1
656: BPF_PROG_TYPE_TRACING         test_subprog2
657: BPF_PROG_TYPE_TRACING         test_subprog3
658: BPF_PROG_TYPE_EXT             new_get_skb_len
659: BPF_PROG_TYPE_EXT             new_get_skb_ifindex
660: BPF_PROG_TYPE_EXT             new_get_constant
                                linked:[650->25: BPF_TRAMP_FEXIT test_pkt_access->tes
                                linked:[650->29: BPF_TRAMP_FEXIT test_pkt_access->tes
                                linked:[650->31: BPF_TRAMP_FEXIT test_pkt_access->tes
                                linked:[650->21: BPF_TRAMP_FEXIT test_pkt_access->tes
                                linked:[650->16: BPF_TRAMP_REPLACE test_pkt_access->g
                                linked:[650->23: BPF_TRAMP_REPLACE test_pkt_access->g
                                linked:[650->19: BPF_TRAMP_REPLACE test_pkt_access->g
```

It can be seen that there is a program `test_pkt_access`, id 650 and there are multiple other tracing and ext programs attached to functions in `test_pkt_access`.

For example the line:

```
658: BPF_PROG_TYPE_EXT             new_get_skb_len
                                linked:[650->16: BPF_TRAMP_REPLACE test_pkt_access->get_
```

, means that BPF program id 658, type `BPF_PROG_TYPE_EXT`, name `new_get_skb_len` replaces (`BPF_TRAMP_REPLACE`) function `get_skb_len()` that has BTf id 16 in BPF program id 650, name `test_pkt_access`.

Getting help:

```
System Message: WARNING/2 (D:\onboarding-resources\sample-onboarding-resources\linux-master\Documentation\bpf\ [linux-master] [Documentation] [bpf] drgn.rst, line 75)
```

Cannot analyze code. No Pygments lexer found for "none".

```
.. code-block:: none

% sudo bpf_inspect.py
usage: bpf_inspect.py [-h] {prog,p,map,m} ...

drgn script to list BPF programs or maps and their properties
unavailable via kernel API.

See https://github.com/osandov/drgn/ for more details on drgn.

optional arguments:
  -h, --help            show this help message and exit

subcommands:
  {prog,p,map,m}
    prog (p)            list BPF programs
    map (m)             list BPF maps
```

Customization

The script is intended to be customized by developers to print relevant information about BPF programs, maps and other objects.

For example, to print `struct bpf_prog_aux` for BPF program id 53077:

```
System Message: WARNING/2 (D:\onboarding-resources\sample-onboarding-resources\linux-master\Documentation\bpf\ [linux-master] [Documentation] [bpf] drgn.rst, line 101)
```

Cannot analyze code. No Pygments lexer found for "none".

```
.. code-block:: none

% git diff
diff --git a/tools/bpf_inspect.py b/tools/bpf_inspect.py
index 650e228..aea2357 100755
--- a/tools/bpf_inspect.py
+++ b/tools/bpf_inspect.py
@@ -112,7 +112,9 @@ def list_bpf_progs(args):
     if linked:
         linked = f" linked:[{linked}]"

-    print(f"{id:>6}: {type_:32} {name:32} {linked}")
+    if id == 53077:
+        print(f"{id:>6}: {type_:32} {name:32}")
+        print(f"{bpf_prog_aux}")
```

```
def list_bpf_maps(args):
```

It produces the output:

```
% sudo bpf_inspect.py p
53077: BPF_PROG_TYPE_XDP tw_xdp_policer
*(struct bpf_prog_aux *)0xffff8893fad4b400 = {
    .refcnt = (atomic64_t){
        .counter = (long)58,
    },
    .used_map_cnt = (u32)1,
    .max_ctx_offset = (u32)8,
    .max_pkt_offset = (u32)15,
    .max_tp_access = (u32)0,
    .stack_depth = (u32)8,
    .id = (u32)53077,
    .func_cnt = (u32)0,
    .func_idx = (u32)0,
    .attach_btf_id = (u32)0,
    .linked_prog = (struct bpf_prog *)0x0,
    .verifier_text = (bool)0,
    .offload_requested = (bool)0,
    .attach_btf_trace = (bool)0,
    .func_proto_unreliable = (bool)0,
    .trampoline_prog_type = (enum bpf_trampoline_type)BPF_TRAMP_FENTRY,
    .trampoline = (struct bpf_trampoline *)0x0,
    .tramp_hlist = (struct hlist_node){
        .next = (struct hlist_node *)0x0,
        .pprev = (struct hlist_node **)0x0,
    },
    .attach_func_proto = (const struct btf_type *)0x0,
    .attach_func_name = (const char *)0x0,
    .func = (struct bpf_prog **)0x0,
    .jit_data = (void *)0x0,
    .poke_tab = (struct bpf_jit_poke_descriptor *)0x0,
    .size_poke_tab = (u32)0,
    .ksym_tnode = (struct latch_tree_node){
        .node = (struct rb_node [2]){
            {
                .__rb_parent_color = (unsigned long)18446612956263126665,
                .rb_right = (struct rb_node *)0x0,
                .rb_left = (struct rb_node *)0xffff88a0be3d0088,
            },
            {
                .__rb_parent_color = (unsigned long)18446612956263126689,
                .rb_right = (struct rb_node *)0x0,
                .rb_left = (struct rb_node *)0xffff88a0be3d00a0,
            },
        },
    },
    .ksym_lnode = (struct list_head){
        .next = (struct list_head *)0xffff88bf481830b8,
        .prev = (struct list_head *)0xffff888309f536b8,
    },
    .ops = (const struct bpf_prog_ops *)xdp_prog_ops+0x0 = 0xffffffff820fa350,
    .used_maps = (struct bpf_map **)0xffff889ff795de98,
    .prog = (struct bpf_prog *)0xffff9000cf2d000,
    .user = (struct user_struct *)root_user+0x0 = 0xffffffff82444820,
    .load_time = (u64)2408348759285319,
    .cgroup_storage = (struct bpf_map *[2]){},
    .name = (char [16])"tw_xdp_policer",
    .security = (void *)0xffff889ff795d548,
    .offload = (struct bpf_prog_offload *)0x0,
    .btf = (struct btf *)0xffff8890ce6d0580,
    .func_info = (struct bpf_func_info *)0xffff889ff795d240,
    .func_info_aux = (struct bpf_func_info_aux *)0xffff889ff795de20,
    .linfo = (struct bpf_line_info *)0xffff888a707afc00,
    .jited_linfo = (void **)0xffff8893fad48600,
    .func_info_cnt = (u32)1,
    .nr_linfo = (u32)37,
    .linfo_idx = (u32)0,
    .num_exentries = (u32)0,
    .extable = (struct exception_table_entry *)0xfffffffffa032d950,
    .stats = (struct bpf_prog_stats *)0x603fe3alf6d0,
    .work = (struct work_struct){
        .data = (atomic_long_t){
            .counter = (long)0,
        },
        .entry = (struct list_head){
            .next = (struct list_head *)0x0,
            .prev = (struct list_head *)0x0,
        },
        .func = (work_func_t)0x0,
    },
    .rcu = (struct callback_head){
        .next = (struct callback_head *)0x0,
        .func = (void (*)(struct callback_head *))0x0,
    },
}
```