

S390 Debug Feature

files:

- arch/s390/kernel/debug.c
- arch/s390/include/asm/debug.h

Description:

The goal of this feature is to provide a kernel debug logging API where log records can be stored efficiently in memory, where each component (e.g. device drivers) can have one separate debug log. One purpose of this is to inspect the debug logs after a production system crash in order to analyze the reason for the crash.

If the system still runs but only a subcomponent which uses dbf fails, it is possible to look at the debug logs on a live system via the Linux debugfs filesystem.

The debug feature may also very useful for kernel and driver development.

Design:

Kernel components (e.g. device drivers) can register themselves at the debug feature with the function call `:c:func:'debug_register()'`. This function initializes a debug log for the caller. For each debug log exists a number of debug areas where exactly one is active at one time. Each debug area consists of contiguous pages in memory. In the debug areas there are stored debug entries (log records) which are written by event- and exception-calls.

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\linux-master\Documentation\s390\[linux-master] [Documentation] [s390] s390dbf.rst, line 25); [backlink](#)

Unknown interpreted text role "c:func".

An event-call writes the specified debug entry to the active debug area and updates the log pointer for the active area. If the end of the active debug area is reached, a wrap around is done (ring buffer) and the next debug entry will be written at the beginning of the active debug area.

An exception-call writes the specified debug entry to the log and switches to the next debug area. This is done in order to be sure that the records which describe the origin of the exception are not overwritten when a wrap around for the current area occurs.

The debug areas themselves are also ordered in form of a ring buffer. When an exception is thrown in the last debug area, the following debug entries are then written again in the very first area.

There are four versions for the event- and exception-calls: One for logging raw data, one for text, one for numbers (unsigned int and long), and one for sprintf-like formatted strings.

Each debug entry contains the following data:

- Timestamp
- Cpu-Number of calling task
- Level of debug entry (0...6)
- Return Address to caller
- Flag, if entry is an exception or not

The debug logs can be inspected in a live system through entries in the debugfs-filesystem. Under the toplevel directory "s390dbf" there is a directory for each registered component, which is named like the corresponding component. The debugfs normally should be mounted to `/sys/kernel/debug` therefore the debug feature can be accessed under `/sys/kernel/debug/s390dbf`.

The content of the directories are files which represent different views to the debug log. Each component can decide which views should be used through registering them with the function `:c:func:'debug_register_view()'`. Predefined views for hex/ascii and sprintf data are provided. It is also possible to define other views. The content of a view can be inspected simply by reading the corresponding debugfs file.

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\linux-master\Documentation\s390\[linux-master] [Documentation] [s390] s390dbf.rst, line 67); [backlink](#)

Unknown interpreted text role "c:func".

All debug logs have an actual debug level (range from 0 to 6). The default level is 3. Event and Exception functions have a `:c:data:'level'` parameter. Only debug entries with a level that is lower or equal than the actual level are written to the log. This means, when writing events, high priority log entries should have a low level value whereas low priority entries should have a high one. The actual debug level can be changed with the help of the debugfs-filesystem through writing a number string "x" to the `level` debugfs

file which is provided for every debug log. Debugging can be switched off completely by using "-" on the `level` debugfs file.

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\linux-master\Documentation\s390\[linux-master] [Documentation] [s390]s390dbf.rst, line 74); [backlink](#)
Unknown interpreted text role "c:data".

Example:

```
> echo "-" > /sys/kernel/debug/s390dbf/dasd/level
```

It is also possible to deactivate the debug feature globally for every debug log. You can change the behavior using 2 sysctl parameters in `/proc/sys/s390dbf`:

There are currently 2 possible triggers, which stop the debug feature globally. The first possibility is to use the `debug_active` sysctl. If set to 1 the debug feature is running. If `debug_active` is set to 0 the debug feature is turned off.

The second trigger which stops the debug feature is a kernel oops. That prevents the debug feature from overwriting debug information that happened before the oops. After an oops you can reactivate the debug feature by piping 1 to `/proc/sys/s390dbf/debug_active`. Nevertheless, it's not suggested to use an oopsed kernel in a production environment.

If you want to disallow the deactivation of the debug feature, you can use the `debug_stoppable` sysctl. If you set `debug_stoppable` to 0 the debug feature cannot be stopped. If the debug feature is already stopped, it will stay deactivated.

Kernel Interfaces:

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\linux-master\Documentation\s390\[linux-master] [Documentation] [s390]s390dbf.rst, line 112)
Unknown directive type "kernel-doc".

```
.. kernel-doc:: arch/s390/kernel/debug.c
```

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\linux-master\Documentation\s390\[linux-master] [Documentation] [s390]s390dbf.rst, line 113)
Unknown directive type "kernel-doc".

```
.. kernel-doc:: arch/s390/include/asm/debug.h
```

Predefined views:

```
extern struct debug_view debug_hex_ascii_view;  
extern struct debug_view debug_sprintf_view;
```

Examples

```
/*  
 * hex_ascii-view Example  
 */  
  
#include <linux/init.h>  
#include <asm/debug.h>  
  
static debug_info_t *debug_info;  
  
static int init(void)  
{  
    /* register 4 debug areas with one page each and 4 byte data field */  
  
    debug_info = debug_register("test", 1, 4, 4 );  
    debug_register_view(debug_info, &debug_hex_ascii_view);  
  
    debug_text_event(debug_info, 4 , "one ");  
    debug_int_exception(debug_info, 4, 4711);  
    debug_event(debug_info, 3, &debug_info, 4);  
  
    return 0;  
}  
  
static void cleanup(void)
```

```

{
    debug_unregister(debug_info);
}

module_init(init);
module_exit(cleanup);

/*
 * sprintf-view Example
 */

#include <linux/init.h>
#include <asm/debug.h>

static debug_info_t *debug_info;

static int init(void)
{
    /* register 4 debug areas with one page each and data field for */
    /* format string pointer + 2 varargs (= 3 * sizeof(long))          */

    debug_info = debug_register("test", 1, 4, sizeof(long) * 3);
    debug_register_view(debug_info, &debug_sprintf_view);

    debug_sprintf_event(debug_info, 2, "first event in %s:%i\n", __FILE__, __LINE__);
    debug_sprintf_exception(debug_info, 1, "pointer to debug info: %p\n", &debug_info);

    return 0;
}

static void cleanup(void)
{
    debug_unregister(debug_info);
}

module_init(init);
module_exit(cleanup);

```

Debugfs Interface

Views to the debug logs can be investigated through reading the corresponding debugfs-files:

Example:

```

> ls /sys/kernel/debug/s390dbf/dasd
flush hex_ascii level pages
> cat /sys/kernel/debug/s390dbf/dasd/hex_ascii | sort -k2,2 -s
00 00974733272:680099 2 - 02 0006ad7e 07 ea 4a 90 | ....
00 00974733272:682210 2 - 02 0006ade6 46 52 45 45 | FREE
00 00974733272:682213 2 - 02 0006adf6 07 ea 4a 90 | ....
00 00974733272:682281 1 * 02 0006ab08 41 4c 4c 43 | EXCP
01 00974733272:682284 2 - 02 0006ab16 45 43 4b 44 | ECKD
01 00974733272:682287 2 - 02 0006ab28 00 00 00 04 | ....
01 00974733272:682289 2 - 02 0006ab3e 00 00 00 20 | ...
01 00974733272:682297 2 - 02 0006ad7e 07 ea 4a 90 | ....
01 00974733272:684384 2 - 00 0006ade6 46 52 45 45 | FREE
01 00974733272:684388 2 - 00 0006adf6 07 ea 4a 90 | ....

```

See section about predefined views for explanation of the above output!

Changing the debug level

Example:

```

> cat /sys/kernel/debug/s390dbf/dasd/level
3
> echo "5" > /sys/kernel/debug/s390dbf/dasd/level
> cat /sys/kernel/debug/s390dbf/dasd/level
5

```

Flushing debug areas

Debug areas can be flushed with piping the number of the desired area (0...n) to the debugfs file "flush". When using "-" all debug areas are flushed.

Examples:

1. Flush debug area 0:

```
> echo "0" > /sys/kernel/debug/s390dbf/dasd/flush
```

2. Flush all debug areas:

```
> echo "-" > /sys/kernel/debug/s390dbf/dasd/flush
```

Changing the size of debug areas

It is possible to change the size of debug areas through piping the number of pages to the debugfs file "pages". The resize request will also flush the debug areas.

Example:

Define 4 pages for the debug areas of debug feature "dasd":

```
> echo "4" > /sys/kernel/debug/s390dbf/dasd/pages
```

Stopping the debug feature

Example:

1. Check if stopping is allowed:

```
> cat /proc/sys/s390dbf/debug_stoppable
```

2. Stop debug feature:

```
> echo 0 > /proc/sys/s390dbf/debug_active
```

crash Interface

The `crash` tool since v5.1.0 has a built-in command `s390dbf` to display all the debug logs or export them to the file system. With this tool it is possible to investigate the debug logs on a live system and with a memory dump after a system crash.

Investigating raw memory

One last possibility to investigate the debug logs at a live system and after a system crash is to look at the raw memory under VM or at the Service Element. It is possible to find the anchor of the debug-logs through the `debug_area_first` symbol in the System map. Then one has to follow the correct pointers of the data-structures defined in `debug.h` and find the debug-areas in memory. Normally modules which use the debug feature will also have a global variable with the pointer to the debug-logs. Following this pointer it will also be possible to find the debug logs in memory.

For this method it is recommended to use `'16 * x + 4'` byte ($x = 0..n$) for the length of the data field in `:c:func:'debug_register()'` in order to see the debug entries well formatted.

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\linux-master\Documentation\s390\linux-master\Documentation\s390\s390dbf.rst, line 290);
[backlink](#)

Unknown interpreted text role "c:func".

Predefined Views

There are two predefined views: `hex_ascii` and `sprintf`. The `hex_ascii` view shows the data field in hex and ascii representation (e.g. 45 43 4b 44 | ECKD).

The `sprintf` view formats the debug entries in the same way as the `sprintf` function would do. The `sprintf` event/exception functions write to the debug entry a pointer to the format string (size = `sizeof(long)`) and for each vararg a long value. So e.g. for a debug entry with a format string plus two varargs one would need to allocate a ($3 * \text{sizeof(long)}$) byte data area in the `debug_register()` function.

IMPORTANT:

Using `"%s"` in `sprintf` event functions is dangerous. You can only use `"%s"` in the `sprintf` event functions, if the memory for the passed string is available as long as the debug feature exists. The reason behind this is that due to performance considerations only a pointer to the string is stored in the debug feature. If you log a string that is freed afterwards, you will get an OOPS when inspecting the debug feature, because then the debug feature will access the already freed memory.

NOTE:

If using the `sprintf` view do NOT use other event/exception functions than the `sprintf-event` and `-exception` functions.

The format of the `hex_ascii` and `sprintf` view is as follows:

- Number of area
- Timestamp (formatted as seconds and microseconds since 00:00:00 Coordinated Universal Time (UTC), January 1, 1970)

- level of debug entry
- Exception flag (* = Exception)
- Cpu-Number of calling task
- Return Address to caller
- data field

A typical line of the hex_ascii view will look like the following (first line is only for explanation and will not be displayed when 'cating' the view):

```
area  time                level exception cpu caller      data (hex + ascii)
-----
00    00964419409:440690 1 -          00  88023fe
```

Defining views

Views are specified with the 'debug_view' structure. There are defined callback functions which are used for reading and writing the debugfs files:

```
struct debug_view {
    char name[DEBUG_MAX_PROCF_LEN];
    debug_prolog_proc_t* prolog_proc;
    debug_header_proc_t* header_proc;
    debug_format_proc_t* format_proc;
    debug_input_proc_t* input_proc;
    void* private_data;
};
```

where:

```
typedef int (debug_header_proc_t) (debug_info_t* id,
                                   struct debug_view* view,
                                   int area,
                                   debug_entry_t* entry,
                                   char* out_buf);

typedef int (debug_format_proc_t) (debug_info_t* id,
                                   struct debug_view* view, char* out_buf,
                                   const char* in_buf);

typedef int (debug_prolog_proc_t) (debug_info_t* id,
                                   struct debug_view* view,
                                   char* out_buf);

typedef int (debug_input_proc_t) (debug_info_t* id,
                                   struct debug_view* view,
                                   struct file* file, const char* user_buf,
                                   size_t in_buf_size, loff_t* offset);
```

The "private_data" member can be used as pointer to view specific data. It is not used by the debug feature itself.

The output when reading a debugfs file is structured like this:

```
"prolog_proc output"

"header_proc output 1" "format_proc output 1"
"header_proc output 2" "format_proc output 2"
"header_proc output 3" "format_proc output 3"
...
```

When a view is read from the debugfs, the Debug Feature calls the 'prolog_proc' once for writing the prolog. Then 'header_proc' and 'format_proc' are called for each existing debug entry.

The input_proc can be used to implement functionality when it is written to the view (e.g. like with `echo "0" > /sys/kernel/debug/s390dbf/dasd/level`).

For header_proc there can be used the default function `c:func:'debug_dflt_header_fn()` which is defined in debug.h. and which produces the same header output as the predefined views. E.g:

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\linux-master\Documentation\s390\linux-master\Documentation\s390\s390dbf.rst, line 400);
[backlink](#)

Unknown interpreted text role "c:func".

```
00 00964419409:440761 2 - 00 88023ec
```

In order to see how to use the callback functions check the implementation of the default views!

Example:

```

#include <asm/debug.h>

#define UNKNOWNSTR "data: %08x"

const char* messages[] =
{
    "This error.....\n",
    "That error.....\n",
    "Problem.....\n",
    "Something went wrong.\n",
    "Everything ok.....\n",
    NULL
};

static int debug_test_format_fn(
    debug_info_t *id, struct debug_view *view,
    char *out_buf, const char *in_buf
)
{
    int i, rc = 0;

    if (id->buf_size >= 4) {
        int msg_nr = *((int*)in_buf);
        if (msg_nr < sizeof(messages) / sizeof(char*) - 1)
            rc += sprintf(out_buf, "%s", messages[msg_nr]);
        else
            rc += sprintf(out_buf, UNKNOWNSTR, msg_nr);
    }
    return rc;
}

struct debug_view debug_test_view = {
    "myview", /* name of view */
    NULL, /* no prolog */
    &debug_dflt_header_fn, /* default header for each entry */
    &debug_test_format_fn, /* our own format function */
    NULL, /* no input function */
    NULL /* no private data */
};

```

test:

```

debug_info_t *debug_info;
int i;
...
debug_info = debug_register("test", 0, 4, 4);
debug_register_view(debug_info, &debug_test_view);
for (i = 0; i < 10; i++)
    debug_int_event(debug_info, 1, i);

```

```

> cat /sys/kernel/debug/s390dbf/test/myview
00 00964419734:611402 1 - 00 88042ca This error.....
00 00964419734:611405 1 - 00 88042ca That error.....
00 00964419734:611408 1 - 00 88042ca Problem.....
00 00964419734:611411 1 - 00 88042ca Something went wrong.
00 00964419734:611414 1 - 00 88042ca Everything ok.....
00 00964419734:611417 1 - 00 88042ca data: 00000005
00 00964419734:611419 1 - 00 88042ca data: 00000006
00 00964419734:611422 1 - 00 88042ca data: 00000007
00 00964419734:611425 1 - 00 88042ca data: 00000008
00 00964419734:611428 1 - 00 88042ca data: 00000009

```