

# Generic vcpu interface

The virtual cpu "device" also accepts the ioctls KVM\_SET\_DEVICE\_ATTR, KVM\_GET\_DEVICE\_ATTR, and KVM\_HAS\_DEVICE\_ATTR. The interface uses the same struct kvm\_device\_attr as other devices, but targets VCPU-wide settings and controls.

The groups and attributes per virtual cpu, if any, are architecture specific.

## 1. GROUP: KVM\_ARM\_VCPU\_PMU\_V3\_CTRL

**Architectures:** ARM64

### 1.1. ATTRIBUTE: KVM\_ARM\_VCPU\_PMU\_V3\_IRQ

**Parameters:** in kvm\_device\_attr.addr the address for PMU overflow interrupt is a pointer to an int

**Returns:**

-EBUSY	The PMU overflow interrupt is already set
-EFAULT	Error reading interrupt number
-ENXIO	PMUv3 not supported or the overflow interrupt not set when attempting to get it
-ENODEV	KVM_ARM_VCPU_PMU_V3 feature missing from VCPU
-EINVAL	Invalid PMU overflow interrupt number supplied or trying to set the IRQ number without using an in-kernel irqchip.

A value describing the PMUv3 (Performance Monitor Unit v3) overflow interrupt number for this vcpu. This interrupt could be a PPI or SPI, but the interrupt type must be same for each vcpu. As a PPI, the interrupt number is the same for all vcpus, while as an SPI it must be a separate number per vcpu.

### 1.2 ATTRIBUTE: KVM\_ARM\_VCPU\_PMU\_V3\_INIT

**Parameters:** no additional parameter in kvm\_device\_attr.addr

**Returns:**

-EEXIST	Interrupt number already used
-ENODEV	PMUv3 not supported or GIC not initialized
-ENXIO	PMUv3 not supported, missing VCPU feature or interrupt number not set
-EBUSY	PMUv3 already initialized

Request the initialization of the PMUv3. If using the PMUv3 with an in-kernel virtual GIC implementation, this must be done after initializing the in-kernel irqchip.

### 1.3 ATTRIBUTE: KVM\_ARM\_VCPU\_PMU\_V3\_FILTER

**Parameters:** in kvm\_device\_attr.addr the address for a PMU event filter is a pointer to a struct kvm\_pmu\_event\_filter

**Returns:**

-ENODEV	PMUv3 not supported or GIC not initialized
-ENXIO	PMUv3 not properly configured or in-kernel irqchip not configured as required prior to calling this attribute
-EBUSY	PMUv3 already initialized or a VCPU has already run
-EINVAL	Invalid filter range

Request the installation of a PMU event filter described as follows:

```
struct kvm_pmu_event_filter {
    __u16      base_event;
    __u16      nevents;

#define KVM_PMU_EVENT_ALLOW 0
#define KVM_PMU_EVENT_DENY 1

    __u8      action;
    __u8      pad[3];
};
```

A filter range is defined as the range [*@base\_event*, *@base\_event* + *@nevents*), together with an *@action* (KVM\_PMU\_EVENT\_ALLOW or KVM\_PMU\_EVENT\_DENY). The first registered range defines the global policy (global ALLOW if the first *@action* is DENY, global DENY if the first *@action* is ALLOW). Multiple ranges can be programmed, and must

fit within the event space defined by the PMU architecture (10 bits on ARMv8.0, 16 bits from ARMv8.1 onwards).

Note: "Cancelling" a filter by registering the opposite action for the same range doesn't change the default action. For example, installing an ALLOW filter for event range [0:10) as the first filter and then applying a DENY action for the same range will leave the whole range as disabled.

Restrictions: Event 0 (SW\_INCR) is never filtered, as it doesn't count a hardware event. Filtering event 0x1E (CHAIN) has no effect either, as it isn't strictly speaking an event. Filtering the cycle counter is possible using event 0x11 (CPU\_CYCLES).

## 1.4 ATTRIBUTE: KVM\_ARM\_VCPU\_PMU\_V3\_SET\_PMU

**Parameters:** in `kvm_device_attr.addr` the address to an int representing the PMU identifier.

<b>Returns:</b>	-EBUSY	PMUv3 already initialized, a VCPU has already run or an event filter has already been set
	-EFAULT	Error accessing the PMU identifier
	-ENXIO	PMU not found
	-ENODEV	PMUv3 not supported or GIC not initialized
	-ENOMEM	Could not allocate memory

Request that the VCPU uses the specified hardware PMU when creating guest events for the purpose of PMU emulation. The PMU identifier can be read from the "type" file for the desired PMU instance under `/sys/devices` (or, equivalent, `/sys/bus/event_source`). This attribute is particularly useful on heterogeneous systems where there are at least two CPU PMUs on the system. The PMU that is set for one VCPU will be used by all the other VCPUs. It isn't possible to set a PMU if a PMU event filter is already present.

Note that KVM will not make any attempts to run the VCPU on the physical CPUs associated with the PMU specified by this attribute. This is entirely left to userspace. However, attempting to run the VCPU on a physical CPU not supported by the PMU will fail and KVM\_RUN will return with `exit_reason = KVM_EXIT_FAIL_ENTRY` and populate the `fail_entry` struct by setting `hardware_entry_failure_reason` field to `KVM_EXIT_FAIL_ENTRY_CPU_UNSUPPORTED` and the `cpu` field to the processor id.

## 2. GROUP: KVM\_ARM\_VCPU\_TIMER\_CTRL

**Architectures:** ARM64

### 2.1. ATTRIBUTES: KVM\_ARM\_VCPU\_TIMER\_IRQ\_VTIMER, KVM\_ARM\_VCPU\_TIMER\_IRQ\_PTIMER

**Parameters:** in `kvm_device_attr.addr` the address for the timer interrupt is a pointer to an int

Returns:

-EINVAL	Invalid timer interrupt number
-EBUSY	One or more VCPUs has already run

A value describing the architected timer interrupt number when connected to an in-kernel virtual GIC. These must be a PPI (16 <= intid < 32). Setting the attribute overrides the default values (see below).

KVM_ARM_VCPU_TIMER_IRQ_VTIMER	The EL1 virtual timer intid (default: 27)
KVM_ARM_VCPU_TIMER_IRQ_PTIMER	The EL1 physical timer intid (default: 30)

Setting the same PPI for different timers will prevent the VCPUs from running. Setting the interrupt number on a VCPU configures all VCPUs created at that time to use the number provided for a given timer, overwriting any previously configured values on other VCPUs. Userspace should configure the interrupt numbers on at least one VCPU after creating all VCPUs and before running any VCPUs.

## 3. GROUP: KVM\_ARM\_VCPU\_PVTIME\_CTRL

**Architectures:** ARM64

### 3.1 ATTRIBUTE: KVM\_ARM\_VCPU\_PVTIME\_IPA

**Parameters:** 64-bit base address

Returns:

-ENXIO	Stolen time not implemented
-EEXIST	Base address already set for this VCPU
-EINVAL	Base address not 64 byte aligned

Specifies the base address of the stolen time structure for this VCPU. The base address must be 64 byte aligned and exist within a valid guest memory region. See `Documentation/virt/kvm/arm/pvtime.rst` for more information including the layout of the stolen time

structure.

## 4. GROUP: KVM\_VCPU\_TSC\_CTRL

**Architectures:** x86

4.1 ATTRIBUTE: KVM\_VCPU\_TSC\_OFFSET

**Parameters:** 64-bit unsigned TSC offset

Returns:

-EFAULT	Error reading/writing the provided parameter address.
-ENXIO	Attribute not supported

Specifies the guest's TSC offset relative to the host's TSC. The guest's TSC is then derived by the following equation:

$$\text{guest\_tsc} = \text{host\_tsc} + \text{KVM\_VCPU\_TSC\_OFFSET}$$

This attribute is useful to adjust the guest's TSC on live migration, so that the TSC counts the time during which the VM was paused. The following describes a possible algorithm to use for this purpose.

From the source VMM process:

1. Invoke the KVM\_GET\_CLOCK ioctl to record the host TSC (tsc\_src), kvmclock nanoseconds (guest\_src), and host CLOCK\_REALTIME nanoseconds (host\_src).
2. Read the KVM\_VCPU\_TSC\_OFFSET attribute for every vCPU to record the guest TSC offset (ofs\_src[i]).
3. Invoke the KVM\_GET\_TSC\_KHZ ioctl to record the frequency of the guest's TSC (freq).

From the destination VMM process:

4. Invoke the KVM\_SET\_CLOCK ioctl, providing the source nanoseconds from kvmclock (guest\_src) and CLOCK\_REALTIME (host\_src) in their respective fields. Ensure that the KVM\_CLOCK\_REALTIME flag is set in the provided structure.

KVM will advance the VM's kvmclock to account for elapsed time since recording the clock values. Note that this will cause problems in the guest (e.g., timeouts) unless CLOCK\_REALTIME is synchronized between the source and destination, and a reasonably short time passes between the source pausing the VMs and the destination executing steps 4-7.

5. Invoke the KVM\_GET\_CLOCK ioctl to record the host TSC (tsc\_dest) and kvmclock nanoseconds (guest\_dest).
6. Adjust the guest TSC offsets for every vCPU to account for (1) time elapsed since recording state and (2) difference in TSCs between the source and destination machine:

$$\text{ofs\_dst}[i] = \text{ofs\_src}[i] - (\text{guest\_src} - \text{guest\_dest}) * \text{freq} + (\text{tsc\_src} - \text{tsc\_dest})$$

("ofs[i] + tsc - guest \* freq" is the guest TSC value corresponding to a time of 0 in kvmclock. The above formula ensures that it is the same on the destination as it was on the source).

7. Write the KVM\_VCPU\_TSC\_OFFSET attribute for every vCPU with the respective value derived in the previous step.