

Split page table lock

Originally, `mm->page_table_lock` spinlock protected all page tables of the `mm_struct`. But this approach leads to poor page fault scalability of multi-threaded applications due high contention on the lock. To improve scalability, split page table lock was introduced.

With split page table lock we have separate per-table lock to serialize access to the table. At the moment we use split lock for PTE and PMD tables. Access to higher level tables protected by `mm->page_table_lock`.

There are helpers to lock/unlock a table and other accessor functions:

- `pte_offset_map_lock()`
maps pte and takes PTE table lock, returns pointer to the taken lock;
- `pte_unmap_unlock()`
unlocks and unmaps PTE table;
- `pte_alloc_map_lock()`
allocates PTE table if needed and take the lock, returns pointer to taken lock or NULL if allocation failed;
- `pte_lockptr()`
returns pointer to PTE table lock;
- `pmd_lock()`
takes PMD table lock, returns pointer to taken lock;
- `pmd_lockptr()`
returns pointer to PMD table lock;

Split page table lock for PTE tables is enabled compile-time if `CONFIG_SPLIT_PTLOCK_CPUS` (usually 4) is less or equal to `NR_CPUS`. If split lock is disabled, all tables are guarded by `mm->page_table_lock`.

Split page table lock for PMD tables is enabled, if it's enabled for PTE tables and the architecture supports it (see below).

Hugetlb and split page table lock

Hugetlb can support several page sizes. We use split lock only for PMD level, but not for PUD.

Hugetlb-specific helpers:

- `huge_pte_lock()`
takes pmd split lock for `PMD_SIZE` page, `mm->page_table_lock` otherwise;
- `huge_pte_lockptr()`
returns pointer to table lock;

Support of split page table lock by an architecture

There's no need in special enabling of PTE split page table lock: everything required is done by `pgtable_pte_page_ctor()` and `pgtable_pte_page_dtor()`, which must be called on PTE table allocation / freeing.

Make sure the architecture doesn't use slab allocator for page table allocation: slab uses `page->slab_cache` for its pages. This field shares storage with `page->ptl`.

PMD split lock only makes sense if you have more than two page table levels.

PMD split lock enabling requires `pgtable_pmd_page_ctor()` call on PMD table allocation and `pgtable_pmd_page_dtor()` on freeing.

Allocation usually happens in `pmd_alloc_one()`, freeing in `pmd_free()` and `pmd_free_tlb()`, but make sure you cover all PMD table allocation / freeing paths: i.e X86_PAE preallocate few PMDs on `pgd_alloc()`.

With everything in place you can set `CONFIG_ARCH_ENABLE_SPLIT_PMD_PTLOCK`.

NOTE: `pgtable_pte_page_ctor()` and `pgtable_pmd_page_ctor()` can fail -- it must be handled properly.

page->ptl

`page->ptl` is used to access split page table lock, where 'page' is struct page of page containing the table. It shares storage with `page->private` (and few other fields in union).

To avoid increasing size of struct page and have best performance, we use a trick:

- if `spinlock_t` fits into long, we use `page->ptr` as spinlock, so we can avoid indirect access and save a cache line.
- if size of `spinlock_t` is bigger then size of long, we use `page->ptl` as pointer to `spinlock_t` and allocate it dynamically. This allows to use split lock with enabled `DEBUG_SPINLOCK` or `DEBUG_LOCK_ALLOC`, but costs one more

cache line for indirect access;

The `spinlock_t` allocated in `pgtable_pte_page_ctor()` for PTE table and in `pgtable_pmd_page_ctor()` for PMD table.

Please, never access `page->ptl` directly -- use appropriate helper.