

This folder contains a [custom training loop \(CTL\)](#) implementation for ResNet50.

Before you begin

Please refer to the [README](#) in the parent directory for information on setup and preparing the data.

ResNet (custom training loop)

Similar to the [estimator implementation](#), the Keras implementation has code for the ImageNet dataset. The ImageNet version uses a ResNet50 model implemented in [resnet_model.py](#).

Pretrained Models

- [ResNet50 Checkpoints](#)
- ResNet50 TFHub: [feature vector](#) and [classification](#)

Again, if you did not download the data to the default directory, specify the location with the `--data_dir` flag:

```
python3 resnet_ctl_imagenet_main.py --data_dir=/path/to/imagenet
```

There are more flag options you can specify. Here are some examples:

- `--use_synthetic_data` : when set to true, synthetic data, rather than real data, are used;
- `--batch_size` : the batch size used for the model;
- `--model_dir` : the directory to save the model checkpoint;
- `--train_epochs` : number of epoches to run for training the model;
- `--train_steps` : number of steps to run for training the model. We now only support a number that is smaller than the number of batches in an epoch.
- `--skip_eval` : when set to true, evaluation as well as validation during training is skipped

For example, this is a typical command line to run with ImageNet data with batch size 128 per GPU:

```
python3 -m resnet_ctl_imagenet_main.py \  
  --model_dir=/tmp/model_dir/something \  
  --num_gpus=2 \  
  --batch_size=128 \  
  --train_epochs=90 \  
  --train_steps=10 \  
  --use_synthetic_data=false
```

See [common.py](#) for full list of options.

Using multiple GPUs

You can train these models on multiple GPUs using `tf.distribute.Strategy` API. You can read more about them in this [guide](#).

In this example, we have made it easier to use with just a command line flag `--num_gpus`. By default this flag is 1 if TensorFlow is compiled with CUDA, and 0 otherwise.

- `--num_gpus=0`: Uses `tf.distribute.OneDeviceStrategy` with CPU as the device.

- `--num_gpus=1`: Uses `tf.distribute.OneDeviceStrategy` with GPU as the device.
- `--num_gpus=2+`: Uses `tf.distribute.MirroredStrategy` to run synchronous distributed training across the GPUs.

If you wish to run without `tf.distribute.Strategy`, you can do so by setting `--distribution_strategy=off`.

Running on multiple GPU hosts

You can also train these models on multiple hosts, each with GPUs, using `tf.distribute.Strategy`.

The easiest way to run multi-host benchmarks is to set the [TF_CONFIG](#) appropriately at each host. e.g., to run using `MultiWorkerMirroredStrategy` on 2 hosts, the `cluster` in `TF_CONFIG` should have 2 `host:port` entries, and host `i` should have the `task` in `TF_CONFIG` set to `{"type": "worker", "index": i}`. `MultiWorkerMirroredStrategy` will automatically use all the available GPUs at each host.

Running on Cloud TPUs

Note: This model will **not** work with TPUs on Colab.

You can train the ResNet CTL model on Cloud TPUs using `tf.distribute.TPUStrategy`. If you are not familiar with Cloud TPUs, it is strongly recommended that you go through the [quickstart](#) to learn how to create a TPU and GCE VM.

To run ResNet model on a TPU, you must set `--distribution_strategy=tpu` and `--tpu=$TPU_NAME`, where `$TPU_NAME` the name of your TPU in the Cloud Console. From a GCE VM, you can run the following command to train ResNet for one epoch on a v2-8 or v3-8 TPU by setting `TRAIN_EPOCHS` to 1:

```
python3 resnet_ctl_imagenet_main.py \
  --tpu=$TPU_NAME \
  --model_dir=$MODEL_DIR \
  --data_dir=$DATA_DIR \
  --batch_size=1024 \
  --steps_per_loop=500 \
  --train_epochs=$TRAIN_EPOCHS \
  --use_synthetic_data=false \
  --dtype=fp32 \
  --enable_eager=true \
  --enable_tensorboard=true \
  --distribution_strategy=tpu \
  --log_steps=50 \
  --single_l2_loss_op=true \
  --use_tf_function=true
```

To train the ResNet to convergence, run it for 90 epochs by setting `TRAIN_EPOCHS` to 90.

Note: `$MODEL_DIR` and `$DATA_DIR` must be GCS paths.