

Evaluation Metrics for Whole Image Parsing

Whole Image Parsing [1], also known as Panoptic Segmentation [2], generalizes the tasks of semantic segmentation for "stuff" classes and instance segmentation for "thing" classes, assigning both semantic and instance labels to every pixel in an image.

Previous works evaluate the parsing result with separate metrics (e.g., one for semantic segmentation result and one for object detection result). Recently, Kirillov et al. propose the unified instance-based Panoptic Quality (PQ) metric [2] into several benchmarks [3, 4].

However, we notice that the instance-based PQ metric often places disproportionate emphasis on small instance parsing, as well as on "thing" over "stuff" classes. To remedy these effects, we propose an alternative region-based Parsing Covering (PC) metric [5], which adapts the Covering metric [6], previously used for class-agnostic segmentation quality evaluation, to the task of image parsing.

Here, we provide implementation of both PQ and PC for evaluating the parsing results. We briefly explain both metrics below for reference.

Panoptic Quality (PQ)

Given a groundtruth segmentation S and a predicted segmentation S' , PQ is defined as follows:



where R and R' are groundtruth regions and predicted regions respectively, and $|TP|$, $|FP|$, and $|FN|$ are the number of true positives, false positives, and false negatives. The matching is determined by a threshold of 0.5 Intersection-Over-Union (IOU).

PQ treats all regions of the same 'stuff' class as one instance, and the size of instances is not considered. For example, instances with 10×10 pixels contribute equally to the metric as instances with 1000×1000 pixels. Therefore, PQ is sensitive to false positives with small regions and some heuristics could improve the performance, such as removing those small regions (as also pointed out in the open-sourced evaluation code from [2]). Thus, we argue that PQ is suitable in applications where one cares equally for the parsing quality of instances irrespective of their sizes.

Parsing Covering (PC)

We notice that there are applications where one pays more attention to large objects, e.g., autonomous driving (where nearby objects are more important than far away ones). Motivated by this, we propose to also evaluate the quality of image parsing results by extending the existing Covering metric [5], which accounts for instance sizes. Specifically, our proposed metric, Parsing Covering (PC), is defined as follows:



where S_i and S'_i are the groundtruth segmentation and predicted segmentation for the i -th semantic class respectively, and N_i is the total number of pixels of groundtruth regions from S_i . The Covering for class i , Cov_i , is computed in the same way as the original Covering metric except that only groundtruth regions from S_i and predicted regions from S'_i are considered. PC is then obtained by computing the average of Cov_i over C semantic classes.

A notable difference between PQ and the proposed PC is that there is no matching involved in PC and hence no matching threshold. As an attempt to treat equally "thing" and "stuff", the segmentation of "stuff" classes still receives partial PC score if the segmentation is only partially correct. For example, if one out of three equally-sized trees is perfectly segmented, the model will get the same partial score by using PC regardless of considering "tree" as "stuff" or "thing".

Tutorial

To evaluate the parsing results with PQ and PC, we provide two options:

1. Python off-line evaluation with results saved in the [COCO format](#).
2. TensorFlow on-line evaluation.

Below, we explain each option in detail.

1. Python off-line evaluation with results saved in COCO format

[COCO result format](#) has been adopted by several benchmarks [3, 4]. Therefore, we provide a convenient function, `eval_coco_format`, to evaluate the results saved in COCO format in terms of PC and re-implemented PQ.

Before using the provided function, the users need to download the official COCO panoptic segmentation task API. Please see [installation](#) for reference.

Once the official COCO panoptic segmentation task API is downloaded, the users should be able to run the `eval_coco_format.py` to evaluate the parsing results in terms of both PC and reimplemented PQ.

To be concrete, let's take a look at the function, `eval_coco_format` in `eval_coco_format.py`:

```
eval_coco_format(gt_json_file,
                 pred_json_file,
                 gt_folder=None,
                 pred_folder=None,
                 metric='pq',
                 num_categories=201,
                 ignored_label=0,
                 max_instances_per_category=256,
                 intersection_offset=None,
                 normalize_by_image_size=True,
                 num_workers=0,
                 print_digits=3):
```

where

1. `gt_json_file`: Path to a JSON file giving ground-truth annotations in COCO format.
2. `pred_json_file`: Path to a JSON file for the predictions to evaluate.
3. `gt_folder`: Folder containing panoptic-format ID images to match ground-truth annotations to image regions.
4. `pred_folder`: Path to a folder containing ID images for predictions.
5. `metric`: Name of a metric to compute. Set to `pc`, `pq` for evaluation in PC or PQ, respectively.
6. `num_categories`: The number of segmentation categories (or "classes") in the dataset.
7. `ignored_label`: A category id that is ignored in evaluation, e.g. the "void" label in COCO panoptic segmentation dataset.

8. `max_instances_per_category` : The maximum number of instances for each category to ensure unique instance labels.
9. `intersection_offset` : The maximum number of unique labels.
10. `normalize_by_image_size` : Whether to normalize groundtruth instance region areas by image size when using PC.
11. `num_workers` : If set to a positive number, will spawn child processes to compute parts of the metric in parallel by splitting the images between the workers. If set to -1, will use the value of `multiprocessing.cpu_count()`.
12. `print_digits` : Number of significant digits to print in summary of computed metrics.

The input arguments have default values set for the COCO panoptic segmentation dataset. Thus, users only need to provide the `gt_json_file` and the `pred_json_file` (following the COCO format) to run the evaluation on COCO with PQ. If users want to evaluate the results on other datasets, they may need to change the default values.

As an example, the interested users could take a look at the provided unit test,

```
test_compare_pq_with_reference_eval, in eval_coco_format_test.py .
```

2. TensorFlow on-line evaluation

Users may also want to run the TensorFlow on-line evaluation, similar to the [tf.contrib.metrics.streaming_mean_iou](https://www.tensorflow.org/api_guides/python/contrib.metrics_streaming_mean_iou).

Below, we provide a code snippet that shows how to use the provided `streaming_panoptic_quality` and `streaming_parsing_covering`.

```
metric_map = {}
metric_map['panoptic_quality'] = streaming_metrics.streaming_panoptic_quality(
    category_label,
    instance_label,
    category_prediction,
    instance_prediction,
    num_classes=201,
    max_instances_per_category=256,
    ignored_label=0,
    offset=256*256)
metric_map['parsing_covering'] = streaming_metrics.streaming_parsing_covering(
    category_label,
    instance_label,
    category_prediction,
    instance_prediction,
    num_classes=201,
    max_instances_per_category=256,
    ignored_label=0,
    offset=256*256,
    normalize_by_image_size=True)
metrics_to_values, metrics_to_updates = slim.metrics.aggregate_metric_map(
    metric_map)
```

where `metric_map` is a dictionary storing the streamed results of PQ and PC.

The `category_label` and the `instance_label` are the semantic segmentation and instance segmentation groundtruth, respectively. That is, in the panoptic segmentation format: `panoptic_label = category_label *`

`max_instances_per_category + instance_label`. Similarly, the `category_prediction` and the `instance_prediction` are the predicted semantic segmentation and instance segmentation, respectively.

Below, we provide a code snippet about how to summarize the results in the context of `tf.summary`.

```
summary_ops = []
for metric_name, metric_value in metrics_to_values.iteritems():
    if metric_name == 'panoptic_quality':
        [pq, sq, rq, total_tp, total_fn, total_fp] = tf.unstack(
            metric_value, 6, axis=0)
        panoptic_metrics = {
            # Panoptic quality.
            'pq': pq,
            # Segmentation quality.
            'sq': sq,
            # Recognition quality.
            'rq': rq,
            # Total true positives.
            'total_tp': total_tp,
            # Total false negatives.
            'total_fn': total_fn,
            # Total false positives.
            'total_fp': total_fp,
        }
        # Find the valid classes that will be used for evaluation. We will
        # ignore the `ignore_label` class and other classes which have (tp + fn
        # + fp) equal to 0.
        valid_classes = tf.logical_and(
            tf.not_equal(tf.range(0, num_classes), void_label),
            tf.not_equal(total_tp + total_fn + total_fp, 0))
        for target_metric, target_value in panoptic_metrics.iteritems():
            output_metric_name = '{}_{}'.format(metric_name, target_metric)
            op = tf.summary.scalar(
                output_metric_name,
                tf.reduce_mean(tf.boolean_mask(target_value, valid_classes)))
            op = tf.Print(op, [target_value], output_metric_name + '_classwise: ',
                           summarize=num_classes)
            op = tf.Print(
                op,
                [tf.reduce_mean(tf.boolean_mask(target_value, valid_classes))],
                output_metric_name + '_mean: ',
                summarize=1)
            summary_ops.append(op)
    elif metric_name == 'parsing_covering':
        [per_class_covering,
         total_per_class_weighted_iious,
         total_per_class_gt_areas] = tf.unstack(metric_value, 3, axis=0)
        # Find the valid classes that will be used for evaluation. We will
        # ignore the `void_label` class and other classes which have
        # total_per_class_weighted_iious + total_per_class_gt_areas equal to 0.
        valid_classes = tf.logical_and(
            tf.not_equal(tf.range(0, num_classes), void_label),
```

```

        tf.not_equal(
            total_per_class_weighted_iou + total_per_class_gt_areas, 0))
op = tf.summary.scalar(
    metric_name,
    tf.reduce_mean(tf.boolean_mask(per_class_covering, valid_classes)))
op = tf.Print(op, [per_class_covering], metric_name + '_classwise: ',
              summarize=num_classes)
op = tf.Print(
    op,
    [tf.reduce_mean(
        tf.boolean_mask(per_class_covering, valid_classes))],
    metric_name + '_mean: ',
    summarize=1)
summary_ops.append(op)
else:
    raise ValueError('The metric_name "%s" is not supported.' % metric_name)

```

Afterwards, the users could use the following code to run the evaluation in TensorFlow.

Users can take a look at `eval.py` for reference which provides a simple example to run the streaming evaluation of mIOU for semantic segmentation.

```

metric_values = slim.evaluation.evaluation_loop(
    master=FLAGS.master,
    checkpoint_dir=FLAGS.checkpoint_dir,
    logdir=FLAGS.eval_logdir,
    num_evals=num_batches,
    eval_op=metrics_to_updates.values(),
    final_op=metrics_to_values.values(),
    summary_op=tf.summary.merge(summary_ops),
    max_number_of_evaluations=FLAGS.max_number_of_evaluations,
    eval_interval_secs=FLAGS.eval_interval_secs)

```

References

1. Image Parsing: Unifying Segmentation, Detection, and Recognition

Zhuowen Tu, Xiangrong Chen, Alan L. Yuille, and Song-Chun Zhu
IJCV, 2005.

2. Panoptic Segmentation

Alexander Kirillov, Kaiming He, Ross Girshick, Carsten Rother and Piotr Dollár
arXiv:1801.00868, 2018.

3. Microsoft COCO: Common Objects in Context

Tsung-Yi Lin, Michael Maire, Serge Belongie, Lubomir Bourdev, Ross Girshick, James Hays, Pietro Perona, Deva Ramanan, C. Lawrence Zitnick, Piotr Dollar
In the Proc. of ECCV, 2014.

4. The Mapillary Vistas Dataset for Semantic Understanding of Street Scenes

Gerhard Neuhold, Tobias Ollmann, Samuel Rota Bulò, and Peter Kotschieder
In the Proc. of ICCV, 2017.

5. DeeperLab: Single-Shot Image Parser

Tien-Ju Yang, Maxwell D. Collins, Yukun Zhu, Jyh-Jing Hwang, Ting Liu, Xiao Zhang, Vivienne Sze, George Papandreou, Liang-Chieh Chen
arXiv: 1902.05093, 2019.

6. Contour Detection and Hierarchical Image Segmentation

Pablo Arbelaez, Michael Maire, Charles Fowlkes, and Jitendra Malik
PAMI, 2011