

Example scripts for User-space, Statically Defined Tracing (USDT)

This directory contains scripts showcasing User-space, Statically Defined Tracing (USDT) support for Bitcoin Core on Linux using. For more information on USDT support in Bitcoin Core see the USDT documentation.

Examples for the two main eBPF front-ends, bpftool and BPF Compiler Collection (BCC), with support for USDT, are listed. BCC is used for complex tools and daemons and **bpftool** is preferred for one-liners and shorter scripts.

To develop and run bpftool and BCC scripts you need to install the corresponding packages. See installing bpftool and installing BCC for more information. For development there exist a bpftool Reference Guide, a BCC Reference Guide, and a bcc Python Developer Tutorial.

Examples

The bpftool examples contain a relative path to the **bitcoind** binary. By default, the scripts should be run from the repository-root and assume a self-compiled **bitcoind** binary. The paths in the examples can be changed, for example, to point to release builds if needed. See the Bitcoin Core USDT documentation on how to list available tracepoints in your **bitcoind** binary.

WARNING: eBPF programs require root privileges to be loaded into a Linux kernel VM. This means the bpftool and BCC examples must be executed with root privileges. Make sure to carefully review any scripts that you run with root privileges first!

log_p2p_traffic.bt

A bpftool script logging information about inbound and outbound P2P network messages. Based on the **net:inbound_message** and **net:outbound_message** tracepoints.

By default, **bpftool** limits strings to 64 bytes due to the limited stack size in the eBPF VM. For example, Tor v3 addresses exceed the string size limit which results in the port being cut off during logging. The string size limit can be increased with the **BPFTRACE_STRLLEN** environment variable (**BPFTRACE_STRLLEN=70** works fine).

```
$ bpftool contrib/tracing/log_p2p_traffic.bt
```

Output

```
outbound 'ping' msg to peer 11 (outbound-full-relay, [2a02:b10c:f747:1:ef:fake:ipv6:addr]:80)
inbound 'pong' msg from peer 11 (outbound-full-relay, [2a02:b10c:f747:1:ef:fake:ipv6:addr]:80)
inbound 'inv' msg from peer 16 (outbound-full-relay, XX.XX.XXX.121:8333) with 37 bytes
outbound 'getdata' msg to peer 16 (outbound-full-relay, XX.XX.XXX.121:8333) with 37 bytes
```

```

inbound 'tx' msg from peer 16 (outbound-full-relay, XX.XX.XXX.121:8333) with 222 bytes
outbound 'inv' msg to peer 9 (outbound-full-relay, faketorv3addressa2ufa6odvoi3s77j4uegey0x)
outbound 'inv' msg to peer 7 (outbound-full-relay, XX.XX.XXX.242:8333) with 37 bytes
...

```

p2p_monitor.py

A BCC Python script using curses for an interactive P2P message monitor.
Based on the `net:inbound_message` and `net:outbound_message` tracepoints.

Inbound and outbound traffic is listed for each peer together with information about the connection. Peers can be selected individually to view recent P2P messages.

```
$ python3 contrib/tracing/p2p_monitor.py ./src/bitcoind
```

Lists selectable peers and traffic and connection information.

P2P Message Monitor

Navigate with UP/DOWN or J/K and select a peer with ENTER or SPACE to see individual P2P messages.

PEER	OUTBOUND	INBOUND	TYPE	ADDR		
0	46	398 byte	61	1407590 byte	block-relay-only	XX.XX.XXX.196:8333
11	1156	253570 byte	3431	2394924 byte	outbound-full-relay	XXX.X.XX.179:8333
13	3425	1809620 byte	1236	305458 byte	inbound	XXX.X.X.X:60380
16	1046	241633 byte	1589	1199220 byte	outbound-full-relay	4faketorv2pbfu7x.onion:8333
19	577	181679 byte	390	148951 byte	outbound-full-relay	kfake4vctorjv2o2.onion:8333
20	11	1248 byte	13	1283 byte	block-relay-only	[2600:fake:64d9:b::1]:8333
21	11	1248 byte	13	1299 byte	block-relay-only	XX.XXX.X.155:8333
22	5	103 byte	1	102 byte	feeler	XX.XX.XXX.173:8333
23	11	1248 byte	12	1255 byte	block-relay-only	XX.XXX.XXX.220:8333
24	3	103 byte	1	102 byte	feeler	XXX.XXX.XXX.64:8333

...

Showing recent P2P messages between our node and a selected peer.

```

-----
|                                PEER 16 (4faketorv2pbfu7x.onion:8333)                                |
| OUR NODE                                outbound-full-relay                                PEER |
|                                <--- sendcmpct (9 bytes)                                |
| inv (37 byte) --->                                |
|                                <--- ping (8 bytes)                                |
| pong (8 byte) --->                                |
| inv (37 byte) --->                                |
|                                <--- addr (31 bytes)                                |
| inv (37 byte) --->                                |
|                                <--- getheaders (1029 bytes)                        |
| headers (1 byte) --->                                |
|                                <--- feefilter (8 bytes)                        |

```

```

|                                     <--- pong (8 bytes) |
|                                     <--- headers (82 bytes) |
|                                     <--- addr (30003 bytes) |
| inv (1261 byte) --->                                     |
|                                     ...                       |

```

log_raw_p2p_msgs.py

A BCC Python script showcasing eBPF and USDT limitations when passing data larger than about 32kb. Based on the `net:inbound_message` and `net:outbound_message` tracepoints.

Bitcoin P2P messages can be larger than 32kb (e.g. `tx`, `block`, ...). The eBPF VM's stack is limited to 512 bytes, and we can't allocate more than about 32kb for a P2P message in the eBPF VM. The **message data is cut off** when the message is larger than `MAX_MSG_DATA_LENGTH` (see script). This can be detected in user-space by comparing the data length to the message length variable. The message is cut off when the data length is smaller than the message length. A warning is included with the printed message data.

Data is submitted to user-space (i.e. to this script) via a ring buffer. The throughput of the ring buffer is limited. Each `p2p_message` is about 32kb in size. In- or outbound messages submitted to the ring buffer in rapid succession fill the ring buffer faster than it can be read. Some messages are lost. BCC prints: `Possibly lost 2 samples` on lost messages.

```
$ python3 contrib/tracing/log_raw_p2p_msgs.py ./src/bitcoind
```

Logging raw P2P messages.

Messages larger than about 32kb will be cut off!

Some messages might be lost!

```
  outbound msg 'inv' from peer 4 (outbound-full-relay, XX.XXX.XX.4:8333) with 253 bytes: 070
```

```
...
```

```
Warning: incomplete message (only 32568 out of 53552 bytes)! inbound msg 'tx' from peer 32
```

```
...
```

```
Possibly lost 2 samples
```

connectblock_benchmark.bt

A `bpftrace` script to benchmark the `ConnectBlock()` function during, for example, a blockchain re-index. Based on the `validation:block_connected` USDT tracepoint.

The script takes three positional arguments. The first two arguments, the start, and end height indicate between which blocks the benchmark should be run. The third acts as a duration threshold in milliseconds. When the `ConnectBlock()` function takes longer than the threshold, information about the block, is printed. For more details, see the header comment in the script.

The following command can be used to benchmark, for example, `ConnectBlock()` between height 20000 and 38000 on SigNet while logging all blocks that take longer than 25ms to connect.

```
$ bpftrace contrib/tracing/connectblock_benchmark.bt 20000 38000 25
```

In a different terminal, starting Bitcoin Core in SigNet mode and with re-indexing enabled.

```
$ ./src/bitcoind -signet -reindex
```

This produces the following output.

Attaching 5 probes...

ConnectBlock Benchmark between height 20000 and 38000 inclusive

Logging blocks taking longer than 25 ms to connect.

Starting Connect Block Benchmark between height 20000 and 38000.

BENCH	39 blk/s	59 tx/s	59 inputs/s	20 sigops/s (height 20038)		
Block	20492	(000000f555653bb05e2f3c6e79925e01a20dd57033f4dc7c354b46e34735d32b)			20 tx	23
BENCH	1840 blk/s	2117 tx/s	4478 inputs/s	2471 sigops/s (height 21879)		
BENCH	1816 blk/s	4972 tx/s	4982 inputs/s	125 sigops/s (height 23695)		
BENCH	2095 blk/s	2890 tx/s	2910 inputs/s	152 sigops/s (height 25790)		
BENCH	1684 blk/s	3979 tx/s	4053 inputs/s	288 sigops/s (height 27474)		
BENCH	1155 blk/s	3216 tx/s	3252 inputs/s	115 sigops/s (height 28629)		
BENCH	1797 blk/s	2488 tx/s	2503 inputs/s	111 sigops/s (height 30426)		
BENCH	1849 blk/s	6318 tx/s	6569 inputs/s	12189 sigops/s (height 32275)		
BENCH	946 blk/s	20209 tx/s	20775 inputs/s	83809 sigops/s (height 33221)		
Block	33406	(0000002adfe4a15cfcd53bd890a89bbae836e5bb7f38bac566f61ad4548c87f6)			25 tx	20
Block	33687	(00000073231307a9828e5607ceb8156b402efe56747271a4442e75eb5b77cd36)			52 tx	17
BENCH	582 blk/s	21581 tx/s	27673 inputs/s	60345 sigops/s (height 33803)		
BENCH	1035 blk/s	19735 tx/s	19776 inputs/s	51355 sigops/s (height 34838)		
Block	35625	(0000006b00b347390c4768ea9df2655e9ff4b120f29d78594a2a702f8a02c997)			20 tx	33
BENCH	887 blk/s	17857 tx/s	22191 inputs/s	24404 sigops/s (height 35725)		
Block	35937	(000000d816d13d6e39b471cd4368db60463a764ba1f29168606b04a22b81ea57)			75 tx	39
BENCH	823 blk/s	16298 tx/s	21031 inputs/s	18440 sigops/s (height 36548)		
Block	36583	(000000c3e260556dbf42968aae3f904dba8b8c1ff96a6f6e3aa5365d2e3ad317)			24 tx	27
Block	36700	(000000b3b173de9e65a3cfa738d976af6347aaf83fa17ab3f2a4d2ede3ddfac4)			73 tx	16
Block	36832	(0000007859578c02c1ac37dabd1b9ec19b98f350b56935f5dd3a41e9f79f836e)			34 tx	14
BENCH	613 blk/s	16718 tx/s	25074 inputs/s	23022 sigops/s (height 37161)		
Block	37870	(000000f5c1086291ba2d943fb0c3bc82e71c5ee341ee117681d1456fbf6c6c38)			25 tx	18
BENCH	811 blk/s	16031 tx/s	20921 inputs/s	18696 sigops/s (height 37972)		

Took 14055 ms to connect the blocks between height 20000 and 38000.

Histogram of block connection times in milliseconds (ms).

@durations:

```
[0]          16838 |@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@|
[1]           882  |@@
```

[2, 4)	236		
[4, 8)	23		
[8, 16)	9		
[16, 32)	9		
[32, 64)	4		

log_utxocache_flush.py

A BCC Python script to log the UTXO cache flushes. Based on the `utxocache:flush` tracepoint.

```
$ python3 contrib/tracing/log_utxocache_flush.py ./src/bitcoind
```

Logging utxocache flushes. Ctrl-C to end...

Duration (µs)	Mode	Coins Count	Memory Usage	Prune
730451	IF_NEEDED	22990	3323.54 kB	True
637657	ALWAYS	122320	17124.80 kB	False
81349	ALWAYS	0	1383.49 kB	False

log_utxos.bt

A `bpfftrace` script to log information about the coins that are added, spent, or uncached from the UTXO set. Based on the `utxocache:add`, `utxocache:spend` and `utxocache:uncache` tracepoints.

```
$ bpfftrace contrib/tracing/log_utxos.bt
```

This should produce an output similar to the following. If you see `bpfftrace` warnings like `Lost 24 events`, the eBPF perf ring-buffer is filled faster than it is being read. You can increase the ring-buffer size by setting the ENV variable `BPFTRACE_PERF_RB_PAGES` (default 64) at a cost of higher memory usage. See the `bpfftrace` reference guide for more information.

Attaching 4 probes...

OP	Outpoint	
Added	6ba9ad857e1ef2eb2a2c94f06813c414c7ab273e3d6bd7ad64e000315a887e7c:1	1
Spent	fa7dc4db56637a151f6649d8f26732956d1c5424c82aae400a83d02b2cc2c87b:0	1822
Added	eeb2f099b1af6a2a12e6ddd2eeb16fc5968582241d7f08ba202d28b60ac264c7:0	1
Added	eeb2f099b1af6a2a12e6ddd2eeb16fc5968582241d7f08ba202d28b60ac264c7:1	1822
Added	a0c7f4ec9cccef2d89672a624a4e6c8237a17572efdd4679eea9e9ee70d2db04:0	1007
Spent	25e0df5cc1aeb1b78e6056bf403e5e8b7e41f138060ca0a50a50134df0549a5e:2	
Spent	42f383c04e09c26a2378272ec33aa0c1bf4883ca5ab739e8b7e06be5a5787d61:1	38
Added	f85e3b4b89270863a389395cc9a4123e417ab19384cef96533c6649abd6b0561:0	378
Added	f85e3b4b89270863a389395cc9a4123e417ab19384cef96533c6649abd6b0561:2	
Spent	a05880b8c77971ed0b9f73062c7c4cdb0ff3856ab14cbf8bc481ed571cd34b83:1	55912
Added	eb689865f7d957938978d6207918748f74e6aa074f47874724327089445b0960:0	55896
Added	eb689865f7d957938978d6207918748f74e6aa074f47874724327089445b0960:1	15