

Go 1.11 release introduces AVX-512 support.

This page describes how to use new features as well as some important encoder details.

Terminology

Most terminology comes from Intel Software Developer’s manual.

Suffixes originate from Go assembler syntax, which is close to AT&T, which also uses size suffixes.

Some terms are listed to avoid ambiguity (for example, opcode can have different meanings).

Term

Description

Operand

Same as “instruction argument”.

Opcode

Name that refers to instruction group. For example, VADDPD is an opcode. It refers to both VEX and EVEX encoded forms and all operand combinations. Most Go assembler opcodes for AVX-512 match Intel manual entries, with exceptions for cases where additional size suffix is used (e.g. VCVTTPD2DQY is VCVTTPD2DQ).

Opcode suffix

Suffix that overrides some opcode properties. Listed after “.” (dot). For example, VADDPD.Z has “Z” opcode suffix. There can be multiple dot-separated opcode suffixes.

Size suffix

Suffix that specifies instruction operand size if it can’t be inferred from operands alone. For example, VCVTSS2USIL has “L” size suffix.

Opmask

Used for both {k1} notation and to describe instructions that have K registers operands. Related to masking support in EVEX prefix.

Register block

Multi-source operand that encodes register range. Intel manual uses +n notation for register blocks. For example, +3 is a register block of 4 registers.

FP

Floating-point

New registers

EVEX-enabled instructions can access additional 16 X (128-bit xmm) and Y (256-bit ymm) registers, plus 32 new Z (512-bit zmm) registers in 64-bit mode. 32-bit mode only gets Z0–Z7.

New opmask registers are named K0–K7.

They can be used for both masking and for special opmask instructions (like KADDB).

Masking support

Instructions that support masking can omit K register operand.

In this case, K0 register is implied (“all ones”) and merging-masking is performed. This is effectively “no masking”.

K1–K7 registers can be used to override default opmask.

K register should be placed right before destination operand.

Zeroing-masking can be activated with Z opcode suffix.

For example, `VADDPD.Z (AX), Z30, K3, Z10` uses zeroing-masking and explicit K register.

- If Z opcode suffix is removed, it's merging-masking. - If K3 operand is removed, K0 operand is implied.

It's compile-time error to use K0 register for {k1} operands (consult manuals for details).

EVEX broadcast/rounding/SAE support

Embedded broadcast, rounding and SAE activated through opcode suffixes.

For reg-reg FP instructions with {er} enabled, rounding opcode suffix can be specified:

- RU_SAE to round towards +Inf
- RD_SAE to round towards -Inf
- RZ_SAE to round towards zero
- RN_SAE to round towards nearest

To read more about rounding modes, see MXCSR.RC info.

For reg-reg FP instructions with {sae} enabled, exception suppression can be specified with SAE opcode suffix.

For reg-mem instructions with `m32bcst/m64bcst` operand, broadcasting can be turned on with BCST opcode suffix.

Zeroing opcode suffix can be combined with any of these.

For example, `VMAXPD.SAE.Z Z3, Z2, Z1` uses both Z and SAE opcode suffixes.

It is important to put zeroing opcode suffix last, otherwise it is a compilation error.

Register block (multi-source) operands

Register blocks are specified using register range syntax.

It would be enough to specify just first (low) register, but Go assembler requires explicit range with both ends for readability reasons.

For example, instructions with +3 range can be used like `VP4DPWSSD Z25, [Z0-Z3], (AX)`.

Range `[Z0-Z3]` reads like “register block of Z0, Z1, Z2, Z3”.

Invalid ranges result in compilation error.

AVX1 and AVX2 instructions with EVEX prefix

Previously existed opcodes that can be encoded using EVEX prefix now can access AVX-512 features like wider register file, zeroing/merging masking, etc. For example, `VADDPD` can now use 512-bit vector registers.

See encoder details for more info.

Supported extensions

Best way to get up-to-date list of supported extensions is to do `ls -1` inside test suite directory.

Latest list includes:

```
aes_avx512f
avx512_4fmaps
avx512_4vnniw
avx512_bitalg
avx512_ifma
avx512_vbmi
avx512_vbmi2
avx512_vnni
avx512_vpopcntdq
avx512bw
avx512cd
avx512dq
avx512er
avx512f
avx512pf
gfni_avx512f
vpclmulqdq_avx512f
```

128-bit and 256-bit instructions additionally require `avx512vl`.

That is, if `VADDPD` is available in `avx512f`, you can't use `X` and `Y` arguments without `avx512vl`.

Filenames follow GNU `as` (gas) conventions.

`avx512extmap.csv` can make naming scheme more apparent.

Instructions with size suffix

Some opcodes do not match Intel manual entries.

This section is provided for search convenience.

Intel opcode	Go assembler opcodes
<code>VCVTPD2DQ</code>	<code>VCVTPD2DQX</code> , <code>VCVTPD2DQY</code>
<code>VCVTPD2PS</code>	<code>VCVTPD2PSX</code> , <code>VCVTPD2PSY</code>
<code>VCVTTPD2DQ</code>	<code>VCVTTPD2DQX</code> , <code>VCVTTPD2DQY</code>
<code>VCVTQQ2PS</code>	<code>VCVTQQ2PSX</code> , <code>VCVTQQ2PSY</code>
<code>VCVTUQQ2PS</code>	<code>VCVTUQQ2PSX</code> , <code>VCVTUQQ2PSY</code>
<code>VCVTPD2UDQ</code>	<code>VCVTPD2UDQX</code> , <code>VCVTPD2UDQY</code>
<code>VCVTTPD2UDQ</code>	<code>VCVTTPD2UDQX</code> , <code>VCVTTPD2UDQY</code>
<code>VFPCLASSPD</code>	<code>VFPCLASSPDX</code> , <code>VFPCLASSPDY</code> , <code>VFPCLASSPDZ</code>
<code>VFPCLASSPS</code>	<code>VFPCLASSPSX</code> , <code>VFPCLASSPSY</code> , <code>VFPCLASSPSZ</code>
<code>VCVTSD2SI</code>	<code>VCVTSD2SI</code> , <code>VCVTSD2SIQ</code>
<code>VCVTSS2SI</code>	<code>VCVTSD2SI</code> , <code>VCVTSD2SIQ</code>
<code>VCVTSS2SI</code>	<code>VCVTSD2SI</code> , <code>VCVTSD2SIQ</code>
<code>VCVTSS2SI</code>	<code>VCVTSD2SI</code> , <code>VCVTSD2SIQ</code>
<code>VCVTSD2USI</code>	<code>VCVTSD2USIL</code> , <code>VCVTSD2USIQ</code>
<code>VCVTSS2USI</code>	<code>VCVTSS2USIL</code> , <code>VCVTSS2USIQ</code>
<code>VCVTSS2USI</code>	<code>VCVTSS2USIL</code> , <code>VCVTSS2USIQ</code>
<code>VCVTSS2USI</code>	<code>VCVTSS2USIL</code> , <code>VCVTSS2USIQ</code>
<code>VCVTSS2USI</code>	<code>VCVTSS2USIL</code> , <code>VCVTSS2USIQ</code>
<code>VCVTUSI2SD</code>	<code>VCVTUSI2SDL</code> , <code>VCVTUSI2SDQ</code>
<code>VCVTUSI2SS</code>	<code>VCVTUSI2SSL</code> , <code>VCVTUSI2SSQ</code>
<code>VCVTSI2SD</code>	<code>VCVTSI2SDL</code> , <code>VCVTSI2SDQ</code>
<code>VCVTSI2SS</code>	<code>VCVTSI2SSL</code> , <code>VCVTSI2SSQ</code>
<code>ANDN</code>	<code>ANDNL</code> , <code>ANDNQ</code>
<code>BEXTR</code>	<code>BEXTRL</code> , <code>BEXTRQ</code>
<code>BLSI</code>	<code>BLSIL</code> , <code>BLSIQ</code>
<code>BLSMSK</code>	<code>BLSMSKL</code> , <code>BLSMSKQ</code>
<code>BLSR</code>	<code>BLSRL</code> , <code>BLSRQ</code>
<code>BZHI</code>	<code>BZHIL</code> , <code>BZHIQ</code>
<code>MULX</code>	<code>MULXL</code> , <code>MULXQ</code>
<code>PDEP</code>	<code>PDEPL</code> , <code>PDEPQ</code>
<code>PEXT</code>	<code>PEXTL</code> , <code>PEXTQ</code>
<code>RORX</code>	<code>RORXL</code> , <code>RORXQ</code>
<code>SARX</code>	<code>SARXL</code> , <code>SARXQ</code>
<code>SHLX</code>	<code>SHLXL</code> , <code>SHLXQ</code>

Intel opcode	Go assembler opcodes
SHRX	SHRXL, SHRXQ

Encoder details

Bitwise comparison with older encoder may fail for VEX-encoded instructions due to slightly different encoder tables order.

This difference may arise for instructions with both `{reg, reg/mem}` and `{reg/mem, reg}` forms for reg-reg case. One of such instructions is `VMOVUPS`.

This does not affect code behavior, nor makes it bigger/less efficient. New encoding selection scheme is borrowed from Intel XED.

EVEX encoding is used when any of the following is true:

- Instruction uses new registers (High 16 X/Y, Z or K registers)
- Instruction uses EVEX-related opcode suffixes like `BCST`
- Instruction uses operands combination that is only available for AVX-512

In all other cases VEX encoding is used.

This means that VEX is used whenever possible, and EVEX whenever required.

Compressed `disp8` is applied whenever possible for EVEX-encoded instructions. This also covers broadcasting `disp8` which sometimes has different N multiplier.

Experienced readers can inspect `avx_optabs.go` to learn about N multipliers for any instruction.

For example, `VADDPD` has these: * `N=64` for 512-bit form; `N=8` when broadcasting * `N=32` for 256-bit form; `N=8` when broadcasting * `N=16` for 128-bit form; `N=8` when broadcasting

Examples

Exhaustive amount of examples can be found in Go assembler test suite.

Each file provides several examples for every supported instruction form in particular AVX-512 extension.

Every example also includes generated machine code.

Here is adopted “Vectorized Histogram Update Using AVX-512CD” from Intel® Optimization Manual:

```
for i := 0; i < 512; i++ {
    histo[key[i]] += 1
}

top:
    VMOVUPS    0x40(SP)(DX*4), Z4    ;; vmovups    zmm4, [rsp+rdx*4+0x40]
    VPXORD     Z1, Z1, Z1           ;; vpxord     zmm1, zmm1, zmm1
```

```

KMOVW      K1, K2                ;; kmovw      k2, k1
VPCONFLICTD Z4, Z2              ;; vpconflictd zmm2, zmm4
VPGATHERDD (AX)(Z4*4), K2, Z1   ;; vpgatherdd zmm1{k2}, [rax+zmm4*4]
VPTESTMD   histo<>(SB), Z2, K0  ;; vptestmd   k0, zmm2, [rip+0x185c]
KMOVW      K0, CX                ;; kmovw      ecx, k0
VPADD      Z0, Z1, Z3            ;; vpadd      zmm3, zmm1, zmm0
TESTL      CX, CX                ;; test       ecx, ecx
JZ         noConflicts           ;; jz         noConflicts
VMOVUPS    histo<>(SB), Z1       ;; vmovups    zmm1, [rip+0x1884]
VPTESTMD   histo<>(SB), Z2, K0  ;; vptestmd   k0, zmm2, [rip+0x18ba]
VPLZCNTD   Z2, Z5                ;; vplzcntd   zmm5, zmm2
XORB       BX, BX                ;; xor        bl, bl
KMOVW      K0, CX                ;; kmovw      ecx, k0
VPSUBD     Z5, Z1, Z1            ;; vpsubd    zmm1, zmm1, zmm5
VPSUBD     Z5, Z1, Z1            ;; vpsubd    zmm1, zmm1, zmm5

```

resolveConflicts:

```

VPBROADCASTD CX, Z5             ;; vpbroadcastd zmm5, ecx
KMOVW      CX, K2                ;; kmovw      k2, ecx
VPERMD     Z3, Z1, K2, Z3       ;; vpermd     zmm3{k2}, zmm1, zmm3
VPADD      Z0, Z3, K2, Z3       ;; vpadd      zmm3{k2}, zmm3, zmm0
VPTESTMD   Z2, Z5, K2, K0      ;; vptestmd   k0{k2}, zmm5, zmm2
KMOVW      K0, SI                ;; kmovw      esi, k0
ANDL       SI, CX                ;; and        ecx, esi
JZ         noConflicts           ;; jz         noConflicts
ADDB       $1, BX                ;; add        bl, 0x1
CMPB       BX, $16              ;; cmp        bl, 0x10
JB         resolveConflicts     ;; jb         resolveConflicts

```

noConflicts:

```

KMOVW      K1, K2                ;; kmovw      k2, k1
VPSCATTERDD Z3, K2, (AX)(Z4*4)  ;; vpscatterdd [rax+zmm4*4]{k2}, zmm3
ADDL       $16, DX               ;; add        edx, 0x10
CMPL       DX, $1024             ;; cmp        edx, 0x400
JB         top                   ;; jb         top

```