

# Resource Consumer

## Overview

Resource Consumer is a tool which allows to generate cpu/memory utilization in a container. The reason why it was created is testing kubernetes autoscaling. Resource Consumer can help with autoscaling tests for: - cluster size autoscaling, - horizontal autoscaling of pod - changing the size of replication controller, - vertical autoscaling of pod - changing its resource limits.

## Usage

Resource Consumer starts an HTTP server and handle sent requests. It listens on port given as a flag (default 8080). Action of consuming resources is send to the container by a POST http request. Each http request creates new process. Http request handler is in file `resource_consumer_handler.go`

The container consumes specified amount of resources:

- CPU in millicores,
- Memory in megabytes,
- Fake custom metrics.

### Consume CPU http request

- suffix “ConsumeCPU”,
- parameters “millicores” and “durationSec”.

Consumes specified amount of millicores for durationSec seconds. Consume CPU uses “./consume-cpu/consume-cpu” binary (file `consume-cpu/consume_cpu.go`). When CPU consumption is too low this binary uses cpu by calculating  $\text{math.sqrt}(0) \cdot 10^7$  times and if consumption is too high binary sleeps for 10 millisecond. One replica of Resource Consumer cannot consume more than 1 cpu.

### Consume Memory http request

- suffix “ConsumeMem”,
- parameters “megabytes” and “durationSec”.

Consumes specified amount of megabytes for durationSec seconds. Consume Memory uses stress tool (`stress -m 1 -vm-bytes megabytes -vm-hang 0 -t durationSec`). Request leading to consuming more memory than container limit will be ignored.

### Bump value of a fake custom metric

- suffix “BumpMetric”,
- parameters “metric”, “delta” and “durationSec”.

Bumps metric with given name by delta for durationSec seconds. Custom metrics in Prometheus format are exposed on “/metrics” endpoint.

### **CURL example**

```
kubect1 run resource-consumer --image=gcr.io/k8s-staging-e2e-test-images/resource-consumer:1.9
kubect1 get services resource-consumer
```

There are two IPs. The first one is internal, while the second one is the external load-balanced IP. Both serve port 8080. (Use second one)

```
curl --data "millicores=300&durationSec=600" http://<EXTERNAL-IP>:8080/ConsumeCPU
```

300 millicores will be consumed for 600 seconds.

### **Image**

Docker image of Resource Consumer can be found in Google Container Registry as gcr.io/k8s-staging-e2e-test-images/resource-consumer:1.9

### **Use cases**

#### **Cluster size autoscaling**

1. Consume more resources on each node that is specified for autoscaler
2. Observe that cluster size increased

#### **Horizontal autoscaling of pod**

1. Create consuming RC and start consuming appropriate amount of resources
2. Observe that RC has been resized
3. Observe that usage on each replica decreased

#### **Vertical autoscaling of pod**

1. Create consuming pod and start consuming appropriate amount of resources
2. Observed that limits has been increased