

pwalk: parallel implementation of filepath.Walk

This is a wrapper for `filepath.Walk` which may speed it up by calling multiple callback functions (`WalkFunc`) in parallel, utilizing goroutines.

By default, it utilizes `2*runtime.NumCPU()` goroutines for callbacks. This can be changed by using `WalkN` function which has the additional parameter, specifying the number of goroutines (concurrency).

pwalk vs pwalkdir

This package is deprecated in favor of `pwalkdir`, which is faster, but requires at least Go 1.16.

Caveats

Please note the following limitations of this code:

- Unlike `filepath.Walk`, the order of calls is non-deterministic;
- Only primitive error handling is supported:
 - `filepath.SkipDir` is not supported;
 - no errors are ever passed to `WalkFunc`;
 - once any error is returned from any `WalkFunc` instance, no more new calls to `WalkFunc` are made, and the error is returned to the caller of `Walk`;
 - if more than one `walkFunc` instance will return an error, only one of such errors will be propagated and returned by `Walk`, others will be silently discarded.

Documentation

For the official documentation, see <https://pkg.go.dev/github.com/opencontainers/selinux/pkg/pwalk?tab=doc>

Benchmarks

For a `WalkFunc` that consists solely of the return statement, this implementation is about 10% slower than the standard library's `filepath.Walk`.

Otherwise (if a `WalkFunc` is doing something) this is usually faster, except when the `WalkN(..., 1)` is used.