

BPF LLVM Relocations

This document describes LLVM BPF backend relocation types.

Relocation Record

LLVM BPF backend records each relocation with the following 16-byte ELF structure:

```
typedef struct
{
    Elf64_Addr    r_offset; // Offset from the beginning of section.
    Elf64_Xword   r_info;   // Relocation type and symbol index.
} Elf64_Rel;
```

For example, for the following code:

```
int g1 __attribute__((section("sec")));
int g2 __attribute__((section("sec")));
static volatile int l1 __attribute__((section("sec")));
static volatile int l2 __attribute__((section("sec")));
int test() {
    return g1 + g2 + l1 + l2;
}
```

Compiled with `clang -target bpf -O2 -c test.c`, the following is the code with `llvm-objdump -dr test.o`:

```
0:      18 01 00 00 00 00 00 00 00 00 00 00 00 00 00 00 r1 = 0 11
      0000000000000000: R_BPF_64_64 g1
2:      61 11 00 00 00 00 00 00 r1 = *(u32 *) (r1 + 0)
3:      18 02 00 00 00 00 00 00 00 00 00 00 00 00 00 00 r2 = 0 11
      0000000000000018: R_BPF_64_64 g2
5:      61 20 00 00 00 00 00 00 r0 = *(u32 *) (r2 + 0)
6:      0f 10 00 00 00 00 00 00 r0 += r1
7:      18 01 00 00 08 00 00 00 00 00 00 00 00 00 00 00 r1 = 8 11
      0000000000000038: R_BPF_64_64 sec
9:      61 11 00 00 00 00 00 00 r1 = *(u32 *) (r1 + 0)
10:     0f 10 00 00 00 00 00 00 r0 += r1
11:     18 01 00 00 0c 00 00 00 00 00 00 00 00 00 00 00 r1 = 12 11
      0000000000000058: R_BPF_64_64 sec
13:     61 11 00 00 00 00 00 00 r1 = *(u32 *) (r1 + 0)
14:     0f 10 00 00 00 00 00 00 r0 += r1
15:     95 00 00 00 00 00 00 00 exit
```

There are four relations in the above for four `LD_imm64` instructions. The following `llvm-readelf -r test.o` shows the binary values of the four relocations:

```
Relocation section '.rel.text' at offset 0x190 contains 4 entries:
      Offset             Info             Type             Symbol's Value  Symbol's Name
0000000000000000 0000000600000001 R_BPF_64_64      0000000000000000 g1
0000000000000018 0000000700000001 R_BPF_64_64      0000000000000004 g2
0000000000000038 0000000400000001 R_BPF_64_64      0000000000000000 sec
0000000000000058 0000000400000001 R_BPF_64_64      0000000000000000 sec
```

Each relocation is represented by `Offset` (8 bytes) and `Info` (8 bytes). For example, the first relocation corresponds to the first instruction (Offset 0x0) and the corresponding `Info` indicates the relocation type of `R_BPF_64_64` (type 1) and the entry in the symbol table (entry 6). The following is the symbol table with `llvm-readelf -s test.o`:

```
Symbol table '.symtab' contains 8 entries:
Num:  Value             Size Type  Bind  Vis      Ndx Name
 0: 0000000000000000      0 NOTYPE LOCAL DEFAULT UND
 1: 0000000000000000      0 FILE  LOCAL DEFAULT ABS test.c
 2: 0000000000000008      4 OBJECT LOCAL DEFAULT 4 l1
 3: 000000000000000c      4 OBJECT LOCAL DEFAULT 4 l2
 4: 0000000000000000      0 SECTION LOCAL DEFAULT 4 sec
 5: 0000000000000000    128 FUNC  GLOBAL DEFAULT 2 test
 6: 0000000000000000      4 OBJECT GLOBAL DEFAULT 4 g1
 7: 0000000000000004      4 OBJECT GLOBAL DEFAULT 4 g2
```

The 6th entry is global variable `g1` with value 0.

Similarly, the second relocation is at `.text` offset 0x18, instruction 3, for global variable `g2` which has a symbol value 4, the offset from the start of `.data` section.

The third and fourth relocations refers to static variables `l1` and `l2`. From `.rel.text` section above, it is not clear which symbols they really refers to as they both refers to symbol table entry 4, symbol `sec`, which has `STT_SECTION` type and represents a section. So for static variable or function, the section offset is written to the original insn buffer, which is called A (addend). Looking at above

insn 7 and 11, they have section offset 8 and 12. From symbol table, we can find that they correspond to entries 2 and 3 for 11 and 12.

In general, the `A` is 0 for global variables and functions, and is the section offset or some computation result based on section offset for static variables/functions. The non-section-offset case refers to function calls. See below for more details.

Different Relocation Types

Six relocation types are supported. The following is an overview and `S` represents the value of the symbol in the symbol table:

Enum	ELF Reloc Type	Description	BitSize	Offset	Calculation
0	<code>R_BPF_NONE</code>	None			
1	<code>R_BPF_64_64</code>	<code>ld_imm64</code> insn	32	<code>r_offset + 4</code>	$S + A$
2	<code>R_BPF_64_ABS64</code>	normal data	64	<code>r_offset</code>	$S + A$
3	<code>R_BPF_64_ABS32</code>	normal data	32	<code>r_offset</code>	$S + A$
4	<code>R_BPF_64_NODYLD32</code>	<code>.BTF[.ext]</code> data	32	<code>r_offset</code>	$S + A$
10	<code>R_BPF_64_32</code>	<code>call</code> insn	32	<code>r_offset + 4</code>	$(S + A) / 8 - 1$

For example, `R_BPF_64_64` relocation type is used for `ld_imm64` instruction. The actual to-be-relocated data (0 or section offset) is stored at `r_offset + 4` and the read/write data bitsize is 32 (4 bytes). The relocation can be resolved with the symbol value plus implicit addend. Note that the `BitSize` is 32 which means the section offset must be less than or equal to `UINT32_MAX` and this is enforced by LLVM BPF backend.

In another case, `R_BPF_64_ABS64` relocation type is used for normal 64-bit data. The actual to-be-relocated data is stored at `r_offset` and the read/write data bitsize is 64 (8 bytes). The relocation can be resolved with the symbol value plus implicit addend.

Both `R_BPF_64_ABS32` and `R_BPF_64_NODYLD32` types are for 32-bit data. But `R_BPF_64_NODYLD32` specifically refers to relocations in `.BTF` and `.BTF.ext` sections. For cases like `bcc` where `llvmExecutionEngine RuntimeDyld` is involved, `R_BPF_64_NODYLD32` types of relocations should not be resolved to actual function/variable address. Otherwise, `.BTF` and `.BTF.ext` become unusable by `bcc` and kernel.

Type `R_BPF_64_32` is used for `call` instruction. The call target section offset is stored at `r_offset + 4` (32bit) and calculated as $(S + A) / 8 - 1$.

Examples

Types `R_BPF_64_64` and `R_BPF_64_32` are used to resolve `ld_imm64` and `call` instructions. For example:

```
__attribute__((noinline)) __attribute__((section("sec1")))
int gfunc(int a, int b) {
    return a * b;
}
static __attribute__((noinline)) __attribute__((section("sec1")))
int lfunc(int a, int b) {
    return a + b;
}
int global __attribute__((section("sec2")));
int test(int a, int b) {
    return gfunc(a, b) + lfunc(a, b) + global;
}
```

Compiled with `clang -target bpf -O2 -c test.c`, we will have following code with `llvm-objdump -dr test.o`:

Disassembly of section `.text`:

```
0000000000000000 <test>:
 0:      bf 26 00 00 00 00 00 00 r6 = r2
 1:      bf 17 00 00 00 00 00 00 r7 = r1
 2:      85 10 00 00 ff ff ff ff call -1
                                0000000000000010:  R_BPF_64_32  gfunc
 3:      bf 08 00 00 00 00 00 00 r8 = r0
 4:      bf 71 00 00 00 00 00 00 r1 = r7
 5:      bf 62 00 00 00 00 00 00 r2 = r6
 6:      85 10 00 00 02 00 00 00 call 2
                                0000000000000030:  R_BPF_64_32  sec1
 7:      0f 80 00 00 00 00 00 00 r0 += r8
 8:      18 01 00 00 00 00 00 00 00 00 00 00 00 00 00 00 r1 = 0 ll
                                0000000000000040:  R_BPF_64_64  global
10:      61 11 00 00 00 00 00 00 r1 = *(u32 *) (r1 + 0)
11:      0f 10 00 00 00 00 00 00 r0 += r1
12:      95 00 00 00 00 00 00 00 exit
```

Disassembly of section `sec1`:

```
0000000000000000 <gfunc>:
 0:      bf 20 00 00 00 00 00 00 r0 = r2
 1:      2f 10 00 00 00 00 00 00 r0 *= r1
 2:      95 00 00 00 00 00 00 00 exit
```

```

000000000000000018 <lfunc>:
3:      bf 20 00 00 00 00 00 00 r0 = r2
4:      0f 10 00 00 00 00 00 00 r0 += r1
5:      95 00 00 00 00 00 00 00 exit

```

The first relocation corresponds to `gfunc(a, b)` where `gfunc` has a value of 0, so the call instruction offset is $(0 + 0)/8 - 1 = -1$. The second relocation corresponds to `lfunc(a, b)` where `lfunc` has a section offset `0x18`, so the call instruction offset is $(0 + 0x18)/8 - 1 = 2$. The third relocation corresponds to `ld_imm64` of `global`, which has a section offset 0.

The following is an example to show how `R_BPF_64_ABS64` could be generated:

```

int global() { return 0; }
struct t { void *g; } gbl = { global };

```

Compiled with `clang -target bpf -O2 -g -c test.c`, we will see a relocation below in `.data` section with command `llvm-readelf -r test.o`:

```

Relocation section '.rel.data' at offset 0x458 contains 1 entries:
  Offset      Info      Type           Symbol's Value Symbol's Name
0000000000000000 0000000700000002 R_BPF_64_ABS64 0000000000000000 global

```

The relocation says the first 8-byte of `.data` section should be filled with address of `global` variable.

With `llvm-readelf` output, we can see that dwarf sections have a bunch of `R_BPF_64_ABS32` and `R_BPF_64_ABS64` relocations:

```

Relocation section '.rel.debug_info' at offset 0x468 contains 13 entries:
  Offset      Info      Type           Symbol's Value Symbol's Name
0000000000000006 0000000300000003 R_BPF_64_ABS32 0000000000000000 .debug_abbrev
000000000000000c 0000000400000003 R_BPF_64_ABS32 0000000000000000 .debug_str
0000000000000012 0000000400000003 R_BPF_64_ABS32 0000000000000000 .debug_str
0000000000000016 0000000600000003 R_BPF_64_ABS32 0000000000000000 .debug_line
000000000000001a 0000000400000003 R_BPF_64_ABS32 0000000000000000 .debug_str
000000000000001e 0000000200000002 R_BPF_64_ABS64 0000000000000000 .text
000000000000002b 0000000400000003 R_BPF_64_ABS32 0000000000000000 .debug_str
0000000000000037 0000000800000002 R_BPF_64_ABS64 0000000000000000 gbl
0000000000000040 0000000400000003 R_BPF_64_ABS32 0000000000000000 .debug_str
.....

```

The `.BTF/.BTF.ext` sections has `R_BPF_64_NODYLD32` relocations:

```

Relocation section '.rel.BTF' at offset 0x538 contains 1 entries:
  Offset      Info      Type           Symbol's Value Symbol's Name
0000000000000084 0000000800000004 R_BPF_64_NODYLD32 0000000000000000 gbl

```

```

Relocation section '.rel.BTF.ext' at offset 0x548 contains 2 entries:
  Offset      Info      Type           Symbol's Value Symbol's Name
000000000000002c 0000000200000004 R_BPF_64_NODYLD32 0000000000000000 .text
0000000000000040 0000000200000004 R_BPF_64_NODYLD32 0000000000000000 .text

```