

Coda Kernel-Venus Interface

Note

This is one of the technical documents describing a component of Coda -- this document describes the client kernel-Venus interface.

For more information:

<http://www.coda.cs.cmu.edu>

For user level software needed to run Coda:

<ftp://ftp.coda.cs.cmu.edu>

To run Coda you need to get a user level cache manager for the client, named Venus, as well as tools to manipulate ACLs, to log in, etc. The client needs to have the Coda filesystem selected in the kernel configuration.

The server needs a user level server and at present does not depend on kernel support.

The Venus kernel interface

Peter J. Braam

v1.0, Nov 9, 1997

This document describes the communication between Venus and kernel level filesystem code needed for the operation of the Coda file system. This document version is meant to describe the current interface (version 1.0) as well as improvements we envisage.

1. Introduction

A key component in the Coda Distributed File System is the cache manager, Venus.

When processes on a Coda enabled system access files in the Coda filesystem, requests are directed at the filesystem layer in the operating system. The operating system will communicate with Venus to service the request for the process. Venus manages a persistent client cache and makes remote procedure calls to Coda file servers and related servers (such as authentication servers) to service these requests it receives from the operating system. When Venus has serviced a request it replies to the operating system with appropriate return codes, and other data related to the request. Optionally the kernel support for Coda may maintain a minicache of recently processed requests to limit the number of interactions with Venus. Venus possesses the facility to inform the kernel when elements from its minicache are no longer valid.

This document describes precisely this communication between the kernel and Venus. The definitions of so called upcalls and downcalls will be given with the format of the data they handle. We shall also describe the semantic invariants resulting from the calls.

Historically Coda was implemented in a BSD file system in Mach 2.6. The interface between the kernel and Venus is very similar to the BSD VFS interface. Similar functionality is provided, and the format of the parameters and returned data is very similar to the BSD VFS. This leads to an almost natural environment for implementing a kernel-level filesystem driver for Coda in a BSD system. However, other operating systems such as Linux and Windows 95 and NT have virtual filesystem with different interfaces.

To implement Coda on these systems some reverse engineering of the Venus/Kernel protocol is necessary. Also it came to light that other systems could profit significantly from certain small optimizations and modifications to the protocol. To facilitate this work as well as to make future ports easier, communication between Venus and the kernel should be documented in great detail. This is the aim of this document.

2. Servicing Coda filesystem calls

The service of a request for a Coda file system service originates in a process P which accessing a Coda file. It makes a system call which traps to the OS kernel. Examples of such calls trapping to the kernel are `read`, `write`, `open`, `close`, `create`, `mkdir`, `rmdir`, `chmod` in a Unix context. Similar calls exist in the Win32 environment, and are named `CreateFile`.

Generally the operating system handles the request in a virtual filesystem (VFS) layer, which is named I/O Manager in NT and IFS manager in Windows 95. The VFS is responsible for partial processing of the request and for locating the specific filesystem(s) which will service parts of the request. Usually the information in the path assists in locating the correct FS drivers. Sometimes after extensive pre-processing, the VFS starts invoking exported routines in the FS driver. This is the point where the FS specific processing of the request starts, and here the Coda specific kernel code comes into play.

The FS layer for Coda must expose and implement several interfaces. First and foremost the VFS must be able to make all necessary calls to the Coda FS layer, so the Coda FS driver must expose the VFS interface as applicable in the operating system. These differ very significantly among operating systems, but share features such as facilities to read/write and create and remove objects. The Coda FS layer services such VFS requests by invoking one or more well defined services offered by the cache manager Venus. When the replies from Venus have come back to the FS driver, servicing of the VFS call continues and finishes with a reply to the kernel's VFS. Finally the VFS layer returns to the process.

As a result of this design a basic interface exposed by the FS driver must allow Venus to manage message traffic. In particular Venus must be able to retrieve and place messages and to be notified of the arrival of a new message. The notification must be through a mechanism which does not block Venus since Venus must attend to other tasks even when no messages are waiting or being processed.

Interfaces of the Coda FS Driver

Furthermore the FS layer provides for a special path of communication between a user process and Venus, called the `pioctl` interface. The `pioctl` interface is used for Coda specific services, such as requesting detailed information about the persistent cache managed by Venus. Here the involvement of the kernel is minimal. It identifies the calling process and passes the information on to Venus. When Venus replies the response is passed back to the caller in unmodified form.

Finally Venus allows the kernel FS driver to cache the results from certain services. This is done to avoid excessive context switches and results in an efficient system. However, Venus may acquire information, for example from the network which implies that cached information must be flushed or replaced. Venus then makes a downcall to the Coda FS layer to request flushes or updates in the cache. The kernel FS driver handles such requests synchronously.

Among these interfaces the VFS interface and the facility to place, receive and be notified of messages are platform specific. We will not go into the calls exported to the VFS layer but we will state the requirements of the message exchange mechanism.

3. The message layer

At the lowest level the communication between Venus and the FS driver proceeds through messages. The synchronization between processes requesting Coda file service and Venus relies on blocking and waking up processes. The Coda FS driver processes VFS- and `pioctl`-requests on behalf of a process P, creates messages for Venus, awaits replies and finally returns to the caller. The implementation of the exchange of messages is platform specific, but the semantics have (so far) appeared to be generally applicable. Data buffers are created by the FS Driver in kernel memory on behalf of P and copied to user memory in Venus.

The FS Driver while servicing P makes upcalls to Venus. Such an upcall is dispatched to Venus by creating a message structure. The structure contains the identification of P, the message sequence number, the size of the request and a pointer to the data in kernel memory for the request. Since the data buffer is re-used to hold the reply from Venus, there is a field for the size of the reply. A flags field is used in the message to precisely record the status of the message. Additional platform dependent structures involve pointers to determine the position of the message on queues and pointers to synchronization objects. In the upcall routine the message structure is filled in, flags are set to 0, and it is placed on the *pending* queue. The routine calling upcall is responsible for allocating the data buffer; its structure will be described in the next section.

A facility must exist to notify Venus that the message has been created, and implemented using available synchronization objects in the OS. This notification is done in the upcall context of the process P. When the message is on the pending queue, process P cannot proceed in upcall. The (kernel mode) processing of P in the filesystem request routine must be suspended until Venus has replied. Therefore the calling thread in P is blocked in upcall. A pointer in the message structure will locate the synchronization object on which P is sleeping.

Venus detects the notification that a message has arrived, and the FS driver allow Venus to retrieve the message with a `getmsg_from_kernel` call. This action finishes in the kernel by putting the message on the queue of processing messages and setting flags to READ. Venus is passed the contents of the data buffer. The `getmsg_from_kernel` call now returns and Venus processes the request.

At some later point the FS driver receives a message from Venus, namely when Venus calls `sendmsg_to_kernel`. At this moment the Coda FS driver looks at the contents of the message and decides if

- the message is a reply for a suspended thread P. If so it removes the message from the processing queue and marks the message as WRITTEN. Finally, the FS driver unblocks P (still in the kernel mode context of Venus) and the `sendmsg_to_kernel` call returns to Venus. The process P will be scheduled at some point and continues processing its upcall with the data buffer replaced with the reply from Venus.
- The message is a `downcall`. A downcall is a request from Venus to the FS Driver. The FS driver processes the request immediately (usually a cache eviction or replacement) and when it finishes `sendmsg_to_kernel` returns.

Now P awakes and continues processing upcall. There are some subtleties to take account of. First P will determine if it was woken up in upcall by a signal from some other source (for example an attempt to terminate P) or as is normally the case by Venus in its `sendmsg_to_kernel` call. In the normal case, the upcall routine will deallocate the message structure and return. The FS routine can proceed with its processing.

Sleeping and IPC arrangements

In case P is woken up by a signal and not by Venus, it will first look at the flags field. If the message is not yet READ, the process P can handle its signal without notifying Venus. If Venus has READ, and the request should not be processed, P can send Venus a signal message to indicate that it should disregard the previous message. Such signals are put in the queue at the head, and read first by Venus. If the message is already marked as WRITTEN it is too late to stop the processing. The VFS routine will now continue. (-- If a VFS request involves more than one upcall, this can lead to complicated state, an extra field "handle_signals" could be added in the message structure to indicate points of no return have been passed.--)

3.1. Implementation details

The Unix implementation of this mechanism has been through the implementation of a character device associated with Coda. Venus retrieves messages by doing a read on the device, replies are sent with a write and notification is through the select system call on the file descriptor for the device. The process P is kept waiting on an interruptible wait queue object.

In Windows NT and the DPMI Windows 95 implementation a DeviceIoControl call is used. The DeviceIoControl call is designed to copy buffers from user memory to kernel memory with OPCODES. The sendmsg_to_kernel is issued as a synchronous call, while the getmsg_from_kernel call is asynchronous. Windows EventObjects are used for notification of message arrival. The process P is kept waiting on a KernelEvent object in NT and a semaphore in Windows 95.

4. The interface at the call level

This section describes the upcalls a Coda FS driver can make to Venus. Each of these upcalls make use of two structures: inputArgs and outputArgs. In pseudo BNF form the structures take the following form:

```
struct inputArgs {
    u_long opcode;
    u_long unique;          /* Keep multiple outstanding msgs distinct */
    u_short pid;            /* Common to all */
    u_short pgid;           /* Common to all */
    struct CodaCred cred;    /* Common to all */

    <union "in" of call dependent parts of inputArgs>
};

struct outputArgs {
    u_long opcode;
    u_long unique;          /* Keep multiple outstanding msgs distinct */
    u_long result;

    <union "out" of call dependent parts of inputArgs>
};
```

Before going on let us elucidate the role of the various fields. The inputArgs start with the opcode which defines the type of service requested from Venus. There are approximately 30 upcalls at present which we will discuss. The unique field labels the inputArg with a unique number which will identify the message uniquely. A process and process group id are passed. Finally the credentials of the caller are included.

Before delving into the specific calls we need to discuss a variety of data structures shared by the kernel and Venus.

4.1. Data structures shared by the kernel and Venus

The CodaCred structure defines a variety of user and group ids as they are set for the calling process. The void_t and vgid_t are 32 bit unsigned integers. It also defines group membership in an array. On Unix the CodaCred has proven sufficient to implement good security semantics for Coda but the structure may have to undergo modification for the Windows environment when these mature:

```
struct CodaCred {
    void_t cr_uid, cr_euid, cr_suid, cr_fsuid; /* Real, effective, set, fs uid */
    vgid_t cr_gid, cr_egid, cr_sgid, cr_fsgid; /* same for groups */
    vgid_t cr_groups[NGROUPS];                /* Group membership for caller */
};
```

Note

It is questionable if we need CodaCreds in Venus. Finally Venus doesn't know about groups, although it does create files with the default uid/gid. Perhaps the list of group membership is superfluous.

The next item is the fundamental identifier used to identify Coda files, the ViceFid. A fid of a file uniquely defines a file or directory in the Coda filesystem within a cell [1]:

```
typedef struct ViceFid {
```

```

    VolumeId Volume;
    VnodeId Vnode;
    Unique_t Unique;
} ViceFid;

```

- [1] A cell is a group of Coda servers acting under the aegis of a single system control machine or SCM. See the Coda Administration manual for a detailed description of the role of the SCM.

Each of the constituent fields: VolumeId, VnodeId and Unique_t are unsigned 32 bit integers. We envisage that a further field will need to be prefixed to identify the Coda cell; this will probably take the form of a IPv6 size IP address naming the Coda cell through DNS.

The next important structure shared between Venus and the kernel is the attributes of the file. The following structure is used to exchange information. It has room for future extensions such as support for device files (currently not present in Coda):

```

struct coda_timespec {
    int64_t      tv_sec;          /* seconds */
    long         tv_nsec;        /* nanoseconds */
};

struct coda_vattr {
    enum coda_vtype va_type;      /* vnode type (for create) */
    u_short         va_mode;      /* files access mode and type */
    short           va_nlink;     /* number of references to file */
    void_t          va_uid;       /* owner user id */
    void_t          va_gid;       /* owner group id */
    long            va_fsid;      /* file system id (dev for now) */
    long            va_fileid;    /* file id */
    u_quad_t        va_size;      /* file size in bytes */
    long            va_blocksize; /* blocksize preferred for i/o */
    struct coda_timespec va_atime; /* time of last access */
    struct coda_timespec va_mtime; /* time of last modification */
    struct coda_timespec va_ctime; /* time file changed */
    u_long          va_gen;       /* generation number of file */
    u_long          va_flags;     /* flags defined for file */
    dev_t           va_rdev;      /* device special file represents */
    u_quad_t        va_bytes;     /* bytes of disk space held by file */
    u_quad_t        va_filerev;   /* file modification number */
    u_int           va_vaflags;    /* operations flags, see below */
    long            va_spare;     /* remain quad aligned */
};

```

4.2. The pioctl interface

Coda specific requests can be made by application through the pioctl interface. The pioctl is implemented as an ordinary ioctl on a fictitious file /coda/.CONTROL. The pioctl call opens this file, gets a file handle and makes the ioctl call. Finally it closes the file.

The kernel involvement in this is limited to providing the facility to open and close and pass the ioctl message and to verify that a path in the pioctl data buffers is a file in a Coda filesystem.

The kernel is handed a data packet of the form:

```

struct {
    const char *path;
    struct ViceIoctl vidata;
    int follow;
} data;

```

where:

```

struct ViceIoctl {
    caddr_t in, out;          /* Data to be transferred in, or out */
    short in_size;            /* Size of input buffer <= 2K */
    short out_size;           /* Maximum size of output buffer, <= 2K */
};

```

The path must be a Coda file, otherwise the ioctl upcall will not be made.

Note

The data structures and code are a mess. We need to clean this up.

We now proceed to document the individual calls:

4.3. root

Arguments

in

empty

out:

```
struct cfs_root_out {  
    ViceFid VFid;  
} cfs_root;
```

Description

This call is made to Venus during the initialization of the Coda filesystem. If the result is zero, the `cfs_root` structure contains the `ViceFid` of the root of the Coda filesystem. If a non-zero result is generated, its value is a platform dependent error code indicating the difficulty Venus encountered in locating the root of the Coda filesystem.

4.4. lookup

Summary

Find the `ViceFid` and type of an object in a directory if it exists.

Arguments

in:

```
struct cfs_lookup_in {  
    ViceFid    VFid;  
    char       *name;           /* Place holder for data. */  
} cfs_lookup;
```

out:

```
struct cfs_lookup_out {  
    ViceFid VFid;  
    int vtype;  
} cfs_lookup;
```

Description

This call is made to determine the `ViceFid` and filetype of a directory entry. The directory entry requested carries name 'name' and Venus will search the directory identified by `cfs_lookup_in.VFid`. The result may indicate that the name does not exist, or that difficulty was encountered in finding it (e.g. due to disconnection). If the result is zero, the field `cfs_lookup_out.VFid` contains the targets `ViceFid` and `cfs_lookup_out.vtype` the `coda_vtype` giving the type of object the name designates.

The name of the object is an 8 bit character string of maximum length `CFS_MAXNAMLEN`, currently set to 256 (including a 0 terminator.)

It is extremely important to realize that Venus bitwise ors the field `cfs_lookup.vtype` with `CFS_NOCACHE` to indicate that the object should not be put in the kernel name cache.

Note

The type of the `vtype` is currently wrong. It should be `coda_vtype`. Linux does not take note of `CFS_NOCACHE`. It should.

4.5. getattr

Summary Get the attributes of a file.

Arguments

in:

```
struct cfs_getattr_in {  
    ViceFid VFid;  
    struct coda_vattr attr; /* XXXXX */  
} cfs_getattr;
```

out:

```
struct cfs_getattr_out {  
    struct coda_vattr attr;  
} cfs_getattr;
```

Description

This call returns the attributes of the file identified by fid.

Errors

Errors can occur if the object with fid does not exist, is inaccessible or if the caller does not have permission to fetch attributes.

Note

Many kernel FS drivers (Linux, NT and Windows 95) need to acquire the attributes as well as the Fid for the instantiation of an internal "inode" or "FileHandle". A significant improvement in performance on such systems could be made by combining the lookup and getattr calls both at the Venus/kernel interaction level and at the RPC level.

The vattr structure included in the input arguments is superfluous and should be removed.

4.6. setattr

Summary

Set the attributes of a file.

Arguments

in:

```
struct cfs_setattr_in {
    ViceFid VFid;
    struct coda_vattr attr;
} cfs_setattr;
```

out

empty

Description

The structure attr is filled with attributes to be changed in BSD style. Attributes not to be changed are set to -1, apart from vtype which is set to VNON. Other are set to the value to be assigned. The only attributes which the FS driver may request to change are the mode, owner, groupid, atime, mtime and ctime. The return value indicates success or failure.

Errors

A variety of errors can occur. The object may not exist, may be inaccessible, or permission may not be granted by Venus.

4.7. access

Arguments

in:

```
struct cfs_access_in {
    ViceFid VFid;
    int flags;
} cfs_access;
```

out

empty

Description

Verify if access to the object identified by VFid for operations described by flags is permitted. The result indicates if access will be granted. It is important to remember that Coda uses ACLs to enforce protection and that ultimately the servers, not the clients enforce the security of the system. The result of this call will depend on whether a token is held by the user.

Errors

The object may not exist, or the ACL describing the protection may not be accessible.

4.8. create

Summary

Invoked to create a file

Arguments

in:

```
struct cfs_create_in {
    ViceFid VFid;
    struct coda_vattr attr;
    int excl;
    int mode;
    char      *name;          /* Place holder for data. */
} cfs_create;
```

out:

```
struct cfs_create_out {
    ViceFid VFid;
    struct coda_vattr attr;
} cfs_create;
```

Description

This upcall is invoked to request creation of a file. The file will be created in the directory identified by VFid, its name will be name, and the mode will be mode. If excl is set an error will be returned if the file already exists. If the size field in attr is set to zero the file will be truncated. The uid and gid of the file are set by converting the CodaCred to a uid using a macro CRTOUID (this macro is platform dependent). Upon success the VFid and attributes of the file are returned. The Coda FS Driver will normally instantiate a vnode, inode or file handle at kernel level for the new object.

Errors

A variety of errors can occur. Permissions may be insufficient. If the object exists and is not a file the error EISDIR is returned under Unix.

Note

The packing of parameters is very inefficient and appears to indicate confusion between the system call creat and the VFS operation create. The VFS operation create is only called to create new objects. This create call differs from the Unix one in that it is not invoked to return a file descriptor. The truncate and exclusive options, together with the mode, could simply be part of the mode as it is under Unix. There should be no flags argument; this is used in open (2) to return a file descriptor for READ or WRITE mode.

The attributes of the directory should be returned too, since the size and mtime changed.

4.9. mkdir

Summary

Create a new directory.

Arguments

in:

```
struct cfs_mkdir_in {
    ViceFid VFid;
    struct coda_vattr attr;
    char      *name;          /* Place holder for data. */
} cfs_mkdir;
```

out:

```
struct cfs_mkdir_out {
    ViceFid VFid;
    struct coda_vattr attr;
} cfs_mkdir;
```

Description

This call is similar to create but creates a directory. Only the mode field in the input parameters is used for creation. Upon successful creation, the attr returned contains the attributes of the new directory.

Errors

As for create.

Note

The input parameter should be changed to mode instead of attributes.

The attributes of the parent should be returned since the size and mtime changes.

4.10. link**Summary**

Create a link to an existing file.

Arguments

in:

```
struct cfs_link_in {
    ViceFid sourceFid;      /* cnode to link *to* */
    ViceFid destFid;        /* Directory in which to place link */
    char      *tname;       /* Place holder for data. */
} cfs_link;
```

out

empty

Description

This call creates a link to the sourceFid in the directory identified by destFid with name tname. The source must reside in the target's parent, i.e. the source must be have parent destFid, i.e. Coda does not support cross directory hard links. Only the return value is relevant. It indicates success or the type of failure.

Errors

The usual errors can occur.

4.11. symlink**Summary**

create a symbolic link

Arguments

in:

```
struct cfs_symlink_in {
    ViceFid    VFid;        /* Directory to put symlink in */
    char      *srcname;
    struct coda_vattr attr;
    char      *tname;
} cfs_symlink;
```

out

none

Description

Create a symbolic link. The link is to be placed in the directory identified by VFid and named tname. It should point to the pathname srcname. The attributes of the newly created object are to be set to attr.

Note

The attributes of the target directory should be returned since its size changed.

4.12. remove**Summary**

Remove a file

Arguments

in:

```
struct cfs_remove_in {
    ViceFid    VFid;
```



```

        char      *name;          /* Place holder for data. */
    } cfs_remove;

    out

    none

```

Description

Remove file named `cfs_remove_in.name` in directory identified by VFid.

Note

The attributes of the directory should be returned since its mtime and size may change.

4.13. rmdir

Summary

Remove a directory

Arguments

in:

```

    struct cfs_rmdir_in {
        ViceFid VFid;
        char      *name;          /* Place holder for data. */
    } cfs_rmdir;

```

out

none

Description

Remove the directory with name 'name' from the directory identified by VFid.

Note

The attributes of the parent directory should be returned since its mtime and size may change.

4.14. readlink

Summary

Read the value of a symbolic link.

Arguments

in:

```

    struct cfs_readlink_in {
        ViceFid VFid;
    } cfs_readlink;

```

out:

```

    struct cfs_readlink_out {
        int count;
        caddr_t data;          /* Place holder for data. */
    } cfs_readlink;

```

Description

This routine reads the contents of symbolic link identified by VFid into the buffer data. The buffer data must be able to hold any name up to CFS_MAXNAMLEN (PATH or NAM??).

Errors

No unusual errors.

4.15. open

Summary

Open a file.

Arguments

in:

```
struct cfs_open_in {
    ViceFid    VFid;
    int flags;
} cfs_open;
```

out:

```
struct cfs_open_out {
    dev_t      dev;
    ino_t      inode;
} cfs_open;
```

Description

This request asks Venus to place the file identified by VFid in its cache and to note that the calling process wishes to open it with flags as in open(2). The return value to the kernel differs for Unix and Windows systems. For Unix systems the Coda FS Driver is informed of the device and inode number of the container file in the fields dev and inode. For Windows the path of the container file is returned to the kernel.

Note

Currently the cfs_open_out structure is not properly adapted to deal with the Windows case. It might be best to implement two upcalls, one to open aiming at a container file name, the other at a container file inode.

4.16. close

Summary

Close a file, update it on the servers.

Arguments

in:

```
struct cfs_close_in {
    ViceFid    VFid;
    int flags;
} cfs_close;
```

out

none

Description

Close the file identified by VFid.

Note

The flags argument is bogus and not used. However, Venus' code has room to deal with an execp input field, probably this field should be used to inform Venus that the file was closed but is still memory mapped for execution. There are comments about fetching versus not fetching the data in Venus vproc_vfscalls. This seems silly. If a file is being closed, the data in the container file is to be the new data. Here again the execp flag might be in play to create confusion: currently Venus might think a file can be flushed from the cache when it is still memory mapped. This needs to be understood.

4.17. ioctl

Summary

Do an ioctl on a file. This includes the pioctl interface.

Arguments

in:

```
struct cfs_ioctl_in {
    ViceFid VFid;
    int cmd;
    int len;
    int rwflag;
    char *data;
} cfs_ioctl; /* Place holder for data. */
```

out:

```
struct cfs_ioctl_out {
    int len;
    caddr_t data;          /* Place holder for data. */
} cfs_ioctl;
```

Description

Do an ioctl operation on a file. The command, len and data arguments are filled as usual. flags is not used by Venus.

Note

Another bogus parameter. flags is not used. What is the business about PREFETCHING in the Venus code?

4.18. rename

Summary

Rename a fid.

Arguments

in:

```
struct cfs_rename_in {
    ViceFid sourceFid;
    char *srcname;
    ViceFid destFid;
    char *destname;
} cfs_rename;
```

out

none

Description

Rename the object with name srcname in directory sourceFid to destname in destFid. It is important that the names srcname and destname are 0 terminated strings. Strings in Unix kernels are not always null terminated.

4.19. readdir

Summary

Read directory entries.

Arguments

in:

```
struct cfs_readdir_in {
    ViceFid VFid;
    int count;
    int offset;
} cfs_readdir;
```

out:

```
struct cfs_readdir_out {
    int size;
    caddr_t data;          /* Place holder for data. */
} cfs_readdir;
```

Description

Read directory entries from VFid starting at offset and read at most count bytes. Returns the data in data and returns the size in size.

Note

This call is not used. Readdir operations exploit container files. We will re-evaluate this during the directory revamp which is about to take place.

4.20. vget

Summary

instructs Venus to do an FSDB->Get.

Arguments

in:

```
struct cfs_vget_in {  
    ViceFid VFid;  
} cfs_vget;
```

out:

```
struct cfs_vget_out {  
    ViceFid VFid;  
    int vtype;  
} cfs_vget;
```

Description

This upcall asks Venus to do a get operation on an fsobj labelled by VFid.

Note

This operation is not used. However, it is extremely useful since it can be used to deal with read/write memory mapped files. These can be "pinned" in the Venus cache using vget and released with inactive.

4.21. fsync

Summary

Tell Venus to update the RVM attributes of a file.

Arguments

in:

```
struct cfs_fsync_in {  
    ViceFid VFid;  
} cfs_fsync;
```

out

none

Description

Ask Venus to update RVM attributes of object VFid. This should be called as part of kernel level fsync type calls. The result indicates if the syncing was successful.

Note

Linux does not implement this call. It should.

4.22. inactive

Summary

Tell Venus a vnode is no longer in use.

Arguments

in:

```
struct cfs_inactive_in {  
    ViceFid VFid;  
} cfs_inactive;
```

out

none

Description

This operation returns EOPNOTSUPP.

Note

This should perhaps be removed.

4.23. rdwr

Summary

Read or write from a file

Arguments

in:

```
struct cfs_rdwr_in {
    ViceFid VFid;
    int rwflag;
    int count;
    int offset;
    int ioflag;
    caddr_t data;          /* Place holder for data. */
} cfs_rdwr;
```

out:

```
struct cfs_rdwr_out {
    int rwflag;
    int count;
    caddr_t data;          /* Place holder for data. */
} cfs_rdwr;
```

Description

This upcall asks Venus to read or write from a file.

Note

It should be removed since it is against the Coda philosophy that read/write operations never reach Venus. I have been told the operation does not work. It is not currently used.

4.24. ody mount

Summary

Allows mounting multiple Coda "filesystems" on one Unix mount point.

Arguments

in:

```
struct ody_mount_in {
    char *name;            /* Place holder for data. */
} ody_mount;
```

out:

```
struct ody_mount_out {
    ViceFid VFid;
} ody_mount;
```

Description

Asks Venus to return the rootfid of a Coda system named name. The fid is returned in VFid.

Note

This call was used by David for dynamic sets. It should be removed since it causes a jungle of pointers in the VFS mounting area. It is not used by Coda proper. Call is not implemented by Venus.

4.25. ody_lookup

Summary

Looks up something.

Arguments

in

irrelevant

out

irrelevant

Note

Gut it. Call is not implemented by Venus.

4.26. ody_expand

Summary

expands something in a dynamic set.

Arguments

in

irrelevant

out

irrelevant

Note

Gut it. Call is not implemented by Venus.

4.27. prefetch

Summary

Prefetch a dynamic set.

Arguments

in

Not documented.

out

Not documented.

Description

Venus worker.cc has support for this call, although it is noted that it doesn't work. Not surprising, since the kernel does not have support for it. (ODY_PREFETCH is not a defined operation).

Note

Gut it. It isn't working and isn't used by Coda.

4.28. signal

Summary

Send Venus a signal about an upcall.

Arguments

in

none

out

not applicable.

Description

This is an out-of-band upcall to Venus to inform Venus that the calling process received a signal after Venus read the message from the input queue. Venus is supposed to clean up the operation.

Errors

No reply is given.

Note

We need to better understand what Venus needs to clean up and if it is doing this correctly. Also we need to handle multiple upcall per system call situations correctly. It would be important to know what state changes in Venus take place after an upcall for which the kernel is responsible for notifying Venus to clean up (e.g. open definitely is such a state change, but many others are maybe not).

5. The minicache and downcalls

The Coda FS Driver can cache results of lookup and access upcalls, to limit the frequency of upcalls. Upcalls carry a price since a process context switch needs to take place. The counterpart of caching the information is that Venus will notify the FS Driver that cached entries must be flushed or renamed.

The kernel code generally has to maintain a structure which links the internal file handles (called vnodes in BSD, inodes in Linux and FileHandles in Windows) with the ViceFid's which Venus maintains. The reason is that frequent translations back and forth are needed in order to make upcalls and use the results of upcalls. Such linking objects are called cnodes.

The current minicache implementations have cache entries which record the following:

1. the name of the file
2. the cnode of the directory containing the object
3. a list of CodaCred's for which the lookup is permitted.
4. the cnode of the object

The lookup call in the Coda FS Driver may request the cnode of the desired object from the cache, by passing its name, directory and the CodaCred's of the caller. The cache will return the cnode or indicate that it cannot be found. The Coda FS Driver must be careful to invalidate cache entries when it modifies or removes objects.

When Venus obtains information that indicates that cache entries are no longer valid, it will make a downcall to the kernel. Downcalls are intercepted by the Coda FS Driver and lead to cache invalidations of the kind described below. The Coda FS Driver does not return an error unless the downcall data could not be read into kernel memory.

5.1. INVALIDATE

No information is available on this call.

5.2. FLUSH

Arguments

None

Summary

Flush the name cache entirely.

Description

Venus issues this call upon startup and when it dies. This is to prevent stale cache information being held. Some operating systems allow the kernel name cache to be switched off dynamically. When this is done, this downcall is made.

5.3. PURGEUSER

Arguments

```
struct cfs_purgeuser_out { /* CFS_PURGEUSER is a venus->kernel call */
    struct CodaCred cred;
} cfs_purgeuser;
```

Description

Remove all entries in the cache carrying the Cred. This call is issued when tokens for a user expire or are flushed.

5.4. ZAPFILE

Arguments

```
struct cfs_zapfile_out { /* CFS_ZAPFILE is a venus->kernel call */
    ViceFid CodaFid;
} cfs_zapfile;
```

Description

Remove all entries which have the (dir vnode, name) pair. This is issued as a result of an invalidation of cached

attributes of a vnode.

Note

Call is not named correctly in NetBSD and Mach. The minicache zapfile routine takes different arguments. Linux does not implement the invalidation of attributes correctly.

5.5. ZAPDIR

Arguments

```
struct cfs_zapdir_out { /* CFS_ZAPDIR is a venus->kernel call */
    ViceFid CodaFid;
} cfs_zapdir;
```

Description

Remove all entries in the cache lying in a directory CodaFid, and all children of this directory. This call is issued when Venus receives a callback on the directory.

5.6. ZAPVNODE

Arguments

```
struct cfs_zapvnode_out { /* CFS_ZAPVNODE is a venus->kernel call */
    struct CodaCred cred;
    ViceFid VFid;
} cfs_zapvnode;
```

Description

Remove all entries in the cache carrying the cred and VFid as in the arguments. This downcall is probably never issued.

5.7. PURGEFID

Arguments

```
struct cfs_purgefid_out { /* CFS_PURGEFID is a venus->kernel call */
    ViceFid CodaFid;
} cfs_purgefid;
```

Description

Flush the attribute for the file. If it is a dir (odd vnode), purge its children from the namecache and remove the file from the namecache.

5.8. REPLACE

Summary

Replace the Fid's for a collection of names.

Arguments

```
struct cfs_replace_out { /* cfs_replace is a venus->kernel call */
    ViceFid NewFid;
    ViceFid OldFid;
} cfs_replace;
```

Description

This routine replaces a ViceFid in the name cache with another. It is added to allow Venus during reintegration to replace locally allocated temp fids while disconnected with global fids even when the reference counts on those fids are not zero.

6. Initialization and cleanup

This section gives brief hints as to desirable features for the Coda FS Driver at startup and upon shutdown or Venus failures. Before entering the discussion it is useful to repeat that the Coda FS Driver maintains the following data:

1. message queues
2. cnodes
3. name cache entries

The name cache entries are entirely private to the driver, so they can easily be manipulated. The message queues will generally have clear points of initialization and destruction. The cnodes are much more delicate. User processes hold reference counts in Coda filesystems and it can be difficult to clean up the cnodes.

It can expect requests through:

1. the message subsystem
2. the VFS layer
3. pioctl interface

Currently the pioctl passes through the VFS for Coda so we can treat these similarly.

6.1. Requirements

The following requirements should be accommodated:

1. The message queues should have open and close routines. On Unix the opening of the character devices are such routines.
 - Before opening, no messages can be placed.
 - Opening will remove any old messages still pending.
 - Close will notify any sleeping processes that their upcall cannot be completed.
 - Close will free all memory allocated by the message queues.
2. At open the namecache shall be initialized to empty state.
3. Before the message queues are open, all VFS operations will fail. Fortunately this can be achieved by making sure that mounting the Coda filesystem cannot succeed before opening.
4. After closing of the queues, no VFS operations can succeed. Here one needs to be careful, since a few operations (lookup, read/write, readdir) can proceed without upcalls. These must be explicitly blocked.
5. Upon closing the namecache shall be flushed and disabled.
6. All memory held by cnodes can be freed without relying on upcalls.
7. Unmounting the file system can be done without relying on upcalls.
8. Mounting the Coda filesystem should fail gracefully if Venus cannot get the rootfid or the attributes of the rootfid. The latter is best implemented by Venus fetching these objects before attempting to mount.

Note

NetBSD in particular but also Linux have not implemented the above requirements fully. For smooth operation this needs to be corrected.