

# Fiemap Ioctl

The fiemap ioctl is an efficient method for userspace to get file extent mappings. Instead of block-by-block mapping (such as bmap), fiemap returns a list of extents.

## Request Basics

A fiemap request is encoded within struct fiemap:

```
struct fiemap {
    __u64    fm_start;          /* logical offset (inclusive) at
                                * which to start mapping (in) */
    __u64    fm_length;        /* logical length of mapping which
                                * userspace cares about (in) */
    __u32    fm_flags;          /* FIEMAP_FLAG_* flags for request (in/out) */
    __u32    fm_mapped_extents; /* number of extents that were
                                * mapped (out) */
    __u32    fm_extent_count;   /* size of fm_extents array (in) */
    __u32    fm_reserved;
    struct fiemap_extent fm_extents[0]; /* array of mapped extents (out) */
};
```

fm\_start, and fm\_length specify the logical range within the file which the process would like mappings for. Extents returned mirror those on disk - that is, the logical offset of the 1st returned extent may start before fm\_start, and the range covered by the last returned extent may end after fm\_length. All offsets and lengths are in bytes.

Certain flags to modify the way in which mappings are looked up can be set in fm\_flags. If the kernel doesn't understand some particular flags, it will return EBADR and the contents of fm\_flags will contain the set of flags which caused the error. If the kernel is compatible with all flags passed, the contents of fm\_flags will be unmodified. It is up to userspace to determine whether rejection of a particular flag is fatal to its operation. This scheme is intended to allow the fiemap interface to grow in the future but without losing compatibility with old software.

fm\_extent\_count specifies the number of elements in the fm\_extents[] array that can be used to return extents. If fm\_extent\_count is zero, then the fm\_extents[] array is ignored (no extents will be returned), and the fm\_mapped\_extents count will hold the number of extents needed in fm\_extents[] to hold the file's current mapping. Note that there is nothing to prevent the file from changing between calls to FIEMAP.

The following flags can be set in fm\_flags:

FIEMAP\_FLAG\_SYNC

If this flag is set, the kernel will sync the file before mapping extents.

FIEMAP\_FLAG\_XATTR

If this flag is set, the extents returned will describe the inodes extended attribute lookup tree, instead of its data tree.

## Extent Mapping

Extent information is returned within the embedded fm\_extents array which userspace must allocate along with the fiemap structure. The number of elements in the fiemap\_extents[] array should be passed via fm\_extent\_count. The number of extents mapped by kernel will be returned via fm\_mapped\_extents. If the number of fiemap\_extents allocated is less than would be required to map the requested range, the maximum number of extents that can be mapped in the fm\_extents[] array will be returned and fm\_mapped\_extents will be equal to fm\_extent\_count. In that case, the last extent in the array will not complete the requested range and will not have the FIEMAP\_EXTENT\_LAST flag set (see the next section on extent flags).

Each extent is described by a single fiemap\_extent structure as returned in fm\_extents:

```
struct fiemap_extent {
    __u64    fe_logical; /* logical offset in bytes for the start of
                        * the extent */
    __u64    fe_physical; /* physical offset in bytes for the start
                        * of the extent */
    __u64    fe_length;   /* length in bytes for the extent */
    __u64    fe_reserved64[2];
    __u32    fe_flags;     /* FIEMAP_EXTENT_* flags for this extent */
    __u32    fe_reserved[3];
};
```

All offsets and lengths are in bytes and mirror those on disk. It is valid for an extents logical offset to start before the request or its logical length to extend past the request. Unless FIEMAP\_EXTENT\_NOT\_ALIGNED is returned, fe\_logical, fe\_physical, and fe\_length will be aligned to the block size of the file system. With the exception of extents flagged as FIEMAP\_EXTENT\_MERGED, adjacent extents will not be merged.

The fe\_flags field contains flags which describe the extent returned. A special flag, FIEMAP\_EXTENT\_LAST is always set on the

last extent in the file so that the process making fiemap calls can determine when no more extents are available, without having to call the ioctl again.

Some flags are intentionally vague and will always be set in the presence of other more specific flags. This way a program looking for a general property does not have to know all existing and future flags which imply that property.

For example, if FIEMAP\_EXTENT\_DATA\_INLINE or FIEMAP\_EXTENT\_DATA\_TAIL are set, FIEMAP\_EXTENT\_NOT\_ALIGNED will also be set. A program looking for inline or tail-packed data can key on the specific flag. Software which simply cares not to try operating on non-aligned extents however, can just key on FIEMAP\_EXTENT\_NOT\_ALIGNED, and not have to worry about all present and future flags which might imply unaligned data. Note that the opposite is not true - it would be valid for FIEMAP\_EXTENT\_NOT\_ALIGNED to appear alone.

#### FIEMAP\_EXTENT\_LAST

This is generally the last extent in the file. A mapping attempt past this extent may return nothing. Some implementations set this flag to indicate this extent is the last one in the range queried by the user (via fiemap->fm\_length).

#### FIEMAP\_EXTENT\_UNKNOWN

The location of this extent is currently unknown. This may indicate the data is stored on an inaccessible volume or that no storage has been allocated for the file yet.

#### FIEMAP\_EXTENT\_DEALLOC

This will also set FIEMAP\_EXTENT\_UNKNOWN.

Delayed allocation - while there is data for this extent, its physical location has not been allocated yet.

#### FIEMAP\_EXTENT\_ENCODED

This extent does not consist of plain filesystem blocks but is encoded (e.g. encrypted or compressed). Reading the data in this extent via I/O to the block device will have undefined results.

Note that it is *always* undefined to try to update the data in-place by writing to the indicated location without the assistance of the filesystem, or to access the data using the information returned by the FIEMAP interface while the filesystem is mounted. In other words, user applications may only read the extent data via I/O to the block device while the filesystem is unmounted, and then only if the FIEMAP\_EXTENT\_ENCODED flag is clear; user applications must not try reading or writing to the filesystem via the block device under any other circumstances.

#### FIEMAP\_EXTENT\_DATA\_ENCRYPTED

This will also set FIEMAP\_EXTENT\_ENCODED. The data in this extent has been encrypted by the file system.

#### FIEMAP\_EXTENT\_NOT\_ALIGNED

Extent offsets and length are not guaranteed to be block aligned.

#### FIEMAP\_EXTENT\_DATA\_INLINE

This will also set FIEMAP\_EXTENT\_NOT\_ALIGNED. Data is located within a meta data block.

#### FIEMAP\_EXTENT\_DATA\_TAIL

This will also set FIEMAP\_EXTENT\_NOT\_ALIGNED. Data is packed into a block with data from other files.

#### FIEMAP\_EXTENT\_UNWRITTEN

Unwritten extent - the extent is allocated but its data has not been initialized. This indicates the extent's data will be all zero if read through the filesystem but the contents are undefined if read directly from the device.

#### FIEMAP\_EXTENT\_MERGED

This will be set when a file does not support extents, i.e., it uses a block based addressing scheme. Since returning an extent for each block back to userspace would be highly inefficient, the kernel will try to merge most adjacent blocks into 'extents'.

## VFS -> File System Implementation

File systems wishing to support fiemap must implement a ->fiemap callback on their inode\_operations structure. The fs ->fiemap call is responsible for defining its set of supported fiemap flags, and calling a helper function on each discovered extent:

```
struct inode_operations {
    ...

    int (*fiemap)(struct inode *, struct fiemap_extent_info *, u64 start,
                  u64 len);
};
```

->fiemap is passed struct fiemap\_extent\_info which describes the fiemap request:

```
struct fiemap_extent_info {
    unsigned int fi_flags; /* Flags as passed from user */
    unsigned int fi_extents_mapped; /* Number of mapped extents */
    unsigned int fi_extents_max; /* Size of fiemap_extent array */
    struct fiemap_extent *fi_extents_start; /* Start of fiemap_extent array */
};
```

It is intended that the file system should not need to access any of this structure directly. Filesystem handlers should be tolerant to signals and return EINTR once fatal signal received.

Flag checking should be done at the beginning of the ->fiemap callback via the fiemap\_prep() helper:

```
int fiemap_prep(struct inode *inode, struct fiemap_extent_info *fieinfo,
                u64 start, u64 *len, u32 supported_flags);
```

The struct fieinfo should be passed in as received from ioctl\_fiemap(). The set of fiemap flags which the fs understands should be passed via fs\_flags. If fiemap\_prep finds invalid user flags, it will place the bad values in fieinfo->fi\_flags and return -EBADR. If the file system gets -EBADR, from fiemap\_prep(), it should immediately exit, returning that error back to ioctl\_fiemap(). Additionally the range is validate against the supported maximum file size.

For each extent in the request range, the file system should call the helper function, fiemap\_fill\_next\_extent():

```
int fiemap_fill_next_extent(struct fiemap_extent_info *info, u64 logical,
                           u64 phys, u64 len, u32 flags, u32 dev);
```

fiemap\_fill\_next\_extent() will use the passed values to populate the next free extent in the fin\_extents array. 'General' extent flags will automatically be set from specific flags on behalf of the calling file system so that the userspace API is not broken.

fiemap\_fill\_next\_extent() returns 0 on success, and 1 when the user-supplied fin\_extents array is full. If an error is encountered while copying the extent to user memory, -EFAULT will be returned.