

Optimized MPEG Filesystem (OMFS)

Overview

OMFS is a filesystem created by SonicBlue for use in the ReplayTV DVR and Rio Karma MP3 player. The filesystem is extent-based, utilizing block sizes from 2k to 8k, with hash-based directories. This filesystem driver may be used to read and write disks from these devices.

Note, it is not recommended that this FS be used in place of a general filesystem for your own streaming media device. Native Linux filesystems will likely perform better.

More information is available at:

<http://linux-karma.sf.net/>

Various utilities, including `mkomfs` and `omfsck`, are included with `omfsprogs`, available at:

<https://bobcopeland.com/karma/>

Instructions are included in its README.

Options

OMFS supports the following mount-time options:

<code>uid=n</code>	make all files owned by specified user
<code>gid=n</code>	make all files owned by specified group
<code>umask=xxx</code>	set permission umask to xxx
<code>fmask=xxx</code>	set umask to xxx for files
<code>dmask=xxx</code>	set umask to xxx for directories

Disk format

OMFS discriminates between "sysblocks" and normal data blocks. The sysblock group consists of super block information, file metadata, directory structures, and extents. Each sysblock has a header containing CRCs of the entire sysblock, and may be mirrored in successive blocks on the disk. A sysblock may have a smaller size than a data block, but since they are both addressed by the same 64-bit block number, any remaining space in the smaller sysblock is unused.

Sysblock header information:

```
struct omfs_header {
    __be64 h_self;                /* FS block where this is located */
    __be32 h_body_size;           /* size of useful data after header */
    __be16 h_crc;                 /* crc-ccitt of body_size bytes */
    char h_fill1[2];
    u8 h_version;                 /* version, always 1 */
    char h_type;                  /* OMFS_INODE_X */
    u8 h_magic;                   /* OMFS_IMAGIC */
    u8 h_check_xor;               /* XOR of header bytes before this */
    __be32 h_fill2;
};
```

Files and directories are both represented by `omfs_inode`:

```
struct omfs_inode {
    struct omfs_header i_head;    /* header */
    __be64 i_parent;              /* parent containing this inode */
    __be64 i_sibling;             /* next inode in hash bucket */
    __be64 i_ctime;               /* ctime, in milliseconds */
    char i_fill1[35];
    char i_type;                  /* OMFS_[DIR,FILE] */
    __be32 i_fill2;
    char i_fill3[64];
    char i_name[OMFS_NAMELEN];    /* filename */
    __be64 i_size;               /* size of file, in bytes */
};
```

Directories in OMFS are implemented as a large hash table. Filenames are hashed then prepended into the bucket list beginning at `OMFS_DIR_START`. Lookup requires hashing the filename, then seeking across `i_sibling` pointers until a match is found on `i_name`. Empty buckets are represented by block pointers with all-1s (~0).

A file is an `omfs_inode` structure followed by an extent table beginning at `OMFS_EXTENT_START`:

```

struct omfs_extent_entry {
    __be64 e_cluster;          /* start location of a set of blocks */
    __be64 e_blocks;          /* number of blocks after e_cluster */
};

struct omfs_extent {
    __be64 e_next;             /* next extent table location */
    __be32 e_extent_count;     /* total # extents in this table */
    __be32 e_fill;
    struct omfs_extent_entry e_entry; /* start of extent entries */
};

```

Each extent holds the block offset followed by number of blocks allocated to the extent. The final extent in each table is a terminator with `e_cluster` being `~0` and `e_blocks` being ones'-complement of the total number of blocks in the table.

If this table overflows, a continuation inode is written and pointed to by `e_next`. These have a header but lack the rest of the inode structure.