

Kernel driver lm93

Supported chips:

- National Semiconductor LM93
Prefix 'lm93'
Addresses scanned: I2C 0x2c-0x2e
Datasheet: <http://www.national.com/ds/cgi/LM/LM93.pdf>
- National Semiconductor LM94
Prefix 'lm94'
Addresses scanned: I2C 0x2c-0x2e
Datasheet: <http://www.national.com/ds/cgi/LM/LM94.pdf>

Authors:

- Mark M. Hoffman <mhoffman@lightlink.com>
- Ported to 2.6 by Eric J. Bowersox <ericb@aspsys.com>
- Adapted to 2.6.20 by Carsten Emde <ce@osadl.org>
- Modified for mainline integration by Hans J. Koch <hjk@hansjkoeh.de>

Module Parameters

- `init`: integer Set to non-zero to force some initializations (default is 0).
- `disable_block`: integer A "0" allows SMBus block data transactions if the host supports them. A "1" disables SMBus block data transactions. The default is 0.
- `vccp_limit_type`: integer array (2) Configures in7 and in8 limit type, where 0 means absolute and non-zero means relative. "Relative" here refers to "Dynamic Vccp Monitoring using VID" from the datasheet. It greatly simplifies the interface to allow only one set of limits (absolute or relative) to be in operation at a time (even though the hardware is capable of enabling both). There's not a compelling use case for enabling both at once, anyway. The default is "0,0".
- `vid_agtl`: integer A "0" configures the VID pins for $V(ih) = 2.1V$ min, $V(il) = 0.8V$ max. A "1" configures the VID pins for $V(ih) = 0.8V$ min, $V(il) = 0.4V$ max. (The latter setting is referred to as AGTL+ Compatible in the datasheet.) I.e. this parameter controls the VID pin input thresholds; if your VID inputs are not working, try changing this. The default value is "0".

Hardware Description

(from the datasheet)

The LM93 hardware monitor has a two wire digital interface compatible with SMBus 2.0. Using an 8-bit ADC, the LM93 measures the temperature of two remote diode connected transistors as well as its own die and 16 power supply voltages. To set fan speed, the LM93 has two PWM outputs that are each controlled by up to four temperature zones. The fancontrol algorithm is lookup table based. The LM93 includes a digital filter that can be invoked to smooth temperature readings for better control of fan speed. The LM93 has four tachometer inputs to measure fan speed. Limit and status registers for all measured values are included. The LM93 builds upon the functionality of previous motherboard management ASICs and uses some of the LM85's features (i.e. smart tachometer mode). It also adds measurement and control support for dynamic Vccp monitoring and PROCHOT. It is designed to monitor a dual processor Xeon class motherboard with a minimum of external components.

LM94 is also supported in LM93 compatible mode. Extra sensors and features of LM94 are not supported.

User Interface

#PROCHOT

The LM93 can monitor two #PROCHOT signals. The results are found in the sysfs files `prochot1`, `prochot2`, `prochot1_avg`, `prochot2_avg`, `prochot1_max`, and `prochot2_max`. `prochot1_max` and `prochot2_max` contain the user limits for #PROCHOT1 and #PROCHOT2, respectively. `prochot1` and `prochot2` contain the current readings for the most recent complete time interval. The value of `prochot1_avg` and `prochot2_avg` is something like a 2 period exponential moving average (but not quite - check the datasheet). Note that this third value is calculated by the chip itself. All values range from 0-255 where 0 indicates no throttling, and 255 indicates > 99.6%.

The monitoring intervals for the two #PROCHOT signals is also configurable. These intervals can be found in the sysfs files `prochot1_interval` and `prochot2_interval`. The values in these files specify the intervals for #P1_PROCHOT and #P2_PROCHOT, respectively. Selecting a value not in this list will cause the driver to use the next largest interval. The available intervals are (in seconds):

#PROCHOT intervals:

0.73, 1.46, 2.9, 5.8, 11.7, 23.3, 46.6, 93.2, 186, 372

It is possible to configure the LM93 to logically short the two #PROCHOT signals. I.e. when #P1_PROCHOT is asserted, the LM93 will automatically assert #P2_PROCHOT, and vice-versa. This mode is enabled by writing a non-zero integer to the sysfs file `prochot_short`.

The LM93 can also override the #PROCHOT pins by driving a PWM signal onto one or both of them. When overridden, the signal has a period of 3.56 ms, a minimum pulse width of 5 clocks (at 22.5kHz => 6.25% duty cycle), and a maximum pulse width of 80 clocks (at 22.5kHz => 99.88% duty cycle).

The sysfs files `prochot1_override` and `prochot2_override` contain boolean integers which enable or disable the override function for #P1_PROCHOT and #P2_PROCHOT, respectively. The sysfs file `prochot_override_duty_cycle` contains a value controlling the duty cycle for the PWM signal used when the override function is enabled. This value ranges from 0 to 15, with 0 indicating minimum duty cycle and 15 indicating maximum.

#VRD_HOT

The LM93 can monitor two #VRD_HOT signals. The results are found in the sysfs files `vrshot1` and `vrshot2`. There is one value per file: a boolean for which 1 indicates #VRD_HOT is asserted and 0 indicates it is negated. These files are read-only.

Smart Tach Mode (from the datasheet):

If a fan is driven using a low-side drive PWM, the tachometer output of the fan is corrupted. The LM93 includes smart tachometer circuitry that allows an accurate tachometer reading to be achieved despite the signal corruption. In smart tach mode all four signals are measured within 4 seconds.

Smart tach mode is enabled by the driver by writing 1 or 2 (associating the fan tachometer with a pwm) to the sysfs file `fan<n>_smart_tach`. A zero will disable the function for that fan. Note that Smart tach mode cannot be enabled if the PWM output frequency is 22500 Hz (see below).

Manual PWM

The LM93 has a fixed or override mode for the two PWM outputs (although, there are still some conditions that will override even this mode - see section 15.10.6 of the datasheet for details.) The sysfs files `pwm1_override` and `pwm2_override` are used to enable this mode; each is a boolean integer where 0 disables and 1 enables the manual control mode. The sysfs files `pwm1` and `pwm2` are used to set the manual duty cycle; each is an integer (0-255) where 0 is 0% duty cycle, and 255 is 100%. Note that the duty cycle values are constrained by the hardware. Selecting a value which is not available will cause the driver to use the next largest value. Also note: when manual PWM mode is disabled, the value of `pwm1` and `pwm2` indicates the current duty cycle chosen by the h/w.

PWM Output Frequency

The LM93 supports several different frequencies for the PWM output channels. The sysfs files `pwm1_freq` and `pwm2_freq` are used to select the frequency. The frequency values are constrained by the hardware. Selecting a value which is not available will cause the driver to use the next largest value. Also note that this parameter has implications for the Smart Tach Mode (see above).

PWM Output Frequencies (in Hz):

12, 36, 48, 60, 72, 84, 96, 22500 (default)

Automatic PWM

The LM93 is capable of complex automatic fan control, with many different points of configuration. To start, each PWM output can be bound to any combination of eight control sources. The final PWM is the largest of all individual control sources to which the PWM output is bound.

The eight control sources are: `temp1-temp4` (aka "zones" in the datasheet), #PROCHOT 1 & 2, and #VRDHOT 1 & 2. The bindings are expressed as a bitmask in the sysfs files `pwm<n>_auto_channels`, where a "1" enables the binding, and a "0" disables it. The h/w default is 0x0f (all temperatures bound).

0x01	Temp 1
0x02	Temp 2
0x04	Temp 3
0x08	Temp 4
0x10	#PROCHOT 1
0x20	#PROCHOT 2
0x40	#VRDHOT 1
0x80	#VRDHOT 2

The function $y = f(x)$ takes a source temperature x to a PWM output y . This function of the LM93 is derived from a base temperature and a table of 12 temperature offsets. The base temperature is expressed in degrees C in the sysfs files `temp<n>_auto_base`. The offsets are expressed in cumulative degrees C, with the value of offset `<i>` for temperature value `<n>` being contained in the file `temp<n>_auto_offset<i>`. E.g. if the base temperature is 40C:

offset #	temp<n>_auto_offset<i>	range	pwm
1	0	•	25.00%
2	0	•	28.57%
3	1	40C - 41C	32.14%
4	1	41C - 42C	35.71%
5	2	42C - 44C	39.29%
6	2	44C - 46C	42.86%
7	2	48C - 50C	46.43%
8	2	50C - 52C	50.00%
9	2	52C - 54C	53.57%
10	2	54C - 56C	57.14%
11	2	56C - 58C	71.43%
12	2	58C - 60C	85.71%
•	•	> 60C	100.00%

Valid offsets are in the range $0C \leq x \leq 7.5C$ in 0.5C increments.

There is an independent base temperature for each temperature channel. Note, however, there are only two tables of offsets: one each for `temp[12]` and `temp[34]`. Therefore, any change to e.g. `temp1_auto_offset<i>` will also affect `temp2_auto_offset<i>`.

The LM93 can also apply hysteresis to the offset table, to prevent unwanted oscillation between two steps in the offsets table. These values are found in the sysfs files `temp<n>_auto_offset_hyst`. The value in this file has the same representation as in `temp<n>_auto_offset<i>`.

If a temperature reading falls below the base value for that channel, the LM93 will use the minimum PWM value. These values are found in the sysfs files `temp<n>_auto_pwm_min`. Note, there are only two minimums: one each for `temp[12]` and `temp[34]`. Therefore, any change to e.g. `temp1_auto_pwm_min` will also affect `temp2_auto_pwm_min`.

PWM Spin-Up Cycle

A spin-up cycle occurs when a PWM output is commanded from 0% duty cycle to some value $> 0\%$. The LM93 supports a minimum duty cycle during spin-up. These values are found in the sysfs files `pwm<n>_auto_spinup_min`. The value in this file has the same representation as other PWM duty cycle values. The duration of the spin-up cycle is also configurable. These values are found in the sysfs files `pwm<n>_auto_spinup_time`. The value in this file is the spin-up time in seconds. The available spin-up times are constrained by the hardware. Selecting a value which is not available will cause the driver to use the next largest value.

Spin-up Durations:

0 (disabled, h/w default), 0.1, 0.25, 0.4, 0.7, 1.0, 2.0, 4.0

#PROCHOT and #VRDHOT PWM Ramping

If the `#PROCHOT` or `#VRDHOT` signals are asserted while bound to a PWM output channel, the LM93 will ramp the PWM output up to 100% duty cycle in discrete steps. The duration of each step is configurable. There are two files, with one value each in seconds: `pwm_auto_prochot_ramp` and `pwm_auto_vrdhot_ramp`. The available ramp times are constrained by the hardware. Selecting a value which is not available will cause the driver to use the next largest value.

Ramp Times:

0 (disabled, h/w default) to 0.75 in 0.05 second intervals

Fan Boost

For each temperature channel, there is a boost temperature: if the channel exceeds this limit, the LM93 will immediately drive both PWM outputs to 100%. This limit is expressed in degrees C in the sysfs files `temp<n>_auto_boost`. There is also a hysteresis temperature for this function: after the boost limit is reached, the temperature channel must drop below this value before the boost function is disabled. This temperature is also expressed in degrees C in the sysfs files `temp<n>_auto_boost_hyst`.

GPIO Pins

The LM93 can monitor the logic level of four dedicated GPIO pins as well as the four tach input pins. GPIO0-GPIO3 correspond to (fan) tach 1-4, respectively. All eight GPIOs are read by reading the bitmask in the sysfs file `gpio`. The LSB is GPIO0, and the MSB is GPIO7.

LM93 Unique sysfs Files

file	description
prochot<n>	current #PROCHOT %
prochot<n>_avg	moving average #PROCHOT %
prochot<n>_max	limit #PROCHOT %
prochot_short	enable or disable logical #PROCHOT pin short
prochot<n>_override	force #PROCHOT assertion as PWM
prochot_override_duty_cycle	duty cycle for the PWM signal used when #PROCHOT is overridden
prochot<n>_interval	#PROCHOT PWM sampling interval
vrldhot<n>	0 means negated, 1 means asserted
fan<n>_smart_tach	enable or disable smart tach mode
pwm<n>_auto_channels	select control sources for PWM outputs
pwm<n>_auto_spinup_min	minimum duty cycle during spin-up
pwm<n>_auto_spinup_time	duration of spin-up
pwm_auto_prochot_ramp	ramp time per step when #PROCHOT asserted
pwm_auto_vrldhot_ramp	ramp time per step when #VRDHOT asserted
temp<n>_auto_base	temperature channel base
temp<n>_auto_offset[1-12]	temperature channel offsets
temp<n>_auto_offset_hyst	temperature channel offset hysteresis
temp<n>_auto_boost	temperature channel boost (PWMs to 100%) limit
temp<n>_auto_boost_hyst	temperature channel boost hysteresis
gpio	input state of 8 GPIO pins; read-only