

Reference-count design for elements of lists/arrays protected by RCU

Please note that the percpu-ref feature is likely your first stop if you need to combine reference counts and RCU. Please see `include/linux/percpu-refcount.h` for more information. However, in those unusual cases where percpu-ref would consume too much memory, please read on.

Reference counting on elements of lists which are protected by traditional reader/writer spinlocks or semaphores are straightforward:

CODE LISTING A:

```
1.
add()
{
    alloc_object
    ...
    atomic_set(&el->rc, 1);
    write_lock(&list_lock);
    add_element
    ...
    write_unlock(&list_lock);
}

2.
search_and_reference()
{
    read_lock(&list_lock);
    search_for_element
    atomic_inc(&el->rc);
    ...
    read_unlock(&list_lock);
    ...
}

3.
release_referenced()
{
    ...
    if(atomic_dec_and_test(&el->rc))
        kfree(el);
    ...
}

4.
delete()
{
    write_lock(&list_lock);
    ...

    remove_element
    write_unlock(&list_lock);
    ...
    if (atomic_dec_and_test(&el->rc))
        kfree(el);
    ...
}
```

If this list/array is made lock free using RCU as in changing the `write_lock()` in `add()` and `delete()` to `spin_lock()` and changing `read_lock()` in `search_and_reference()` to `rcu_read_lock()`, the `atomic_inc()` in `search_and_reference()` could potentially hold reference to an element which has already been deleted from the list/array. Use `atomic_inc_not_zero()` in this scenario as follows:

CODE LISTING B:

```
1.
add()
{
    alloc_object
    ...
    atomic_set(&el->rc, 1);
    spin_lock(&list_lock);

    add_element
    ...
    spin_unlock(&list_lock);
}

2.
search_and_reference()
{
    rcu_read_lock();
    search_for_element
    if (!atomic_inc_not_zero(&el->rc)) {
        rcu_read_unlock();
        return FAIL;
    }
    ...
    rcu_read_unlock();
}

3.
release_referenced()
{
    ...
    if (atomic_dec_and_test(&el->rc))
        call_rcu(&el->head, el_free);
    ...
}

4.
delete()
{
    spin_lock(&list_lock);
    ...
    remove_element
    spin_unlock(&list_lock);
    ...
    if (atomic_dec_and_test(&el->rc))
        call_rcu(&el->head, el_free);
    ...
}
```

Sometimes, a reference to the element needs to be obtained in the update (write) stream. In such cases, `atomic_inc_not_zero()` might be overkill, since we hold the update-side spinlock. One might instead use `atomic_inc()` in such cases.

It is not always convenient to deal with "FAIL" in the `search_and_reference()` code path. In such cases, the `atomic_dec_and_test()` may be moved from `delete()` to `el_free()` as follows:

CODE LISTING C:

```

1.
add()
{
    alloc_object
    ...
    atomic_set(&el->rc, 1);
    spin_lock(&list_lock);

    add_element
    ...
    spin_unlock(&list_lock);
}
3.
release_referenced()
{
    ...
    if (atomic_dec_and_test(&el->rc))
        kfree(el);
    ...
}
5.
void el_free(struct rcu_head *rhp)
{
    release_referenced();
}

2.
search_and_reference()
{
    rcu_read_lock();
    search_for_element
    atomic_inc(&el->rc);
    ...

    rcu_read_unlock();
}
4.
delete()
{
    spin_lock(&list_lock);
    ...
    remove_element
    spin_unlock(&list_lock);
    ...
    call_rcu(&el->head, el_free);
    ...
}

```

The key point is that the initial reference added by `add()` is not removed until after a grace period has elapsed following removal. This means that `search_and_reference()` cannot find this element, which means that the value of `el->rc` cannot increase. Thus, once it reaches zero, there are no readers that can or ever will be able to reference the element. The element can therefore safely be freed. This in turn guarantees that if any reader finds the element, that reader may safely acquire a reference without checking the value of the reference counter.

A clear advantage of the RCU-based pattern in listing C over the one in listing B is that any call to `search_and_reference()` that locates a given object will succeed in obtaining a reference to that object, even given a concurrent invocation of `delete()` for that same object. Similarly, a clear advantage of both listings B and C over listing A is that a call to `delete()` is not delayed even if there are an arbitrarily large number of calls to `search_and_reference()` searching for the same object that `delete()` was invoked on. Instead, all that is delayed is the eventual invocation of `kfree()`, which is usually not a problem on modern computer systems, even the small ones.

In cases where `delete()` can sleep, `synchronize_rcu()` can be called from `delete()`, so that `el_free()` can be subsumed into `delete` as follows:

```

4.
delete()
{
    spin_lock(&list_lock);
    ...
    remove_element
    spin_unlock(&list_lock);
    ...
    synchronize_rcu();
    if (atomic_dec_and_test(&el->rc))
        kfree(el);
    ...
}

```

As additional examples in the kernel, the pattern in listing C is used by reference counting of `struct pid`, while the pattern in listing B is used by `struct posix_acl`.