

Speculation Control

Quite some CPUs have speculation-related misfeatures which are in fact vulnerabilities causing data leaks in various forms even across privilege domains.

The kernel provides mitigation for such vulnerabilities in various forms. Some of these mitigations are compile-time configurable and some can be supplied on the kernel command line.

There is also a class of mitigations which are very expensive, but they can be restricted to a certain set of processes or tasks in controlled environments. The mechanism to control these mitigations is via `manpage:prctl(2)`.

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\linux-master\Documentation\userspace-api\linux-master) (Documentation) (userspace-api) spec_ctrl.rst, line 13); [backlink](#)

Unknown interpreted text role "manpage".

There are two `prctl` options which are related to this:

- `PR_GET_SPECULATION_CTRL`
- `PR_SET_SPECULATION_CTRL`

PR_GET_SPECULATION_CTRL

`PR_GET_SPECULATION_CTRL` returns the state of the speculation misfeature which is selected with `arg2` of `prctl(2)`. The return value uses bits 0-3 with the following meaning:

Bit	Define	Description
0	<code>PR_SPEC_PRCTL</code>	Mitigation can be controlled per task by <code>PR_SET_SPECULATION_CTRL</code> .
1	<code>PR_SPEC_ENABLE</code>	The speculation feature is enabled, mitigation is disabled.
2	<code>PR_SPEC_DISABLE</code>	The speculation feature is disabled, mitigation is enabled.
3	<code>PR_SPEC_FORCE_DISABLE</code>	Same as <code>PR_SPEC_DISABLE</code> , but cannot be undone. A subsequent <code>prctl(..., PR_SPEC_ENABLE)</code> will fail.
4	<code>PR_SPEC_DISABLE_NOEXEC</code>	Same as <code>PR_SPEC_DISABLE</code> , but the state will be cleared on <code>manpage:execve(2)</code> . <div>System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\linux-master\Documentation\userspace-api\linux-master) (Documentation) (userspace-api) spec_ctrl.rst, line 43); backlink Unknown interpreted text role "manpage".</div>

If all bits are 0 the CPU is not affected by the speculation misfeature.

If `PR_SPEC_PRCTL` is set, then the per-task control of the mitigation is available. If not set, `prctl(PR_SET_SPECULATION_CTRL)` for the speculation misfeature will fail.

PR_SET_SPECULATION_CTRL

`PR_SET_SPECULATION_CTRL` allows to control the speculation misfeature, which is selected by `arg2` of `manpage:prctl(2)` per task. `arg3` is used to hand in the control value, i.e. either `PR_SPEC_ENABLE` or `PR_SPEC_DISABLE` or `PR_SPEC_FORCE_DISABLE`.

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\linux-master\Documentation\userspace-api\linux-master) (Documentation) (userspace-api) spec_ctrl.rst, line 57); [backlink](#)

Unknown interpreted text role "manpage".

Common error codes

Value	Meaning
<code>EINVAL</code>	The <code>prctl</code> is not implemented by the architecture or unused <code>prctl(2)</code> arguments are not 0.

Value	Meaning
-------	---------

ENODEV	arg2 is selecting a not supported speculation misfeature.
--------	---

PR_SET_SPECULATION_CTRL error codes

Value	Meaning
0	Success
ERANGE	arg3 is incorrect, i.e. it's neither PR_SPEC_ENABLE nor PR_SPEC_DISABLE nor PR_SPEC_FORCE_DISABLE.
ENXIO	Control of the selected speculation misfeature is not possible. See PR_GET_SPECULATION_CTRL.
EPERM	Speculation was disabled with PR_SPEC_FORCE_DISABLE and caller tried to enable it again.

Speculation misfeature controls

- PR_SPEC_STORE_BYPASS: Speculative Store Bypass

Invocations:

- prctl(PR_GET_SPECULATION_CTRL, PR_SPEC_STORE_BYPASS, 0, 0, 0);
- prctl(PR_SET_SPECULATION_CTRL, PR_SPEC_STORE_BYPASS, PR_SPEC_ENABLE, 0, 0);
- prctl(PR_SET_SPECULATION_CTRL, PR_SPEC_STORE_BYPASS, PR_SPEC_DISABLE, 0, 0);
- prctl(PR_SET_SPECULATION_CTRL, PR_SPEC_STORE_BYPASS, PR_SPEC_FORCE_DISABLE, 0, 0);
- prctl(PR_SET_SPECULATION_CTRL, PR_SPEC_STORE_BYPASS, PR_SPEC_DISABLE_NOEXEC, 0, 0);

- PR_SPEC_INDIR_BRANCH: Indirect Branch Speculation in User Processes
(Mitigate Spectre V2 style attacks against user processes)

Invocations:

- prctl(PR_GET_SPECULATION_CTRL, PR_SPEC_INDIRECT_BRANCH, 0, 0, 0);
- prctl(PR_SET_SPECULATION_CTRL, PR_SPEC_INDIRECT_BRANCH, PR_SPEC_ENABLE, 0, 0);
- prctl(PR_SET_SPECULATION_CTRL, PR_SPEC_INDIRECT_BRANCH, PR_SPEC_DISABLE, 0, 0);
- prctl(PR_SET_SPECULATION_CTRL, PR_SPEC_INDIRECT_BRANCH, PR_SPEC_FORCE_DISABLE, 0, 0);

- PR_SPEC_L1D_FLUSH: Flush L1D Cache on context switch out of the task
(works only when tasks run on non SMT cores)

Invocations:

- prctl(PR_GET_SPECULATION_CTRL, PR_SPEC_L1D_FLUSH, 0, 0, 0);
- prctl(PR_SET_SPECULATION_CTRL, PR_SPEC_L1D_FLUSH, PR_SPEC_ENABLE, 0, 0);
- prctl(PR_SET_SPECULATION_CTRL, PR_SPEC_L1D_FLUSH, PR_SPEC_DISABLE, 0, 0);