

PAT (Page Attribute Table)

x86 Page Attribute Table (PAT) allows for setting the memory attribute at the page level granularity. PAT is complementary to the MTRR settings which allows for setting of memory types over physical address ranges. However, PAT is more flexible than MTRR due to its capability to set attributes at page level and also due to the fact that there are no hardware limitations on number of such attribute settings allowed. Added flexibility comes with guidelines for not having memory type aliasing for the same physical memory with multiple virtual addresses.

PAT allows for different types of memory attributes. The most commonly used ones that will be supported at this time are:

| | |
|-----|----------------|
| WB | Write-back |
| UC | Uncached |
| WC | Write-combined |
| WT | Write-through |
| UC- | Uncached Minus |

PAT APIs

There are many different APIs in the kernel that allows setting of memory attributes at the page level. In order to avoid aliasing, these interfaces should be used thoughtfully. Below is a table of interfaces available, their intended usage and their memory attribute relationships. Internally, these APIs use a `reserve_memtype()/free_memtype()` interface on the physical address range to avoid any aliasing.

| API | RAM | ACPI,... | Reserved/Holes |
|---|-----|------------------------------------|------------------------------------|
| <code>ioremap</code> | -- | UC- | UC- |
| <code>ioremap_cache</code> | -- | WB | WB |
| <code>ioremap_uc</code> | -- | UC | UC |
| <code>ioremap_wc</code> | -- | -- | WC |
| <code>ioremap_wt</code> | -- | -- | WT |
| <code>set_memory_uc</code> , <code>set_memory_wb</code> | UC- | -- | -- |
| <code>set_memory_wc</code> , <code>set_memory_wb</code> | WC | -- | -- |
| <code>set_memory_wt</code> , <code>set_memory_wb</code> | WT | -- | -- |
| <code>pci sysfs resource</code> | -- | -- | UC- |
| <code>pci sysfs resource_wc</code> is <code>IORESOURCE_PREFETCH</code> | -- | -- | WC |
| <code>pci proc !PCIIOC_WRITE_COMBINE</code> | -- | -- | UC- |
| <code>pci proc PCIIOC_WRITE_COMBINE</code> | -- | -- | WC |
| <code>/dev/mem read-write</code> | -- | WB/WC/UC- | WB/WC/UC- |
| <code>/dev/mem mmap SYNC flag</code> | -- | UC- | UC- |
| <code>/dev/mem mmap !SYNC flag and any alias to this area</code> | -- | WB/WC/UC- (from existing alias) | WB/WC/UC- (from existing alias) |
| <code>/dev/mem mmap !SYNC flag no alias to this area and MTRR says WB</code> | -- | WB | WB |
| <code>/dev/mem mmap !SYNC flag no alias to this area and MTRR says !WB</code> | -- | -- | UC- |

Advanced APIs for drivers

A. Exporting pages to users with `remap_pfn_range`, `io_remap_pfn_range`, `vmf_insert_pfn`.

Drivers wanting to export some pages to userspace do it by using `mmap` interface and a combination of:

1. `pgprot_noncached()`
2. `io_remap_pfn_range()` or `remap_pfn_range()` or `vmf_insert_pfn()`

With PAT support, a new API `pgprot_writecombine` is being added. So, drivers can continue to use the above sequence, with either `pgprot_noncached()` or `pgprot_writecombine()` in step 1, followed by step 2.

In addition, step 2 internally tracks the region as UC or WC in `memtype` list in order to ensure no conflicting mapping.

Note that this set of APIs only works with IO (non RAM) regions. If driver wants to export a RAM region, it has to do `set_memory_uc()` or `set_memory_wc()` as step 0 above and also track the usage of those pages and use `set_memory_wb()` before the page is freed to free pool.

MTRR effects on PAT / non-PAT systems

The following table provides the effects of using write-combining MTRRs when using `ioremap*()` calls on x86 for both non-PAT and PAT systems. Ideally `mtrr_add()` usage will be phased out in favor of `arch_phys_wc_add()` which will be a no-op on PAT enabled systems. The region over which a `arch_phys_wc_add()` is made, should already have been ioremapped with WC attributes or PAT entries, this can be done by using `ioremap_wc()` / `set_memory_wc()`. Devices which combine areas of IO memory desired to remain uncacheable with areas where write-combining is desirable should consider use of `ioremap_uc()` followed by `set_memory_wc()` to white-list effective write-combined areas. Such use is nevertheless discouraged as the effective memory type is considered implementation defined, yet this strategy can be used as last resort on devices with size-constrained regions where otherwise MTRR write-combining would otherwise not be effective.

| MTRR | Non-PAT | PAT | Linux ioremap value | Effective memory type | |
|------|---------|-----|--|-----------------------|-----|
| | | | | Non-PAT | PAT |
| | PAT | | | | |
| | PCD | | | | |
| | PWT | | | | |
| | | | | | |
| WC | 000 | WB | <code>_PAGE_CACHE_MODE_WB</code> | WC | WC |
| WC | 001 | WC | <code>_PAGE_CACHE_MODE_WC</code> | WC* | WC |
| WC | 010 | UC- | <code>_PAGE_CACHE_MODE_UC_MINUS</code> | WC* | UC |
| WC | 011 | UC | <code>_PAGE_CACHE_MODE_UC</code> | UC | UC |

(*) denotes implementation defined and is discouraged

Note

-- in the above table mean "Not suggested usage for the API". Some of the --'s are strictly enforced by the kernel. Some others are not really enforced today, but may be enforced in future.

For `ioremap` and `pci` access through `/sys` or `/proc` - The actual type returned can be more restrictive, in case of any existing aliasing for that address. For example: If there is an existing uncached mapping, a new `ioremap_wc` can return uncached mapping in place of write-combine requested.

`set_memory_[uc|wc|wt]` and `set_memory_wb` should be used in pairs, where driver will first make a region `uc`, `wc` or `wt` and switch it back to `wb` after use.

Over time writes to `/proc/mtrr` will be deprecated in favor of using PAT based interfaces. Users writing to `/proc/mtrr` are suggested to use above interfaces.

Drivers should use `ioremap_[uc|wc]` to access PCI BARs with `[uc|wc]` access types.

Drivers should use `set_memory_[uc|wc|wt]` to set access type for RAM ranges.

PAT debugging

With `CONFIG_DEBUG_FS` enabled, PAT memtype list can be examined by:

```
# mount -t debugfs debugfs /sys/kernel/debug
# cat /sys/kernel/debug/x86/pat_memtype_list
PAT memtype list:
uncached-minus @ 0x7fadb000-0x7fae0000
uncached-minus @ 0x7fb19000-0x7fb1a000
uncached-minus @ 0x7fb1a000-0x7fb1b000
uncached-minus @ 0x7fb1b000-0x7fb1c000
uncached-minus @ 0x7fb1c000-0x7fb1d000
uncached-minus @ 0x7fb1d000-0x7fb1e000
uncached-minus @ 0x7fb1e000-0x7fb25000
uncached-minus @ 0x7fb25000-0x7fb26000
uncached-minus @ 0x7fb26000-0x7fb27000
uncached-minus @ 0x7fb27000-0x7fb28000
uncached-minus @ 0x7fb28000-0x7fb2e000
uncached-minus @ 0x7fb2e000-0x7fb2f000
uncached-minus @ 0x7fb2f000-0x7fb30000
uncached-minus @ 0x7fb31000-0x7fb32000
uncached-minus @ 0x80000000-0x90000000
```

This list shows physical address ranges and various PAT settings used to access those physical address ranges.

Another, more verbose way of getting PAT related debug messages is with "debugpat" boot parameter. With this parameter, various debug messages are printed to `dmesg` log.

PAT Initialization

The following table describes how PAT is initialized under various configurations. The PAT MSR must be updated by Linux in order

to support WC and WT attributes. Otherwise, the PAT MSR has the value programmed in it by the firmware. Note, Xen enables WC attribute in the PAT MSR for guests.

| MTRR | PAT | Call Sequence | PAT State | PAT MSR |
|------|------|--------------------------|-----------|---------|
| E | E | MTRR -> PAT init | Enabled | OS |
| E | D | MTRR -> PAT init | Disabled | • |
| D | E | MTRR -> PAT disable | Disabled | BIOS |
| D | D | MTRR -> PAT disable | Disabled | • |
| • | np/E | PAT -> PAT disable | Disabled | BIOS |
| • | np/D | PAT -> PAT disable | Disabled | • |
| E | !P/E | MTRR -> PAT init | Disabled | BIOS |
| D | !P/E | MTRR -> PAT disable | Disabled | BIOS |
| !M | !P/E | MTRR stub -> PAT disable | Disabled | BIOS |

Legend

| | |
|----------|---|
| E | Feature enabled in CPU |
| D | Feature disabled/unsupported in CPU |
| np | "nopat" boot option specified |
| !P | CONFIG_X86_PAT option unset |
| !M | CONFIG_MTRR option unset |
| Enabled | PAT state set to enabled |
| Disabled | PAT state set to disabled |
| OS | PAT initializes PAT MSR with OS setting |
| BIOS | PAT keeps PAT MSR with BIOS setting |