

## Image Classification

**Warning:** the features in the `image_classification/` directory have been fully integrated into the new code base.

This folder contains TF 2 model examples for image classification:

- MNIST
- Classifier Trainer, a framework that uses the Keras compile/fit methods for image classification models, including:
  - ResNet
  - EfficientNet<sup>1</sup>

### Before you begin

Please make sure that you have the latest version of TensorFlow installed and add the models folder to your Python path.

### ImageNet preparation

**Using TFDS** `classifier_trainer.py` supports ImageNet with TensorFlow Datasets (TFDS).

Please see the following example snippet for more information on how to use TFDS to download and prepare datasets, and specifically the TFDS ImageNet readme for manual download instructions.

**Legacy TFRecords** Download the ImageNet dataset and convert it to TFRecord format. The following script and README provide a few options.

Note that the legacy ResNet runners, e.g. `resnet/resnet_ctl_imagenet_main.py` require TFRecords whereas `classifier_trainer.py` can use both by setting the builder to 'records' or 'tfds' in the configurations.

### Running on Cloud TPUs

Note: These models will **not** work with TPUs on Colab.

You can train image classification models on Cloud TPUs using `tf.distribute.TPUStrategy`. If you are not familiar with Cloud TPUs, it is strongly recommended that you go through the quickstart to learn how to create a TPU and GCE VM.

### Running on multiple GPU hosts

You can also train these models on multiple hosts, each with GPUs, using `tf.distribute.experimental.MultiWorkerMirroredStrategy`.

---

<sup>1</sup>Currently a work in progress. We cannot match “AutoAugment (AA)” in the original version. For more information about other types of models, please refer to this README file.

The easiest way to run multi-host benchmarks is to set the `TF_CONFIG` appropriately at each host. e.g., to run using `MultiWorkerMirroredStrategy` on 2 hosts, the `cluster` in `TF_CONFIG` should have 2 `host:port` entries, and host `i` should have the `task` in `TF_CONFIG` set to `{"type": "worker", "index": i}`. `MultiWorkerMirroredStrategy` will automatically use all the available GPUs at each host.

## MNIST

To download the data and run the MNIST sample model locally for the first time, run one of the following command:

```
python3 mnist_main.py \
  --model_dir=$MODEL_DIR \
  --data_dir=$DATA_DIR \
  --train_epochs=10 \
  --distribution_strategy=one_device \
  --num_gpus=$NUM_GPUS \
  --download
```

To train the model on a Cloud TPU, run the following command:

```
python3 mnist_main.py \
  --tpu=$TPU_NAME \
  --model_dir=$MODEL_DIR \
  --data_dir=$DATA_DIR \
  --train_epochs=10 \
  --distribution_strategy=tpu \
  --download
```

Note: the `--download` flag is only required the first time you run the model.

## Classifier Trainer

The classifier trainer is a unified framework for running image classification models using Keras's `compile/fit` methods. Experiments should be provided in the form of YAML files, some examples are included within the `configs/examples` folder. Please see `configs/examples` for more example configurations.

The provided configuration files use a per replica batch size and is scaled by the number of devices. For instance, if `batch_size = 64`, then for 1 GPU the global batch size would be  $64 * 1 = 64$ . For 8 GPUs, the global batch size would be  $64 * 8 = 512$ . Similarly, for a v3-8 TPU, the global batch size would be  $64 * 8 = 512$ , and for a v3-32, the global batch size is  $64 * 32 = 2048$ .

## ResNet50

On GPU:

```
python3 classifier_trainer.py \
  --mode=train_and_eval \
  --model_type=resnet \
  --dataset=imagenet \
  --model_dir=$MODEL_DIR \
  --data_dir=$DATA_DIR \
  --config_file=configs/examples/resnet/imagenet/gpu.yaml \
  --params_override='runtime.num_gpus=$NUM_GPUS'
```

To train on multiple hosts, each with GPUs attached using MultiWorkerMirroredStrategy please update `runtime` section in `gpu.yaml` (or override using `--params_override`) with:

```
# gpu.yaml
runtime:
  distribution_strategy: 'multi_worker_mirrored'
  worker_hosts: '$HOST1:port,$HOST2:port'
  num_gpus: $NUM_GPUS
  task_index: 0
```

By having `task_index: 0` on the first host and `task_index: 1` on the second and so on. `$HOST1` and `$HOST2` are the IP addresses of the hosts, and `port` can be chosen any free port on the hosts. Only the first host will write TensorBoard Summaries and save checkpoints.

### On TPU:

```
python3 classifier_trainer.py \
  --mode=train_and_eval \
  --model_type=resnet \
  --dataset=imagenet \
  --tpu=$TPU_NAME \
  --model_dir=$MODEL_DIR \
  --data_dir=$DATA_DIR \
  --config_file=configs/examples/resnet/imagenet/tpu.yaml
```

### VGG-16

#### On GPU:

```
python3 classifier_trainer.py \
  --mode=train_and_eval \
  --model_type=vgg \
  --dataset=imagenet \
  --model_dir=$MODEL_DIR \
  --data_dir=$DATA_DIR \
  --config_file=configs/examples/vgg/imagenet/gpu.yaml \
  --params_override='runtime.num_gpus=$NUM_GPUS'
```

## EfficientNet

**Note:** EfficientNet development is a work in progress. ##### On GPU:

```
python3 classifier_trainer.py \  
  --mode=train_and_eval \  
  --model_type=efficientnet \  
  --dataset=imagenet \  
  --model_dir=$MODEL_DIR \  
  --data_dir=$DATA_DIR \  
  --config_file=configs/examples/efficientnet/imagenet/efficientnet-b0-gpu.yaml \  
  --params_override='runtime.num_gpus=$NUM_GPUS'
```

### On TPU:

```
python3 classifier_trainer.py \  
  --mode=train_and_eval \  
  --model_type=efficientnet \  
  --dataset=imagenet \  
  --tpu=$TPU_NAME \  
  --model_dir=$MODEL_DIR \  
  --data_dir=$DATA_DIR \  
  --config_file=configs/examples/efficientnet/imagenet/efficientnet-b0-tpu.yaml
```

Note that the number of GPU devices can be overridden in the command line using `--params_overrides`. The TPU does not need this override as the device is fixed by providing the TPU address or name with the `--tpu` flag.