

Layers

Layers are the fundamental building blocks for NLP models. They can be used to assemble new `tf.keras` layers or models.

- [MultiHeadAttention](#) implements an optionally masked attention between query, key, value tensors as described in "[Attention Is All You Need](#)". If `from_tensor` and `to_tensor` are the same, then this is self-attention.
- [BigBirdAttention](#) implements a sparse attention mechanism that reduces this quadratic dependency to linear described in "[Big Bird: Transformers for Longer Sequences](#)".
- [CachedAttention](#) implements an attention layer with cache used for auto-aggressive decoding.
- [KernelAttention](#) implements a group of attention mechanisms that express the self-attention as a linear dot-product of kernel feature maps and make use of the associativity property of matrix products to reduce the complexity from quadratic to linear. The implementation includes methods described in "[Transformers are RNNs: Fast Autoregressive Transformers with Linear Attention](#)", "[Rethinking Attention with Performers](#)", "[Random Feature Attention](#)".
- [MatMulWithMargin](#) implements a matrix multiplication with margin layer used for training retrieval / ranking tasks, as described in "[Improving Multilingual Sentence Embedding using Bi-directional Dual Encoder with Additive Margin Softmax](#)".
- [MultiChannelAttention](#) implements an variant of multi-head attention which can be used to merge multiple streams for cross-attentions.
- [TalkingHeadsAttention](#) implements the talking heads attention, as described in "[Talking-Heads Attention](#)".
- [Transformer](#) implements an optionally masked transformer as described in "[Attention Is All You Need](#)".
- [TransformerDecoderBlock](#) TransformerDecoderBlock is made up of self multi-head attention, cross multi-head attention and feedforward network.
- [RandomFeatureGaussianProcess](#) implements random feature-based Gaussian process described in "[Random Features for Large-Scale Kernel Machines](#)".
- [ReuseMultiHeadAttention](#) supports passing attention scores to be reused and avoid recomputation described in "[Leveraging redundancy in attention with Reuse Transformers](#)".
- [ReuseTransformer](#) supports reusing attention scores from lower layers in higher layers to avoid recomputing attention scores described in "[Leveraging redundancy in attention with Reuse Transformers](#)".
- [ReZeroTransformer](#) implements Transformer with ReZero described in "[ReZero is All You Need: Fast Convergence at Large Depth](#)".
- [OnDeviceEmbedding](#) implements efficient embedding lookups designed for TPU-based models.
- [PositionalEmbedding](#) creates a positional embedding as described in "[BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding](#)".
- [SelfAttentionMask](#) creates a 3D attention mask from a 2D tensor mask.
- [SpectralNormalization](#) implements a `tf.Wrapper` that applies spectral normalization regularization to a given layer. See [Spectral Norm Regularization for Improving the Generalizability of Deep Learning](#).

- [MaskedSoftmax](#) implements a softmax with an optional masking input. If no mask is provided to this layer, it performs a standard softmax; however, if a mask tensor is applied (which should be 1 in positions where the data should be allowed through, and 0 where the data should be masked), the output will have masked positions set to approximately zero.
- [MaskedLM](#) implements a masked language model. It assumes the embedding table variable is passed to it.
- [ClassificationHead](#) A pooling head over a sequence of embeddings, commonly used by classification tasks.
- [GaussianProcessClassificationHead](#) A spectral-normalized neural Gaussian process (SNGP)-based classification head as described in "[Simple and Principled Uncertainty Estimation with Deterministic Deep Learning via Distance Awareness](#)".
- [GatedFeedforward](#) implements the gated linear layer feedforward as described in "[GLU Variants Improve Transformer](#)".
- [MultiHeadRelativeAttention](#) implements a variant of multi-head attention with support for relative position encodings as described in "[Transformer-XL: Attentive Language Models Beyond a Fixed-Length Context](#)". This also has extended support for segment-based attention, a re-parameterization introduced in "[XLNet: Generalized Autoregressive Pretraining for Language Understanding](#)".
- [TwoStreamRelativeAttention](#) implements a variant of multi-head relative attention as described in ["XLNet: Generalized Autoregressive Pretraining for Language Understanding"] (<https://arxiv.org/abs/1906.08237>). This takes in a query and content stream and applies self attention.
- [TransformerXL](#) implements Transformer XL introduced in ["Transformer-XL: Attentive Language Models Beyond a Fixed-Length Context"] (<https://arxiv.org/abs/1901.02860>). This contains `TransformerXLBlock`, a block containing either one or two stream relative self-attention as well as subsequent feedforward networks. It also contains `TransformerXL`, which contains attention biases as well as multiple `TransformerXLBlocks`.
- [MobileBertEmbedding](#) and [MobileBertTransformer](#) implement the embedding layer and also transformer layer proposed in the [MobileBERT paper](#).
- [BertPackInputs](#) and [BertTokenizer](#) and [SentencepieceTokenizer](#) implements the layer to tokenize raw text and pack them into the inputs for BERT models.
- [TransformerEncoderBlock](#) implements an optionally masked transformer as described in "[Attention Is All You Need](#)".