

DELF/DELG Training Instructions

This README documents the end-to-end process for training a local and/or global image feature model on the Google Landmarks Dataset v2 (GLDv2). This can be achieved following these steps:

1. Install the DELF Python library.
2. Download the raw images of the GLDv2 dataset.
3. Prepare the training data.
4. Run the training.

The next sections will cover each of these steps in greater detail.

Prerequisites

Clone the TensorFlow Model Garden repository and move into the `models/research/delf/delf/python/trainingfolder`.

```
git clone https://github.com/tensorflow/models.git
cd models/research/delf/delf/python/training
```

Install the DELF Library

To be able to use this code, please follow these instructions to properly install the DELF library.

Download the GLDv2 Training Data

The GLDv2 images are grouped in 3 datasets: TRAIN, INDEX, TEST. Images in each dataset are grouped into `*.tar` files and individually referenced in `*.csv` files containing training metadata and licensing information. The number of `*.tar` files per dataset is as follows:

- TRAIN: 500 files.
- INDEX: 100 files.
- TEST: 20 files.

To download the GLDv2 images, run the `download_dataset.sh` script like in the following example:

```
bash download_dataset.sh 500 100 20
```

The script takes the following parameters, in order:

- The number of image files from the TRAIN dataset to download (maximum 500).
- The number of image files from the INDEX dataset to download (maximum 100).
- The number of image files from the TEST dataset to download (maximum 20).

The script downloads the GLDv2 images under the following directory structure:

- `gldv2_dataset/`
 - `train/` - Contains raw images from the TRAIN dataset.
 - `index/` - Contains raw images from the INDEX dataset.
 - `test/` - Contains raw images from the TEST dataset.

Each of the three folders `gldv2_dataset/train/`, `gldv2_dataset/index/` and `gldv2_dataset/test/` contains the following:

- The downloaded `*.tar` files.
- The corresponding MD5 checksum files, `*.txt`.
- The unpacked content of the downloaded files. (*Images are organized in folders and subfolders based on the first, second and third character in their file name.*)
- The CSV files containing training and licensing metadata of the downloaded images.

Please note that due to the large size of the GLDv2 dataset, the download can take up to 12 hours and up to 1 TB of disk space. In order to save bandwidth and disk space, you may want to start by downloading only the TRAIN dataset, the only one required for the training, thus saving approximately ~95 GB, the equivalent of the INDEX and TEST datasets. To further save disk space, the `.tar` files can be deleted after downloading and unpacking them.*

Prepare the Data for Training

Preparing the data for training consists of creating TFRecord files from the raw GLDv2 images grouped into TRAIN and VALIDATION splits. The training set produced contains only the *clean* subset of the GLDv2 dataset. The CVPR'20 paper introducing the GLDv2 dataset contains a detailed description of the *clean* subset.

Generating the TFRecord files containing the TRAIN and VALIDATION splits of the *clean* GLDv2 subset can be achieved by running the `build_image_dataset.py` script. Assuming that the GLDv2 images have been downloaded to the `gldv2_dataset` folder, the script can be run as follows:

```
python3 build_image_dataset.py \
--train_csv_path=gldv2_dataset/train/train.csv \
--train_clean_csv_path=gldv2_dataset/train/train_clean.csv \
--train_directory=gldv2_dataset/train/*/*/* \
--output_directory=gldv2_dataset/tfrecord/ \
--num_shards=128 \
--generate_train_validation_splits \
--validation_split_size=0.2
```

Please refer to the source code of the `build_image_dataset.py` script for a detailed description of its parameters.

The TFRecord files written in the `OUTPUT_DIRECTORY` will be prefixed as follows:

- TRAIN split: `train-*`
- VALIDATION split: `validation-*`

The same script can be used to generate TFRecord files for the TEST split for post-training evaluation purposes. This can be achieved by adding the parameters:

```
--test_csv_path=gldv2_dataset/train/test.csv \
--test_directory=gldv2_dataset/test/*/*/*/ \
```

In this scenario, the TFRecord files of the TEST split written in the `OUTPUT_DIRECTORY` will be named according to the pattern `test-*`.

Please note that due to the large size of the GLDV2 dataset, the generation of the TFRecord files can take up to 12 hours and up to 500 GB of space disk.

Running the Training

For the training to converge faster, it is possible to initialize the ResNet backbone with the weights of a pretrained ImageNet model. The ImageNet checkpoint is available at the following location: http://storage.googleapis.com/delf/resnet50_imagenet_weights.tar.gz. To download and unpack it run the following commands on a Linux box:

```
curl -Os http://storage.googleapis.com/delf/resnet50_imagenet_weights.tar.gz
tar -xzf resnet50_imagenet_weights.tar.gz
```

Training with Local Features

Assuming the TFRecord files were generated in the `gldv2_dataset/tfrecord/` directory, running the following command should start training a model and output the results in the `gldv2_training` directory:

```
python3 train.py \
--train_file_pattern=gldv2_dataset/tfrecord/train* \
--validation_file_pattern=gldv2_dataset/tfrecord/validation* \
--imagenet_checkpoint=resnet50_weights_tf_dim_ordering_tf_kernels_notop.h5 \
--dataset_version=gld_v2_clean \
--logdir=gldv2_training/
```

NOTE: The `--use_autoencoder` parameter is set by default to `True`, therefore the model will be by default trained with the autoencoder.

Training with Local and Global Features

It is also possible to train the model with an improved global features head as introduced in the DELG paper. To do this, specify the additional parameter `--delg_global_features` when launching the training, like in the following example:

```
python3 train.py \
  --train_file_pattern=gldv2_dataset/tfrecord/train* \
  --validation_file_pattern=gldv2_dataset/tfrecord/validation* \
  --imagenet_checkpoint=resnet50_weights_tf_dim_ordering_tf_kernels_notop.h5 \
  --dataset_version=gld_v2_clean \
  --logdir=gldv2_training/ \
  --delg_global_features
```

NOTE: The `--use_autoencoder` parameter is set by default to `True`, therefore the model will be by default trained with the autoencoder.

Hyperparameter Guidelines

In order to improve the convergence of the training, the following hyperparameter values have been tested and validated on the following infrastructures, the remaining `train.py` flags keeping their **default values**: * 8 Tesla P100 GPUs: `--batch_size=256`, `--initial_lr=0.01` * 4 Tesla P100 GPUs: `--batch_size=128`, `--initial_lr=0.005`

Exporting the Trained Model

Assuming the training output, the TensorFlow checkpoint, is in the `gldv2_training` directory, running the following commands exports the model.

DELF local feature-only model

This should be used when you are only interested in having a local feature model.

```
python3 model/export_local_model.py \
  --ckpt_path=gldv2_training/delf_weights \
  --export_path=gldv2_model_local
```

DELG global feature-only model

This should be used when you are only interested in having a global feature model.

```
python3 model/export_global_model.py \
  --ckpt_path=gldv2_training/delf_weights \
  --export_path=gldv2_model_global \
  --delg_global_features
```

DELG local+global feature model

This should be used when you are interested in jointly extracting local and global features.

```
python3 model/export_local_and_global_model.py \
  --ckpt_path=gldv2_training/delf_weights \
```

```
--export_path=gldv2_model_local_and_global \
--delg_global_features
```

Kaggle-compatible global feature model

To export a global feature model in the format required by the 2020 Landmark Retrieval challenge, you can use the following command:

NOTE: this command is helpful to use the model directly in the above-mentioned Kaggle competition; however, this is a different format than the one required in this DELF/DELG codebase (ie, if you export the model this way, the commands found in the DELG instructions would not work). To export the model in a manner compatible to this codebase, use a similar command as the “DELG global feature-only model” above.

```
python3 model/export_global_model.py \
--ckpt_path=gldv2_training/delf_weights \
--export_path=gldv2_model_global \
--input_scales_list=0.70710677,1.0,1.4142135 \
--multi_scale_pool_type=sum \
--normalize_global_descriptor
```

Testing the trained model

Testing the trained local feature model

After the trained model has been exported, it can be used to extract DELF features from 2 images of the same landmark and to perform a matching test between the 2 images based on the extracted features to validate they represent the same landmark.

Start by downloading the Oxford buildings dataset:

```
mkdir data && cd data
wget http://www.robots.ox.ac.uk/~vgg/data/oxbuildings/oxbuild_images.tgz
mkdir oxford5k_images oxford5k_features
tar -xvzf oxbuild_images.tgz -C oxford5k_images/
cd ../
echo data/oxford5k_images/hertford_000056.jpg >> list_images.txt
echo data/oxford5k_images/oxford_000317.jpg >> list_images.txt
```

Make a copy of the `delf_config_example.pbtxt` protobuffer file which configures the DELF feature extraction. Update the file by making the following changes:

- set the `model_path` attribute to the directory containing the exported model, `gldv2_model_local` in this example
- add at the root level the attribute `is_tf2_exported` with the value `true`
- set to `false` the `use_pca` attribute inside `delf_local_config`

The ensuing file should resemble the following:

```
model_path: "gldv2_model_local"
image_scales: .25
image_scales: .3536
image_scales: .5
image_scales: .7071
image_scales: 1.0
image_scales: 1.4142
image_scales: 2.0
is_tf2_exported: true
delf_local_config {
  use_pca: false
  max_feature_num: 1000
  score_threshold: 100.0
}
```

Run the following command to extract DELF features for the images hertford_000056.jpg and oxford_000317.jpg:

```
python3 ../examples/extract_features.py \
  --config_path delf_config_example.pbtxt \
  --list_images_path list_images.txt \
  --output_dir data/oxford5k_features
```

Run the following command to perform feature matching between the images hertford_000056.jpg and oxford_000317.jpg:

```
python3 ../examples/match_images.py \
  --image_1_path data/oxford5k_images/hertford_000056.jpg \
  --image_2_path data/oxford5k_images/oxford_000317.jpg \
  --features_1_path data/oxford5k_features/hertford_000056.delf \
  --features_2_path data/oxford5k_features/oxford_000317.delf \
  --output_image matched_images.png
```

The generated image matched_images.png should look similar to this one:

MatchedImagesDemo

Testing the trained global (or global+local) feature model

Please follow the DELG instructions. The only modification should be to pass a different `delf_config_path` when doing feature extraction, which should point to the newly-trained model. As described in the `DelfConfig`, you should set the `use_local_features` and `use_global_features` in the right way, depending on which feature modalities you are using. Note also that you should set `is_tf2_exported` to `true`.