

Frequently Asked Questions

Q: How can I ensure that all the groundtruth boxes are used during train and eval?

A: For the object detection framework to be TPU-compliant, we must pad our input tensors to static shapes. This means that we must pad to a fixed number of bounding boxes, configured by `InputReader.max_number_of_boxes`. It is important to set this value to a number larger than the maximum number of groundtruth boxes in the dataset. If an image is encountered with more bounding boxes, the excess boxes will be clipped.

Q: AttributeError: ‘module’ object has no attribute ‘BackupHandler’

A: This BackupHandler (`tf_slim.tfexample_decoder.BackupHandler`) was introduced in tensorflow 1.5.0 so running with earlier versions may cause this issue. It now has been replaced by `object_detection.data_decoders.tf_example_decoder.BackupHandler`. Whoever sees this issue should be able to resolve it by syncing your fork to HEAD. Same for LookupTensor.

Q: AttributeError: ‘module’ object has no attribute ‘LookupTensor’

A: Similar to BackupHandler, syncing your fork to HEAD should make it work.

Q: Why can’t I get the inference time as reported in model zoo?

A: The inference time reported in model zoo is mean time of testing hundreds of images with an internal machine. As mentioned in TensorFlow detection model zoo, this speed depends highly on one’s specific hardware configuration and should be treated more as relative timing.