

Intel's Deep Learning Inference Engine (DL IE) is a part of Intel® OpenVINO™ toolkit. You can use it as a computational backend for OpenCV deep learning module.

To use OpenCV with Inference Engine, choose one of the options: \* Intel® OpenVINO™ toolkit - includes ready to use build of OpenCV \* OpenCV+OpenVINO Windows package (community version) \* Build from source \* Linux \* Microsoft Windows \* Raspbian

---

### Intel® OpenVINO™ toolkit

- Download and install Intel® OpenVINO™ toolkit.

**Important note:** if you want to transfer the installed Inference Engine binaries to another machine w/o running OpenVINO installer there, you need the redistributable files of Intel C++ compiler (use the latest update, 64-bit version), otherwise the Inference Engine or some of its essential plugins will refuse to load and run, which may result in an app crash.

- To perform deep neural networks inference on ARM CPUs, build ARM CPU Plugin from `openvino_contrib`. ARM CPU Plugin isn't distributed in the package. This way requires building OpenVINO and OpenCV from source. ARM CPU Plugin build instructions.
  - OpenVINO release tag 2021.3
  - OpenVINO contrib release tag 2021.3

---

### ### OpenCV+OpenVINO Windows package (community version)

This community package uses open source version of Inference Engine from Deep Learning Deployment Toolkit repository (distributed under Apache 2 license).

Hardware requirements: - CPU with support of AVX2 instruction set - *[optional]* Intel® Integrated Graphics - *[optional]* Intel® Neural Compute Stick 2

Software requirements: - Windows\* 10, 64-bit - *[recommended]* Microsoft\* Visual Studio\* 2019 - or MSVS\* 2019 redistributables: `vc_redist.x64.exe` - Python 3.x (64-bit) to use OpenCV Python bindings - CMake 3.5.1+ (MSVS 2015+) or CMake 3.14+ (MSVS 2019+)

### OpenCV 4.5.5

OpenVINO version	GitHub releases	SourceForge
2021.4.2 (release binaries)	.7z (164.5 Mb) / .zip (253.4 Mb)	.7z (164.5 Mb) / .zip (253.4 Mb)
2021.4.2 (debug binaries)	.7z (196.8 Mb)	.7z (196.8 Mb)

DLDT components included in this release:

- IE MKLDNN plugin (CPU)
- IE clDNN plugin (GPU-OpenCL)
- IE Myriad plugin (VPU)
- IE Hetero plugin
- Intel® Threading Building Blocks (TBB)
- nGraph

Both package versions include .PDB files for better debugging experience (generated by MSVS 2019). Package with debug binaries is intended for development of C++ applications only (missing Python support, etc).

Archive

#### OpenCV 4.5.4

OpenVINO version	GitHub releases	SourceForge
2021.4.1 (release binaries)	.7z (~164.2 Mb) / .zip (~252.7Mb)	.7z (~164.2Mb) / .zip (~252.7Mb)
2021.4.1 (debug binaries)	.7z (~195.6Mb)	.7z (~195.6Mb)

DLDT components included in this release:

- IE MKLDNN plugin (CPU)
- IE clDNN plugin (GPU-OpenCL)
- IE Myriad plugin (VPU)
- IE Hetero plugin
- Intel® Threading Building Blocks (TBB)
- nGraph

Both package versions include .PDB files for better debugging experience (generated by MSVS 2019). Package with debug binaries is intended for development of C++ applications only (missing Python support, etc).

### OpenCV 4.5.3

OpenVINO version	GitHub releases	SourceForge
2021.4 (release binaries)	.7z (~166.9Mb) / .zip (~257.4Mb)	.7z (~166.9Mb) / .zip (~257.4Mb)
2021.4 (debug binaries)	.7z (~199.7Mb)	.7z (~199.7Mb)

DLDT components included in this release:

- IE MKLDNN plugin (CPU)
- IE cldNN plugin (GPU-OpenCL)
- IE Myriad plugin (VPU)
- IE Hetero plugin
- Intel® Threading Building Blocks (TBB)
- nGraph

Both package versions include .PDB files for better debugging experience (generated by MSVS 2019). Package with debug binaries is intended for development of C++ applications only (missing Python support, etc).

### OpenCV 4.5.2

DLDT version	GitHub releases	SourceForge
2021.3 (release binaries)	.7z (~161.1Mb) / .zip (~246Mb)	.7z (~161.1Mb) / .zip (~246Mb)
2021.3 (debug binaries)	.7z (~192.5Mb)	.7z (~192.5Mb)

DLDT components included in this release:

- IE MKLDNN plugin (CPU)
- IE cldNN plugin (GPU-OpenCL)
- IE Myriad plugin (VPU)
- Intel® Threading Building Blocks (TBB)
- nGraph

Both package versions include .PDB files for better debugging experience (generated by MSVS 2019). Package with debug binaries is intended for development of C++ applications only (missing Python support, etc).

## OpenCV 4.5.1

DLDT version	GitHub releases	SourceForge
2021.2 (release binaries)	.7z (~152.9Mb) / .zip (~230Mb)	.7z (~152.9Mb) / .zip (~230Mb)
2021.2 (debug binaries)	.7z (~185.5Mb)	.7z (~185.5Mb)

DLDT components included in this release:

- IE MKLDNN plugin (CPU)
- IE cldNN plugin (GPU-OpenCL)
- IE Myriad plugin (VPU)
- Intel® Threading Building Blocks (TBB)
- nGraph

Both package versions include .PDB files for better debugging experience (generated by MSVS 2019). Package with debug binaries is intended for development of C++ applications only (missing Python support, etc).

## OpenCV 4.5.0

DLDT version	GitHub releases	SourceForge
2021.1 (release binaries)	.7z (~154Mb) / .zip (~232Mb)	.7z (~154Mb) / .zip (~232Mb)
2021.1 (debug binaries)	.7z (~183Mb)	.7z (~183Mb)

DLDT components included in this release:

- IE MKLDNN plugin (CPU)
- IE cldNN plugin (GPU-OpenCL)
- IE Myriad plugin (VPU)
- Intel® Threading Building Blocks (TBB)
- nGraph

Both package versions include .PDB files for better debugging experience (generated by MSVS 2019). Package with debug binaries is intended for development of C++ applications only (missing Python support, etc).

## OpenCV 4.4.0

DLDT version	GitHub releases	SourceForge
2020.4 (release binaries)	.7z (~155Mb) / .zip (~236Mb)	.7z (~155Mb) / .zip (~236Mb)
2020.4 (debug binaries)	.7z (~184Mb)	.7z (~184Mb)
2020.3 (release binaries)	.7z (~152Mb) / .zip (~231Mb)	.7z (~152Mb) / .zip (~231Mb)
2020.3 (debug binaries)	.7z (~179Mb)	.7z (~179Mb)

DLDT components included in this release:

- IE MKLDNN plugin (CPU)
- IE cldNN plugin (GPU-OpenCL)
- IE Myriad plugin (VPU)
- Intel® Threading Building Blocks (TBB)
- nGraph

Both package versions include .PDB files for better debugging experience (generated by MSVS 2019). Package with debug binaries is intended for development of C++ applications only (missing Python support, etc).

## OpenCV 4.3.0

DLDT version	GitHub releases	SourceForge
2020.2 (release binaries)	.7z (~150Mb) / .zip (~220Mb)	.7z (~150Mb) / .zip (~220Mb)
2020.2 (debug binaries)	.7z (~175Mb)	.7z (~175Mb)
2020.1 (release binaries)	.7z (~150Mb) / .zip (~220Mb)	.7z (~150Mb) / .zip (~220Mb)
2020.1 (debug binaries)	.7z (~175Mb)	.7z (~175Mb)

DLDT components included in this release:

- IE MKLDNN plugin (CPU)
- IE cldNN plugin (GPU-OpenCL)
- IE Myriad plugin (VPU)
- Intel® Threading Building Blocks (TBB)

- nGraph

Both package versions include .PDB files for better debugging experience (generated by MSVS 2019). Package with debug binaries is intended for development of C++ applications only (missing Python support, etc).

**First steps** Steps below require package with release binaries. Steps for debug binaries might differ.

1. Extract package contents

- Unpack .7z archive using 7-Zip: <https://www.7-zip.org/download.html>
- .zip archive can be extracted using Windows 10 builtin tools (check content menu in explorer)

2. Run OpenCV Python samples (can be launched from explorer):

```
<package_root>\src\opencv\samples\python\_run_winpack_demo.cmd
```

**Note:** OpenCV DNN Python samples are located in `samples\dnn` directory

3. Open console terminal and configure OpenCV environment:

```
CALL <package_root>\build\setup_vars_opencv4.cmd
opencv_version
```

For better experience consider running MSVS instance from the terminal with configured environment (type `start MyProject.sln`)

4. Run `opencv_version`:

```
<package_root>\build\bin\opencv_version_win32.exe
```

or execute in the terminal:

```
<package_root>\build\setup_vars_opencv4.cmd opencv_version -v
```

5. Run Python interpreter (python.exe should be in PATH):

```
<package_root>\build\setup_vars_opencv4.cmd python
> import cv2 as cv
> print(cv.getBuildInformation())
```

**Note:** Python support is missing in the package with “Debug” binaries. Use package with “Release” binaries instead.

6. Build and run C++ sample:

- open `src\opencv\samples` in explorer (unmaximize window)
- open `src\opencv\samples\cpp` in explorer (unmaximize window)
- drag&drop `drawing.cpp` file from `cpp` directory onto `_winpack_build_sample.cmd` script

- or execute in the terminal:

```
<package_root>\src\opencv\samples\_winpack_build_sample.cmd <package_root>\src\oper
```

---

## Build OpenCV from source

Build OpenCV with pre-built Inference Engine binaries from OpenVINO toolkit.

**Linux** Setup environment variables to detect Inference Engine:

```
source /opt/intel/openvino/bin/setupvars.sh
export ngraph_DIR=/opt/intel/openvino/deployment_tools/ngraph/cmake/
```

Build OpenCV with extra flags:

```
cmake \
  -DWITH_INF_ENGINE=ON \
  -DENABLE_CXX11=ON \
  ...
```

**Microsoft Windows** Setup environment variables to detect Inference Engine:

```
"C:\Program Files (x86)\IntelSWTools\openvino\bin\setupvars.bat"
```

Build OpenCV with extra flags:

```
cmake ^
  -DWITH_INF_ENGINE=ON ^
  -DENABLE_CXX11=ON ^
  ...
```

**Raspbian Buster** Use Docker to cross-compile OpenCV for Raspberry Pi. Check that `uname -m` detects `armv7l` CPU architecture (starts from Raspberry Pi 2 model B).

1. Create a folder named `debian_armhf` with a file `Dockerfile` with the following content:

```
FROM debian:buster
```

```
USER root
```

```
RUN dpkg --add-architecture armhf && \
  apt-get update && \
  apt-get install -y --no-install-recommends \
  crossbuild-essential-armhf \
  cmake \
  pkg-config \
```

```
wget \
xz-utils \
libgtk2.0-dev:armhf \
libpython-dev:armhf \
libpython3-dev:armhf \
python-numpy \
python3-numpy \
libgstreamer1.0-dev:armhf \
libgstreamer-plugins-base1.0-dev:armhf
```

# Install Inference Engine

```
RUN wget --no-check-certificate https://download.01.org/opencv/2019/openvinotoolkit/R2/l_op
tar -xf l_openvino_toolkit_runtime_raspbian_p_2019.2.242.tgz
```

2. Build a Docker image

```
docker image build -t debian_armhf debian_armhf
```

3. Run Docker container mounting source code folder from host.

```
docker run -it -v /absolute/path/to/opencv:/opencv debian_armhf /bin/bash
```

4. Build

```
cd opencv && mkdir opencv_build && mkdir opencv_install && cd opencv_build
```

```
cmake -DCMAKE_BUILD_TYPE=Release \
-DMAKE_INSTALL_PREFIX="./opencv_install" \
-DOPENCV_CONFIG_INSTALL_PATH="cmake" \
-DCMAKE_TOOLCHAIN_FILE="./platforms/linux/arm-gnueabi.toolchain.cmake" \
-DWITH_IPP=OFF \
-DBUILD_TESTS=OFF \
-DBUILD_PERF_TESTS=OFF \
-DOPENCV_ENABLE_PKG_CONFIG=ON \
-DPKG_CONFIG_EXECUTABLE="/usr/bin/arm-linux-gnueabi-hf-pkg-config" \
-DPYTHON2_INCLUDE_PATH="/usr/include/python2.7" \
-DPYTHON2_NUMPY_INCLUDE_DIRS="/usr/local/lib/python2.7/dist-packages/numpy/core/include" \
-DPYTHON3_INCLUDE_PATH="/usr/include/python3.7" \
-DPYTHON3_NUMPY_INCLUDE_DIRS="/usr/local/lib/python3.7/dist-packages/numpy/core/include" \
-DPYTHON3_CVPY_SUFFIX=".cpython-37m-arm-linux-gnueabi-hf.so" \
-DENABLE_NEON=ON \
-DCPU_BASELINE="NEON" \
-DWITH_INF_ENGINE=ON \
-DINF_ENGINE_LIB_DIRS="/l_openvino_toolkit_runtime_raspbian_p_2019.2.242/inference_engine" \
-DINF_ENGINE_INCLUDE_DIRS="/l_openvino_toolkit_runtime_raspbian_p_2019.2.242/inference_engine" \
-DCMAKE_FIND_ROOT_PATH="/l_openvino_toolkit_runtime_raspbian_p_2019.2.242" \
-DENABLE_CXX11=ON ..
make -j4 && make install
```

5. Copy opencv\_install to the board. Follow [http://docs.openvinotoolkit.org/latest/\\_docs\\_install\\_guides\\_](http://docs.openvinotoolkit.org/latest/_docs_install_guides_)



to install OpenVINO distribution for Raspberry Pi. Then type the following commands to specify new location of OpenCV:

```
export PYTHONPATH=/path/to/opencv_install/lib/python2.7/dist-packages/:$PYTHONPATH
export PYTHONPATH=/path/to/opencv_install/lib/python3.7/dist-packages/:$PYTHONPATH
export LD_LIBRARY_PATH=/path/to/opencv_install/lib/:$LD_LIBRARY_PATH
```

---

## Usage

- Enable Intel's Inference Engine backend right after `cv::dnn::readNet` invocation:

```
net.setPreferableBackend(DNN_BACKEND_INFERENCE_ENGINE);
```

**NOTE:** Starts from OpenCV 3.4.2 (OpenVINO 2018.R2) this backend is used by default if OpenCV is built with the Inference Engine support. To switch to origin implementation, use `DNN_BACKEND_OPENCV`. Also, the Inference engine backend is the only available option when the loaded model is represented in OpenVINO™ Model Optimizer format (`.bin` and `.xml`).

- Then, optionally you can also set the device to use for the inference (by default it will use CPU):

```
net.setPreferableTarget(DNN_TARGET_OPENCL);
// the possible options are
// DNN_TARGET_CPU,
// DNN_TARGET_OPENCL,
// DNN_TARGET_OPENCL_FP16
// (fall back to OPENCL if the hardware does not support FP16),
// DNN_TARGET_MYRIAD
```

- You may also import pre-trained models from Open Model Zoo passing paths to `.bin` and `.xml` files to `cv::dnn::readNet` function.
- 

- Other names and brands may be claimed as the property of others.