

Distributed Data Parallel Benchmark

This tool is used to measure distributed training iteration time. This is helpful for evaluating the performance impact of code changes to `torch.nn.parallel.DistributedDataParallel`, `torch.distributed`, or anything in between.

It optionally produces a JSON file with all measurements, allowing for an easy A/B comparison of code, configuration, or environment. This comparison can be produced by `diff.py`.

Requirements

This benchmark depends on PyTorch and torchvision.

How to run

Run as many copies of this script as you have model replicas.

If you launch a single task per machine with multiple GPUs, consider using `torch.distributed.launch` to spawn multiple processes per machine.

Example output (only on rank 0):

```
-----  
PyTorch distributed benchmark suite  
-----
```

```
* PyTorch version: 1.4.0a0+05140f0  
* CUDA version: 10.0  
* Distributed backend: nccl
```

```
--- nvidia-smi topo -m ---
```

	GPU0	GPU1	GPU2	GPU3	GPU4	GPU5	GPU6	GPU7	mlx5_2	mlx5_0	mlx5_3
GPU0	X	NV1	NV1	NV2	NV2	SYS	SYS	SYS	SYS	PIX	SYS
GPU1	NV1	X	NV2	NV1	SYS	NV2	SYS	SYS	SYS	PIX	SYS
GPU2	NV1	NV2	X	NV2	SYS	SYS	NV1	SYS	SYS	PHB	SYS
GPU3	NV2	NV1	NV2	X	SYS	SYS	SYS	NV1	SYS	PHB	SYS
GPU4	NV2	SYS	SYS	SYS	X	NV1	NV1	NV2	PIX	SYS	PHB
GPU5	SYS	NV2	SYS	SYS	NV1	X	NV2	NV1	PIX	SYS	PHB
GPU6	SYS	SYS	NV1	SYS	NV1	NV2	X	NV2	PHB	SYS	PIX
GPU7	SYS	SYS	SYS	NV1	NV2	NV1	NV2	X	PHB	SYS	PIX
mlx5_2	SYS	SYS	SYS	SYS	PIX	PIX	PHB	PHB	X	SYS	PHB
mlx5_0	PIX	PIX	PHB	PHB	SYS	SYS	SYS	SYS	SYS	X	SYS
mlx5_3	SYS	SYS	SYS	SYS	PHB	PHB	PIX	PIX	PHB	SYS	X
mlx5_1	PHB	PHB	PIX	PIX	SYS	SYS	SYS	SYS	SYS	PHB	SYS

Legend:

X = Self
 SYS = Connection traversing PCIe as well as the SMP interconnect between NUMA nodes (e.g.
 NODE = Connection traversing PCIe as well as the interconnect between PCIe Host Bridges w
 PHB = Connection traversing PCIe as well as a PCIe Host Bridge (typically the CPU)
 PXB = Connection traversing multiple PCIe switches (without traversing the PCIe Host Brid
 PIX = Connection traversing a single PCIe switch
 NV# = Connection traversing a bonded set of # NVLinks

 Benchmark: resnet50 with batch size 32

			sec/iter	ex/sec		sec/iter	ex/sec		sec/iter	ex/sec
1 GPUs --	no ddp:	p50:	0.097s	329/s	p75:	0.097s	329/s	p90:	0.097s	329/s
1 GPUs --	1M/1G:	p50:	0.100s	319/s	p75:	0.100s	318/s	p90:	0.100s	318/s
2 GPUs --	1M/2G:	p50:	0.103s	310/s	p75:	0.103s	310/s	p90:	0.103s	310/s
4 GPUs --	1M/4G:	p50:	0.103s	310/s	p75:	0.103s	310/s	p90:	0.103s	310/s
8 GPUs --	1M/8G:	p50:	0.104s	307/s	p75:	0.104s	307/s	p90:	0.104s	306/s
16 GPUs --	2M/8G:	p50:	0.104s	306/s	p75:	0.104s	306/s	p90:	0.104s	306/s

Benchmark: resnet101 with batch size 32

			sec/iter	ex/sec		sec/iter	ex/sec		sec/iter	ex/sec
1 GPUs --	no ddp:	p50:	0.162s	197/s	p75:	0.162s	197/s	p90:	0.162s	197/s
1 GPUs --	1M/1G:	p50:	0.171s	187/s	p75:	0.171s	186/s	p90:	0.171s	186/s
2 GPUs --	1M/2G:	p50:	0.176s	182/s	p75:	0.176s	181/s	p90:	0.176s	181/s
4 GPUs --	1M/4G:	p50:	0.176s	182/s	p75:	0.176s	181/s	p90:	0.176s	181/s
8 GPUs --	1M/8G:	p50:	0.179s	179/s	p75:	0.179s	178/s	p90:	0.180s	178/s
16 GPUs --	2M/8G:	p50:	0.179s	178/s	p75:	0.180s	177/s	p90:	0.183s	174/s

Benchmark: resnext50_32x4d with batch size 32

			sec/iter	ex/sec		sec/iter	ex/sec		sec/iter	ex/sec
1 GPUs --	no ddp:	p50:	0.145s	220/s	p75:	0.145s	220/s	p90:	0.145s	220/s
1 GPUs --	1M/1G:	p50:	0.147s	217/s	p75:	0.147s	217/s	p90:	0.148s	216/s
2 GPUs --	1M/2G:	p50:	0.153s	209/s	p75:	0.153s	209/s	p90:	0.153s	209/s
4 GPUs --	1M/4G:	p50:	0.153s	208/s	p75:	0.153s	208/s	p90:	0.154s	208/s
8 GPUs --	1M/8G:	p50:	0.157s	204/s	p75:	0.157s	204/s	p90:	0.157s	203/s
16 GPUs --	2M/8G:	p50:	0.157s	203/s	p75:	0.157s	203/s	p90:	0.158s	203/s

Benchmark: resnext101_32x8d with batch size 32

			sec/iter	ex/sec		sec/iter	ex/sec		sec/iter	ex/sec
--	--	--	----------	--------	--	----------	--------	--	----------	--------

1 GPUs	--	no ddp:	p50:	0.415s	77/s	p75:	0.415s	77/s	p90:	0.416s	76
1 GPUs	--	1M/1G:	p50:	0.425s	75/s	p75:	0.426s	75/s	p90:	0.426s	75
2 GPUs	--	1M/2G:	p50:	0.438s	73/s	p75:	0.439s	72/s	p90:	0.439s	72
4 GPUs	--	1M/4G:	p50:	0.439s	72/s	p75:	0.439s	72/s	p90:	0.440s	72
8 GPUs	--	1M/8G:	p50:	0.447s	71/s	p75:	0.447s	71/s	p90:	0.448s	71
16 GPUs	--	2M/8G:	p50:	0.450s	71/s	p75:	0.451s	70/s	p90:	0.451s	70

How to diff

Run the benchmark with the `--json PATH_TO_REPORT_FILE` argument to produce the JSON file that the diff script can consume.

Then, run the diff script as follows:

```
$ python3 diff.py PATH_TO_BASELINE_FILE PATH_TO_TEST_FILE
```

	baseline		test
	-----		-----
bucket_size:	25	vs	1
cuda_version:	10.0	vs	10.0
distributed_backend:	nccl	vs	nccl
pytorch_version:	1.4.0a0+05140f0	vs	1.4.0a0+05140f0

Benchmark: resnet50 with batch size 32

		sec/iter	ex/sec	diff		sec/iter	ex/sec	diff
1 GPUs:	p75:	0.101s	317/s	-0.3%	p95:	0.101s	317/s	-0.4%
2 GPUs:	p75:	0.104s	306/s	-1.0%	p95:	0.104s	306/s	-1.0%
4 GPUs:	p75:	0.105s	305/s	-1.6%	p95:	0.105s	304/s	-1.8%
8 GPUs:	p75:	0.107s	299/s	-2.6%	p95:	0.107s	298/s	-2.7%
16 GPUs:	p75:	0.108s	294/s	-3.8%	p95:	0.122s	262/s	-16.4%

Benchmark: resnet101 with batch size 32

		sec/iter	ex/sec	diff		sec/iter	ex/sec	diff
1 GPUs:	p75:	0.172s	185/s	-1.2%	p95:	0.172s	185/s	-1.3%
2 GPUs:	p75:	0.179s	178/s	-2.1%	p95:	0.179s	178/s	-2.0%
4 GPUs:	p75:	0.180s	177/s	-2.6%	p95:	0.180s	177/s	-2.6%
8 GPUs:	p75:	0.184s	173/s	-3.5%	p95:	0.184s	173/s	-3.5%
16 GPUs:	p75:	0.187s	170/s	-0.1%	p95:	0.204s	157/s	-7.9%

Benchmark: resnext50_32x4d with batch size 32

		sec/iter	ex/sec	diff		sec/iter	ex/sec	diff
1 GPUs:	p75:	0.149s	214/s	-1.0%	p95:	0.149s	214/s	-0.9%
2 GPUs:	p75:	0.156s	205/s	-1.5%	p95:	0.156s	205/s	-1.6%
4 GPUs:	p75:	0.156s	204/s	-1.6%	p95:	0.157s	204/s	-1.8%
8 GPUs:	p75:	0.159s	200/s	-1.5%	p95:	0.159s	200/s	-1.5%

16 GPUs:	p75:	0.161s	198/s	-1.9%	p95:	0.162s	197/s	-2.3%
----------	------	--------	-------	-------	------	--------	-------	-------

Benchmark: resnext101_32x8d with batch size 32

		sec/iter	ex/sec	diff		sec/iter	ex/sec	diff
1 GPUs:	p75:	0.427s	74/s	-0.8%	p95:	0.428s	74/s	-0.7%
2 GPUs:	p75:	0.444s	72/s	-1.3%	p95:	0.445s	71/s	-0.7%
4 GPUs:	p75:	0.444s	72/s	-1.1%	p95:	0.445s	71/s	-0.8%
8 GPUs:	p75:	0.452s	70/s	-1.3%	p95:	0.452s	70/s	-1.3%
16 GPUs:	p75:	0.455s	70/s	-0.7%	p95:	0.456s	70/s	-0.6%

This compares throughput between `bucket_cap_mb=25` (the default) and `bucket_cap_mb=1` on 8 DGX machines with V100 GPUs. It confirms that even for a relatively small model on machines with a very fast interconnect (4x 100Gb InfiniBand per machine), it still pays off to batch allreduce calls.