

# torch.optim

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\pytorch-master) (docs) (source) optim.rst, line 4)**

Unknown directive type "automodule".

```
.. automodule:: torch.optim
```

## How to use an optimizer

To use `mod:torch.optim` you have to construct an optimizer object, that will hold the current state and will update the parameters based on the computed gradients.

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\pytorch-master) (docs) (source) optim.rst, line 9); [backlink](#)**

Unknown interpreted text role "mod".

## Constructing it

To construct an `class:Optimizer` you have to give it an iterable containing the parameters (all should be `class:~torch.autograd.Variable` s) to optimize. Then, you can specify optimizer-specific options such as the learning rate, weight decay, etc.

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\pytorch-master) (docs) (source) optim.rst, line 15); [backlink](#)**

Unknown interpreted text role "class".

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\pytorch-master) (docs) (source) optim.rst, line 15); [backlink](#)**

Unknown interpreted text role "class".

### Note

If you need to move a model to GPU via `.cuda()`, please do so before constructing optimizers for it. Parameters of a model after `.cuda()` will be different objects with those before the call.

In general, you should make sure that optimized parameters live in consistent locations when optimizers are constructed and used.

Example:

```
optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.9)
optimizer = optim.Adam([var1, var2], lr=0.0001)
```

## Per-parameter options

`class:Optimizer` s also support specifying per-parameter options. To do this, instead of passing an iterable of `class:~torch.autograd.Variable` s, pass in an iterable of `class:dict` s. Each of them will define a separate parameter group, and should contain a `params` key, containing a list of parameters belonging to it. Other keys should match the keyword arguments accepted by the optimizers, and will be used as optimization options for this group.

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\pytorch-master) (docs) (source) optim.rst, line 36); [backlink](#)**

Unknown interpreted text role "class".

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\pytorch-master) (docs) (source) optim.rst, line 36); [backlink](#)**

Unknown interpreted text role "class".

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\ (pytorch-master) (docs) (source) optim.rst, line 36); [backlink](#)**

Unknown interpreted text role "class".

#### Note

You can still pass options as keyword arguments. They will be used as defaults, in the groups that didn't override them. This is useful when you only want to vary a single option, while keeping all others consistent between parameter groups.

For example, this is very useful when one wants to specify per-layer learning rates:

```
optim.SGD([
    {'params': model.base.parameters()},
    {'params': model.classifier.parameters(), 'lr': 1e-3}
], lr=1e-2, momentum=0.9)
```

This means that `model.base`'s parameters will use the default learning rate of `1e-2`, `model.classifier`'s parameters will use a learning rate of `1e-3`, and a momentum of `0.9` will be used for all parameters.

## Taking an optimization step

All optimizers implement a `.func:~Optimizer.step` method, that updates the parameters. It can be used in two ways:

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\ (pytorch-master) (docs) (source) optim.rst, line 65); [backlink](#)**

Unknown interpreted text role "func".

`optimizer.step()`

This is a simplified version supported by most optimizers. The function can be called once the gradients are computed using e.g. `.func:~torch.autograd.Variable.backward`.

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\ (pytorch-master) (docs) (source) optim.rst, line 71); [backlink](#)**

Unknown interpreted text role "func".

Example:

```
for input, target in dataset:
    optimizer.zero_grad()
    output = model(input)
    loss = loss_fn(output, target)
    loss.backward()
    optimizer.step()
```

`optimizer.step(closure)`

Some optimization algorithms such as Conjugate Gradient and LBFGS need to reevaluate the function multiple times, so you have to pass in a closure that allows them to recompute your model. The closure should clear the gradients, compute the loss, and return it.

Example:

```
for input, target in dataset:
    def closure():
        optimizer.zero_grad()
        output = model(input)
        loss = loss_fn(output, target)
        loss.backward()
        return loss
    optimizer.step(closure)
```

## Base class

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\ (pytorch-master) (docs) (source) optim.rst, line 108)**

Unknown directive type "autoclass".

```
.. autoclass:: Optimizer
```

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\ (pytorch-master) (docs) (source) optim.rst, line 110)**

Unknown directive type "autosummary".

```
.. autosummary::
    :toctree: generated
    :nosignatures:

    Optimizer.add_param_group
    Optimizer.load_state_dict
    Optimizer.state_dict
    Optimizer.step
    Optimizer.zero_grad
```

## Algorithms

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\ (pytorch-master) (docs) (source) optim.rst, line 123)**

Unknown directive type "autosummary".

```
.. autosummary::
    :toctree: generated
    :nosignatures:

    Adadelta
    Adagrad
    Adam
    AdamW
    SparseAdam
    Adamax
    ASGD
    LBFGS
    NAdam
    RAdam
    RMSprop
    Rprop
    SGD
```

## How to adjust learning rate

`mod:torch.optim.lr_scheduler` provides several methods to adjust the learning rate based on the number of epochs.  
`class:torch.optim.lr_scheduler.ReduceLROnPlateau` allows dynamic learning rate reducing based on some validation measurements.

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\ (pytorch-master) (docs) (source) optim.rst, line 144); [backlink](#)**

Unknown interpreted text role "mod".

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\ (pytorch-master) (docs) (source) optim.rst, line 144); [backlink](#)**

Unknown interpreted text role "class".

Learning rate scheduling should be applied after optimizer's update; e.g., you should write your code this way:

Example:

```
model = [Parameter(torch.randn(2, 2, requires_grad=True))]
optimizer = SGD(model, 0.1)
scheduler = ExponentialLR(optimizer, gamma=0.9)

for epoch in range(20):
    for input, target in dataset:
        optimizer.zero_grad()
        output = model(input)
        loss = loss_fn(output, target)
        loss.backward()
```

```
optimizer.step()
scheduler.step()
```

Most learning rate schedulers can be called back-to-back (also referred to as chaining schedulers). The result is that each scheduler is applied one after the other on the learning rate obtained by the one preceding it.

Example:

```
model = [Parameter(torch.randn(2, 2, requires_grad=True))]
optimizer = SGD(model, 0.1)
scheduler1 = ExponentialLR(optimizer, gamma=0.9)
scheduler2 = MultiStepLR(optimizer, milestones=[30,80], gamma=0.1)

for epoch in range(20):
    for input, target in dataset:
        optimizer.zero_grad()
        output = model(input)
        loss = loss_fn(output, target)
        loss.backward()
        optimizer.step()
    scheduler1.step()
    scheduler2.step()
```

In many places in the documentation, we will use the following template to refer to schedulers algorithms.

```
>>> scheduler = ...
>>> for epoch in range(100):
>>>     train(...)
>>>     validate(...)
>>>     scheduler.step()
```

### Warning

Prior to PyTorch 1.1.0, the learning rate scheduler was expected to be called before the optimizer's update; 1.1.0 changed this behavior in a BC-breaking way. If you use the learning rate scheduler (calling `scheduler.step()`) before the optimizer's update (calling `optimizer.step()`), this will skip the first value of the learning rate schedule. If you are unable to reproduce results after upgrading to PyTorch 1.1.0, please check if you are calling `scheduler.step()` at the wrong time.

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\ (pytorch-master) (docs) (source) optim.rst, line 205)**

Unknown directive type "autosummary".

```
.. autosummary::
   :toctree: generated
   :nosignatures:

   lr_scheduler.LambdaLR
   lr_scheduler.MultiplicativeLR
   lr_scheduler.StepLR
   lr_scheduler.MultiStepLR
   lr_scheduler.ConstantLR
   lr_scheduler.LinearLR
   lr_scheduler.ExponentialLR
   lr_scheduler.CosineAnnealingLR
   lr_scheduler.ChainedScheduler
   lr_scheduler.SequentialLR
   lr_scheduler.ReduceLROnPlateau
   lr_scheduler.CyclicLR
   lr_scheduler.OneCycleLR
   lr_scheduler.CosineAnnealingWarmRestarts
```

## Stochastic Weight Averaging

`mod:torch.optim.swa_utils` implements Stochastic Weight Averaging (SWA). In particular, `class:torch.optim.swa_utils.AveragedModel` class implements SWA models, `class:torch.optim.swa_utils.SWALR` implements the SWA learning rate scheduler and `func:torch.optim.swa_utils.update_bn` is a utility function used to update SWA batch normalization statistics at the end of training.

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\ (pytorch-master) (docs) (source) optim.rst, line 227); [backlink](#)**

Unknown interpreted text role "mod".

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\ (pytorch-master) (docs) (source) optim.rst, line 227); [backlink](#)**

Unknown interpreted text role "class".

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\ (pytorch-master) (docs) (source) optim.rst, line 227); [backlink](#)**

Unknown interpreted text role "class".

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\ (pytorch-master) (docs) (source) optim.rst, line 227); [backlink](#)**

Unknown interpreted text role "func".

SWA has been proposed in [Averaging Weights Leads to Wider Optima and Better Generalization](#).

## Constructing averaged models

`AveragedModel` class serves to compute the weights of the SWA model. You can create an averaged model by running:

```
>>> swa_model = AveragedModel(model)
```

Here the `model` can be an arbitrary `:class:'torch.nn.Module'` object. `swa_model` will keep track of the running averages of the parameters of the `model`. To update these averages, you can use the `:func:'update_parameters'` function:

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\ (pytorch-master) (docs) (source) optim.rst, line 245); [backlink](#)**

Unknown interpreted text role "class".

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\ (pytorch-master) (docs) (source) optim.rst, line 245); [backlink](#)**

Unknown interpreted text role "func".

```
>>> swa_model.update_parameters(model)
```

## SWA learning rate schedules

Typically, in SWA the learning rate is set to a high constant value. `:class:'SWALR'` is a learning rate scheduler that anneals the learning rate to a fixed value, and then keeps it constant. For example, the following code creates a scheduler that linearly anneals the learning rate from its initial value to 0.05 in 5 epochs within each parameter group:

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\ (pytorch-master) (docs) (source) optim.rst, line 255); [backlink](#)**

Unknown interpreted text role "class".

```
>>> swa_scheduler = torch.optim.swa_utils.SWALR(optimizer, \
>>>     anneal_strategy="linear", anneal_epochs=5, swa_lr=0.05)
```

You can also use cosine annealing to a fixed value instead of linear annealing by setting `anneal_strategy="cos"`.

## Taking care of batch normalization

`:func:'update_bn'` is a utility function that allows to compute the batchnorm statistics for the SWA model on a given dataloader loader at the end of training:

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\ (pytorch-master) (docs) (source) optim.rst, line 270); [backlink](#)**

Unknown interpreted text role "func".

```
>>> torch.optim.swa_utils.update_bn(loader, swa_model)
```

`:func:'update_bn'` applies the `swa_model` to every element in the dataloader and computes the activation statistics for each batch normalization layer in the model.

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\ (pytorch-master) (docs) (source) optim.rst, line 275); [backlink](#)**

Unknown interpreted text role "func".

### Warning

`:func:`update_bn`` assumes that each batch in the dataloader `loader` is either a tensors or a list of tensors where the first element is the tensor that the network `swa_model` should be applied to. If your dataloader has a different structure, you can update the batch normalization statistics of the `swa_model` by doing a forward pass with the `swa_model` on each element of the dataset.

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\ (pytorch-master) (docs) (source) optim.rst, line 279); [backlink](#)**

Unknown interpreted text role "func".

## Custom averaging strategies

By default, `:class:`torch.optim.swa_utils.AveragedModel`` computes a running equal average of the parameters that you provide, but you can also use custom averaging functions with the `avg_fn` parameter. In the following example `ema_model` computes an exponential moving average.

**System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\ (pytorch-master) (docs) (source) optim.rst, line 288); [backlink](#)**

Unknown interpreted text role "class".

Example:

```
>>> ema_avg = lambda averaged_model_parameter, model_parameter, num_averaged:\
>>>     0.1 * averaged_model_parameter + 0.9 * model_parameter
>>> ema_model = torch.optim.swa_utils.AveragedModel(model, avg_fn=ema_avg)
```

## Putting it all together

In the example below, `swa_model` is the SWA model that accumulates the averages of the weights. We train the model for a total of 300 epochs and we switch to the SWA learning rate schedule and start to collect SWA averages of the parameters at epoch 160:

```
>>> loader, optimizer, model, loss_fn = ...
>>> swa_model = torch.optim.swa_utils.AveragedModel(model)
>>> scheduler = torch.optim.lr_scheduler.CosineAnnealingLR(optimizer, T_max=300)
>>> swa_start = 160
>>> swa_scheduler = SWALR(optimizer, swa_lr=0.05)
>>>
>>> for epoch in range(300):
>>>     for input, target in loader:
>>>         optimizer.zero_grad()
>>>         loss_fn(model(input), target).backward()
>>>         optimizer.step()
>>>     if epoch > swa_start:
>>>         swa_model.update_parameters(model)
>>>         swa_scheduler.step()
>>>     else:
>>>         scheduler.step()
>>>
>>> # Update bn statistics for the swa_model at the end
>>> torch.optim.swa_utils.update_bn(loader, swa_model)
>>> # Use swa_model to make predictions on test data
>>> preds = swa_model(test_input)
```