

orphan:

Distributed Autograd Design

This note will present the detailed design for distributed autograd and walk through the internals of the same. Make sure you're familiar with [ref`autograd-mechanics`](#) and the [ref`distributed-rpc-framework`](#) before proceeding.

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\rpc\pytorch-master) (docs) (source) (rpc) distributed_autograd.rst, line 8); [backlink](#)

Unknown interpreted text role "ref".

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\rpc\pytorch-master) (docs) (source) (rpc) distributed_autograd.rst, line 8); [backlink](#)

Unknown interpreted text role "ref".

Background

Let's say you have two nodes and a very simple model partitioned across two nodes. This can be implemented using `mod`torch.distributed.rpc`` as follows:

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\rpc\pytorch-master) (docs) (source) (rpc) distributed_autograd.rst, line 16); [backlink](#)

Unknown interpreted text role "mod".

```
import torch
import torch.distributed.rpc as rpc

def my_add(t1, t2):
    return torch.add(t1, t2)

# On worker 0:
t1 = torch.rand((3, 3), requires_grad=True)
t2 = torch.rand((3, 3), requires_grad=True)

# Perform some computation remotely.
t3 = rpc.rpc_sync("worker1", my_add, args=(t1, t2))

# Perform some computation locally based on remote result.
t4 = torch.rand((3, 3), requires_grad=True)
t5 = torch.mul(t3, t4)

# Compute some loss.
loss = t5.sum()
```

The main motivation behind distributed autograd is to enable running a backward pass on such distributed models with the `loss` that we've computed and record appropriate gradients for all tensors that require gradients.

Autograd recording during the forward pass

PyTorch builds the autograd graph during the forward pass and this graph is used to execute the backward pass. For more details see [ref`how-autograd-encodes-history`](#).

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\rpc\pytorch-master) (docs) (source) (rpc) distributed_autograd.rst, line 50); [backlink](#)

Unknown interpreted text role "ref".

For distributed autograd, we need to keep track of all RPCs during the forward pass to ensure the backward pass is executed appropriately. For this purpose, we attach `send` and `recv` functions to the autograd graph when we perform an RPC.

- The `send` function is attached to the source of the RPC and its output edges point to the autograd function for the input tensors of the RPC. The input for this function during the backward pass is received from the destination as the output of the appropriate `recv` function.

- The `recv` function is attached to the destination of the RPC and its inputs are retrieved from operators executed on the destination using the input tensors. The output gradients of this function are sent to the source node to the appropriate `send` function during the backward pass.
- Each `send-recv` pair is assigned a globally unique `autograd_message_id` to uniquely identify the pair. This is useful to look up the corresponding function on a remote node during the backward pass.
- For `ref`, whenever we call `meth:torch.distributed.rpc.RRef.to_here` we attach an appropriate `send-recv` pair for the tensors involved.

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\rpc\ (pytorch-master) (docs) (source) (rpc) distributed_autograd.rst, line 70); [backlink](#)

Unknown interpreted text role "ref".

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\rpc\ (pytorch-master) (docs) (source) (rpc) distributed_autograd.rst, line 70); [backlink](#)

Unknown interpreted text role "meth".

As an example, this is what the autograd graph for our example above would look like (`t5.sum()` excluded for simplicity):



Distributed Autograd Context

Each forward and backward pass that uses distributed autograd is assigned a unique `class:torch.distributed.autograd.context` and this context has a globally unique `autograd_context_id`. This context is created on each node as needed.

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\rpc\ (pytorch-master) (docs) (source) (rpc) distributed_autograd.rst, line 83); [backlink](#)

Unknown interpreted text role "class".

This context serves the following purpose:

1. Multiple nodes running distributed backward passes might accumulate gradients on the same tensor and as a result the `.grad` field of the tensor would have gradients from a variety of distributed backward passes before we have the opportunity to run the optimizer. This is similar to calling `meth:torch.autograd.backward` multiple times locally. In order to provide a way of separating out the gradients for each backward pass, the gradients are accumulated in the `class:torch.distributed.autograd.context` for each backward pass.

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\rpc\ (pytorch-master) (docs) (source) (rpc) distributed_autograd.rst, line 90); [backlink](#)

Unknown interpreted text role "meth".

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\rpc\ (pytorch-master) (docs) (source) (rpc) distributed_autograd.rst, line 90); [backlink](#)

Unknown interpreted text role "class".

2. During the forward pass we store the `send` and `recv` functions for each autograd pass in this context. This ensures we hold references to the appropriate nodes in the autograd graph to keep it alive. In addition to this, it is easy to look up the appropriate `send` and `recv` functions during the backward pass.
3. In general we also use this context to store some metadata for each distributed autograd pass.

From the user's perspective the autograd context is setup as follows:

```
import torch.distributed.autograd as dist_autograd
with dist_autograd.context() as context_id:
    loss = model.forward()
    dist_autograd.backward(context_id, loss)
```

It is important to note that your model's forward pass must be invoked within the distributed autograd context manager, as a valid context is needed in order to ensure that all `send` and `recv` functions are stored properly to run the backward pass across all participating nodes.

Distributed Backward Pass

In this section we outline the challenge of computing dependencies accurately during a distributed backward pass and describe a couple of algorithms (with tradeoffs) on how we can execute a distributed backward pass.

Computing dependencies

Consider the following piece of code being run on a single machine

```
import torch
a = torch.rand((3, 3), requires_grad=True)
b = torch.rand((3, 3), requires_grad=True)
c = torch.rand((3, 3), requires_grad=True)
d = a + b
e = b * c
d.sum().backward()
```

This is what the autograd graph for the code above would look like:



The first step the autograd engine performs as part of the backward pass is computing the number of dependencies for each node in the autograd graph. This helps the autograd engine know when a node in the graph is ready for execution. The numbers in brackets for `add(1)` and `mul(0)` denote the number of dependencies. As you can see, this means during the backward pass the `add` node needs 1 input and the `mul` node doesn't need any inputs (in other words doesn't need to be executed). The local autograd engine computes these dependencies by traversing the graph from the root nodes (`d` in this case).

The fact that certain nodes in the autograd graph might not be executed in the backward pass poses a challenge for distributed autograd. Consider this piece of code which uses RPC.

```
import torch
import torch.distributed.rpc as rpc

a = torch.rand((3, 3), requires_grad=True)
b = torch.rand((3, 3), requires_grad=True)
c = torch.rand((3, 3), requires_grad=True)

d = rpc.rpc_sync("worker1", torch.add, args=(a, b))
e = rpc.rpc_sync("worker1", torch.mul, args=(b, c))
loss = d.sum()
```

The associated autograd graph for the code above would be:



Computing dependencies of this distributed autograd graph is much more challenging and requires some overhead (either in terms of computation or network communication).

For performance sensitive applications we can avoid a lot of overhead by assuming every `send` and `recv` function are valid as part of the backward pass (most applications don't perform RPCs that aren't used). This simplifies the distributed autograd algorithm and is much more efficient, but at the cost that the application needs to be aware of the limitations. This algorithm is called the [FAST mode algorithm](#) and is described in detail below.

In the general case it might not be necessary that every `send` and `recv` function is valid as part of the backward pass. To address this, we have proposed a [SMART mode algorithm](#) which is described in a later section. Please note that currently, only the *FAST* mode algorithm is implemented.

FAST mode algorithm

The key assumption of this algorithm is that each `send` function has a dependency of 1 when we run a backward pass. In other words, we assume we'll receive a gradient over RPC from another node.

The algorithm is as follows:

1. We start from the worker which has the roots for the backward pass (all roots must be local).
2. Lookup all the `send` functions for the current [Distributed Autograd Context](#).
3. Compute dependencies locally starting from the provided roots and all the `send` functions we retrieved.
4. After computing dependencies, kick off the local autograd engine with the provided roots.
5. When the autograd engine executes the `recv` function, the `recv` function sends the input gradients via RPC to the appropriate worker. Each `recv` function knows the destination worker id since it is recorded as part of the forward pass. The `recv` function also sends over the `autograd_context_id` and `autograd_message_id` to the remote host.

- When this request is received on the remote host, we use the `autograd_context_id` and `autograd_message_id` to look up the appropriate `send` function.
- If this is the first time a worker has received a request for the given `autograd_context_id`, it will compute dependencies locally as described in points 1-3 above.
- The `send` function retrieved in 6. is then enqueued for execution on the local autograd engine for that worker.
- Finally, instead of accumulating the gradients on the `.grad` field of the Tensor, we accumulate the gradients separately per [Distributed Autograd Context](#). The gradients are stored in a `Dict[Tensor, Tensor]`, which is basically a map from Tensor to its associated gradient and this map can be retrieved using the `meth:~torch.distributed.autograd.get_gradients` API.

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\rpc\ (pytorch-master) (docs) (source) (rpc)distributed_autograd.rst, line 228); [backlink](#)

Unknown interpreted text role "meth".

As an example the complete code with distributed autograd would be as follows:

```
import torch
import torch.distributed.autograd as dist_autograd
import torch.distributed.rpc as rpc

def my_add(t1, t2):
    return torch.add(t1, t2)

# On worker 0:

# Setup the autograd context. Computations that take
# part in the distributed backward pass must be within
# the distributed autograd context manager.
with dist_autograd.context() as context_id:
    t1 = torch.rand((3, 3), requires_grad=True)
    t2 = torch.rand((3, 3), requires_grad=True)

    # Perform some computation remotely.
    t3 = rpc.rpc_sync("worker1", my_add, args=(t1, t2))

    # Perform some computation locally based on remote result.
    t4 = torch.rand((3, 3), requires_grad=True)
    t5 = torch.mul(t3, t4)

    # Compute some loss.
    loss = t5.sum()

    # Run the backward pass.
    dist_autograd.backward(context_id, [loss])

    # Retrieve the gradients from the context.
    dist_autograd.get_gradients(context_id)
```

The distributed autograd graph with dependencies would be as follows (`t5.sum()` excluded for simplicity):

The [FAST mode algorithm](#) applied to the above example would be as follows:

- On Worker 0 we start from the roots `loss` and `send1` to compute dependencies. As a result `send1` is marked with a dependency of 1 and `mul` on Worker 0 is marked with a dependency of 1.
- Now, we kickoff the local autograd engine on Worker 0. We first execute the `mul` function, accumulate its output in the autograd context as the gradient for `t4`. Then, we execute `recv2` which sends the gradients to Worker 1.
- Since this is the first time Worker 1 has heard about this backward pass, it starts dependency computation and marks the dependencies for `send2`, `add` and `recv1` appropriately.
- Next, we enqueue `send2` on the local autograd engine of Worker 1, which in turn executes `add` and `recv1`.
- When `recv1` is executed it sends the gradients over to Worker 0.
- Since Worker 0 has already computed dependencies for this backward pass, it just enqueues and executes `send1` locally.
- Finally, gradients for `t1`, `t2` and `t4` are accumulated in the [Distributed Autograd Context](#).

SMART mode algorithm

Full details of this algorithm are still in the works, but for the general idea you can refer to [Distributed Autograd Algorithm Smart mode](#) section in the [RFC](#).

Distributed Optimizer

The `:class:`~torch.distributed.optim.DistributedOptimizer`` operates as follows:

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\rpc\ (pytorch-master) (docs) (source) (rpc) distributed_autograd.rst, line 306); [backlink](#)

Unknown interpreted text role "class".

1. Takes a list of remote parameters (`:class:`~torch.distributed.rpc.RRef``) to optimize. These could also be local parameters wrapped within a local `RRef`.

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\rpc\ (pytorch-master) (docs) (source) (rpc) distributed_autograd.rst, line 308); [backlink](#)

Unknown interpreted text role "class".

2. Takes a `:class:`~torch.optim.Optimizer`` class as the local optimizer to run on all distinct `RRef` owners.

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\rpc\ (pytorch-master) (docs) (source) (rpc) distributed_autograd.rst, line 311); [backlink](#)

Unknown interpreted text role "class".

3. The distributed optimizer creates an instance of the local `Optimizer` on each of the worker nodes and holds an `RRef` to them.
4. When `:meth:`~torch.distributed.optim.DistributedOptimizer.step`` is invoked, the distributed optimizer uses RPC to remotely execute all the local optimizers on the appropriate remote workers. A distributed autograd `context_id` must be provided as input to `:meth:`~torch.distributed.optim.DistributedOptimizer.step``. This is used by local optimizers to apply gradients stored in the corresponding context.

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\rpc\ (pytorch-master) (docs) (source) (rpc) distributed_autograd.rst, line 315); [backlink](#)

Unknown interpreted text role "meth".

System Message: ERROR/3 (D:\onboarding-resources\sample-onboarding-resources\pytorch-master\docs\source\rpc\ (pytorch-master) (docs) (source) (rpc) distributed_autograd.rst, line 315); [backlink](#)

Unknown interpreted text role "meth".

5. If multiple concurrent distributed optimizers are updating the same parameters on a worker, these updates are serialized via a lock.

Simple end to end example

Putting it all together, the following is a simple end to end example using distributed autograd and the distributed optimizer. If the code is placed into a file called "dist_autograd_simple.py", it can be run with the command `MASTER_ADDR="localhost"`

`MASTER_PORT=29500 python dist_autograd_simple.py:`

```
import torch
import torch.multiprocessing as mp
import torch.distributed.autograd as dist_autograd
from torch.distributed import rpc
from torch import optim
from torch.distributed.optim import DistributedOptimizer

def random_tensor():
    return torch.rand((3, 3), requires_grad=True)

def _run_process(rank, dst_rank, world_size):
    name = "worker{}".format(rank)
    dst_name = "worker{}".format(dst_rank)
```

```

# Initialize RPC.
rpc.init_rpc(
    name=name,
    rank=rank,
    world_size=world_size
)

# Use a distributed autograd context.
with dist_autograd.context() as context_id:
    # Forward pass (create references on remote nodes).
    rref1 = rpc.remote(dst_name, random_tensor)
    rref2 = rpc.remote(dst_name, random_tensor)
    loss = rref1.to_here() + rref2.to_here()

    # Backward pass (run distributed autograd).
    dist_autograd.backward(context_id, [loss.sum()])

    # Build DistributedOptimizer.
    dist_optim = DistributedOptimizer(
        optim.SGD,
        [rref1, rref2],
        lr=0.05,
    )

    # Run the distributed optimizer step.
    dist_optim.step(context_id)

def run_process(rank, world_size):
    dst_rank = (rank + 1) % world_size
    _run_process(rank, dst_rank, world_size)
    rpc.shutdown()

if __name__ == '__main__':
    # Run world_size workers
    world_size = 2
    mp.spawn(run_process, args=(world_size,), nprocs=world_size)

```