

In this section, we look at the timeline of simple tensor. The content is extracted from a live presentation. It reflects the PyTorch callstacks as a snapshot on July 10, 2019. All the code refers to PyTorch location inside FB, but the opensource version points to similar locations.

Let's start with a simple tensor:

```
import torch
r = torch.rand(3,4)[0] + torch.rand(3 , 4)
```

output:

```
tensor([[0.3091, 0.5503, 1.0780, 0.9044],
        [0.5770, 0.5245, 0.3225, 1.4672],
        [0.1581, 1.0439, 0.3313, 0.9924]])
```

The code is equivalent to:

```
_t1 = torch.rand(3, 4)
_t2 = _t1.__getitem__(0)
del _t1
_t3 = torch.rand(3, 4)
r = _t2.__add__(_t3)
del _t2
del _t3
# only r remains at this point
```

Looking at them one by one:

```
_t1 = torch.rand(3, 4) # <--- here
_t2 = _t1.__getitem__(0)
del _t1
_t3 = torch.rand(3, 4)
r = _t2.__add__(_t3)
del _t2
del _t3
```

The Python code for torch.rand doesn't exist. It all comes from

aten/src/ATen/native/native_functions.yaml

```
- func: scalar_tensor(Scalar s, *, ScalarType? dtype=None, Layout? layout=None,
    Device? device=None, bool? pin_memory=None) -> Tensor

- func: rand(int[] size, *, ScalarType? dtype=None, Layout? layout=None,
    Device? device=None, bool? pin_memory=None) -> Tensor

- func: rand(int[] size, *, Generator? generator, ScalarType? dtype=None,
    Layout? layout=None, Device? device=None, bool? pin_memory=None) -> Tensor

- func: rand(int[] size, *, Tensor(a!) out) -> Tensor(a!)
```

```

- func: rand(int[] size, *, Generator? generator, Tensor(a!) out) -> Tensor(a!)

- func: rand_like(Tensor self) -> Tensor

- func: rand_like(Tensor self, *, ScalarType dtype, Layout layout,
    Device device, bool pin_memory=False) -> Tensor

```

tools/autograd/templates/python_torch_functions.cpp

```

static PyMethodDef torch_functions[] = {
    {"arange", (PyCFunction)THPVariable_arange, METH_VARARGS | METH_KEYWORDS | METH_STATIC, NULL},
    {"as_tensor", (PyCFunction)THPVariable_as_tensor, METH_VARARGS | METH_KEYWORDS | METH_STATIC, NULL},
    {"dsmm", (PyCFunction)THPVariable_mm, METH_VARARGS | METH_KEYWORDS | METH_STATIC, NULL},
    {"from_numpy", (PyCFunction)THPVariable_from_numpy, METH_STATIC | METH_O, NULL},
    {"hsmm", (PyCFunction)THPVariable_hspmm, METH_VARARGS | METH_KEYWORDS | METH_STATIC, NULL},
    {"_promote_types", (PyCFunction)THPVariable__promote_types, METH_VARARGS | METH_KEYWORDS | METH_STATIC, NULL},
    {"nonzero", (PyCFunction)THPVariable_nonzero, METH_VARARGS | METH_KEYWORDS | METH_STATIC, NULL},
    {"randint", (PyCFunction)THPVariable_randint, METH_VARARGS | METH_KEYWORDS | METH_STATIC, NULL},
    {"range", (PyCFunction)THPVariable_range, METH_VARARGS | METH_KEYWORDS | METH_STATIC, NULL},
    {"saddmm", (PyCFunction)THPVariable_sspaddmm, METH_VARARGS | METH_KEYWORDS | METH_STATIC, NULL},
    {"sparse_coo_tensor", (PyCFunction)THPVariable_sparse_coo_tensor, METH_VARARGS | METH_KEYWORDS | METH_STATIC, NULL},
    {"spmm", (PyCFunction)THPVariable_mm, METH_VARARGS | METH_KEYWORDS | METH_STATIC, NULL},
    {"tensor", (PyCFunction)THPVariable_tensor, METH_VARARGS | METH_KEYWORDS | METH_STATIC, NULL},
    {"get_device", (PyCFunction)THPVariable_get_device, METH_VARARGS | METH_KEYWORDS | METH_STATIC, NULL},
    {py_method_defs},
    {NULL}
};

```

gen/generate-code-outputs/generate-code-outputs/python_torch_functions.cpp

```

{"quantized_gru_cell", (PyCFunction)THPVariable_quantized_gru_cell, METH_VARARGS | METH_KEYWORDS | METH_STATIC, NULL},
{"quantized_lstm", (PyCFunction)THPVariable_quantized_lstm, METH_VARARGS | METH_KEYWORDS | METH_STATIC, NULL},
{"quantized_lstm_cell", (PyCFunction)THPVariable_quantized_lstm_cell, METH_VARARGS | METH_KEYWORDS | METH_STATIC, NULL},
{"quantized_rnn_relu_cell", (PyCFunction)THPVariable_quantized_rnn_relu_cell, METH_VARARGS | METH_KEYWORDS | METH_STATIC, NULL},
{"quantized_rnn_tanh_cell", (PyCFunction)THPVariable_quantized_rnn_tanh_cell, METH_VARARGS | METH_KEYWORDS | METH_STATIC, NULL},
{"rand", (PyCFunction)THPVariable_rand, METH_VARARGS | METH_KEYWORDS | METH_STATIC, NULL},
{"rand_like", (PyCFunction)THPVariable_rand_like, METH_VARARGS | METH_KEYWORDS | METH_STATIC, NULL},
{"randint_like", (PyCFunction)THPVariable_randint_like, METH_VARARGS | METH_KEYWORDS | METH_STATIC, NULL},
{"randn", (PyCFunction)THPVariable_randn, METH_VARARGS | METH_KEYWORDS | METH_STATIC, NULL},
{"randn_like", (PyCFunction)THPVariable_randn_like, METH_VARARGS | METH_KEYWORDS | METH_STATIC, NULL},
{"randperm", (PyCFunction)THPVariable_randperm, METH_VARARGS | METH_KEYWORDS | METH_STATIC, NULL}

```

tools/autograd/templates/python_torch_functions.cpp

```

static PyTypeObject THPVariableFunctions = {
    PyVarObject_HEAD_INIT(NULL, 0)
    "torch._C._VariableFunctions", /* tp_name */
    0, /* tp_basicsize */

```

```

0, /* tp_itemsize */
0, /* tp_dealloc */
0, /* tp_print */
0, /* tp_getattr */
0, /* tp_setattr */
0, /* tp_reserved */
0, /* tp_repr */
0, /* tp_as_number */
0, /* tp_as_sequence */
0, /* tp_as_mapping */
0, /* tp_hash */
0, /* tp_call */
0, /* tp_str */
0, /* tp_getattro */
0, /* tp_setattro */
0, /* tp_as_buffer */
Py_TPFLAGS_DEFAULT, /* tp_flags */
NULL, /* tp_doc */
0, /* tp_traverse */
0, /* tp_clear */
0, /* tp_richcompare */
0, /* tp_weaklistoffset */
0, /* tp_iter */
0, /* tp_iternext */
torch_functions, /* tp_methods */
0, /* tp_members */
0, /* tp_getset */
0, /* tp_base */
0, /* tp_dict */
0, /* tp_descr_get */
0, /* tp_descr_set */
0, /* tp_dictoffset */
0, /* tp_init */
0, /* tp_alloc */
0 /* tp_new */
};

```

tools/autograd/templates/python_torch_functions.cpp

```

void initTorchFunctions(PyObject* module) {
    if (PyType_Ready(&THPVariableFunctions) < 0) {
        throw python_error();
    }
    Py_INCREF(&THPVariableFunctions);
    if (PyModule_AddObject(module, "_VariableFunctions",
        (PyObject*)&THPVariableFunctions) < 0) {
        throw python_error();
    }
}

```

```

    }
}

torch/init.py

for name in dir(_C._VariableFunctions):
    if name.startswith('__'):
        continue
    globals()[name] = getattr(_C._VariableFunctions, name)

gen/generate-code-outputs/generate-code-outputs/python_torch_functions.cpp

{"rand", (PyCFunction)THPVariable_rand,
    METH_VARARGS | METH_KEYWORDS | METH_STATIC, NULL},

static PyObject * THPVariable_rand(PyObject* self_, PyObject* args,
    PyObject* kwargs)
{
    HANDLE_TH_ERRORS
    static PythonArgParser parser({
        "rand(IntArrayRef size, *, Generator generator, Tensor out=None, ScalarType dtype=None,
        "rand(IntArrayRef size, *, Tensor out=None, ScalarType dtype=None, Layout layout=torch.
    }, /*traceable=*/true);

    ParsedArgs<8> parsed_args;
    auto r = parser.parse(args, kwargs, parsed_args);

    if (r.idx == 0) {
        if (r.isNone(2)) {
            auto size = r.intlist(0);
            auto generator = r.generator(1);
            auto dtype = r.scalar_type(3);
            auto device = r.device(5);
            const auto options = TensorOptions()
                .dtype(dtype)
                .device(device)
                .layout(r.layout(4).layout)
                .requires_grad(r.toBool(7))
                .pinned_memory(r.toBool(6));
            return wrap(dispatch_rand(size, generator, options));
        } else {
            check_out_type_matches(r.tensor(2), r.scalar_type(3), r.isNone(3),
                r.layout(4), r.isNone(4),
                r.device(5), r.isNone(5));
            return wrap(dispatch_rand(r.intlist(0), r.generator(1),
                r.tensor(2)).set_requires_grad(r.toBool(7)));
        }
    }
    ...
}

gen/generate-code-outputs/generate-code-outputs/python_torch_functions_dispatch.h

```

```

inline Tensor dispatch_rand(IntArrayRef size, Generator * generator,
    const TensorOptions & options) {
    maybe_initialize_cuda(options);
    AutoNoGIL no_gil;
    return torch::rand(size, generator, options);
}

gen/generate-code-outputs/generate-code-outputs/variable_factories.h

inline at::Tensor rand(at::IntArrayRef size, at::Generator * generator,
    const at::TensorOptions & options = {}) {
    torch::jit::Node* node = nullptr;
    std::shared_ptr<jit::tracer::TracingState> tracer_state;
    if (jit::tracer::isTracing()) {
        tracer_state = jit::tracer::getTracingState();
        at::Symbol op_name;
        op_name = jit::Symbol::fromQualString("aten::rand");
        node = tracer_state->graph->create(op_name, /*num_outputs=*/0);
        jit::tracer::recordSourceLocation(node);
        jit::tracer::addInputs(node, "size", size);
        jit::tracer::addInputs(node, "generator", generator);
        jit::tracer::addInputs(node, "options", options);
        tracer_state->graph->insertNode(node);

        jit::tracer::setTracingState(nullptr);
    }
    at::Tensor tensor = at::rand(size, generator,
        at::TensorOptions(options).is_variable(false));
    at::Tensor result =
        autograd::make_variable_consuming(std::move(tensor),
            /*requires_grad=*/options.requires_grad());
    if (tracer_state) {
        jit::tracer::setTracingState(std::move(tracer_state));
        jit::tracer::addOutput(node, result);
    }
    return result;
}

gen/aten/gen_aten-outputs/gen_aten-outputs/Functions.h

static inline Tensor rand(IntArrayRef size, Generator * generator,
    const TensorOptions & options) {
    globalLegacyTypeDispatch().initForBackend(options.backend());
    static auto table = globalATenDispatch().getOpTable(
        "aten::rand(int[] size, *, Generator? generator, "
        "ScalarType? dtype=None, Layout? layout=None, Device? device=None, "
        "bool? pin_memory=None) -> Tensor");
    return table->getOp<Tensor>(IntArrayRef, Generator *, const TensorOptions &)

```

```

        >(options.backend(), options.is_variable())(size, generator, options);
    }

gen/aten/gen_aten-outputs/gen_aten-outputs/TypeDefault.cpp
static auto& registerer = globalATenDispatch()
    .registerOp<Tensor (const Tensor &, bool)>(Backend::Undefined, "aten::_cast_Byte(Tensor se
    .registerOp<Tensor (const Tensor &, bool)>(Backend::Undefined, "aten::_cast_Char(Tensor se
    .registerOp<Tensor (const Tensor &, bool)>(Backend::Undefined, "aten::_cast_Double(Tensor
    .registerOp<Tensor (const Tensor &, bool)>(Backend::Undefined, "aten::_cast_Float(Tensor s
    .registerOp<Tensor (const Tensor &, bool)>(Backend::Undefined, "aten::_cast_Int(Tensor sel
    .registerOp<Tensor (IntArrayRef, Generator *, const TensorOptions &)>(Backend::Undefined,
        "aten::rand(int[] size, *, Generator? generator, ScalarType? dtype=None, "
        "Layout? layout=None, Device? device=None, bool? pin_memory=None) -> Tensor",
        &TypeDefault::rand)

Tensor TypeDefault::rand(IntArrayRef size, Generator * generator, const TensorOptions & opti
    const DeviceGuard device_guard(options.device());
    return at::native::rand(size, generator, options);
}

aten/src/ATen/native/TensorFactories.cpp
Tensor rand(IntArrayRef size, Generator* generator, const TensorOptions& options) {
    auto result = at::empty(size, options);
    return result.uniform_(0, 1, generator);
}

aten/src/ATen/native/native_functions.yaml
- func: empty(int[] size, *, ScalarType? dtype=None, Layout? layout=None,
    Device? device=None, bool? pin_memory=None) -> Tensor
dispatch:
    CPU: empty_cpu
    CUDA: empty_cuda
    MklDnnCPU: empty_mkldnn
    SparseCPU: empty_sparse
    SparseCUDA: empty_sparse

aten/src/ATen/native/TensorFactories.cpp
Tensor empty_cpu(IntArrayRef size, const TensorOptions& options) {
    AT_ASSERT(options.backend() == Backend::CPU);
    AT_ASSERT(!options.is_variable()); // is_variable should have been 'unpacked' // TODO:
    check_size_nonnegative(size);

    c10::Allocator* allocator;
    if (options.pinned_memory()) {
        allocator = detail::getCUDataHooks().getPinnedMemoryAllocator();
    } else {

```

```

        allocator = at::getCPUAllocator();
    }

    int64_t nelements = prod_intlist(size);
    auto dtype = options.dtype();
    auto storage_impl = c10::make_intrusive<StorageImpl>(
        dtype,
        nelements,
        allocator->allocate(nelements * dtype.itemsize()),
        allocator,
        /*resizeable=*/true);

    auto tensor = detail::make_tensor<TensorImpl>(storage_impl, at::CPUTensorId());
    // Default TensorImpl has size [0]
    if (size.size() != 1 || size[0] != 0) {
        tensor.unsafeGetTensorImpl()->set_sizes_contiguous(size);
    }
    return tensor;
}

aten/src/ATen/Context.cpp

Allocator* getCPUAllocator() {
    return getTHDefaultAllocator();
}

aten/src/TH/THAllocator.cpp

at::Allocator* getTHDefaultAllocator() {
    return c10::GetCPUAllocator();
}

c10/core/CPUAllocator.cpp

at::Allocator* GetCPUAllocator() {
    return GetAllocator(DeviceType::CPU);
}

c10/core/Allocator.cpp

at::Allocator* GetAllocator(const at::DeviceType& t) {
    auto* alloc = allocator_array[static_cast<int>(t)];
    AT_ASSERTM(alloc, "Allocator for ", t, " is not set.");
    return alloc;
}

c10/core/Allocator.h

template <DeviceType t>
struct AllocatorRegisterer {
    explicit AllocatorRegisterer(Allocator* alloc) {

```

```

        SetAllocator(t, alloc);
    }
};

#define REGISTER_ALLOCATOR(t, f) \
    namespace { \
        static AllocatorRegisterer<t> g_allocator_d(f); \
    }

c10/core/CPUAllocator.cpp

REGISTER_ALLOCATOR(DeviceType::CPU, &g_cpu_alloc);

static DefaultCPUAllocator g_cpu_alloc;

struct C10_API DefaultCPUAllocator final : at::Allocator {
    DefaultCPUAllocator() {}
    ~DefaultCPUAllocator() override {}
    at::DataPtr allocate(size_t nbytes) const override {
        void* data = alloc_cpu(nbytes);
        if (FLAGS_caffe2_report_cpu_memory_usage && nbytes > 0) {
            getMemoryAllocationReporter().New(data, nbytes);
            return {data, data, &ReportAndDelete, at::Device(at::DeviceType::CPU)};
        }
        return {data, data, &free_cpu, at::Device(at::DeviceType::CPU)};
    }

void* alloc_cpu(size_t nbytes) {
    void* data;
#ifdef __ANDROID__
    data = memalign(gAlignment, nbytes);
#elif defined(_MSC_VER)
    data = _aligned_malloc(nbytes, gAlignment);
#else
    int err = posix_memalign(&data, gAlignment, nbytes);
#endif

    NUMAMove(data, nbytes, GetCurrentNUMANode());

    if (FLAGS_caffe2_cpu_allocator_do_zero_fill) {
        memset(data, 0, nbytes);
    } else if (FLAGS_caffe2_cpu_allocator_do_junk_fill) {
        memset_junk(data, nbytes);
    }

constexpr size_t gAlignment = 64;

void free_cpu(void* data) {
#ifdef _MSC_VER
    _aligned_free(data);

```



```

#else
    free(data);
#endif
}

aten/src/ATen/native/TensorFactories.cpp
Tensor empty_cpu(IntArrayRef size, const TensorOptions& options) {
    .....
    int64_t nelements = prod_intlist(size);
    auto dtype = options.dtype();
    auto storage_impl = c10::make_intrusive<StorageImpl>(
        dtype,
        nelements,
        allocator->allocate(nelements * dtype.itemsize()),
        allocator,
        /*resizeable=*/true);
c10/util/intrusive_ptr.h
template <
    class TTarget,
    class NullType = detail::intrusive_target_default_null_type<TTarget>,
    class... Args>
inline intrusive_ptr<TTarget, NullType> make_intrusive(Args&&... args) {
    return intrusive_ptr<TTarget, NullType>::make(std::forward<Args>(args)...);
}

template <
    class TTarget,
    class NullType = detail::intrusive_target_default_null_type<TTarget>>
class intrusive_ptr final {
public:
    intrusive_ptr(const intrusive_ptr& rhs) : target_(rhs.target_) {
        retain_();
    }

    ~intrusive_ptr() noexcept {
        reset_();
    }

private:
    TTarget* target_;

    void retain_() {
        size_t new_refcount = ++target_->refcount_;
    }
}

```

```

void reset_() noexcept {
    if (target_ != NullType::singleton() && --target_>refcount_ == 0) {
        auto weak_count = --target_>weakcount_;
        const_cast<c10::guts::remove_const_t<TTarget>*>(target_)->release_resources();
        if (weak_count == 0) {
            delete target_;
        }
    }
}

struct C10_API StorageImpl final : public c10::intrusive_ptr_target {

class C10_API intrusive_ptr_target {
    mutable std::atomic<size_t> refcount_;
    mutable std::atomic<size_t> weakcount_;

c10/core/Allocator.h

class C10_API DataPtr {
private:
    c10::detail::UniqueVoidPtr ptr_;
    Device device_;

public:
    DataPtr() : ptr_(), device_(DeviceType::CPU) {}
    DataPtr(void* data, Device device) : ptr_(data), device_(device) {}
    DataPtr(void* data, void* ctx, DeleterFnPtr ctx_deleter, Device device)
        : ptr_(data, ctx, ctx_deleter), device_(device) {}

c10/util/UniqueVoidPtr.h

class UniqueVoidPtr {
private:
    // Lifetime tied to ctx_
    void* data_;
    std::unique_ptr<void, DeleterFnPtr> ctx_;

public:
    UniqueVoidPtr(void* data, void* ctx, DeleterFnPtr ctx_deleter)
        : data_(data), ctx_(ctx, ctx_deleter ? ctx_deleter : &deleteNothing) {}

c10/core/StorageImpl.h

struct C10_API StorageImpl final : public c10::intrusive_ptr_target {
public:
    StorageImpl(caffe2::TypeMeta data_type, int64_t numel, at::DataPtr data_ptr,
        at::Allocator* allocator, bool resizable);

private:
    caffe2::TypeMeta data_type_;
    DataPtr data_ptr_;

```

```

        int64_t numel_;
        bool resizable_;
        bool received_cuda_;
        Allocator* allocator_;
aten/src/ATen/native/TensorFactories.cpp
Tensor empty_cpu(IntArrayRef size, const TensorOptions& options) {
    .....
    auto tensor = detail::make_tensor<TensorImpl>(storage_impl, at::CPUTensorId());
aten/src/ATen/core/Tensor.h
class CAFFE2_API Tensor {
protected:
    c10::intrusive_ptr<TensorImpl, UndefinedTensorImpl> impl_;

public:
    int64_t dim() const {
        return impl_->dim();
    }
    int64_t storage_offset() const {
        return impl_->storage_offset();
    }

    Tensor abs() const;
    Tensor& abs_();
    Tensor add(const Tensor & other, Scalar alpha=1) const;
c10/core/TensorImpl.h
struct C10_API TensorImpl : public c10::intrusive_ptr_target {
public:
    virtual int64_t dim() const;
    virtual int64_t storage_offset() const;

private:
    Storage storage_;
#ifdef NAMEDTENSOR_ENABLED
    std::unique_ptr<c10::NamedTensorMetaInterface> named_tensor_meta_ = nullptr;
#endif
    c10::VariableVersion version_counter_;
    PyObject* pyobj_ = nullptr; // weak reference
    SmallVector<int64_t,5> sizes_;
    SmallVector<int64_t,5> strides_;
    int64_t storage_offset_ = 0;
    int64_t numel_ = 1;
    caffe2::TypeMeta data_type_;
    c10::optional<c10::Device> device_opt_;

```

```

    TensorTypeId type_id_;
    bool is_contiguous_ = true;
    bool is_wrapped_number_ = false;
    bool allow_tensor_metadata_change_ = true;
    bool reserved_ = false;

class Caffe2_API Tensor {
    c10::intrusive_ptr<TensorImpl, UndefinedTensorImpl> impl_;
struct C10_API TensorImpl : public c10::intrusive_ptr_target {
    Storage storage_;

struct C10_API Storage {
protected:
    c10::intrusive_ptr<StorageImpl> storage_impl_;

struct C10_API StorageImpl final : public c10::intrusive_ptr_target {
    DataPtr data_ptr_;

class C10_API DataPtr {
    c10::detail::UniqueVoidPtr ptr_;

class UniqueVoidPtr {
    std::unique_ptr<void, DeleterFnPtr> ctx_;

aten/src/ATen/native/TensorFactories.cpp

Tensor rand(IntArrayRef size, Generator* generator, const TensorOptions& options) {
    auto result = at::empty(size, options);
    return result.uniform_(0, 1, generator);
}

aten/src/ATen/core/TensorMethods.h

inline Tensor & Tensor::uniform_(double from, double to, Generator * generator) {
    static auto table = globalATenDispatch().getOpTable(
        "aten::uniform_(Tensor(a!) self, float from=0, float to=1, *, "
        "Generator? generator=None) -> Tensor(a!)");
    return table->getOp<Tensor & (Tensor &, double, double, Generator *)>(
        tensorTypeIdToBackend(type_id()),
        is_variable())(*this, from, to, generator);
}

aten/src/ATen/native/native_functions.yaml

- func: uniform_(Tensor(a!) self, float from=0, float to=1, *, Generator? generator=None) ->
  variants: method
  dispatch:
    CPU: legacy::cpu::_th_uniform_
    CUDA: uniform_cuda_

gen/aten/gen_aten=CPUType.cpp/CPUType.cpp

```

```

Tensor & CPUPType::uniform_(Tensor & self, double from, double to, Generator * generator) {
#ifdef NAMEDTENSOR_ENABLED
    if (self.is_named()) {
        AT_ERROR("uniform_: no named inference rule implemented.");
    }
#endif
    const OptionalDeviceGuard device_guard(device_of(self));
    return at::native::legacy::cpu::_th_uniform_(self, from, to, generator);
}

aten/src/ATen/Declarations.cwrap

name: _th_uniform_
types:
    - floating_point
backends:
    - CPU
cname: uniform
variants: function
return: self
arguments:
    - THTensor* self
    - arg: THGenerator* GeneratorExit

gen/aten/gen_aten-outputs/gen_aten-outputs/LegacyTHFunctionsCPU.cpp

Tensor & _th_uniform_(Tensor & self, double from, double to, Generator * generator) {
    auto dispatch_scalar_type = infer_scalar_type(self);
    switch (dispatch_scalar_type) {
        case ScalarType::Float: {
            auto self_ = checked_tensor_unwrap(self, "self", 1, false, Backend::CPU, ScalarType::Float);
            THFloatTensor_uniform(self_, generator, from, to);
            return self;
        }
        break;
    }
}

aten/src/TH/generic/THTensorRandom.cpp

void THTensor_(uniform)(THTensor *self, at::Generator *_generator, double a, double b)
{
    auto gen = at::get_generator_or_default<at::CPUGenerator>(_generator, at::detail::getDefaultGenerator());
    at::uniform_real_distribution<float> uniform((float)a, (float)b);
    TH_TENSOR_APPLY(scalar_t, self, *self_data = (scalar_t)uniform(gen));
}

aten/src/ATen/native/TensorFactories.cpp

Tensor rand(IntArrayRef size, Generator* generator, const TensorOptions& options) {
    auto result = at::empty(size, options);
    return result.uniform_(0, 1, generator);
}

```

Move to slicing:

```
_t1 = torch.rand(3, 4)
_t2 = _t1.__getitem__(0)    # <--- here
del _t1
_t3 = torch.rand(3, 4)
r = _t2.__add__(_t3)
del _t2
del _t3
```

torch/tensor.py

```
class Tensor(torch._C._TensorBase):
```

torch/csrc/autograd/python_variable.cpp

```
PyTypeObject THPVariableType = {
    PyVarObject_HEAD_INIT(NULL, 0)
    "torch._C._TensorBase",          /* tp_name */
    sizeof(THPVariable),              /* tp_basicsize */
    (destructor)THPVariable_dealloc,  /* tp_dealloc */
    &THPVariable_as_mapping,          /* tp_as_mapping */
    Py_TPFLAGS_DEFAULT | Py_TPFLAGS_BASETYPE | Py_TPFLAGS_HAVE_GC, /* tp_flags */
    (traverseproc)THPVariable_traverse, /* tp_traverse */
    (inquiry)THPVariable_clear,       /* tp_clear */
    THPVariable_properties,           /* tp_getset */
    THPVariable_pynew                 /* tp_new */
};
```

```
static PyMappingMethods THPVariable_as_mapping = {
    THPVariable_length,
    THPVariable_getitem,
    THPVariable_setitem,
};
```

```
bool THPVariable_initModule(PyObject *module)
{
```

```
    PyModule_AddObject(module, "_TensorBase", (PyObject *)&THPVariableType);
```

torch/csrc/autograd/python_variable_indexing.cpp

```
PyObject* THPVariable_getitem(PyObject* self, PyObject* index) {
    if (index == Py_None) {
        return wrap(self._unsqueeze(0));
    } else if (index == Py_Ellipsis) {
        return wrap(at::alias(self_));
    } else if (THPUtility_checkLong(index)) {
        return wrap(applySelect(self_, 0, THPUtility_unpackLong(index)));
    } else if (PySlice_Check(index)) {
        return wrap(applySlice(self_, 0, index, true));
    }
}
```

```

    }

    // wrap index in a tuple if it's not already one
    THPObjPtr holder = wrapTuple(index);

    variable_list variableIndices;
    Variable sliced = applySlicing(self_, holder.get(), variableIndices);

static Variable applySelect(const Variable& self, int64_t dim, int64_t index,
    int64_t real_dim=0) {
    int64_t size = self.size(dim);
    return self.select(dim, index);
}

aten/src/ATen/core/TensorMethods.h

inline Tensor Tensor::select(int64_t dim, int64_t index) const {
    static auto table = globalATenDispatch().getOpTable("aten::select(Tensor(a) self, int d
    return table->getOp<Tensor (const Tensor &, int64_t, int64_t)>(tensorTypeIdToBackend(ty
}

aten/src/ATen/native/native_functions.yaml

- func: select(Tensor(a) self, int dim, int index) -> Tensor(a)
  variants: function, method
  device_guard: False
  named_guard: False

gen/aten/gen_aten-outputs/gen_aten-outputs/TypeDefault.cpp

.registerOp<Tensor (const Tensor &, int64_t, int64_t)>(Backend::Undefined,
    "aten::select(Tensor(a) self, int dim, int index) -> Tensor(a)",
    &TypeDefault::select)

Tensor TypeDefault::select(const Tensor & self, int64_t dim, int64_t index) {
    return at::native::select(self, dim, index);
}

aten/src/ATen/native/TensorShape.cpp

Tensor select(const Tensor& self, int64_t dim, int64_t index) {
    auto sizes = self.sizes().vec();
    auto strides = self.strides().vec();
    auto storage_offset = self.storage_offset() + index * strides[dim];
    sizes.erase(sizes.begin() + dim);
    strides.erase(strides.begin() + dim);
    auto result = self.as_strided(sizes, strides, storage_offset);
}

aten/src/ATen/core/TensorMethods.h

inline Tensor Tensor::as_strided(IntArrayRef size, IntArrayRef stride, c10::optional<int64_t
    static auto table = globalATenDispatch().getOpTable("aten::as_strided(Tensor(a) self, in

```

```

        return table->getOp<Tensor>(const Tensor &, IntArrayRef, IntArrayRef, c10::optional<int>)
    }
}

aten/src/ATen/native/native_functions.yaml
- func: as_strided(Tensor(a) self, int[] size, int[] stride, int? storage_offset=None) -> Tensor
  variants: function, method
  dispatch:
    CPU: as_strided_tensorimpl
    CUDA: as_strided_tensorimpl

aten/src/ATen/native/TensorShape.cpp
Tensor as_strided_tensorimpl(const Tensor& self, IntArrayRef size,
    IntArrayRef stride, optional<int64_t> storage_offset_) {
    auto storage_offset = storage_offset_.value_or(self.storage_offset());
    auto tid = self.type_id();
    auto result = detail::make_tensor<TensorImpl>(Storage(self.storage()), tid);
    setStrided(result, size, stride, storage_offset);
    return result;
}

c10/core/Storage.h
struct C10_API Storage {
protected:
    c10::intrusive_ptr<StorageImpl> storage_impl_;

    _t1 = torch.rand(3, 4)
    _t2 = _t1.__getitem__(0)
    del _t1 # <--- here
    _t3 = torch.rand(3, 4)
    r = _t2.__add__(_t3)
    del _t2
    del _t3

torch/tensor.py
class Tensor(torch._C._TensorBase):

torch/csrc/autograd/python_variable.cpp
PyTypeObject THPVariableType = {
    PyVarObject_HEAD_INIT(nullptr, 0)
    "torch._C._TensorBase", /* tp_name */
    sizeof(THPVariable), /* tp_basicsize */
    (destructor)THPVariable_dealloc, /* tp_dealloc */
    &THPVariable_as_mapping, /* tp_as_mapping */
    Py_TPFLAGS_DEFAULT | Py_TPFLAGS_BASETYPE | Py_TPFLAGS_HAVE_GC, /* tp_flags */
    (traverseproc)THPVariable_traverse, /* tp_traverse */
    (inquiry)THPVariable_clear, /* tp_clear */

```



```

    THPVariable_properties,          /* tp_getset */
    THPVariable_pynew                /* tp_new */
};

static void THPVariable_dealloc(THPVariable* self)
{
    PyObject_GC_UnTrack(self);
    THPVariable_clear(self);
    self->cdata.~Variable();
    Py_TYPE(self)->tp_free((PyObject*)self);
}

torch/csrc/autograd/python_variable.h

struct THPVariable {
    PyObject_HEAD
    torch::autograd::Variable cdata;
    PyObject* backward_hooks = nullptr;
};

torch/csrc/autograd/variable.h

struct TORCH_API Variable : public at::Tensor {
class CAFE2_API Tensor {
    c10::intrusive_ptr<TensorImpl, UndefinedTensorImpl> impl_;
struct C10_API TensorImpl : public c10::intrusive_ptr_target {
    Storage storage_;
struct C10_API Storage {
    protected:
    c10::intrusive_ptr<StorageImpl> storage_impl_;
struct C10_API StorageImpl final : public c10::intrusive_ptr_target {
    DataPtr data_ptr_;
class C10_API DataPtr {
    c10::detail::UniqueVoidPtr ptr_;
class UniqueVoidPtr {
    std::unique_ptr<void, DeleterFnPtr> ctx_;
void free_cpu(void* data) {
#ifdef _MSC_VER
    _aligned_free(data);
#else
    free(data);
#endif
}

```

The last step: addition

```

_t1 = torch.rand(3, 4)
_t2 = _t1.__getitem__(0)
del _t1
_t3 = torch.rand(3, 4)
r = _t2.__add__(_t3)      # <--- here
del _t2
del _t3

tools/autograd/templates/python_variable_methods.cpp

PyMethodDef variable_methods[] = {
    {"__add__", (PyCFunction)THPVariable_add, METH_VARARGS | METH_KEYWORDS, NULL},
    {"__radd__", (PyCFunction)THPVariable_add, METH_VARARGS | METH_KEYWORDS, NULL},
    {"__iadd__", (PyCFunction)THPVariable_add_, METH_VARARGS | METH_KEYWORDS, NULL},
};

bool THPVariable_initModule(PyObject *module)
{
    static std::vector<PyMethodDef> methods;
    THPUUtils_addPyMethodDefs(methods, torch::autograd::variable_methods);
    PyModule_AddObject(module, "_TensorBase", (PyObject *)&THPVariableType);
}

aten/src/ATen/native/native_functions.yaml

- func: add(Tensor self, Tensor other, *, Scalar alpha=1) -> Tensor
  variants: function, method
  dispatch:
    CPU: add
    CUDA: add
    SparseCPU: add
    SparseCUDA: add
    MklDnnCPU: mklDnn_add

gen/generate-code-outputs/generate-code-outputs/python_variable_methods.cpp

static PyObject * THPVariable_add(PyObject* self_, PyObject* args, PyObject* kwargs)
{
    static PythonArgParser parser({
        "add(Scalar alpha, Tensor other)|deprecated",
        "add(Tensor other, *, Scalar alpha=1)",
    });
    ParsedArgs<3> parsed_args;
    auto r = parser.parse(args, kwargs, parsed_args);

    if (r.idx == 0) {
        return wrap(dispatch_add(self, r.scalar(0), r.tensor(1)));
    } else if (r.idx == 1) {
        return wrap(dispatch_add(self, r.tensor(0), r.scalar(1)));
    }
}

```

```

gen/generate-code=python_torch_functions_dispatch.h/python_torch_functions_dispatch.h

inline Tensor dispatch_add(const Tensor & self, const Tensor & other, Scalar alpha) {
    return self.add(other, alpha);
}

aten/src/ATen/core/TensorMethods.h

inline Tensor Tensor::add(const Tensor & other, Scalar alpha) const {
    static auto table = globalATenDispatch().getOpTable(
        "aten::add(Tensor self, Tensor other, *, Scalar alpha=1) -> Tensor");
    return table->getOp<Tensor (const Tensor &, const Tensor &, Scalar)>(
        tensorTypeIdToBackend(type_id()), is_variable())(*this, other, alpha);
}

aten/src/ATen/native/BinaryOps.cpp

namespace at {
namespace native {
Tensor add(const Tensor& self, const Tensor& other, Scalar alpha) {
    Tensor result;
    auto iter = TensorIterator::binary_op(result, self, other);
    add_stub(iter->device_type(), *iter, alpha);
    return iter->output();
}

aten/src/ATen/native/TensorIterator.cpp

std::unique_ptr<TensorIterator> TensorIterator::binary_op(Tensor& out,
    const Tensor& a, const Tensor& b) {
    auto builder = TensorIterator::Builder();
    builder.add_output(out);
    builder.add_input(a);
    builder.add_input(b);
    return builder.build();

std::unique_ptr<TensorIterator> TensorIterator::Builder::build() {
    iter_->mark_outputs();
    iter_->compute_shape();
    iter_->compute_strides();
    iter_->reorder_dimensions();
    iter_->compute_types();
    iter_->allocate_outputs();

void TensorIterator::allocate_outputs() {
    for (int i = 0; i < num_outputs_; i++) {
        op.tensor = at::empty_strided(tensor_shape, tensor_stride, op.options());
    }
}

aten/src/ATen/native/BinaryOps.h

```

```

using binary_fn_alpha = void(*) (TensorIterator&, Scalar alpha);
DECLARE_DISPATCH(binary_fn_alpha, add_stub);

aten/src/ATen/native/cpu/BinaryOpsKernel.cpp

REGISTER_DISPATCH(add_stub, &add_kernel);

void add_kernel(TensorIterator& iter, Scalar alpha_scalar) {
    if (iter.dtype() == ScalarType::Bool) {
        cpu_kernel(iter, [=](bool a, bool b) -> bool { return a + b; });
    } else {
        AT_DISPATCH_ALL_TYPES(iter.dtype(), "add_cpu", [&]() {
            auto alpha = alpha_scalar.to<scalar_t>();
            auto alpha_vec = Vec256<scalar_t>(alpha);
            cpu_kernel_vec(iter,
                [=](scalar_t a, scalar_t b) -> scalar_t { return a + alpha * b; },
                [=](Vec256<scalar_t> a, Vec256<scalar_t> b) {
                    return vec256::fmadd(b, alpha_vec, a);
                });
        });
    }
}

```