

PyTorch uses `linrunner` to execute lints locally and in CI. This provides developers with a single command to run all the linters, and ensures consistency between the CI and local development environments. See the `linrunner` repo for more info.

To get started, run the following to install `linrunner`. Make sure you are at the root of the PyTorch repo:

```
pip install linrunner
linrunner init
```

This will install `linrunner` on your system and download all the necessary dependencies to run linters locally. Note that this will install new packages with `pip` and download binaries maintained by the PyTorch team. If you want to see what `linrunner init` will install, run `linrunner init --dry-run`.

After that, just run `linrunner` to lint your changes! For more docs, see the `linrunner` repo.