

Welcome to the PyTorch developer's wiki!

Please read our [\[\[best practices|Where or how should I add documentation\]\]](#) if you're interested in adding a page or making edits

User docs

- [Release notes](#)
- [PyTorch Versions](#)
- [Public API definition and documentation](#)

Onboarding

New to PyTorch? Don't know where to start?

- [\[\[Core Onboarding|Core Frontend Onboarding\]\]](#)

Developer docs

- [Developer FAQ](#)
- [\[\[Where should I add documentation?|Where or how should I add documentation\]\]](#)
- [PyTorch Data Flow and Interface Diagram](#)
- [Multiprocessing Technical Notes](#)
- [Software Architecture for c10](#)
- [PyTorch JIT IR format](#) (slightly out of date now)
- [TH to ATen porting guide](#)
- [Writing Python in C++ \(a manifesto\)](#)
- [Introducing Quantized Tensor](#)
- [Life of a Tensor](#)
- [How to use `TensorIterator`](#)
- [Running and writing tests](#)
- [Writing memory format aware operators](#)
- [Guide for adding type annotations to PyTorch](#)
- [The torch.fft module in PyTorch 1.7](#)
- [PyTorch-ONNX exporter](#)

Notes

- [Automatic Mixed Precision package](#)
- [Automatic Mixed Precision examples](#)
- [Autograd mechanics](#)
- [Broadcasting semantics](#)
- [CPU threading and TorchScript inference](#)
- [CUDA semantics](#)
- [Frequently Asked Questions](#)
- [Extending PyTorch](#)
- [Features for large-scale deployments](#)
- [Multiprocessing best practices](#)
- [Reproducibility](#)
- [Serialization semantics](#)
- [Windows FAQ](#)
- [Python Language Reference Coverage](#)
- [Complex Numbers](#)

- [Android](#)
- [iOS](#)
- [How-to: Writing PyTorch & Caffe2 Operators](#)
- [CUDA IPC Refcounting implementation explained](#)
- [Autograd](#)
- [Code Coverage Tool for Pytorch](#)
- [How to write tests using FileCheck](#)
- [PyTorch Release Scripts](#)
- [Serialized operator test framework](#)
- [Observers](#)
- [Snapdragon NPE Support](#)
- [Using TensorBoard in ifbpy](#)

Named Tensors

- [Named Tensors](#)
- [Named Tensors operator coverage](#)

Quantization

- [Introduction to Quantization](#)
- [Quantization Operation coverage](#)
- [Implementing native quantized ops](#)
- [Extend PyTorch Quantization to Custom Backends](#)

JIT/TorchScript

- [JIT Technical Overview](#)
- [Current workflow](#)
- [Static Runtime](#)
- [TorchScript serialization](#)
- [PyTorch Fuser](#)
- [Implementation reference for the CUDA PyTorch JIT Fuser](#)
- [TorchScript](#)
- [TorchScript Language Reference](#)
- [TorchScript Unsupported Pytorch Constructs](#)

Distributed

- [Distributed RPC Framework](#)
- [Distributed Autograd Design](#)
- [Remote Reference Protocol](#)
- [Distributed Data Parallel](#)
- [Distributed communication package](#)
- [Contributing to PyTorch Distributed](#)

C++

- [PyTorch with C++](#)
- [The C++ Frontend](#)
- [PyTorch C++ API](#)
- [Tensor basics](#)
- [Tensor Creation API](#)
- [Tensor Indexing API](#)
- [MaybeOwned<Tensor>](#)

- [Installing C++ Distributions of PyTorch](#)
- [Torch Library API](#)
- [libtorch](#)
- [C++ / Python API parity tracker](#)
- [TensorExpr C++ Tests](#)
- [JIT C++ Tests](#)
- [C++ Frontend Tests](#)
- [FAQ](#)
- [Best Practices to Edit and Compile Pytorch Source Code On Window](#)

Benchmarks

- [Distributed Data Parallel Benchmark](#)
- [Fast RNN benchmarks](#)
- [PyTorch/Caffe2 Operator Micro-benchmarks](#)
- [torch function micro-benchmarks](#)
- [Benchmarking tool for the autograd AP](#)
- [Modular Benchmarking Components](#)

DataLoader

- [DataPipe](#)
- [DataPipe test requirements](#)

Workflow docs

- [Continuous Integration](#)
- [Bot commands](#)
- [Code review values](#)
- [Lint as you type](#)
- [Pull request review etiquette](#)
- [Debugging with SSH on Github Actions](#)
- [Using hud.pytorch.org](#)

Community

- [Code of Conduct](#)
- [Contributing](#)
- [PyTorch Contribution Guide](#)
- [PyTorch Governance](#)

Archived

- [Breaking Changes from Variable and Tensor merge](#) (from 0.4 release)
- [Tensor API changes for Caffe2 developers](#) (from 1.0 release, plus some stuff on master)
- [Autograd and Fork](#)

Caffe2

- [Caffe2](#)
- [Building Caffe2](#)
- [Doxygen Notes](#)
- [Docker & Caffe2](#)
- [Caffe2 implementation of Open Neural Network Exchange \(ONNX\)](#)
- [Caffe2 ONNX op coverage](#)

- [nomnigraph](#)
- [Caffe2 & TensorRT integration](#)
- [Playground for Caffe2 Models](#)
- [How to run FakeLowP vs Glow tests](#)
- [Using ONNX and ATen to export models from PyTorch to Caffe2](#)
- [An ATen operator for Caffe2](#)