

## Changelog

### 2.0.8

- Remove unimplemented functions (`mz_zip_locate_file` and `mz_zip_locate_file_v2`)
- Add license, changelog, readme and example files to release zip
- Fix heap overflow to user buffer in `tinfl_status` `tinfl_decompress`
- Fix corrupt archive if uncompressed file smaller than 4 byte and file is added by `mz_zip_writer_add_mem*`

### 2.0.7

- Removed need in C++ compiler in cmake build
- Fixed loads of uninitialized value errors found with Valgrind by memsetting `m_dict` to 0 in `tdefl_init`.
- Fix resource leak in `mz_zip_reader_init_file_v2`
- Fix assert with `mz_zip_writer_add_mem*` w/`MZ_DEFAULT_COMPRESSION`
- cmake build: install library and headers
- Remove `_LARGEFILE64_SOURCE` requirement from apple defines for large files

### 2.0.6

- Improve `MZ_ZIP_FLAG_WRITE_ZIP64` documentation
- Remove check for `cur_archive_file_ofs > UINT_MAX`, because `cur_archive_file_ofs` is not used after this point
- Add cmake debug configuration
- Fix PNG height when creating png files
- Add “iterative” file extraction method based on `mz_zip_reader_extract_to_callback`.
- Option to use `memcpy` for unaligned data access
- Define processor/arch macros as zero if not set to one

### 2.0.4/2.0.5

- Fix compilation with the various omission compile definitions

### 2.0.3

- Fix GCC/clang compile warnings
- Added callback for periodic flushes (for ZIP file streaming)
- Use UTF-8 for file names in ZIP files per default

### 2.0.2

- Fix source backwards compatibility with 1.x
- Fix a ZIP bit not being set correctly

### 2.0.1

- Added some tests
- Added CI
- Make source code ANSI C compatible

### 2.0.0 beta

- Matthew Sitton merged to vogl ZIP64 changes. Miniz is now licensed as MIT since the vogl code base is MIT licensed
- Miniz is now split into several files
- Miniz does now not seek backwards when creating ZIP files. That is the ZIP files can be streamed
- Miniz automatically switches to the ZIP64 format the created ZIP files goes over ZIP file limits
- Similar to SQLite the Miniz source code is amalgamated into one `miniz.c/miniz.h` pair in a build step (`amalgamate.sh`). Please use `miniz.c/miniz.h` in your projects

### v1.16 BETA Oct 19, 2013

Still testing, this release is downloadable from [here](#). Two key inflator-only robustness and streaming related changes. Also merged in `tdefl_compressor_alloc()`, `tdefl_compressor_free()` helpers to make script bindings easier for `rustyzip`. I would greatly appreciate any help with testing or any feedback.

The inflator in raw (non-zlib) mode is now usable on gzip or similar streams that have a bunch of bytes following the raw deflate data (problem discovered by `rustyzip` author `williamw520`). This version should never read beyond the last byte of the raw deflate data independent of how many bytes you pass into the input buffer.

The inflator now has a new failure status `TINFL_STATUS_FAILED_CANNOT_MAKE_PROGRESS` (-4). Previously, if the inflator was starved of bytes and could not make progress (because the input buffer was empty and the caller did not set the `TINFL_FLAG_HAS_MORE_INPUT` flag - say on truncated or corrupted compressed data stream) it would append all 0's to the input and try to soldier on. This is scary behavior if the caller didn't know when to stop accepting output (because it didn't know how much uncompressed data was expected, or didn't enforce a sane maximum). v1.16 will instead return `TINFL_STATUS_FAILED_CANNOT_MAKE_PROGRESS` immediately if it needs 1 or more bytes to make progress, the input buf is empty, and the caller has indicated that no more input is available. This is a "soft" failure, so you can call the inflator again with more input and it will try to continue, or you can give up and fail. This could be very useful in network streaming scenarios.

- The inflator coroutine func. is subtle and complex so I'm being cautious about this release. I would greatly appreciate any help with testing or

any feedback. I feel good about these changes, and they've been through several hours of automated testing, but they will probably not fix anything for the majority of prev. users so I'm going to mark this release as beta for a few weeks and continue testing it at work/home on various things.

- The inflator in raw (non-zlib) mode is now usable on gzip or similar data streams that have a bunch of bytes following the raw deflate data (problem discovered by rustzip author williamw520). This version should *never* read beyond the last byte of the raw deflate data independent of how many bytes you pass into the input buffer. This issue was caused by the various Huffman bitbuffer lookahead optimizations, and would not be an issue if the caller knew and enforced the precise size of the raw compressed data *or* if the compressed data was in zlib format (i.e. always followed by the byte aligned zlib Adler32). So in other words, you can now call the inflator on deflate streams that are followed by arbitrary amounts of data and it's guaranteed that decompression will stop exactly on the last byte.
- The inflator now has a new failure status: `TINFL_STATUS_FAILED_CANNOT_MAKE_PROGRESS` (-4). Previously, if the inflator was starved of bytes and could not make progress (because the input buffer was empty and the caller did not set the `TINFL_FLAG_HAS_MORE_INPUT` flag - say on truncated or corrupted compressed data stream) it would append all 0's to the input and try to soldier on. This is scary, because in the worst case, I believe it was possible for the prev. inflator to start outputting large amounts of literal data. If the caller didn't know when to stop accepting output (because it didn't know how much uncompressed data was expected, or didn't enforce a sane maximum) it could continue forever. v1.16 cannot fall into this failure mode, instead it'll return `TINFL_STATUS_FAILED_CANNOT_MAKE_PROGRESS` immediately if it needs 1 or more bytes to make progress, the input buf is empty, and the caller has indicated that no more input is available. This is a "soft" failure, so you can call the inflator again with more input and it will try to continue, or you can give up and fail. This could be very useful in network streaming scenarios.
- Added documentation to all the tinfl return status codes, fixed miniz\_tester so it accepts double minus params for Linux, tweaked example1.c, added a simple "follower bytes" test to miniz\_tester.cpp. ### v1.15 r4 STABLE - Oct 13, 2013

Merged over a few very minor bug fixes that I fixed in the zip64 branch. This is downloadable from [here](#) and also in SVN head (as of 10/19/13).

### **v1.15 - Oct. 13, 2013**

Interim bugfix release while I work on the next major release with zip64 and streaming compression/decompression support. Fixed the `MZ_ZIP_FLAG_DO_NOT_SORT_CENTRAL_DIRECTORY` bug (thanks kahmyong.moon@hp.com), which could cause the locate files

func to not find files when this flag was specified. Also fixed a bug in `mz_zip_reader_extract_to_mem_no_alloc()` with user provided read buffers (thanks kymoon). I also merged lots of compiler fixes from various github repo branches and Google Code issue reports. I finally added cmake support (only tested under for Linux so far), compiled and tested with clang v3.3 and gcc 4.6 (under Linux), added `defl_write_image_to_png_file_in_memory_ex()` (supports Y flipping for OpenGL use, real-time compression), added a new PNG example (`example6.c` - Mandelbrot), and I added 64-bit file I/O support (`stat64()`, etc.) for glibc.

- Critical fix for the `MZ_ZIP_FLAG_DO_NOT_SORT_CENTRAL_DIRECTORY` bug (thanks kahmyong.moon@hp.com) which could cause locate files to not find files. This bug would only have occurred in earlier versions if you explicitly used this flag, OR if you used `mz_zip_extract_archive_file_to_heap()` or `mz_zip_add_mem_to_archive_file_in_place()` (which used this flag). If you can't switch to v1.15 but want to fix this bug, just remove the uses of this flag from both helper funcs (and of course don't use the flag).
- Bugfix in `mz_zip_reader_extract_to_mem_no_alloc()` from kymoon when `pUser_read_buf` is not NULL and compressed size is > uncompressed size
- Fixing `mz_zip_reader_extract_*`() funcs so they don't try to extract compressed data from directory entries, to account for weird zipfiles which contain zero-size compressed data on dir entries. Hopefully this fix won't cause any issues on weird zip archives, because it assumes the low 16-bits of zip external attributes are DOS attributes (which I believe they always are in practice).
- Fixing `mz_zip_reader_is_file_a_directory()` so it doesn't check the internal attributes, just the filename and external attributes
- `mz_zip_reader_init_file()` - missing `MZ_FCLOSE()` call if the seek failed
- Added cmake support for Linux builds which builds all the examples, tested with clang v3.3 and gcc v4.6.
- Clang fix for `tdefl_write_image_to_png_file_in_memory()` from toffaletti
- Merged `MZ_FORCEINLINE` fix from hdeanclark
- Fix `<time.h>` include before config `#ifdef`, thanks emil.brink
- Added `tdefl_write_image_to_png_file_in_memory_ex()`: supports Y flipping (super useful for OpenGL apps), and explicit control over the compression level (so you can set it to 1 for real-time compression).
- Merged in some compiler fixes from paulharris's github repro.
- Retested this build under Windows (VS 2010, including static analysis), tcc 0.9.26, gcc v4.6 and clang v3.3.
- Added `example6.c`, which dumps an image of the mandelbrot set to a PNG file.
- Modified `example2` to help test the `MZ_ZIP_FLAG_DO_NOT_SORT_CENTRAL_DIRECTORY` flag more.
- In r3: Bugfix to `mz_zip_writer_add_file()` found during merge: Fix possible src file `fclose()` leak if alignment bytes+local header file write

failed

- In r4: Minor bugfix to `mz_zip_writer_add_from_zip_reader()`: Was pushing the wrong central dir header offset, appears harmless in this release, but it became a problem in the zip64 branch

#### **v1.14 - May 20, 2012**

(SVN Only) Minor tweaks to get `miniz.c` compiling with the Tiny C Compiler, added `#ifndef MINIZ_NO_TIME` guards around `utime.h` includes. Adding `mz_free()` function, so the caller can free heap blocks returned by `miniz` using whatever heap functions it has been configured to use, MSVC specific fixes to use “safe” variants of several functions (`localtime_s`, `fopen_s`, `freopen_s`).

MinGW32/64 GCC 4.6.1 compiler fixes: added `MZ_FORCEINLINE`, `#include <time.h>` (thanks `fermtect`).

Compiler specific fixes, some from `fermtect`. I upgraded to TDM GCC 4.6.1 and now static `__forceinline` is giving it fits, so I’m changing all usage of `__forceinline` to `MZ_FORCEINLINE` and forcing gcc to use `attribute((always_inline))` (and MSVC to use `__forceinline`). Also various fixes from `fermtect` for MinGW32: added `#include`, 64-bit `ftell/fseek` fixes.

#### **v1.13 - May 19, 2012**

From `jason@cornsyrup.org` and `kelwert@mtu.edu` - Most importantly, fixed `mz_crc32()` so it doesn’t compute the wrong CRC-32’s when `mz_ulong` is 64-bits. Temporarily/locally slammed in “`typedef unsigned long mz_ulong`” and re-ran a randomized regression test on ~500k files. Other stuff:

Eliminated a bunch of warnings when compiling with GCC 32-bit/64. Ran all examples, `miniz.c`, and `tinfl.c` through MSVC 2008’s `/analyze` (static analysis) option and fixed all warnings (except for the silly “Use of the comma-operator in a tested expression.” analysis warning, which I purposely use to work around a MSVC compiler warning).

Created 32-bit and 64-bit Codeblocks projects/workspace. Built and tested Linux executables. The codeblocks workspace is compatible with Linux+Win32/x64. Added `miniz_tester` solution/project, which is a useful little app derived from LZHAM’s tester app that I use as part of the regression test. Ran `miniz.c` and `tinfl.c` through another series of regression testing on ~500,000 files and archives. Modified `example5.c` so it purposely disables a bunch of high-level functionality (`MINIZ_NO_STDIO`, etc.). (Thanks to `corysama` for the `MINIZ_NO_STDIO` bug report.)

Fix `ftell()` usage in a few of the examples so they exit with an error on files which are too large (a limitation of the examples, not `miniz` itself). Fix fail logic handling in `mz_zip_add_mem_to_archive_file_in_place()` so it always calls `mz_zip_writer_finalize_archive()` and `mz_zip_writer_end()`, even if the file add fails.

- From jason@cornsyrup.org and kelwert@mtu.edu - Fix `mz_crc32()` so it doesn't compute the wrong CRC-32's when `mz_ulong` is 64-bit.
- Temporarily/locally slammed in "typedef unsigned long `mz_ulong`" and re-ran a randomized regression test on ~500k files.
- Eliminated a bunch of warnings when compiling with GCC 32-bit/64.
- Ran all examples, `miniz.c`, and `tinfl.c` through MSVC 2008's `/analyze` (static analysis) option and fixed all warnings (except for the silly "Use of the comma-operator in a tested expression.." analysis warning, which I purposely use to work around a MSVC compiler warning).
- Created 32-bit and 64-bit Codeblocks projects/workspace. Built and tested Linux executables. The codeblocks workspace is compatible with Linux+Win32/x64.
- Added `miniz_tester` solution/project, which is a useful little app derived from LZHAM's tester app that I use as part of the regression test.
- Ran `miniz.c` and `tinfl.c` through another series of regression testing on ~500,000 files and archives.
- Modified `example5.c` so it purposely disables a bunch of high-level functionality (`MINIZ_NO_STDIO`, etc.). (Thanks to corysama for the `MINIZ_NO_STDIO` bug report.)
- Fix `ftell()` usage in examples so they exit with an error on files which are too large (a limitation of the examples, not `miniz` itself).

#### **v1.12 - 4/12/12**

More comments, added low-level `example5.c`, fixed a couple minor `level_and_flags` issues in the archive API's. `level_and_flags` can now be set to `MZ_DEFAULT_COMPRESSION`. Thanks to Bruce Dawson [bruced@valvesoftware.com](mailto:bruced@valvesoftware.com) for the feedback/bug report.

#### **v1.11 - 5/28/11**

Added statement from [unlicense.org](http://unlicense.org)

#### **v1.10 - 5/27/11**

- Substantial compressor optimizations:
- Level 1 is now ~4x faster than before. The L1 compressor's throughput now varies between 70-110MB/sec. on a Core i7 (actual throughput varies depending on the type of data, and x64 vs. x86).
- Improved baseline L2-L9 compression perf. Also, greatly improved compression perf. issues on some file types.
- Refactored the compression code for better readability and maintainability.
- Added level 10 compression level (L10 has slightly better ratio than level 9, but could have a potentially large drop in throughput on some files).

**v1.09 - 5/15/11**

Initial stable release.