

DO NOT READ THIS FILE ON GITHUB, GUIDES ARE PUBLISHED ON <https://guides.rubyonrails.org>.

Active Support Instrumentation

Active Support is a part of core Rails that provides Ruby language extensions, utilities, and other things. One of the things it includes is an instrumentation API that can be used inside an application to measure certain actions that occur within Ruby code, such as that inside a Rails application or the framework itself. It is not limited to Rails, however. It can be used independently in other Ruby scripts if it is so desired.

In this guide, you will learn how to use the instrumentation API inside of Active Support to measure events inside of Rails and other Ruby code.

After reading this guide, you will know:

- What instrumentation can provide.
- How to add a subscriber to a hook.
- The hooks inside the Rails framework for instrumentation.
- How to build a custom instrumentation implementation.

Introduction to instrumentation

The instrumentation API provided by Active Support allows developers to provide hooks which other developers may hook into. There are several of these within the [Rails framework](#). With this API, developers can choose to be notified when certain events occur inside their application or another piece of Ruby code.

For example, there is a hook provided within Active Record that is called every time Active Record uses an SQL query on a database. This hook could be **subscribed** to, and used to track the number of queries during a certain action. There's another hook around the processing of an action of a controller. This could be used, for instance, to track how long a specific action has taken.

You are even able to [create your own events](#) inside your application which you can later subscribe to.

Subscribing to an event

Subscribing to an event is easy. Use `ActiveSupport::Notifications.subscribe` with a block to listen to any notification.

The block receives the following arguments:

- The name of the event
- Time when it started
- Time when it finished
- A unique ID for the instrumenter that fired the event
- The payload (described in future sections)

```
ActiveSupport::Notifications.subscribe "process_action.action_controller" do |name,
  started, finished, unique_id, data|
  # your own custom stuff
  Rails.logger.info "#{name} Received! (started: #{started}, finished: #{finished})"
  # process_action.action_controller Received (started: 2019-05-05 13:43:57 -0800,
  finished: 2019-05-05 13:43:58 -0800)
end
```

If you are concerned about the accuracy of `started` and `finished` to compute a precise elapsed time then use `ActiveSupport::Notifications.monotonic_subscribe`. The given block would receive the same arguments as above but the `started` and `finished` will have values with an accurate monotonic time instead of wall-clock time.

```
ActiveSupport::Notifications.monotonic_subscribe "process_action.action_controller"
do |name, started, finished, unique_id, data|
  # your own custom stuff
  Rails.logger.info "#{name} Received! (started: #{started}, finished: #{finished})"
  # process_action.action_controller Received (started: 1560978.425334, finished:
  1560979.429234)
end
```

Defining all those block arguments each time can be tedious. You can easily create an `ActiveSupport::Notifications::Event` from block arguments like this:

```
ActiveSupport::Notifications.subscribe "process_action.action_controller" do |*args|
  event = ActiveSupport::Notifications::Event.new *args

  event.name      # => "process_action.action_controller"
  event.duration  # => 10 (in milliseconds)
  event.payload   # => {:extra=>information}

  Rails.logger.info "#{event} Received!"
end
```

You may also pass a block that accepts only one argument, and it will receive an event object:

```
ActiveSupport::Notifications.subscribe "process_action.action_controller" do |event|
  event.name      # => "process_action.action_controller"
  event.duration  # => 10 (in milliseconds)
  event.payload   # => {:extra=>information}

  Rails.logger.info "#{event} Received!"
end
```

Most times you only care about the data itself. Here is a shortcut to just get the data.

```
ActiveSupport::Notifications.subscribe "process_action.action_controller" do |*args|
  data = args.extract_options!
  data # { extra: :information }
end
```

You may also subscribe to events matching a regular expression. This enables you to subscribe to multiple events at once. Here's how to subscribe to everything from `ActionController`.

```
ActiveSupport::Notifications.subscribe /action_controller/ do |*args|
  # inspect all ActionController events
end
```

Rails framework hooks

Within the Ruby on Rails framework, there are a number of hooks provided for common events. These are detailed below.

Action Controller

write_fragment.action_controller

Key	Value
:key	The complete key

```
{
  key: 'posts/1-dashboard-view'
}
```

read_fragment.action_controller

Key	Value
:key	The complete key

```
{
  key: 'posts/1-dashboard-view'
}
```

expire_fragment.action_controller

Key	Value
:key	The complete key

```
{
  key: 'posts/1-dashboard-view'
}
```

exist_fragment?.action_controller

Key	Value
:key	The complete key

```
{
  key: 'posts/1-dashboard-view'
```

```
}
```

start_processing.action_controller

Key	Value
:controller	The controller name
:action	The action
:params	Hash of request parameters without any filtered parameter
:headers	Request headers
:format	html/js/json/xml etc
:method	HTTP request verb
:path	Request path

```
{
  controller: "PostsController",
  action: "new",
  params: { "action" => "new", "controller" => "posts" },
  headers: #<ActionDispatch::Http::Headers:0x0055a67a519b88>,
  format: :html,
  method: "GET",
  path: "/posts/new"
}
```

process_action.action_controller

Key	Value
:controller	The controller name
:action	The action
:params	Hash of request parameters without any filtered parameter
:headers	Request headers
:format	html/js/json/xml etc
:method	HTTP request verb
:path	Request path
:request	The ActionDispatch::Request
:response	The ActionDispatch::Response
:status	HTTP status code
:view_runtime	Amount spent in view in ms
:db_runtime	Amount spent executing database queries in ms

```
{
  controller: "PostsController",
  action: "index",
  params: {"action" => "index", "controller" => "posts"},
  headers: #<ActionDispatch::Http::Headers:0x0055a67a519b88>,
  format: :html,
  method: "GET",
  path: "/posts",
  request: #<ActionDispatch::Request:0x00007ff1cb9bd7b8>,
  response: #<ActionDispatch::Response:0x00007f8521841ec8>,
  status: 200,
  view_runtime: 46.848,
  db_runtime: 0.157
}
```

send_file.action_controller

Key	Value
:path	Complete path to the file

INFO. Additional keys may be added by the caller.

send_data.action_controller

`ActionController` does not add any specific information to the payload. All options are passed through to the payload.

redirect_to.action_controller

Key	Value
:status	HTTP response code
:location	URL to redirect to
:request	The <code>ActionDispatch::Request</code>

```
{
  status: 302,
  location: "http://localhost:3000/posts/new",
  request: #<ActionDispatch::Request:0x00007ff1cb9bd7b8>
}
```

halted_callback.action_controller

Key	Value
:filter	Filter that halted the action

```
{
  filter: ":halting_filter"
}
```

unpermitted_parameters.action_controller

Key	Value
:keys	The unpermitted keys
:context	Hash with the following keys: :controller, :action, :params, :request

Action Dispatch

process_middleware.action_dispatch

Key	Value
:middleware	Name of the middleware

Action View

render_template.action_view

Key	Value
:identifier	Full path to template
:layout	Applicable layout

```
{
  identifier: "/Users/adam/projects/notifications/app/views/posts/index.html.erb",
  layout: "layouts/application"
}
```

render_partial.action_view

Key	Value
:identifier	Full path to template

```
{
  identifier: "/Users/adam/projects/notifications/app/views/posts/_form.html.erb"
}
```

render_collection.action_view

Key	Value
:identifier	Full path to template
:count	Size of collection

:cache_hits	Number of partials fetched from cache
-------------	---------------------------------------

:cache_hits is only included if the collection is rendered with cached: true .

```
{
  identifier: "/Users/adam/projects/notifications/app/views/posts/_post.html.erb",
  count: 3,
  cache_hits: 0
}
```

render_layout.action_view

Key	Value
:identifier	Full path to template

```
{
  identifier:
"/Users/adam/projects/notifications/app/views/layouts/application.html.erb"
}
```

Active Record

sql.active_record

Key	Value
:sql	SQL statement
:name	Name of the operation
:connection	Connection object
:binds	Bind parameters
:type_casted_binds	Typecasted bind parameters
:statement_name	SQL Statement name
:cached	true is added when cached queries used

INFO. The adapters will add their own data as well.

```
{
  sql: "SELECT \"posts\".* FROM \"posts\" ",
  name: "Post Load",
  connection: #
<ActiveRecord::ConnectionAdapters::SQLite3Adapter:0x00007f9f7a838850>,
  binds: [#<ActiveModel::Attribute::WithCastValue:0x00007fe19d15dc00>],
  type_casted_binds: [11],
  statement_name: nil
}
```

instantiation.active_record

Key	Value
:record_count	Number of records that instantiated
:class_name	Record's class

```
{
  record_count: 1,
  class_name: "User"
}
```

Action Mailer

deliver.action_mailer

Key	Value
:mailer	Name of the mailer class
:message_id	ID of the message, generated by the Mail gem
:subject	Subject of the mail
:to	To address(es) of the mail
:from	From address of the mail
:bcc	BCC addresses of the mail
:cc	CC addresses of the mail
:date	Date of the mail
:mail	The encoded form of the mail
:perform_deliveries	Whether delivery of this message is performed or not

```
{
  mailer: "Notification",
  message_id: "4f5b5491f1774_181b23fc3d4434d38138e5@mba.local.mail",
  subject: "Rails Guides",
  to: ["users@rails.com", "dhh@rails.com"],
  from: ["me@rails.com"],
  date: Sat, 10 Mar 2012 14:18:09 +0100,
  mail: "...", # omitted for brevity
  perform_deliveries: true
}
```

process.action_mailer

Key	Value
:mailer	Name of the mailer class

:action	The action
:args	The arguments

```
{
  mailer: "Notification",
  action: "welcome_email",
  args: []
}
```

Active Support

cache_read.active_support

Key	Value
:key	Key used in the store
:store	Name of the store class
:hit	If this read is a hit
:super_operation	:fetch is added when a read is used with #fetch

cache_generate.active_support

This event is only used when `#fetch` is called with a block.

Key	Value
:key	Key used in the store
:store	Name of the store class

INFO. Options passed to fetch will be merged with the payload when writing to the store

```
{
  key: "name-of-complicated-computation",
  store: "ActiveSupport::Cache::MemCacheStore"
}
```

cache_fetch_hit.active_support

This event is only used when `#fetch` is called with a block.

Key	Value
:key	Key used in the store
:store	Name of the store class

INFO. Options passed to fetch will be merged with the payload.

```
{
  key: "name-of-complicated-computation",
  store: "ActiveSupport::Cache::MemCacheStore"
}
```

cache_write.active_support

Key	Value
:key	Key used in the store
:store	Name of the store class

INFO. Cache stores may add their own keys

```
{
  key: "name-of-complicated-computation",
  store: "ActiveSupport::Cache::MemCacheStore"
}
```

cache_delete.active_support

Key	Value
:key	Key used in the store
:store	Name of the store class

```
{
  key: "name-of-complicated-computation",
  store: "ActiveSupport::Cache::MemCacheStore"
}
```

cache_exist?.active_support

Key	Value
:key	Key used in the store
:store	Name of the store class

```
{
  key: "name-of-complicated-computation",
  store: "ActiveSupport::Cache::MemCacheStore"
}
```

Active Job

enqueue_at.active_job

Key	Value
-----	-------

:adapter	QueueAdapter object processing the job
:job	Job object

enqueue.active_job

Key	Value
:adapter	QueueAdapter object processing the job
:job	Job object

enqueue_retry.active_job

Key	Value
:job	Job object
:adapter	QueueAdapter object processing the job
:error	The error that caused the retry
:wait	The delay of the retry

perform_start.active_job

Key	Value
:adapter	QueueAdapter object processing the job
:job	Job object

perform.active_job

Key	Value
:adapter	QueueAdapter object processing the job
:job	Job object

retry_stopped.active_job

Key	Value
:adapter	QueueAdapter object processing the job
:job	Job object
:error	The error that caused the retry

discard.active_job

Key	Value
:adapter	QueueAdapter object processing the job
:job	Job object

<code>:error</code>	The error that caused the discard
---------------------	-----------------------------------

Action Cable

`perform_action.action_cable`

Key	Value
<code>:channel_class</code>	Name of the channel class
<code>:action</code>	The action
<code>:data</code>	A hash of data

`transmit.action_cable`

Key	Value
<code>:channel_class</code>	Name of the channel class
<code>:data</code>	A hash of data
<code>:via</code>	Via

`transmit_subscription_confirmation.action_cable`

Key	Value
<code>:channel_class</code>	Name of the channel class

`transmit_subscription_rejection.action_cable`

Key	Value
<code>:channel_class</code>	Name of the channel class

`broadcast.action_cable`

Key	Value
<code>:broadcasting</code>	A named broadcasting
<code>:message</code>	A hash of message
<code>:coder</code>	The coder

Active Storage

`service_upload.active_storage`

Key	Value
<code>:key</code>	Secure token
<code>:service</code>	Name of the service
<code>:checksum</code>	Checksum to ensure integrity

service_streaming_download.active_storage

Key	Value
:key	Secure token
:service	Name of the service

service_download_chunk.active_storage

Key	Value
:key	Secure token
:service	Name of the service
:range	Byte range attempted to be read

service_download.active_storage

Key	Value
:key	Secure token
:service	Name of the service

service_delete.active_storage

Key	Value
:key	Secure token
:service	Name of the service

service_delete_prefixed.active_storage

Key	Value
:prefix	Key prefix
:service	Name of the service

service_exist.active_storage

Key	Value
:key	Secure token
:service	Name of the service
:exist	File or blob exists or not

service_url.active_storage

Key	Value
:key	Secure token

:service	Name of the service
:url	Generated URL

service_update_metadata.active_storage

Key	Value
:key	Secure token
:service	Name of the service
:content_type	HTTP Content-Type field
:disposition	HTTP Content-Disposition field

INFO. The only ActiveStorage service that provides this hook so far is GCS.

preview.active_storage

Key	Value
:key	Secure token

transform.active_storage

analyze.active_storage

Key	Value
:analyzer	Name of analyzer e.g., ffprobe

Action Mailbox

process.action_mailbox

Key	Value
:mailbox	Instance of the Mailbox class inheriting from ActionMailbox::Base
:inbound_email	Hash with data about the inbound email being processed

```
{
  mailbox: #<RepliesMailbox:0x00007f9f7a8388>,
  inbound_email: {
    id: 1,
    message_id: "0CB459E0-0336-41DA-BC88-E6E28C697DDB@37signals.com",
    status: "processing"
  }
}
```

Railties

load_config_initializer.railties

--	--

Key	Value
<code>:initializer</code>	Path to loaded initializer from <code>config/initializers</code>

Rails

`deprecation.rails`

Key	Value
<code>:message</code>	The deprecation warning
<code>:callstack</code>	Where the deprecation came from

Exceptions

If an exception happens during any instrumentation the payload will include information about it.

Key	Value
<code>:exception</code>	An array of two elements. Exception class name and the message
<code>:exception_object</code>	The exception object

Creating custom events

Adding your own events is easy as well. `ActiveSupport::Notifications` will take care of all the heavy lifting for you. Simply call `instrument` with a `name`, `payload` and a block. The notification will be sent after the block returns. `ActiveSupport` will generate the start and end times and add the instrumenter's unique ID. All data passed into the `instrument` call will make it into the payload.

Here's an example:

```
ActiveSupport::Notifications.instrument "my.custom.event", this: :data do
  # do your custom stuff here
end
```

Now you can listen to this event with:

```
ActiveSupport::Notifications.subscribe "my.custom.event" do |name, started,
  finished, unique_id, data|
  puts data.inspect # {:this=>:data}
end
```

You also have the option to call `instrument` without passing a block. This lets you leverage the instrumentation infrastructure for other messaging uses.

```
ActiveSupport::Notifications.instrument "my.custom.event", this: :data

ActiveSupport::Notifications.subscribe "my.custom.event" do |name, started,
  finished, unique_id, data|
```

```
puts data.inspect # {:this=>:data}
end
```

You should follow Rails conventions when defining your own events. The format is: `event.library` . If your application is sending Tweets, you should create an event named `tweet.twitter` .