

Line number information

Line number information maps between machine code instructions and the source level location.

Encoding

The line number information is stored in the `.debug_line` section for ELF and `__debug_line` section of the `__DWARF` segment for Mach-O object files. The line number information contains a header followed by the line program. The line program is a program for a virtual machine with instructions like set line number for the current machine code instruction and advance the current machine code instruction.

Tips

You need to set either `DW_AT_low_pc` and `DW_AT_high_pc` or `DW_AT_ranges` of a `DW_TAG_compilation_unit` to the range of addresses in the compilation unit. After that you need to set `DW_AT_stmt_list` to the `.debug_line` section offset of the line program. Otherwise a debugger won't find the line number information. On macOS the debuginfo relocations **must** be section relative and not symbol relative. See #303 (comment) for more information.

Function debuginfo

Tips

`DW_TAG_subprogram` requires `DW_AT_name`, `DW_AT_low_pc` and `DW_AT_high_pc` or `DW_AT_ranges`. Otherwise gdb will silently skip it. When `DW_AT_high_pc` is a length instead of an address, the DWARF version must be at least 4.

IRC log of #gdb on irc.freenode.org at 2020-04-23

```
(13:46:11) bjorn3: i am writing a backend for a compiler that uses DWARF for debuginfo. for
(13:47:49) bjorn3: this is the output of llvm-dwarfdump: https://gist.github.com/bjorn3/8a34
(13:47:50) osa1: luispm: why is that problem not exists in 'commands'? (the target vs. host)
(13:52:16) luispm: osa1, commands is a bit more high level. It executes isolated commands.
(13:52:36) luispm: osa1, Oh, i see your point now. Commands is only executed on the host.
(13:53:18) luispm: osa1, The commands are not tied to the execution context of the debugged
(13:55:00) luispm: bjorn3, Likely something GDB thinks is wrong. Does enabling "set debug dw
(13:56:01) bjorn3: luispm: no
(13:56:12) bjorn3: for more context: https://github.com/bjorn3/rustc_codegen_cranelift/pull/
(13:58:16) osa1 verliet de ruimte (quit: Quit: osa1).
(13:58:28) bjorn3: luispm: wait, for b m<TAB> it shows nothing, but when stepping into a new
(13:58:45) bjorn3: it still doesn't show anything for `info args` though
(13:58:50) bjorn3: No symbol table info available.
(14:00:50) luispm: bjorn3, Is that expected given the nature of the binary?
```

(14:01:17) bjorn3: b main<TAB> may show nothing as I only set DW_AT_linkage_name and not DW
 (14:01:24) bjorn3: info args should work though
 (14:03:26) luispm: Sorry, I'm not sure what's up. There may be a genuine bug there.
 (14:03:41) luispm: tromeey (not currently in the channel, but maybe later today) may have mor
 (14:04:08) bjorn3: okay, thanks luispm!
 (14:04:27) luispm: In the worst case, reporting a bug may prompt someone to look into that a
 (14:04:48) luispm: Or send an e-mail to the gdb@sourceware.org mailing list.
 (14:05:11) bjorn3: I don't know if it is a bug in gdb, or just me producing (slightly) wrong
 (14:39:40) irker749: gdb: tom binutils-gdb.git:master * 740480b88af / gdb/ChangeLog gdb/darv
 (15:22:45) irker749: gdb: tromeey binutils-gdb.git:master * ecc6c6066b5 / gdb/ChangeLog gdb/o
 (15:23:13) bjorn3: tromeey: ping
 (15:23:29) tromeey: bjorn3: hey
 (15:24:16) bjorn3: I am writing a backend for a compiler which uses DWARF for debuginfo. I u
 (15:25:13) bjorn3: it just says: No symbol table info available.
 (15:25:21) bjorn3: any idea what it could be?
 (15:25:34) bjorn3: dwarfdump output: <https://gist.github.com/bjorn3/8a34e333c80f13cb048381e>
 (15:26:48) bjorn3: more context: https://github.com/bjorn3/rustc_codegen_cranefift/pull/978
 (15:28:05) tromeey: offhand I don't know, but if you can send me an executable I can look
 (15:28:17) bjorn3: how should I send it?
 (15:29:26) tromeey: good question
 (15:29:41) tromeey: you could try emailing it to tromeey at adacore.com
 (15:29:47) tromeey: dunno if that will work or not
 (15:30:26) bjorn3: i will try
 (15:37:27) bjorn3: tromeey: i sent an email with the subject "gdb args not showing"
 (15:38:29) tromeey: will check now
 (15:38:40) bjorn3: thanks!
 (15:42:51) irker749: gdb: tdevries binutils-gdb.git:master * de82891ce5b / gdb/ChangeLog gdb
 (15:42:52) irker749: gdb: tdevries binutils-gdb.git:master * 70bc38f5138 / gdb/ChangeLog gdb
 (15:43:36) tromeey: bjorn3: sorry, got distracted. I have the file now
 (15:45:35) tromeey: my first thing when investigating was to enable complaints
 (15:45:37) tromeey: so I did
 (15:45:40) tromeey: set complaints 1000
 (15:45:42) tromeey: then
 (15:45:51) tromeey: file -readnow mini_core_hello_world
 (15:46:00) tromeey: gdb printed just one style of complaint
 (15:46:07) tromeey: During symbol reading: missing name for subprogram DIE at 0x3f7
 (15:46:18) tromeey: (which is really pretty good, most compilers manage to generate a bunch)
 (15:46:29) tromeey: and then the gdb DWARF reader says
 (15:46:34) tromeey: /* Ignore functions with missing or empty names. These are actually
 (15:46:34) tromeey: illegal according to the DWARF standard. */
 (15:46:34) tromeey: if (name == NULL)
 (15:46:34) tromeey: {
 (15:46:37) tromeey: complaint (_("missing name for subprogram DIE at %s"),
 (15:46:40) tromeey: sect_offset_str (die->sect_off));
 (15:46:47) tromeey: I wonder if that comment is correct though
 (15:47:34) tromeey: I guess pedantically maybe it is, DWARF 5 3.3.1 says

(15:47:43) tromeY: The subroutine or entry point entry has a DW_AT_name attribute whose value is a null-terminated string containing the subroutine or entry point name.
 (15:47:43) tromeY: a null-terminated string containing the subroutine or entry point name.
 (15:48:14) bjorn3: i tried set complaints, but it returned complaints for system files. i don't know what to do.
 (15:48:21) tromeY: cool
 (15:48:26) bjorn3: i will try adding DW_AT_name
 (15:48:45) tromeY: without readnow unfortunately you get less stuff, because for whatever reason readnow is broken.
 (15:49:02) tromeY: sort of anyway
 (15:49:43) tromeY: this seems kind of pedantic of gdb, like if there's a linkage name but no DW_AT_name.
 (15:50:01) tromeY: also what about anonymous functions
 (15:50:17) tromeY: but anyway this explains the current situation and if you don't mind adding DW_AT_name.
 (15:51:47) bjorn3: i added DW_AT_name.
 (15:51:54) bjorn3: now it says cannot get low and high bounds for subprogram DIE at ...
 (15:52:01) tromeY: ugh
 (15:52:10) bjorn3: i will add DW_AT_low_pc and DW_AT_high_pc
 (15:52:15) tromeY: /* Ignore functions with missing or invalid low and high pc attributes.
 (15:52:37) tromeY: you can also use DW_AT_ranges
 (15:52:55) tromeY: if you'd prefer
 (15:53:08) bjorn3: already using DW_AT_ranges for DW_TAG_compilation_unit
 (15:53:19) bjorn3: for individual functions, there are no gaps
 (15:57:07) bjorn3: still the same error with DW_AT_low_pc and DW_AT_high_pc
 (15:57:24) bjorn3: tromeY: ^
 (15:58:08) tromeY: hmmm
 (15:58:30) bjorn3: should i send the new executable?
 (15:58:31) tromeY: send me another executable & I will debug
 (15:58:33) tromeY: yep
 (15:59:23) bjorn3: sent as reply of the previous mail
 (16:03:23) tromeY: the low PC has DW_FORM_addr, but the high PC has DW_FORM_udata, which seems odd.
 (16:03:50) mJw: no
 (16:03:54) tromeY: no?
 (16:04:00) mJw: I suggested that for the DWARF standard...
 (16:04:05) mJw: sorry
 (16:04:58) mJw: The idea was that instead of two relocations and two address wide fields, you could use a single relocation and a single address wide field.
 (16:05:05) tromeY: ahh, I see the code now
 (16:05:07) tromeY: I forgot about this
 (16:05:18) tromeY: if (cu->header.version >= 4 && attr_high->form_is_constant ())
 (16:05:18) tromeY: high += low;
 (16:05:36) mJw: that second offset doesn't need a relocation and can often be packed in some way.
 (16:05:51) mJw: using udata might not be ideal though, but is allowed
 (16:05:51) tromeY: bjorn3: the problem is that this CU claims to be DWARF 3 but is using a DWARF 4 attribute.
 (16:05:58) mJw: aha
 (16:05:59) bjorn3: which one?
 (16:06:03) ryoshu: hi
 (16:06:08) tromeY: high_pc (udata) 107 (+0x000000000000011b0 <_ZN2...
 (16:06:08) tromeY: just soft ping, I have a queue of patches :)
 (16:06:12) ryoshu: just soft ping, I have a queue of patches :)
 (16:06:22) tromeY: using this as a length requires DWARF 4

(16:06:36) tromeey: for gdb at least it's fine to always emit DWARF 4
(16:06:44) bjorn3: trying dwarf 4 now
(16:06:48) tromeey: I think there are some DWARF 5 features still in the works but DWARF 4 sh
(16:07:03) tromeey: fini
(16:07:08) tromeey: lol wrong window
(16:07:56) mjlw: Maybe you can accept it for DWARF < 4. But if I remember correctly it might
(16:08:13) tromeey: yeah, I vaguely recall this as well, though I'd expect there to be a comm
(16:08:21) mjlw: Cannot really remember why it needed version >= 4. Maybe there was no good r
(16:08:32) bjorn3: tromeey: it works!!!! thanks for all the help!
(16:08:41) tromeey: my pleasure bjorn3